

Learn coding and testing with puzzles and games

Tiny Python Projects

Ken Youens-Clark





Save 35% at [manning.com!](https://manning.com)

Use the code **humble535** at checkout to save on your first purchase at manning.com.

Manning publishes high-quality books and courses for technology professionals. Here are a few Manning advantages:

Early access

Don't wait to start learning! In MEAP, the [Manning Early Access Program](#), you read books while they're being written.

Access anywhere with liveBook

The [Manning liveBook platform](#) provides instant browser-based access to our content.

Beyond books

Cutting edge liveProjects, liveAudio, and liveVideo courses give you new ways to learn. Only available at manning.com

Impeccable quality

We believe in excellence. Our customers tell us we produce the highest quality content you can buy.

Exclusive eBooks

Manning eBooks are only available from manning.com. You won't find them anywhere else.

Shop now at manning.com



Register to get exclusive offers and updates!

Tiny Python Projects

KEN YOUENS-CLARK



MANNING
SHELTER ISLAND

For online information and ordering of this and other Manning books, please visit www.manning.com. The publisher offers discounts on this book when ordered in quantity. For more information, please contact

Special Sales Department
Manning Publications Co.
20 Baldwin Road
PO Box 761
Shelter Island, NY 11964
Email: orders@manning.com

©2020 by Manning Publications Co. All rights reserved.

No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by means electronic, mechanical, photocopying, or otherwise, without prior written permission of the publisher.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in the book, and Manning Publications was aware of a trademark claim, the designations have been printed in initial caps or all caps.

⊗ Recognizing the importance of preserving what has been written, it is Manning's policy to have the books we publish printed on acid-free paper, and we exert our best efforts to that end. Recognizing also our responsibility to conserve the resources of our planet, Manning books are printed on paper that is at least 15 percent recycled and processed without the use of elemental chlorine.

 Manning Publications Co.
20 Baldwin Road
PO Box 761
Shelter Island, NY 11964

Development editor: Elesha Hyde
Technical development editor: Al Scherer
Review editor: Aleksandar Dragosavljević
Production editor: Deirdre S. Hiam
Copy editor: Andy Carroll
Proofreader: Katie Tennant
Technical proofreader: Mathijs Affourtit
Typesetter: Dennis Dalinnik
Cover designer: Marija Tudor

ISBN: 9781617297519
Printed in the United States of America

brief contents

- | | |
|--|-----|
| Getting started: Introduction and installation guide | 1 |
| 1 ■ How to write and test a Python program | 15 |
| 2 ■ The crow's nest: Working with strings | 35 |
| 3 ■ Going on a picnic: Working with lists | 55 |
| 4 ■ Jump the Five: Working with dictionaries | 76 |
| 5 ■ Howler: Working with files and STDOUT | 92 |
| 6 ■ Words count: Reading files and STDIN, iterating lists,
formatting strings | 107 |
| 7 ■ Gashlycrumb: Looking items up in a dictionary | 118 |
| 8 ■ Apples and Bananas: Find and replace | 128 |
| 9 ■ Dial-a-Curse: Generating random insults from
lists of words | 150 |
| 10 ■ Telephone: Randomly mutating strings | 165 |
| 11 ■ Bottles of Beer Song: Writing and
testing functions | 178 |
| 12 ■ Ransom: Randomly capitalizing text | 195 |
| 13 ■ Twelve Days of Christmas: Algorithm design | 207 |

- 14 ■ Rhymer: Using regular expressions to create rhyming words 225
- 15 ■ The Kentucky Friar: More regular expressions 248
- 16 ■ The Scrambler: Randomly reordering the middles of words 268
- 17 ■ Mad Libs: Using regular expressions 281
- 18 ■ Gematria: Numeric encoding of text using ASCII values 295
- 19 ■ Workout of the Day: Parsing CSV files, creating text table output 311
- 20 ■ Password strength: Generating a secure and memorable password 331
- 21 ■ Tic-Tac-Toe: Exploring state 351
- 22 ■ Tic-Tac-Toe redux: An interactive version with type hints 367

contents

preface xv
acknowledgments xvii
about this book xix
about the author xxii
about the cover xxiii

Getting started: Introduction and installation guide 1

- Writing command-line programs 1
- Using test-driven development 4
- Setting up your environment 4
- Code examples 5
- Getting the code 8
- Installing modules 10
- Code formatters 10
- Code linters 11
- How to start writing new programs 11
- Why not Notebooks? 12
- The scope of topics we'll cover 12
- Why not object-oriented programming? 13
- A note about the lingo 13

1 How to write and test a Python program 15

- 1.1 Creating your first program 15
- 1.2 Comment lines 16
- 1.3 Testing your program 17
- 1.4 Adding the #! (shebang) line 18
- 1.5 Making a program executable 20
- 1.6 Understanding \$PATH 20
 - Altering your \$PATH* 21
- 1.7 Adding a parameter and help 22
- 1.8 Making the argument optional 24
- 1.9 Running our tests 26
- 1.10 Adding the main() function 26
- 1.11 Adding the get_args() function 27
 - Checking style and errors* 28
- 1.12 Testing hello.py 29
- 1.13 Starting a new program with new.py 30
- 1.14 Using template.py as an alternative to new.py 33

2 The crow's nest: Working with strings 35

- 2.1 Getting started 36
 - How to use the tests* 36 ▪ *Creating programs with new.py* 37
 - Write, test, repeat* 38 ▪ *Defining your arguments* 39
 - Concatenating strings* 41 ▪ *Variable types* 42 ▪ *Getting just part of a string* 43 ▪ *Finding help in the REPL* 44
 - String methods* 44 ▪ *String comparisons* 45 ▪ *Conditional branching* 47 ▪ *String formatting* 48 ▪ *Time to write* 49
- 2.2 Solution 49
- 2.3 Discussion 50
 - Defining the arguments with get_args()* 50 ▪ *The main() thing* 51 ▪ *Classifying the first character of a word* 51
 - Printing the results* 52 ▪ *Running the test suite* 52
- 2.4 Going further 53

3 Going on a picnic: Working with lists 55

- 3.1 Starting the program 56
- 3.2 Writing picnic.py 58

3.3	Introducing lists	59
	<i>Adding one element to a list</i>	60
	<i>Adding many elements to a list</i>	61
	<i>Indexing lists</i>	63
	<i>Slicing lists</i>	64
	<i>Finding elements in a list</i>	64
	<i>Removing elements from a list</i>	65
	<i>Sorting and reversing a list</i>	67
	<i>Lists are mutable</i>	69
	<i>Joining a list</i>	70
3.4	Conditional branching with if/elif/else	70
	<i>Time to write</i>	71
3.5	Solution	71
3.6	Discussion	73
	<i>Defining the arguments</i>	73
	<i>Assigning and sorting the items</i>	73
	<i>Formatting the items</i>	73
	<i>Printing the items</i>	74
3.7	Going further	75

4 *Jump the Five: Working with dictionaries* 76

4.1	Dictionaries	77
	<i>Creating a dictionary</i>	78
	<i>Accessing dictionary values</i>	80
	<i>Other dictionary methods</i>	81
4.2	Writing jump.py	82
4.3	Solution	84
4.4	Discussion	85
	<i>Defining the parameters</i>	85
	<i>Using a dict for encoding</i>	85
	<i>Various ways to process items in a series</i>	86
	<i>(Not) using str.replace()</i>	90
4.5	Going further	91

5 *Howler: Working with files and STDOUT* 92

5.1	Reading files	93
5.2	Writing files	97
5.3	Writing howler.py	99
5.4	Solution	101
5.5	Discussion	102
	<i>Defining the arguments</i>	102
	<i>Reading input from a file or the command line</i>	103
	<i>Choosing the output file handle</i>	104
	<i>Printing the output</i>	104
	<i>A low-memory version</i>	104
5.6	Going further	106

6 Words count: Reading files and STDIN, iterating lists, formatting strings 107

6.1 Writing wc.py 109

Defining file inputs 110 ▪ Iterating lists 111 ▪ What you're counting 111 ▪ Formatting your results 112

6.2 Solution 114

6.3 Discussion 115

Defining the arguments 115 ▪ Reading a file using a for loop 115

6.4 Going further 117

7 Gashlycrumb: Looking items up in a dictionary 118

7.1 Writing gashlycrumb.py 119

7.2 Solution 122

7.3 Discussion 123

Handling the arguments 123 ▪ Reading the input file 124
Using a dictionary comprehension 125 ▪ Dictionary lookups 126

7.4 Going further 126

8 Apples and Bananas: Find and replace 128

8.1 Altering strings 130

Using the str.replace() method 131 ▪ Using str.translate() 131
Other ways to mutate strings 132

8.2 Solution 133

8.3 Discussion 134

Defining the parameters 134 ▪ Eight ways to replace the vowels 135

8.4 Refactoring with tests 149

8.5 Going further 149

9 Dial-a-Curse: Generating random insults from lists of words 150

9.1 Writing abuse.py 151

Validating arguments 153 ▪ Importing and seeding the random module 154 ▪ Defining the adjectives and nouns 155 ▪ Taking random samples and choices 156 ▪ Formatting the output 156

9.2	Solution	157
9.3	Discussion	159
	<i>Defining the arguments</i>	159
	<i>Using parser.error()</i>	160
	<i>Program exit values and STDERR</i>	160
	<i>Controlling randomness with random.seed()</i>	161
	<i>Iterating with range() and using throwaway variables</i>	162
	<i>Constructing the insults</i>	162
9.4	Going further	163

10 Telephone: Randomly mutating strings 165

10.1	Writing telephone.py	167
	<i>Calculating the number of mutations</i>	168
	<i>The mutation space</i>	169
	<i>Selecting the characters to mutate</i>	169
	<i>Mutating a string</i>	172
	<i>Time to write</i>	173
10.2	Solution	173
10.3	Discussion	175
	<i>Mutating a string</i>	175
	<i>Using a list instead of a str</i>	176
10.4	Going further	177

11 Bottles of Beer Song: Writing and testing functions 178

11.1	Writing bottles.py	179
	<i>Counting down</i>	180
	<i>Writing a function</i>	181
	<i>Writing a test for verse()</i>	182
	<i>Using the verse() function</i>	186
11.2	Solution	187
11.3	Discussion	189
	<i>Counting down</i>	189
	<i>Test-driven development</i>	189
	<i>The verse() function</i>	190
	<i>Iterating through the verses</i>	191
	<i>1,500 other solutions</i>	194
11.4	Going further	194

12 Ransom: Randomly capitalizing text 195

12.1	Writing ransom.py	197
	<i>Mutating the text</i>	197
	<i>Flipping a coin</i>	198
	<i>Creating a new string</i>	198
12.2	Solution	199
12.3	Discussion	200
	<i>Iterating through elements in a sequence</i>	200
	<i>Writing a function to choose the letter</i>	202
	<i>Another way to write list.append()</i>	202
	<i>Using a str instead of a list</i>	203
	<i>Using a list comprehension</i>	203
	<i>Using a map() function</i>	204

- 12.4 Comparing methods 204
- 12.5 Going further 205

13 *Twelve Days of Christmas: Algorithm design 207*

- 13.1 Writing twelve_days.py 208
 - Counting 209 ▪ Creating the ordinal value 211 ▪ Making the verses 213 ▪ Using the verse() function 215 ▪ Printing 215 Time to write 215
- 13.2 Solution 216
- 13.3 Discussion 218
 - Making one verse 218 ▪ Generating the verses 221
 - Printing the verses 222
- 13.4 Going further 223

14 *Rhymer: Using regular expressions to create rhyming words 225*

- 14.1 Writing rhymer.py 227
 - Breaking a word 228 ▪ Using regular expressions 229
 - Using capture groups 232 ▪ Truthiness 236 ▪ Creating the output 238
- 14.2 Solution 238
- 14.3 Discussion 240
 - Stemming a word 240 ▪ Formatting and commenting the regular expression 242 ▪ Using the stemmer() function outside your program 243 ▪ Creating rhyming strings 244 ▪ Writing stemmer() without regular expressions 245
- 14.4 Going further 246

15 *The Kentucky Friar: More regular expressions 248*

- 15.1 Writing friar.py 250
 - Splitting text using regular expressions 251 ▪ Shorthand classes 252 ▪ Negated shorthand classes 254 ▪ Using re.split() with a captured regex 255 ▪ Writing the fry() function 256
 - Using the fry() function 261
- 15.2 Solution 262
- 15.3 Discussion 263
 - Writing the fry() function manually 264 ▪ Writing the fry() function with regular expressions 266
- 15.4 Going further 266

16 *The Scrambler: Randomly reordering the middles of words* 268

- 16.1 Writing scrambler.py 269
 - Breaking the text into lines and words* 270 ▪ *Capturing, non-capturing, and optional groups* 272 ▪ *Compiling a regex* 272
 - Scrambling a word* 273 ▪ *Scrambling all the words* 275
- 16.2 Solution 276
- 16.3 Discussion 277
 - Processing the text* 277 ▪ *Scrambling a word* 279
- 16.4 Going further 280

17 *Mad Libs: Using regular expressions* 281

- 17.1 Writing mad.py 282
 - Using regular expressions to find the pointy bits* 284
 - Halting and printing errors* 287 ▪ *Getting the values* 288
 - Substituting the text* 289
- 17.2 Solution 289
- 17.3 Discussion 290
 - Substituting with regular expressions* 291 ▪ *Finding the placeholders without regular expressions* 291
- 17.4 Going further 293

18 *Gematria: Numeric encoding of text using ASCII values* 295

- 18.1 Writing gematria.py 296
 - Cleaning a word* 297 ▪ *Ordinal character values and ranges* 298 ▪ *Summing and reducing* 300 ▪ *Using functools.reduce* 302 ▪ *Encoding the words* 303
 - Breaking the text* 304
- 18.2 Solution 304
- 18.3 Discussion 305
 - Writing word2num()* 306 ▪ *Sorting* 308 ▪ *Testing* 309
- 18.4 Going further 309

19 *Workout of the Day: Parsing CSV files, creating text table output* 311

- 19.1 Writing wod.py 312
 - Reading delimited text files* 313 ▪ *Manually reading a CSV file* 315 ▪ *Parsing with the csv module* 318 ▪ *Creating a*

function to read a CSV file 320 ▪ Selecting the exercises 321
Formatting the output 322 ▪ Handling bad data 322
Time to write 323

19.2 Solution 323

19.3 Discussion 325

Reading a CSV file 325 ▪ Potential runtime errors 326
Using `pandas.read_csv()` to parse the file 327 ▪ Formatting
the table 328

19.4 Going further 330

20 *Password strength: Generating a secure and memorable password 331*

20.1 Writing `password.py` 334

Creating a unique list of words 335 ▪ Cleaning the text 337
Using a set 339 ▪ Filtering the words 340 ▪ Titlecasing the
words 341 ▪ Sampling and making a password 341
l33t-ify 342 ▪ Putting it all together 343

20.2 Solution 343

20.3 Discussion 346

Cleaning the text 346 ▪ A king's ransom 347 ▪ How to
l33t() 347 ▪ Processing the files 347 ▪ Sampling and
creating the passwords 348

20.4 Going further 349

21 *Tic-Tac-Toe: Exploring state 351*

21.1 Writing `tictactoe.py` 353

Validating user input 355 ▪ Altering the board 355
Printing the board 356 ▪ Determining a winner 356
Solution 357 ▪ Validating the arguments and mutating
the board 360 ▪ Formatting the board 363 ▪ Finding the
winner 364

21.2 Going further 366

22 *Tic-Tac-Toe redux: An interactive version with type hints 367*

22.1 Writing `itictactoe.py` 368

Tuple talk 369 ▪ Named tuples 371 ▪ Adding type
hints 372 ▪ Type verification with Mypy 373 ▪ Updating

	<i>immutable structures</i>	375	■ <i>Adding type hints to function definitions</i>	376
22.2	Solution	377		
			<i>A version using TypedDict</i>	379
			■ <i>Thinking about state</i>	381
22.3	Going further	381		
			<i>Epilogue</i>	383
appendix	<i>Using argparse</i>	385		
	<i>index</i>	405		

preface

Why write Python?

Python is an excellent, general-purpose programming language. You can write a program to send secret messages to your friends or to play chess. There are Python modules to help you wrangle complex scientific data, explore machine learning algorithms, and generate publication-ready graphics. Many college-level computer science programs have moved away from languages like C and Java to Python as their introductory language because Python is a relatively easy language to learn. We can use Python to study fundamental and powerful ideas from computer science. As I show you ideas like regular expressions and higher-order functions, I hope to encourage you to study further.

Why did I write this book?

Over the years, I've had many opportunities to help people learn programming, and I always find it rewarding. The structure of this book comes from my own experience in the classroom, where I think formal specifications and tests can be useful aids in learning how to break a program into smaller problems that need to be solved to create the whole program.

The biggest barrier to entry I've found when I'm learning a new language is that small concepts of the language are usually presented outside of any useful context. Most programming language tutorials will start with printing "HELLO, WORLD!" (and this book is no exception). Usually that's pretty simple. After that, I usually struggle to write a complete program that will accept some arguments and do something *useful*.

In this book, I'll show you many, many examples of programs that do useful things, in the hopes that you can modify these programs to make more programs for your own use.

More than anything, I think you need to practice. It's like the old joke: "What's the way to Carnegie Hall? Practice, practice, practice." These coding challenges are short enough that you could probably finish each in a few hours or days. This is more material than I could work through in a semester-long university-level class, so I imagine the whole book will take you several months. I hope you will solve the problems, then think about them, and then return later to see if you can solve them differently, maybe using a more advanced technique or making them run faster.

acknowledgments

This being my first book, it has been interesting to note the many people who have helped me create it. It all started with a call with Mike Stephens, the acquisitions editor for Manning, who entertained the idea of a book on learning how to produce serious, tested software by writing silly games and puzzles. That eventually led to a call with Marjan Bace, the publisher, who was enthusiastic about using test-driven development ideas to motivate readers to actively engage with writing the programs.

My first development editor, Susanna Kline, had to help me wrestle the first few chapters of the book into something people would actually want to read. My second development editor, Elesha Hyde, provided patient and thoughtful guidance through months of writing, editing, and reviews. I thank my technical editors, Scott Chaussee, Al Scherer, and Mathijs Affourtit, for carefully checking all my code and text for mistakes. I appreciated the efforts of Manning's MEAP team, especially Mehmed Pasic for producing the PDFs and giving me technical guidance on how to use AsciiDoc. I would also like to thank my project editor Deirdre Hiam, my copyeditor Andy Carroll, my proofreader Katie Tennant, and my review editor Aleksandar Dragosavljević. Also, the readers of the liveBook edition and the many technical reviewers who provided such great feedback: Amanda Debler, Conor Redmond, Drew Leon, Joaquin Beltran, José Apablaza, Kimberly Winston-Jackson, Maciej Jurkowski, Mafinar Khan, Manuel Ricardo Gonzalez Cova, Marcel van den Brink, Marcin Sek, Mathijs Affourtit, Paul R Hendrik, Shayn Cornwell, Víctor M. Pérez.

I especially want to acknowledge the countless people who create the open source software upon which all of this is built. From the people who maintain the Python

language and modules and documentation to the countless hackers who answer questions on the internet, I thank you for all that you do.

Of course, none of this would have ever been possible without the love and support of my family, especially my wife, Lori Kindler, who has been an unbelievable source of love and support for over 27 years. (I'm still really, really sorry about wrecking on my mountain bike and the year it took for me to recover!) Our three children bring me such challenges and joy, and I hope that I am making them proud. They constantly have to feign interest in topics they know and care nothing about, and they have shown such patience for the many hours I've spent writing this book.

about this book

Who should read this book

After you read this book and write all the programs, I would hope that you will be a zealot for creating programs that are documented, tested, and reproducible.

I think my ideal reader is someone who's been trying to learn to code well but isn't quite sure how to level up. Perhaps you are someone who's been playing with Python or some other language that has a similar syntax, like Java(Script) or Perl. Maybe you've cut your teeth on something really different, like Haskell or Scheme, and you're wondering how to translate your ideas to Python. Maybe you've been writing Python for a while and are looking for interesting challenges with enough structure to help you know when you're moving in the right direction.

This is a book that will teach you to write well-structured, documented, testable code in Python. The material introduces best practices from industry such as *test-driven development*—that's when the *tests* for a program exist even before the program itself is written! I will show you how to read documentation and Python Enhancement Proposals (PEPs) and how to write idiomatic code that other Python programmers would immediately recognize and understand.

This is probably not an ideal book for the absolute beginning programmer. I assume no prior knowledge of the Python language specifically, because I'm thinking of someone who is coming from another language. If you've never written a program in *any* language at all, you might do well to come back to this material when you are comfortable with ideas like variables, loops, and functions.

How this book is organized: A roadmap

The book is written with chapters building on previous chapters, so I really recommend you start at the beginning and work sequentially through the material.

- Every program uses command-line arguments, so we start off discussing how to use argparse to handle this. Every program is also tested, so you'll have to learn how to install and use pytest. The introduction and chapter 1 will get you up and running.
- Chapters 2–4 discuss the basic Python structures like strings, lists, and dictionaries.
- Chapters 5 and 6 move into how we can work with files as input and output and how files are related to “standard in” and “standard out” (STDIN/STDOUT).
- Chapters 7 and 8 start combining ideas so you can write more complicated programs.
- Chapters 9 and 10 introduce the random module and how to control and test random events.
- In chapters 11–13 you'll learn more about compartmentalizing code into functions and how to write and run tests for them.
- In chapters 14–18 we'll start digging into denser topics like higher-order functions as well as regular expressions to find patterns of text.
- In chapters 19–22 we'll start writing more complex, “real-world” programs that will put all your skills together while pushing your knowledge of the Python language and testing.

About the code

Every program and test shown in the book can be found at https://github.com/kyclark/tiny_python_projects.

Software/hardware requirements

All the programs were written and tested with Python 3.8, but version 3.6 would be sufficient for almost every program. Several additional modules are required, such as pytest for running the tests. There are instructions for how to use the pip module to install these.

liveBook discussion forum

Purchase of *Tiny Python Projects* includes free access to a private web forum run by Manning Publications where you can make comments about the book, ask technical questions, and receive help from the author and from other users. To access the forum, go to <https://livebook.manning.com/book/tiny-python-projects/welcome/v-6>. You can also learn more about Manning's forums and the rules of conduct at <https://livebook.manning.com/#!/discussion>.

Manning's commitment to our readers is to provide a venue where a meaningful dialogue between individual readers and between readers and the author can take

place. It is not a commitment to any specific amount of participation on the part of the author, whose contribution to the forum remains voluntary (and unpaid). We suggest you try asking him some challenging questions lest his interest stray! The forum and the archives of previous discussions will be accessible from the publisher's website as long as the book is in print.

Other online resources

One element missing from many programming courses is a demonstration of how one can go from having no program to having one that works. In my classroom teaching, I spend a lot of time showing students how to start writing a program and then how to work through the process of adding and testing new features. I've recorded videos for each chapter and shared them at www.youtube.com/user/kyclark. There is a playlist for each chapter, and the videos follow the pattern of each chapter by introducing the problem and the language features you might use to write your program, followed by a discussion of the solution(s).

about the author

My name is Ken Youens-Clark. I work as a Senior Scientific Programmer at the University of Arizona. Most of my career has been spent working in bioinformatics, using computer science ideas to study biological data.

I began my undergraduate degree as a Jazz Studies major on the drum set at the University of North Texas in 1990. I changed my major a few times and eventually ended up with a BA in English literature in 1995. I didn't really have a plan for my career, but I did like computers.

Around 1995, I started tinkering with databases and HTML at my first job out of college, building the company's mailing list and first website. I was definitely hooked! After that, I managed to learn Visual Basic on Windows 3.1 and, during the next few years, I programmed in several languages and companies before landing in a bioinformatics group at Cold Spring Harbor Laboratory in 2001, led by Lincoln Stein, a prominent author of books and modules in Perl and an early advocate for open software, data, and science. In 2014 I moved to Tucson, AZ, to work at the University of Arizona, where I completed my MS in Biosystems Engineering in 2019.

When I'm not coding, I like playing music, riding bikes, cooking, reading, and being with my wife and children.

about the cover

The figure on the cover of *Tiny Python Projects* is captioned “Femme Turc allant par les rues,” or “Turkish woman going through the streets.” The illustration is taken from a collection of dress costumes from various countries by Jacques Grasset de Saint-Sauveur (1757–1810), titled *Costumes de Différents Pays*, published in France in 1788. Each illustration is finely drawn and colored by hand. The rich variety of Grasset de Saint-Sauveur’s collection reminds us vividly of how culturally apart the world’s towns and regions were just 200 years ago. Isolated from each other, people spoke different dialects and languages. In the streets or in the countryside, it was easy to identify where they lived and what their trade or station in life was just by their dress.

The way we dress has changed since then, and the diversity by region, so rich at the time, has faded away. It is now hard to tell apart the inhabitants of different continents, let alone different towns, regions, or countries. Perhaps we have traded cultural diversity for a more varied personal life—certainly for a more varied and fast-paced technological life.

At a time when it is hard to tell one computer book from another, Manning celebrates the inventiveness and initiative of the computer business with book covers based on the rich diversity of regional life of two centuries ago, brought back to life by Grasset de Saint-Sauveur’s pictures.

Getting started: Introduction and installation guide

This book will teach you how to write Python programs that run on the command line. If you have never used the command line before, don't worry! You can use programs like PyCharm (see figure I.1) or Microsoft's VS Code to help you write and run these programs. If you are completely new to programming or to the Python language, I will try to cover everything I think you'll need to know, although you might find it useful to read another book first if you've never heard of things like variables and functions.

In this introduction, we'll discuss

- Why you should learn to write command-line programs
- Tools and environments for writing code
- How and why we test software

Writing command-line programs

Why do I want you to write command-line programs? For one, I think they strip a program down to its most bare essentials. We're not going to try to write complicated programs like an interactive 3D game that requires lots of other software to work. The programs in this book will all work with the barest of inputs and create only text output. We're going to focus on learning the core Python language and how to write *and test* programs.

Another reason for focusing on command-line programs is that I want to show you how to write programs that can run on any computer that has Python installed. I'm writing this book on my Mac laptop, but I can run all the programs on any of

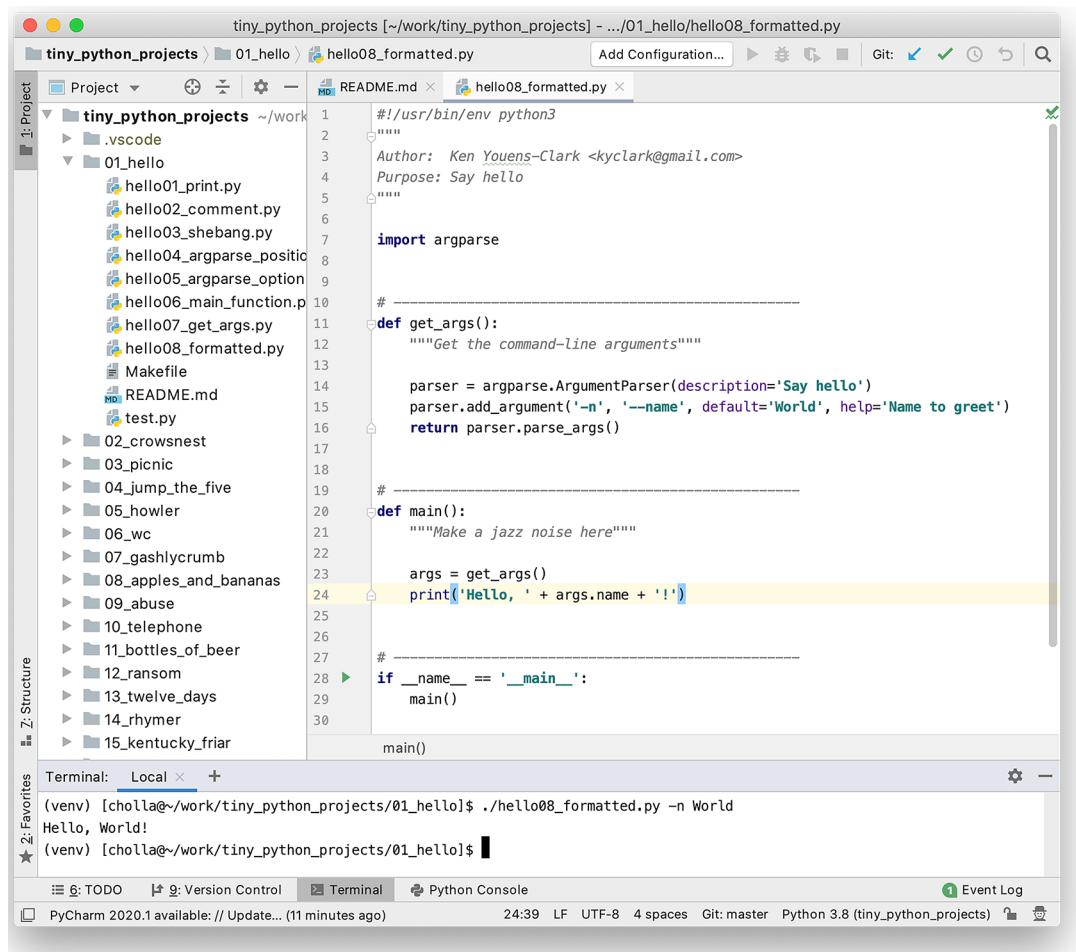


Figure I.1 This is the PyCharm tool being used to edit and run the `hello.py` program from chapter 1. “Hello, World!”

the Linux machines I use in my work or on a friend’s Windows machine. Any computer with the same version of Python can run any of these programs, and that is pretty cool.

The biggest reason I want to show you how to write command-line programs, though, is because I want to show you how to *test* programs to make sure they work. While I don’t think anyone will die if I make a mistake in one of my programs, I still really, really want to be sure that my code is as perfect as possible.

What does it mean to test a program? Well, if my program is supposed to add two numbers together, I’ll need to run it with many pairs of numbers and check that it prints the correct sum. I might also give it a number and a word, to make sure that it doesn’t try to add “3” plus “seahorse” but instead complains that I didn’t give it two

numbers. Testing gives me some measure of confidence in my code, and I hope you will come to see how testing can help you understand programming more deeply.

The exercises in this book are meant to be silly enough to pique your interest, but they each contain lessons that can be applied to all sorts of real-world problems. Almost every program I've ever written needs to accept some input data, whether from the user or from a file, and produce some output—sometimes text on the screen or maybe a new file. These are the kinds of skills you'll learn by writing these programs.

In each chapter, I'll describe some program that I want you to write and the tests you'll use to check if your program is working correctly. Then I'll show you a solution and discuss how it works. As the problems get harder, I'll start suggesting ways you might write your own tests to explore and verify your code.

When you're done with this book, you should be able to

- Write and run command-line Python programs
- Handle arguments to your programs
- Write and run tests for your programs and functions
- Use Python data structures like strings, lists, and dictionaries
- Have your programs read and write text files
- Use regular expressions to find patterns in text
- Use and control randomness to make your programs behave unpredictably

“Codes are a puzzle. A game, just like any other game.”

—Alan Turing

Alan Turing is perhaps most famous for cracking the Enigma code that the Nazis used to encrypt messages during World War II. The fact that the Allies could read enemy messages is credited with shortening the war by years and saving millions of lives. *The Imitation Game* is a fun movie that shows how Turing published puzzles in newspapers to find people who could help him break what was supposed to be an unbreakable code.

I think we can learn tons from writing fun programs that generate random insults or produce verses to “The Twelve Days of Christmas” or play Tic-Tac-Toe. Some of the programs in this book even dabble a bit in cryptography, like in chapter 4 where we encode all the numbers in a piece of text or in chapter 18 where we create signatures for words by summing the numeric representations of their letters. I hope you'll find the programs both amusing and challenging.

The programming techniques in each exercise are not specific to Python. Most every language has variables, loops, functions, strings, lists, and dictionaries, as well as ways to parameterize and test programs. After you write your solutions in Python, I encourage you to write solutions in another language you know and compare what parts of the different languages make it easier or harder to write your programs. If your programs support the same command-line options, you can even use the included tests to verify those programs.

Using test-driven development

Test-driven development is described by Kent Beck in his 2002 book by that title as a method for creating more reliable programs. The basic idea is that we write tests even before we write code. The tests define what it means to say that our program works “correctly.” First we write and run our tests to verify that our code fails. Then we write the code to make each test pass. We always run *all of the tests* so that, as we fix new tests, we ensure we don’t break tests that were passing before. When all the tests pass, we have at least some assurance that the code we’ve written conforms to some manner of specification.

Each program you are asked to write in this book comes with tests that will tell you when the code is working acceptably. The first test in every exercise checks whether the expected program exists. The second test checks that the program will print a help message if we ask for help. After that, your program will be run with various inputs and options.

Since I’ve written around 250 tests for the programs in this book, and you have not yet written one of the programs, you’re going to encounter many failed tests. That’s OK! In fact, it’s a really good thing, because when you pass all the tests, you’ll know that your programs are correct. You’ll learn to read the failed tests carefully to figure out what needs fixing. Then you’ll correct the program and run the tests again. You may get another failed test, in which case you’ll repeat the process until finally all the tests pass. Then you’ll be done.



It doesn’t matter if you solve the problems the same way as in the solution I provide. All that matters is that you figure out a way to pass the tests.

Setting up your environment

If you want to write these programs on your computer, you will need Python version 3.6 or later. It’s quite possible that it’s already installed on your computer.

You’ll also need some way to execute the `python3` command—something we often call a *command line*. If you use a Windows computer, you may want to install Windows Subsystem for Linux (WSL). On a Mac, the default Terminal app is sufficient. You can also use a tool like VS Code (in figure I.2) or PyCharm, which have terminals built into them.

I wrote and tested the programs for this book with Python version 3.8, but they should work with version 3.6 or newer. Python 2 reached its end of life at the end of 2019 and should no longer be used. To see what version of Python you have installed, open a terminal window and type `python3 --version`. If it says something like “command “`python3`” not found,” then you need to install Python. You can download the latest version from the Python site (www.python.org/downloads).

If you are using a computer that doesn’t have Python, and you don’t have any way to install Python, you can do everything in this book using the Repl.it website (<http://repl.it>).

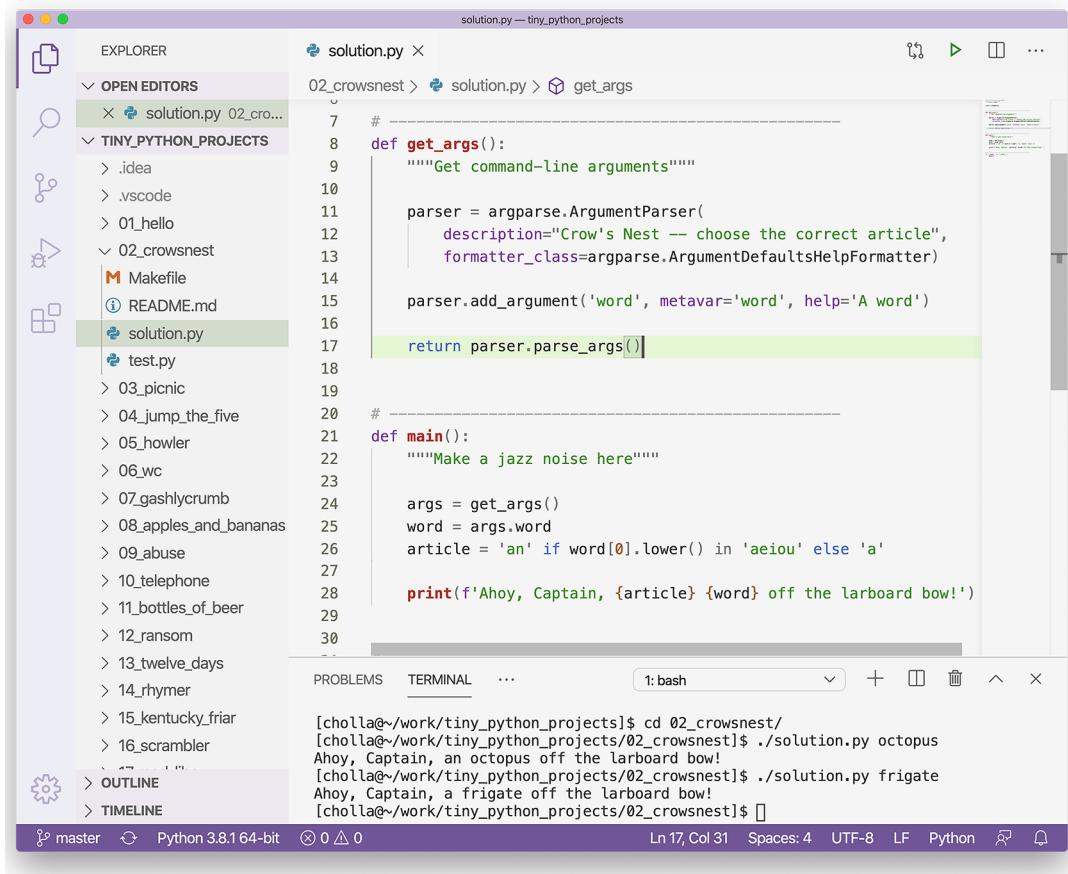


Figure I.2 An IDE like VS Code combines a text editor for writing your code along with a terminal (lower-right window) for running your programs, and many other tools.

Code examples

Throughout the book, I will show commands and code using a fixed-width font. When the text is preceded with a dollar sign (\$), that means it's something you can type on the command line. For instance, there is a program called `cat` (short for “concatenate”) that will print the contents of a file to the screen. Here is how I can run it to print the contents of the `spiders.txt` file that lives in the `inputs` directory:

```
$ cat inputs/spiders.txt
Don't worry, spiders,
I keep house
casually.
```

If you want to run that command, *do not copy* the leading \$, only the text that follows. Otherwise you'll probably get an error like “\$: command not found.”

Python has a really excellent tool called IDLE that allows you to interact directly with the language to try out ideas. You can start it with the command `idle3`. That should open a new window with a prompt that looks like `>>>` (see figure I.3).

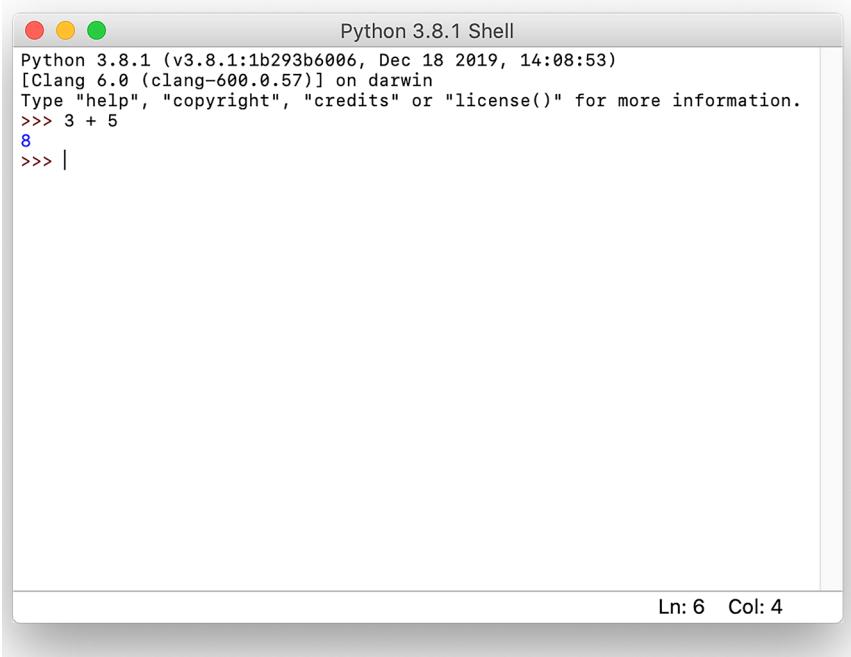


Figure I.3 The IDLE application allows you to interact directly with the Python language. Each statement you type is evaluated when you press Enter, and the results are shown in the window.

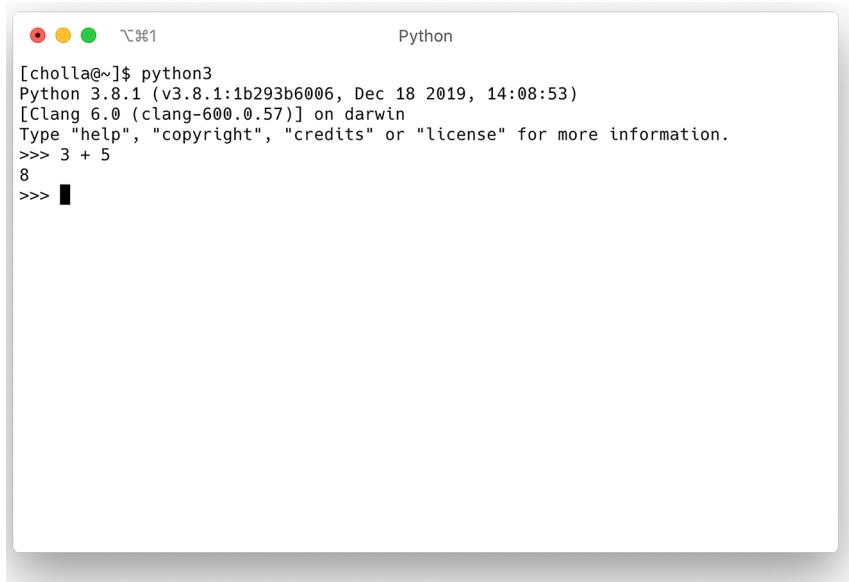
You can type Python statements there, and they will be immediately evaluated and printed. For example, type `3 + 5` and press Enter, and you should see 8:

```
>>> 3 + 5
8
```

This interface is called a *REPL* because it's a Read-Evaluate-Print-Loop. (I pronounce this like “repple” in a way that sort of rhymes with “pebble.”) You can get a similar tool by typing `python3` on the command line (see figure I.4).

The IPython program is yet another “interactive Python” REPL that has many enhancements over IDLE and `python3`. Figure I.5 shows what it looks like on my system.

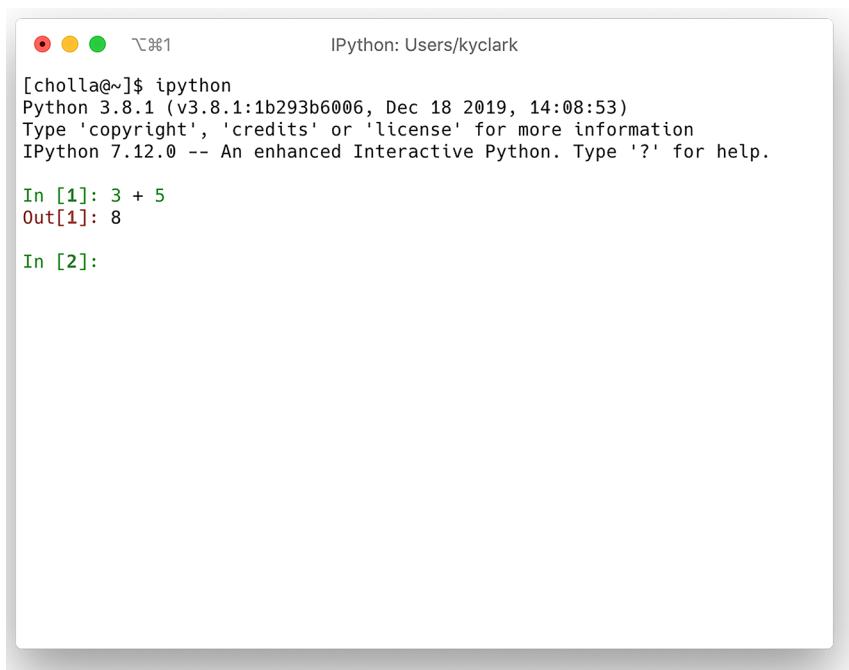
I also recommend you look into using Jupyter Notebooks, as they allow you to interactively run code with the added bonus that you can save a Notebook as a file and share all your code with other people.



A screenshot of a terminal window titled "Python". The window shows a command-line interface where the user has typed "python3" and is presented with the standard Python 3.8.1 startup message. Below this, the user types "3 + 5" and receives the output "8". The window has a dark grey header bar with three colored circles (red, yellow, green) and a small icon.

```
[cholla@~]$ python3
Python 3.8.1 (v3.8.1:1b293b6006, Dec 18 2019, 14:08:53)
[Clang 6.0 (clang-600.0.57)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> 3 + 5
8
>>> █
```

Figure I.4 Typing the command `python3` in the terminal will give you a REPL similar to the IDLE interface.



A screenshot of a terminal window titled "IPython: Users/kyclark". The window shows an IPython 7.12.0 session. The user has typed "ipython" and is presented with the standard IPython startup message. Below this, the user types "3 + 5" and receives the output "8". The window has a dark grey header bar with three colored circles (red, yellow, green) and a small icon.

```
[cholla@~]$ ipython
Python 3.8.1 (v3.8.1:1b293b6006, Dec 18 2019, 14:08:53)
Type 'copyright', 'credits' or 'license' for more information
IPython 7.12.0 -- An enhanced Interactive Python. Type '?' for help.

In [1]: 3 + 5
Out[1]: 8

In [2]:
```

Figure I.5 The IPython application is another REPL interface you can use to try out your ideas with Python.

Whichever REPL interface you use, you can type Python statements like `x = 10` and press Enter to assign the value 10 to the variable `x`:

```
>>> x = 10
```

As with the command-line prompt, \$, do not copy the leading >>> or Python will complain:

```
>>> >>> x = 10
File "<stdin>", line 1
>>> x = 10
^
SyntaxError: invalid syntax
```

The IPython REPL has a magical %paste mode that removes the leading >>> prompts so that you can copy and paste all the code examples:

```
In [1]: >>> x = 10

In [2]: x
Out[2]: 10
```

Whichever way you choose to interact with Python, I suggest you *manually type all the code yourself* in this book, as this builds muscle memory and forces you to interact with the syntax of the language.

Getting the code

All the tests and solutions are available at https://github.com/kyclark/tiny_python_projects. You can use the program Git (which you may need to install) to copy that code to your computer with the following command:

```
$ git clone https://github.com/kyclark/tiny_python_projects
```

Now you should have a new directory called `tiny_python_projects` on your computer.

You may prefer to make a copy of the code into your own repository, so that you can track your changes and share your solutions with others. This is called “forking” because you’re breaking off from my code and adding your own programs to the repository. If you plan to use Repl.it to write the exercises, I recommend you do fork my repo into your own account so that you can configure Repl.it to interact with your own GitHub repositories.

To fork, do the following:

- 1 Create an account on GitHub.com.
- 2 Go to https://github.com/kyclark/tiny_python_projects.
- 3 Click the Fork button (see figure I.6) to make a copy of the repository into your account.

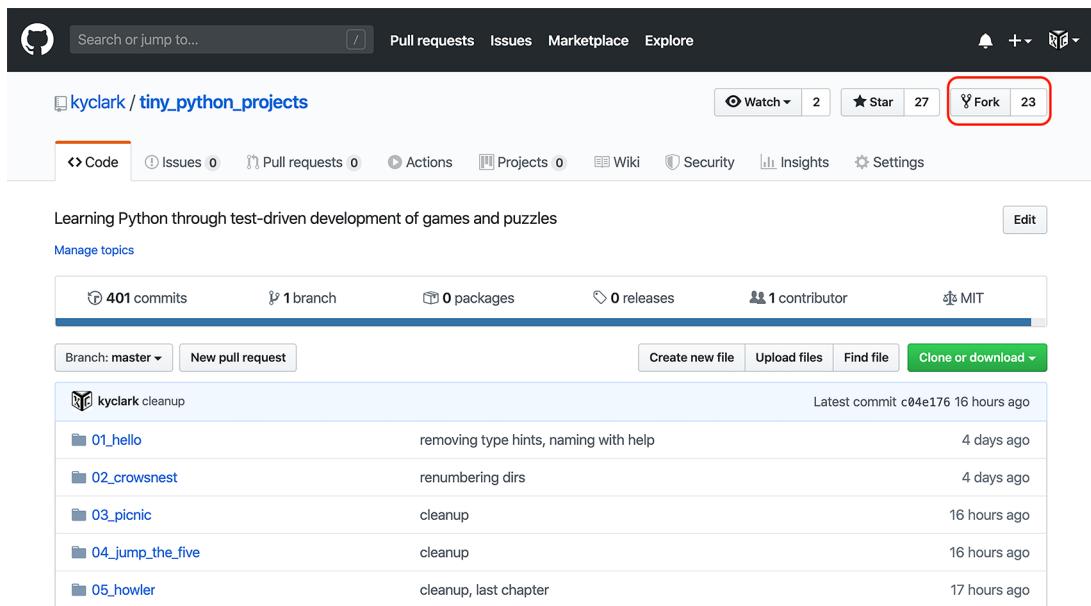


Figure I.6 The Fork button on my GitHub repository will make a copy of the code into your account.

Now you have a copy of my all code in your own repository. You can use Git to copy that code to your computer. Be sure to replace “YOUR_GITHUB_ID” with your actual GitHub ID:

```
$ git clone https://github.com/YOUR_GITHUB_ID/tiny_python_projects
```

I may update the repo after you make your copy. If you would like to be able to get those updates, you will need to configure Git to set my repository as an “upstream” source. To do so, after you have cloned your repository to your computer, go into your tiny_python_projects directory:

```
$ cd tiny_python_projects
```

Then execute this command:

```
$ git remote add upstream https://github.com/kyclark/tiny_python_projects.git
```

Whenever you would like to update your repository from mine, you can execute this command:

```
$ git pull upstream master
```

Installing modules

I recommend using a few tools that may not be installed on your system. You can use the pip module to install them like so:

```
$ python3 -m pip install black flake8 ipython mypy pylint pytest yapf
```

I've also included a requirements.txt file in the top level of the repository. You can use it to install all the modules and tools with this command:

```
$ python3 -m pip install -r requirements.txt
```

If, for example, you wish to write the exercises on Repl.it, you will need to run this command to set up your environment, as the modules are not already installed.

Code formatters

Most IDEs and text editors will have tools to help you format your code so that it's easier to read and find problems. In addition, the Python community has created a standard for writing code so that other Python programmers can readily understand it. The PEP 8 (Python Enhancement Proposal) document at www.python.org/dev/peps/pep-0008/ describes best practices for formatting code, and most editors will automatically apply formatting for you. For instance, the Repl.it interface has an autoformat button (see figure I.7), VS Code has a Format Document command, and PyCharm has a Reformat Code command.

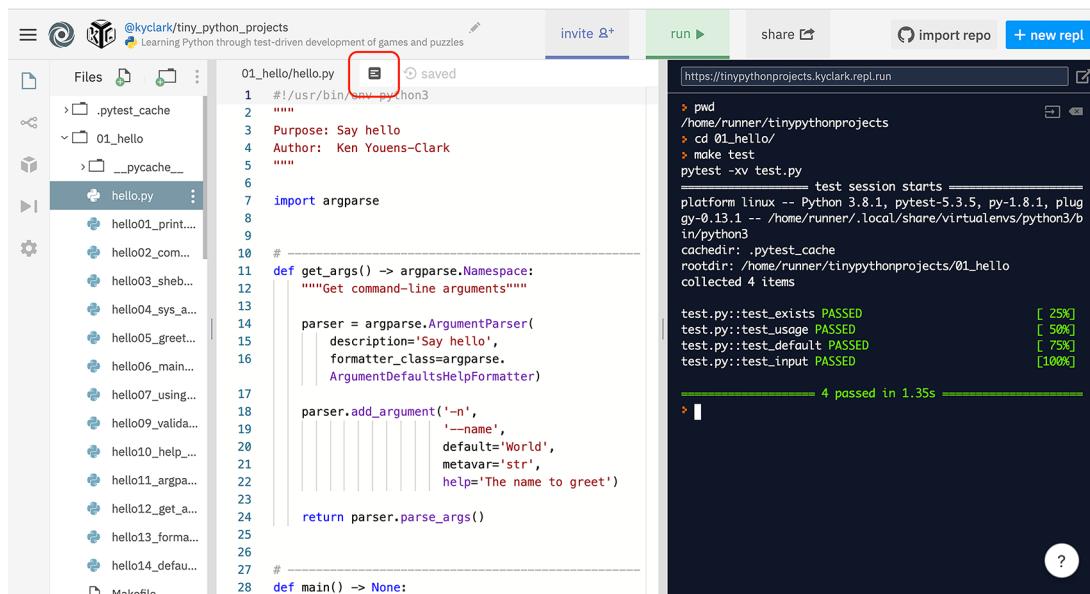


Figure I.7 The Repl.it tool has an autoformat button to reformat your code according to community standards. The interface also includes a command line for running and testing your program.

There are also command-line tools that integrate with your editor. I used YAPF (Yet Another Python Formatter, <https://github.com/google/yapf>) to format every program in the book, but another popular formatter is Black (<https://github.com/psf/black>). Whatever you use, I encourage you to use it *often*. For instance, I can tell YAPF to format the hello.py program that we will write in chapter 1 by running the following command. Note that the `-i` tells YAPF to format the code “in place,” so that the original file will be overwritten with the newly formatted code.

```
$ yapf -i hello.py
```

Code linters

A *code linter* is a tool that will report problems in your code, such as declaring a variable but never using it. Two that I like are Pylint (www.pylint.org/) and Flake8 (<http://flake8.pycqa.org/en/latest/>), and both can find errors in your code that the Python interpreter itself will not complain about.

In the final chapter, I will show you how to incorporate *type hints* into your code that the Mypy tool (<http://mypy-lang.org/>) can use to find problems, such as using text when you should be using a number.

How to start writing new programs

I think it’s much easier to start writing code with a standard template, so I wrote a program called new.py that will help you create new Python programs with boilerplate code that will be expected of every program. It’s located in the bin directory, so if you are in the top directory of the repository, you can run it like this:

```
$ bin/new.py
usage: new.py [-h] [-s] [-n NAME] [-e EMAIL] [-p PURPOSE] [-f] program
new.py: error: the following arguments are required: program
```

Here you can see that new.py is asking you to provide the name of the “program” to create. For each chapter, the program you write needs to live in the directory that has the test.py file for that program.

For example, you can use new.py to start off chapter 2’s crowsnest.py program in the 02_crowsnest directory like so:

```
$ bin/new.py 02_crowsnest/crowsnest.py
Done, see new script "02_crowsnest/crowsnest.py."
```

If you open that file now, you’ll see that it has written a lot of code for you that I’ll explain later. For now, just realize that the resulting crowsnest.py program is one that can be run like so:

```
$ 02_crowsnest/crowsnest.py
usage: crowsnest.py [-h] [-a str] [-i int] [-f FILE] [-o] str
crowsnest.py: error: the following arguments are required: str
```

Later you'll learn how to modify the program to do what the tests expect.

An alternative to running new.py is to copy the file template.py from the template directory to the directory and program name you need to write. You could create the crowsnest.py program file like so:

```
$ cp template/template.py 02_crowsnest/crowsnest.py
```

You do not have to use either new.py or copy the template.py file to start your programs. These are provided to save you time and provide your programs with an initial structure, but you are welcome to write your programs however you please.

Why not Notebooks?

Many people are familiar with Jupyter Notebooks, as they provide a way to integrate Python code and text and images into a document that other people can execute like a program. I really love Notebooks, especially for interactively exploring data, but I find them difficult to use in teaching for the following reasons:

- A Notebook is stored in JavaScript Object Notation (JSON), not as line-oriented text. This makes it really difficult to compare Notebooks to each other to find out how they differ.
- Code and text and images can live mixed together in separate cells. These cells can be interactively run in any order, which can lead to very subtle problems in the logic of a program. The programs we write in this book will always be run from top to bottom in entirety every time, which I think makes them easier to understand.
- There is no way for Notebooks to accept different values when they are run. That is, if you test a program with one input file and then want to change to a different file, you have to change *the program itself*. You will learn how to pass in a file as an *argument* to the program, so that you can change the value without changing the *code*.
- It's difficult to automatically run tests on a Notebook or on the functions they contain. We will use the pytest module to run our programs over and over with different input values and verify that the programs create the correct output.

The scope of topics we'll cover

The purpose of this book is to show you how amazingly useful all the built-in features of the Python language are. The exercises will push you to practice manipulating strings, lists, dictionaries, and files. We'll spend several chapters focusing on regular expressions, and every exercise except for the last requires you to accept and validate command-line arguments of varying types and numbers.

Every author is biased toward some subjects, and I'm no different. I've chosen these topics because they reflect ideas that are fundamental to the work I've done over the last 20 years. For instance, I have spent many more hours than I would care to

admit parsing really messy data from countless Excel spreadsheets and XML files. The world of genomics that has consumed most of my career is based primarily on efficiently parsing text files, and much of my web development work is predicated on understanding how text is encoded and transferred to and from the web browser. For that reason, you'll find many exercises that entail processing text and files, and that will challenge you to think about how to transform inputs into outputs. If you work through every exercise, I believe you'll be a much improved programmer who understands the basic ideas that are common across many languages.

Why not object-oriented programming?

One topic that you'll notice is missing from this book is writing object-oriented code in Python. If you are not familiar with *object-oriented programming* (OOP), you can skip this section.

I think OOP is a somewhat advanced topic that is beyond the scope of this book. I prefer to focus on how to write small functions and their accompanying tests. I think this leads to more transparent code, because the functions should be short, should only use the values explicitly passed as arguments, and should have enough tests that you can completely understand how they will behave under both favorable and unfavorable circumstances.

The Python language is itself inherently object-oriented. Almost everything from strings to the lists and dictionaries that we'll use are actually *objects*, so you'll get plenty of practice using objects. But I don't think it's necessary to create objects to solve any of the problems I present. In fact, even though I spent many years writing object-oriented code, I haven't written in this style for the last few years. I tend to draw my inspiration from the world of purely functional programming, and I hope I can convince you by the end of this book that you can do anything you want by combining functions.

Although I personally avoid OOP, I would recommend you learn about it. There have been several seismic paradigm shifts in the world of programming from procedural to object-oriented and now functional. You can find dozens of books on OOP in general and on programming objects in Python specifically. This is a deep and fascinating topic, and I encourage you to try writing object-oriented solutions and compare them to my solutions.

A note about the lingo

Often in programming books you will see *foobar* used in examples. The word has no real meaning, but its origin probably comes from the military acronym "FUBAR" (Fouled Up Beyond All Recognition). If I use "foobar" in an example, it's because I don't want to talk about any specific thing in the universe, just the idea of a string of characters. If I need a list of items, usually the first item will be "foo" and the next will be "bar." After that, convention uses "baz" and "quux," again because they mean nothing at all. Don't get hung up on "foobar." It's just a placeholder for something that could be more interesting later.

Programmers also tend to call errors in code *bugs*. This comes from the days of computing before the invention of transistors. Early machines used vacuum tubes, and the heat from the machines would attract actual bugs like moths that could cause short circuits. The operators (the people running the machines) would have to hunt through the machinery to find and remove the bugs; hence, the term “to debug.”



How to write and test a Python program

Before you start working on the exercises, I want to discuss how to write programs that are documented and tested. Specifically, we're going to

- Write a Python program to say “Hello, World!”
- Handle command-line arguments using argparse
- Run tests for the code with Pytest.
- Learn about \$PATH
- Use tools like YAPF and Black to format the code
- Use tools like Flake8 and Pylint to find problems in the code
- Use the new.py program to create new programs



1.1 *Creating your first program*

It's pretty common to write “Hello, World!” as your first program in any language, so let's start there. We're going to work toward making a version that will greet whichever name is passed as an argument. It will also print a helpful message when we ask for it, and we'll use tests to make sure it does everything correctly.

In the 01_hello directory, you'll see several versions of the hello program we'll write. There is also a program called test.py that we'll use to test the program.

Start off by creating a text file called hello.py in that directory. If you are working in VS Code or PyCharm, you can use File > Open to open the 01_hello directory as a project. Both tools have something like a File > New menu option that will allow you to create a new file in that directory. It's very important to create the hello.py file *inside* the 01_hello directory so that the test.py program can find it.

Once you've started a new file, add this line:

```
print('Hello, World!')
```

It's time to run your new program! Open a terminal window in VS Code or PyCharm or in some other terminal, and navigate to the directory where your hello.py program is located. You can run it with the command `python3 hello.py`—this causes Python version 3 to execute the commands in the file named hello.py. You should see this:

```
$ python3 hello.py
Hello, World!
```

Figure 1.1 shows how it looks in the Repl.it interface.

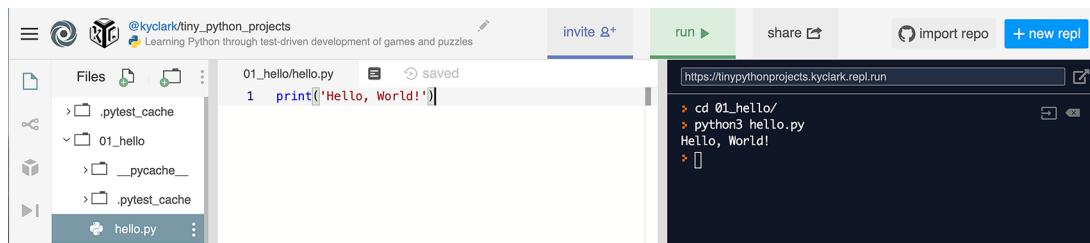


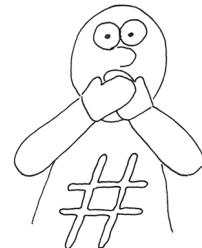
Figure 1.1 Writing and running our first program using Repl.it

If that was your first Python program, congratulations!

1.2 Comment lines

In Python, the `#` character and anything following it is ignored by Python. This is useful for adding comments to your code or temporarily disabling lines of code when testing and debugging. It's always a good idea to document your programs, indicating the purpose of the program or the author's name and email address, or both. We can use a comment for that:

```
# Purpose: Say hello
print('Hello, World!')
```



If you run this program again, you should see the same output as before because the “Purpose” line is ignored. Note that any text to the left of the `#` is executed, so you can add a comment to the end of a line if you like.

1.3 Testing your program

The most fundamental idea I want to teach you is how to test your programs. I've written a test.py program in the 01_hello directory that we can use to test our new hello.py program.

We will use pytest to execute all the commands and tell us how many tests we passed. We'll include the -v option, which tells pytest to create "verbose" output. If you run it like this, you should see the following output as the first several lines. After that will follow many more lines showing you more information about the tests that didn't pass.

NOTE If you get the error "pytest: command not found," you need to install the pytest module. Refer to the "Installing modules" section in the book's introduction.

The second test tries to run the program with python3 hello.py and then checks if the program printed "Hello, World!" If you miss even one character, like forgetting a comma, the test will point out the error, so read carefully!

```
$ pytest -v test.py
=====
test session starts =====
...
collected 5 items

test.py::test_exists PASSED
test.py::test_runnable PASSED
test.py::test_executable FAILED
test.py::test_usage FAILED
test.py::test_input FAILED

=====
FAILURES =====
```

The first test always checks that the expected file exists. Here the test looks for hello.py.

[20%] ←
[40%]
[60%]
[80%]
[100%] ←

The fourth test asks the program for help and doesn't get anything. We're going to add the ability to print a "usage" statement that describes how to use our program.

The third test checks that the program is "executable." This test fails, so next we'll talk about how to make that pass.

The last test checks that the program can greet a name that we'll pass as an argument. Since our program doesn't yet accept arguments, we'll need to add that, too.

I've written the tests in an order that I hope will help you write the program in a logical fashion. If the program doesn't pass one of the tests, there's no reason to continue running the tests after it. I recommend you always run the tests with the flags -x, to stop on the first failing test, and -v, to print verbose output. You can combine these like -xv or -vx. Here's what our tests look like with those options:

```
$ pytest -xv test.py
=====
test session starts =====
...
collected 5 items
```

```

test.py::test_exists PASSED [ 20%]
test.py::testRunnable PASSED [ 40%]
test.py::test_executable FAILED [ 60%]

This test fails. No more tests are run because
we ran pytest with the -x option.

=====
FAILURES =====
test_executable

def test_executable():
    """Says 'Hello, World!' by default"""

    out = getoutput({prg})
>   assert out.strip() == 'Hello, World!' ←
E   AssertionError: assert '/bin/sh: ./h...ission denied' == 'Hello, World!' ←
E       - /bin/sh: ./hello.py: Permission denied ←
E       + Hello, World!

test.py:30: AssertionError
!!!!!!!!!!!!!! stopping after 1 failures !!!!!!!!!!!!!!!
=====
1 failed, 2 passed in 0.09s =====

The plus character (+) shows
that the test expected to get
"Hello, World!" ←

The "E" at the beginning of this line shows that this is an
"Error" you should read. The AssertionError is saying that the
test.py program is trying to execute the command ./hello.py
to see if it will produce the text "Hello, World!" ←

```

Let's talk about how to fix this error.

1.4 Adding the `#! (shebang)` line

One thing you have learned so far is that Python programs live in plain text files that you ask `python3` to execute. Many other programming languages, such as Ruby and Perl, work in the same way—we type Ruby or Perl commands into a text file and run it with the right language. It's common to put a special comment line in programs like these to indicate which language needs to be used to execute the commands in the file.

This comment line starts off with `#!`, and the nickname for this is “shebang” (pronounced “shuh-bang”—I always think of the `#` as the “shuh” and the `!` as the “bang!”). Just as with any other comment, Python will ignore the shebang, but the operating system (like macOS or Windows) will use it to decide which program to use to run the rest of the file.

Here is the shebang you should add:

```
#!/usr/bin/env python3
```

The `env` program will tell you about your “environment.” When I run `env` on my computer, I see many lines of output like `USER=kyclark` and `HOME=/Users/kyclark`. These values are accessible as the variables `$USER` and `$HOME`:

```
$ echo $USER
kyclark
$ echo $HOME
/Users/kyclark
```

If you run `env` on your computer, you should see your login name and your home directory. They will, of course, have different values from mine, but we both (probably) have both of these concepts.

You can use the `env` command to find and run programs. If you run `env python3`, it will run a `python3` program if it can find one. Here's what I see on my computer:

```
$ env python3
Python 3.8.1 (v3.8.1:1b293b6006, Dec 18 2019, 14:08:53)
[Clang 6.0 (clang-600.0.57)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>>
```

The `env` program is looking for `python3` in the environment. If Python has not been installed, it won't be able to find it, but it's also possible that Python has been installed more than once. You can use the `which` command to see which `python3` it finds:

```
$ which python3
/Library/Frameworks/Python.framework/Versions/3.8/bin/python3
```

If I run this on Repl.it, I can see that `python3` exists in a different place. Where does it exist on your computer?

```
$ which python3
/home/runner/.local/share/virtualenvs/python3/bin/python3
```

Just as my `$USER` name is different from yours, my `python3` is probably different from yours. If the `env` command is able to find a `python3`, it will execute it. As shown previously, if you run `python3` by itself, it will open a REPL.

If I were to put my `python3` path as the shebang line, like so,

```
#!/Library/Frameworks/Python.framework/Versions/3.8/bin/python3
```

my program would not work on another computer that has `python3` installed in a different location. I doubt it would work on your computer, either. This is why you should always use the `env` program to find the `python3` that is specific to the machine on which it's running.

Now your program should look like this:

```
#!/usr/bin/env python3
# Purpose: Say hello
print('Hello, World!')
```

The shebang line tells the operating system to use `/usr/bin/env` to find `python3` to interpret this program.

A comment line documenting the purpose of the program

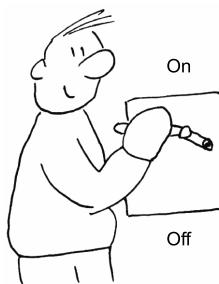
A Python command to print some text to the screen

1.5 Making a program executable

So far we've been explicitly telling `python3` to run our program, but since we added the shebang, we can execute the program directly and let the OS figure out that it should use `python3`. The advantage of this is that we could copy our program to a directory where other programs live and execute it from anywhere on our computer.

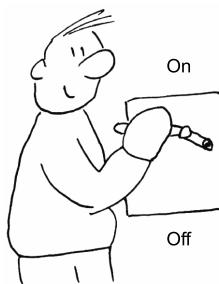
The first step in doing this is to make our program "executable" using the command `chmod` (*change mode*). Think of it as turning your program "on." Run this command to make `hello.py` executable:

```
$ chmod +x hello.py
```

 The `+x` will add an "executable" attribute to the file.

Now you can run the program like so:

```
$ ./hello.py
```

 The `./` is the current directory, and it's necessary to run a program when you are in the same directory as the program.

1.6 Understanding \$PATH

One of the biggest reasons to set the shebang line and make your program executable is so that you can install your Python programs just like other commands and programs. We used the `which` command earlier to find the location of `python3` on the Repl.it instance:

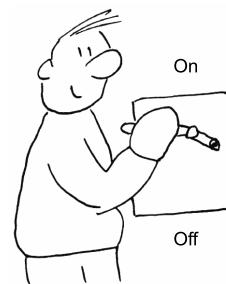
```
$ which python3
/home/runner/.local/share/virtualenvs/python3/bin/python3
```

How was the `env` program able to find it? Windows, macOS, and Linux all have a `$PATH` variable, which is a list of directories the OS will look in to find a program. For instance, here is the `$PATH` for my Repl.it instance:

```
> echo $PATH
/home/runner/.local/share/virtualenvs/python3/bin:/usr/local/bin:\n
/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

The directories are separated by colons (:). Notice that the directory where `python3` lives is the first one in `$PATH`. It's a pretty long string, so I broke it with the \ character to make it easier to read. If you copy your `hello.py` program to any of the directories listed in your `$PATH`, you can execute a program like `hello.py` without the leading `./` and without having to be in the same directory as the program.

Think about `$PATH` like this: If you lose your keys in your house, would you start looking in the upper-left kitchen cabinet and work your way through each cabinet, and then all the drawers where you keep your silverware and kitchen gadgets, and then move on to your bathrooms and bedroom closets? Or would you start by looking in places where you normally put your keys, like the key hooks beside the front door,



and then move on to search the pockets of your favorite jacket and your purse or backpack, and then maybe look under the couch cushions, and so forth?

The \$PATH variable is a way of telling your computer to only look in places where executable programs can be found. The only alternative is for the OS to search *every directory*, and that could take several minutes or possibly even hours! You can control both the names of the directories in the \$PATH variable and their relative order so that the OS will find the programs you need.

It's very common for programs to be installed into /usr/local/bin, so we could try to copy our program there using the cp command. Unfortunately, I do not have permission to do this on Repl.it:

```
> cp 01_hello/hello.py /usr/local/bin
cp: cannot create regular file '/usr/local/bin/hello.py': Permission denied
```

But I can do this on my own laptop:

```
$ cp hello.py /usr/local/bin/
```

I can verify that the program is found:

```
$ which hello.py
/usr/local/bin/hello.py
```

And now I can execute it from any directory on my computer:

```
$ hello.py
Hello, World!
```

1.6.1 Altering your \$PATH

Often you may find yourself working on a computer that won't allow you to install programs into your \$PATH, such as on Repl.it. An alternative is to alter your \$PATH to include a directory where you can put your programs. For instance, I often create a bin directory in my home directory, which can often be written with the tilde (~).

On most computers, ~/bin would mean "the bin directory in my home directory." It's also common to see \$HOME/bin where \$HOME is the name of your home directory. Here is how I create this directory on the Repl.it machine, copy a program to it, and then add it to my \$PATH:

```
$ mkdir ~bin
$ cp 01_hello/hello.py ~bin
$ PATH=~/bin:$PATH
$ which hello.py
/home/runner/bin/hello.py
```

Use the mkdir ("make directory") command to create ~/bin.

Use the cp command to copy the 01_hello/hello.py program to the ~/bin directory.

Put the ~/bin directory first in \$PATH.

Use the which command to look for the hello.py program. If the previous steps worked, the OS should now be able to find the program in one of the directories listed in \$PATH.

Now I can be in any directory,

```
$ pwd
/home/runner/tiny python projects
```

and I can run it:

```
$ hello.py
Hello, World!
```

Although the shebang and the executable stuff may seem like a lot of work, the payoff is that you can create a Python program that can be installed onto your computer or anyone else's and run just like any other program.

1.7 Adding a parameter and help

Throughout the book, I'll use string diagrams to visualize the inputs and outputs of the programs we'll write. If we created one for our program now (as in figure 1.2), there would be no inputs, and the output would always be "Hello, World!"

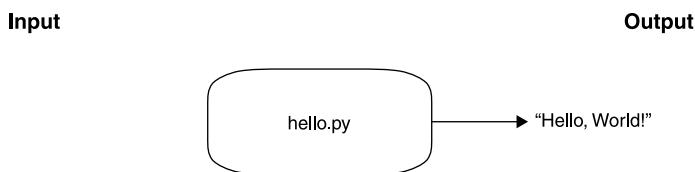


Figure 1.2 A string diagram representing our `hello.py` program that takes no inputs and always produces the same output

It's not terribly interesting for our program to always say "Hello, World!" It would be nice if it could say "Hello" to something else, like the entire universe. We could do this by changing the code as follows:

```
print('Hello, Universe')
```

But that would mean we'd have to change the code every time we wanted to make it greet a different name. It would be better to change the *behavior* of the program without always having to change *the program itself*.

We can do that by finding the parts of the program that we want to change—like the name to greet—and providing that value as an *argument* to our program. That is, we'd like our program to work like this:

```
$ ./hello.py Terra
Hello, Terra!
```

How would the person using our program know to do this? *It's our program's responsibility to provide a help message!* Most command-line programs will respond to arguments

like `-h` and `--help` with helpful messages about how to use the programs. We need our program to print something like this:

```
$ ./hello.py -h
usage: hello.py [-h] name

Say hello

positional arguments:
  name      Name to greet      Note that name is called
                                         a positional argument.

optional arguments:
  -h, --help show this help message and exit
```

To do this, we can use the `argparse` module. Modules are files of code we can bring into our programs. We can also create modules to share our code with other people. There are hundreds to thousands of modules you can use in Python, which is one of the reasons why it's so exciting to use the language.

The `argparse` module will “parse” the “arguments” to the program. To use it, change your program as follows. I recommend you type everything yourself and don't copy and paste.

We must import the `argparse` module to use it.

```
#!/usr/bin/env python3
# Purpose: Say hello
import argparse

parser = argparse.ArgumentParser(description='Say hello')
parser.add_argument('name', help='Name to greet')
args = parser.parse_args()
print('Hello, ' + args.name + '!')

The shebang line tells the OS which
program to use to execute this program.

This comment documents the
purpose of the program.

The parser will figure out
all the arguments. The
description appears in
the help message.

We need to tell the parser to
expect a name that will be the
object of our salutations.

We ask the parser to parse any
arguments to the program.

We print the greeting using
the args.name value.
```

Figure 1.3 shows a string diagram of our program now.

Now when you try to run the program like before, it triggers an error and a “usage” statement (notice that “usage” is the first word of the output):

```
$ ./hello.py
usage: hello.py [-h] name
hello.py: error: the following arguments are required: name
```

Since the program doesn't get the expected argument, it stops and prints a “usage” message to let the user know how to properly invoke the program.

We run the program with no arguments,
but the program now expects a single
argument (a “name”).

The error message tells the user that
they have not supplied a required
parameter called “name.”

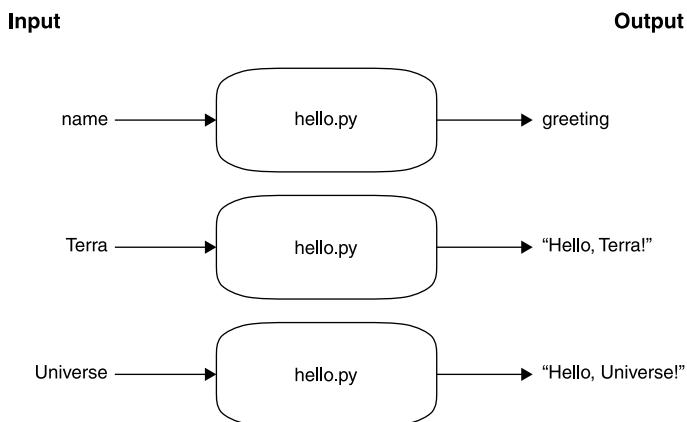


Figure 1.3 Now our string diagram shows that our program can take an argument and produce a message based on that value.

We've changed the program so that it requires a name or it won't run. That's pretty cool! Let's give it a name to greet:

```
$ ./hello.py Universe
Hello, Universe!
```

Try running your program with both the `-h` and `--help` arguments, and verify that you see the help messages.

The program works really well now and has nice documentation, all because we added those few lines using argparse. That's a big improvement.

1.8 Making the argument optional

Suppose we'd like to run the program like before, with no arguments, and have it print "Hello, World!" We can make the name optional by changing the name of the argument to `--name`:

```
#!/usr/bin/env python3
# Purpose: Say hello
import argparse

parser = argparse.ArgumentParser(description='Say hello')
parser.add_argument('-n', '--name', metavar='name',
                    default='World', help='Name to greet')
args = parser.parse_args()
print('Hello, ' + args.name + '!')

The only change to this program is adding -n and
--name for the "short" and "long" option names.
We also indicate a default value. "metavar" will
show up in the usage to describe the argument.
```

Now we can run it like before:

```
$ ./hello.py
Hello, World!
```

Or we can use the `--name` option:

```
$ ./hello.py --name Terra
Hello, Terra!
```

And our help message has changed:

```
$ ./hello.py -h
usage: hello.py [-h] [-n NAME]

Say hello

optional arguments:
  -h, --help            show this help message and exit
  -n name, --name name  Name to greet
```

The argument is now optional and no longer a positional argument. It's common to provide both short and long names to make it easy to type the options. The metavar value of "name" appears here to describe what the value should be.

Figure 1.4 shows a string diagram that describes our program.

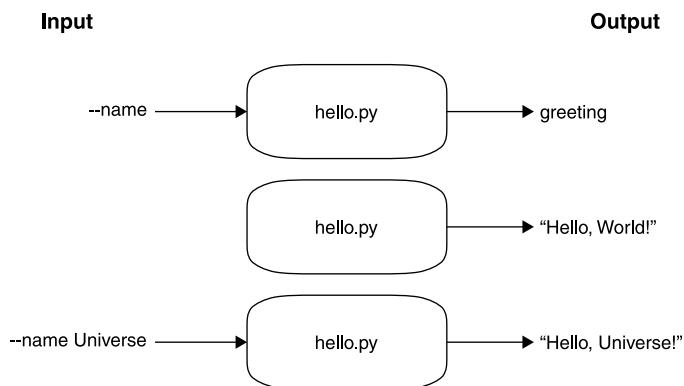


Figure 1.4 The name parameter is now optional. The program will greet a given name or will use a default value when it's missing.

Our program is really flexible now, greeting a default value when run with no arguments or allowing us to say “hi” to something else. Remember that parameters that start with dashes are *optional*, so they can be left out, and they may have default values. Parameters that *don't* start with dashes are *positional* and are usually required, so they do not have default values.

Table 1.1 Two kinds of command-line parameters

Type	Example	Required	Default
Positional	name	Yes	No
Optional	-n (short), --name (long)	No	Yes

1.9 Running our tests

Let's run our tests again to see how we are doing:

```
$ make test
pytest -xv test.py
=====
test session starts =====
...
collected 5 items

test.py::test_exists PASSED [ 20%]
test.py::test_runnable PASSED [ 40%]
test.py::test_executable PASSED [ 60%]
test.py::test_usage PASSED [ 80%]
test.py::test_input PASSED [100%]

===== 5 passed in 0.38s =====
```

Wow, we're passing all our tests! I actually get excited whenever I see my programs pass all their tests, even when I'm the one who wrote the tests. Before we were failing on the usage and input tests. Adding the argparse code fixed both of those because it allows us to accept arguments when our program runs, and it will also create documentation about how to run our program.

1.10 Adding the `main()` function

Our program works really well now, but it's not quite up to community standards and expectations. For instance, it's very common for computer programs—not just ones written in Python—to start at a place called `main()`. Most Python programs define a function called `main()`, and there is an idiom to call the `main()` function at the end of the code, like this:

```
#!/usr/bin/env python3
# Purpose: Say hello

import argparse

def main():
    parser = argparse.ArgumentParser(description='Say hello')
    parser.add_argument('-n', '--name', metavar='name',
                        default='World', help='Name to greet')
    args = parser.parse_args()
    print('Hello, ' + args.name + '!')

if __name__ == '__main__':
    main()
```

def defines a function,
named `main()` in this case.
The empty parentheses
show that this function
accepts no arguments.

Every program or module in Python has a
name that can be accessed through the
variable `__name__`. When the program is
executing, `__name__` is set to "`__main__`".¹

If this is true, call
the `main()` function.

¹ See Python's documentation of `main` for more information: https://docs.python.org/3/library/__main__.html.

As our programs get longer, we'll start creating more functions. Python programmers approach this in different ways, but in this book I will always create and execute a `main()` function to be consistent. To start off, we'll always put the main part of our program inside the `main()` function.

1.11 Adding the `get_args()` function

As a matter of personal taste, I like to put all the `argparse` code into a separate place that I always call `get_args()`. Getting and validating arguments is one concept in my mind, so it belongs by itself. For some programs, this function can get quite long.

I always put `get_args()` as the first function so that I can see it immediately when I read the source code. I usually put `main()` right after it. You are, of course, welcome to structure your programs however you like.

Here is what the program looks like now:

```
#!/usr/bin/env python3
# Purpose: Say hello
import argparse
def get_args():
    parser = argparse.ArgumentParser(description='Say hello')
    parser.add_argument('-n', '--name', metavar='name',
                        default='World', help='Name to greet')
    return parser.parse_args()
def main():
    args = get_args()
    print('Hello, ' + args.name + '!')
if __name__ == '__main__':
    main()
```

The main() function is much shorter now.

The get_args() function is dedicated to getting the arguments. All the argparse code now lives here.

We need to call return to send the results of parsing the arguments back to the main() function.

Call the get_args() function to get parsed arguments. If there is a problem with the arguments or if the user asks for --help, the program never gets to this point because argparse will cause it to exit. If our program does make it this far, the input values must have been OK.

Nothing has changed about the way the program works. We're just organizing the code to group ideas together—the code that deals with `argparse` now lives in the `get_args()` function, and everything else lives in `main()`. Just to be sure, go run the test suite!

1.11.1 Checking style and errors

Our program works really well now. We can use tools like Flake8 and Pylint to check if our program has problems. These tools are called *linters*, and their job is to suggest ways to improve a program. If you haven't installed them yet, you can use the pip module to do so now:

```
$ python3 -m pip install flake8 pylint
```

The Flake8 program wants me to put two blank lines between each of the function def definitions:

```
$ flake8 hello.py
hello.py:6:1: E302 expected 2 blank lines, found 1
hello.py:12:1: E302 expected 2 blank lines, found 1
hello.py:16:1: E305 expected 2 blank lines after class or function definition,
found 1
```



And Pylint says that the functions are missing documentation ("docstrings"):

```
$ pylint hello.py
*****
Module hello
hello.py:1:0: C0114: Missing module docstring (missing-module-docstring)
hello.py:6:0: C0116: Missing function or method docstring (missing-function-
docstring)
hello.py:12:0: C0116: Missing function or method docstring (missing-function-
docstring)

-----
Your code has been rated at 7.00/10 (previous run: -10.00/10, +17.00)
```

A *docstring* is a string that occurs just after the def of the function. It's common to have several lines of documentation for a function, so programmers often will use Python's triple quotes (single or double) to create a multiline string. Following is what the program looks like when I add docstrings. I have also used YAPF to format the program and fix the spacing problems, but you are welcome to use Black or any other tool you like.

```
#!/usr/bin/env python3
"""
Author: Ken Youens-Clark <kyclark@gmail.com>
Purpose: Say hello
"""

import argparse
```

Triple-quoted, multiline docstring for the entire program. It's common practice to write a long docstring just after the shebang to document the overall purpose of the function. I like to include at least my name, email address, and the purpose of the script so that any future person using the program will know who wrote it, how to get in touch with me if they have problems, and what the program is supposed to do.

```

→ # -----
def get_args():
    """Get the command-line arguments"""

    parser = argparse.ArgumentParser(description='Say hello')
    parser.add_argument('-n', '--name',
                        default='World', help='Name to greet')
    return parser.parse_args()

```

A big horizontal “line” comment to help me find the functions. You can omit these if you don’t like them.

```

# -----
def main():
    """Make a jazz noise here"""
    ←
    args = get_args()
    print('Hello, ' + args.name + '!')

```

The `main()` function is simply where the program begins, so there’s not much to say in the docstring. I think it’s (at least a little) funny to always put “Make a jazz noise here,” but you can put whatever you like.

```

# -----
if __name__ == '__main__':
    main()

```

To learn how to use YAPF or Black on the command line, run them with the `-h` or `--help` flag and read the documentation. If you are using an IDE like VS Code or PyCharm, or if you are using the Repl.it interface, there are commands to reformat your code.

1.12 Testing *hello.py*

We’ve made many changes to our program—are we sure it still works correctly? Let’s run our test again.

This is something you will do literally hundreds of times, so I’ve created a shortcut you might like to use. In every directory, you’ll find a file called `Makefile` that looks like this:

```

$ cat Makefile
.PHONY: test

test:
    pytest -xv test.py

```

If you have the program `make` installed on your computer, you can run `make test` when you are in the `01_hello` directory. The `make` program will look for a `Makefile` in your current working directory and then look for a recipe called “`test`.” There it will find that the command to run for the “`test`” target is `pytest -xv test.py`, so it will run that command for you.

```

$ make test
pytest -xv test.py

```

```
===== test session starts =====
...
collected 5 items

test.py::test_exists PASSED [ 20%]
test.py::test_runnable PASSED [ 40%]
test.py::test_executable PASSED [ 60%]
test.py::test_usage PASSED [ 80%]
test.py::test_input PASSED [100%]

===== 5 passed in 0.75s =====
```

If you do not have `make` installed, you might like to install it and learn about how Makefiles can be used to execute complicated sets of commands. If you do not want to install or use `make`, you can always run `pytest -xv test.py` yourself. They both accomplish the same task.

The important point is that we were able to use our tests to verify that our program still does exactly what it is supposed to do. As you write programs, you may want to try different solutions. Tests give you the freedom to rewrite a program (also called “refactoring your code”) and know that it still works.

1.13 Starting a new program with `new.py`

The `argparse` module is a standard module that is always installed with Python. It’s widely used because it can save us so much time in parsing and validating the arguments to our programs. You’ll be using `argparse` in every program for this book, and you’ll learn how you can use it to convert text to numbers, to validate and open files, and much more. There are so many options that I created a Python program called `new.py` that will help you start writing new Python programs that use `argparse`.



I have put this `new.py` program into the `bin` directory of the GitHub repo. I suggest you use it to start every new program you write. For instance, you could create a new version of `hello.py` using `new.py`. Go to the top level of your repository and run this:

```
$ bin/new.py 01_hello/hello.py
"01_hello/hello.py" exists. Overwrite? [yN] n
Will not overwrite. Bye!
```

The `new.py` program will not overwrite an existing file unless you tell it to, so you can use it without worrying that you might erase your work. Try using it to create a program with a different name:

```
$ bin/new.py 01_hello/hello2.py
Done, see new script "01_hello/hello2.py."
```

Now try executing that program:

```
$ 01_hello/hello2.py
usage: hello2.py [-h] [-a str] [-i int] [-f FILE] [-o] str
hello2.py: error: the following arguments are required: str
```

Let's look at the source code of the new program:

```
#!/usr/bin/env python3
"""
Author : Ken Youens-Clark <kyclark@gmail.com>
Date   : 2020-02-28
Purpose: Rock the Casbah
"""

import argparse
import os
import sys

# -----
def get_args():
    """Get command-line arguments"""
    parser = argparse.ArgumentParser(
        description='Rock the Casbah',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('positional',
                        metavar='str',
                        help='A positional argument')
    parser.add_argument('-a',
                        '--arg',
                        help='A named string argument',
                        metavar='str',
                        type=str,
                        default='')

    parser.add_argument('-i',
                        '--int',
                        help='A named integer argument',
                        metavar='int',
                        type=int,
                        default=0)

    parser.add_argument('-f',
                        '--file',
                        help='A readable file',
                        metavar='FILE',
                        type=argparse.FileType('r'),
                        default=None)

    parser.add_argument('-o',
                        '--on',
```

The shebang line should use the env program to find the python3 program.

This docstring is for the program as a whole.

These lines import various modules that the program needs.

The get_args() function is responsible for parsing and validating arguments.

Define a “positional” argument like our first version of hello.py that had a name argument.

Define an “optional” argument like when we changed to use the --name option.

Define an optional argument that must be an integer value.

Define an optional argument that must be a file.

Define a “flag” option that is either “on” when present or “off” when absent. You’ll learn more about these later.

```

        help='A boolean flag',
        action='store_true')

return parser.parse_args()      ← Return the parsed arguments to main(). If there
                                are any problems, like if the --int value is some text
                                rather than a number like 42, argparse will print
                                an error message and the "usage" for the user.

# ----- ← Define the main() function
def main():                      where the program starts.

    """Make a jazz noise here"""

    args = get_args()             ← The first thing our main() functions will always
    str_arg = args.arg            do is call get_args() to get the arguments.

    int_arg = args.int

    file_arg = args.file

    flag_arg = args.on

    pos_arg = args.positional

    print(f'str_arg = "{str_arg}"')
    print(f'int_arg = "{int_arg}"')
    print('file_arg = "{}".format(file_arg.name if file_arg else "")')
    print(f'flag_arg = "{flag_arg}"')
    print(f'positional = "{pos_arg}"')

# ----- ← Each argument's value is accessible through
if __name__ == '__main__':          the long name of the argument. It is not
    main()                         required to have both a short and long name,
                                but it is common and tends to make your
                                program more readable.

                                ← If the condition is true, this
                                calls the main() function.

                                ← When the program is being executed,
                                the __name__ value will be equal to
                                the text "__main__".

```

This program will accept the following arguments:

- A single positional argument of the type `str`. *Positional* means it is not preceded by a flag to name it but has meaning because of its position relative to the command name.
- An automatic `-h` or `--help` flag that will cause `argparse` to print the usage.
- A string option called either `-a` or `--arg`.
- A named option argument called `-i` or `--int`.
- A file option called `-f` or `--file`.
- A Boolean (off/on) flag called `-o` or `--on`.

Looking at the preceding list, you can see that `new.py` has done the following for you:

- Created a new Python program called `hello2.py`
- Used a template to generate a working program complete with docstrings, a `main()` function to start the program, a `get_args()` function to parse and document various kinds of arguments, and code to start the program running in the `main()` function
- Made the new program executable so that it can be run like `./hello2.py`

The result is a program that you can immediately execute and that will produce documentation on how to run it. After you use new.py to create your new program, you should open it with your editor and modify the argument names and types to suit the needs of your program. For instance, in chapter 2 you'll be able to delete everything but the positional argument, which you should rename from 'positional' to something like 'word' (because the argument is going to be a word).

Note that you can control the "name" and "email" values that are used by new.py by creating a file called .new.py (note the leading dot!) in your home directory. Here is mine:

```
$ cat ~/.new.py
name=Ken Youens-Clark
email=kyclark@gmail.com
```

1.14 Using template.py as an alternative to new.py

If you don't want to use new.py, I have included a sample of the preceding program as template/template.py, which you can copy. For instance, in chapter 2 you will need to create the program 02_crowsnest/crowsnest.py.

You can do this with new.py from the top level of the repository:

```
$ bin/new.py 02_crowsnest/crowsnest.py
```

Or you can use cp (copy) command to copy the template to your new program:

```
$ cp template/template.py 02_crowsnest/crowsnest.py
```

The main point is that you won't have to start every program from scratch. I think it's much easier to start with a complete, working program and modify it.

NOTE You can copy new.py to your ~/bin directory. Then you can use it from any directory to create a new program.

Be sure to skim the appendix—it has many examples of programs that use argparse. You can copy many of those examples to help you with the exercises.

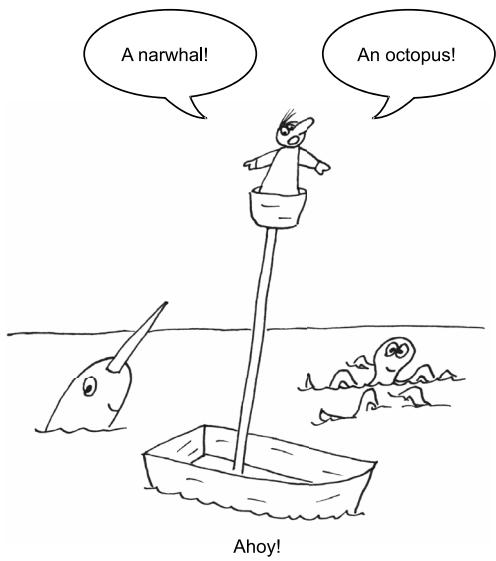
Summary

- A Python program is plain text that lives in a file. You need the python3 program to interpret and execute the program file.
- You can make a program executable and copy it to a location in your \$PATH so that you can run it like any other program on your computer. Be sure to set the shebang to use env to find the correct python3.
- The argparse module will help you document and parse all the parameters to your program. You can validate the types and numbers of arguments, which can be positional, optional, or flags. The usage will be automatically generated.

- We will use the `pytest` program to run the `test.py` programs for each exercise. The `make test` shortcut will execute `pytest -xv test.py`, or you can run this command directly.
- You should run your tests often to ensure that everything works.
- Code formatters like YAPF and Black will automatically format your code to community standards, making it easier to read and debug.
- Code linters like Pylint and Flake8 can help you correct both programmatic and stylistic problems.
- You can use the `new.py` program to generate new Python programs that use `argparse`.

The crow's nest: Working with strings

Avast, you corny-faced gollumpus! Ye are barrelman for this watch. D'ye ken what I mean, ye addlepated blunderbuss?! Ah, landlubber ye be! OK, then, you are the lookout in the crow's nest—the little bucket attached to the top of a mast of a sailing ship. Your job is to keep a lookout for interesting or dangerous things, like a ship to plunder or an iceberg to avoid. When you see something like a narwhal, you are supposed to cry out, “Ahoy, Captain, *a narwhal* off the larboard bow!” If you see an octopus, you’ll shout “Ahoy, Captain, *an octopus* off the larboard bow!” (We’ll assume everything is “off the larboard bow” for this exercise. It’s a great place for things to be.)



From this point on, each chapter will present a coding challenge that you should complete on your own. I will discuss the key ideas you’ll need to solve the problems as well as how to use the provided tests to determine when your program is correct. You should have a copy of the Git repository locally (see the setup instructions in the book’s introduction), and you should write each program in that chapter’s directory. For example, this chapter’s program should be written in the `02_crowsnest` directory, where the tests for the program live.

In this chapter, we’re going to start working with strings. By the end, you will be able to

- Create a program that accepts a positional argument and produces usage documentation
- Create a new output string depending on the inputs to the program
- Run a test suite

Your program should be called `crowsnest.py`. It will accept a single positional argument and will print the given argument inside the “Ahoy” bit, along with the word “a” or “an” depending on whether the argument starts with a consonant or a vowel.

That is, if given “narwhal,” it should do this:

```
$ ./crowsnest.py narwhal
Ahoy, Captain, a narwhal off the larboard bow!
```

And if given “octopus,”

```
$ ./crowsnest.py octopus
Ahoy, Captain, an octopus off the larboard bow!
```

This means you’re going to need to write a program that accepts some input on the command line, decides on the proper article (“a” or “an”) for the input, and prints out a string that puts those two values into the “Ahoy” phrase.

2.1 Getting started

You’re probably ready to start writing the program! Well, hold on just a minute longer, ye duke of limbs. We need to discuss how you can use the tests to know when your program is working and how you might get started programming.

2.1.1 How to use the tests

“The greatest teacher, failure is.”

—Yoda

In the code repository, I’ve included tests that will guide you in the writing of your program. Before you even write the first line of code, I’d like you to run the tests so you can look at the first failed test:

```
$ cd 02_crowsnest
$ make test
```

Instead of `make test` you could also run `pytest -xv test.py`. Among the output, you’ll see this line:

```
$ pytest -xv test.py
===== test session starts =====
...
```

```
collected 6 items

→ test.py::test_exists FAILED [ 16%]
```

This test failed. There are more tests after this, but testing stops here because of the `-x` flag to pytest.

You'll also see lots of other output trying to convince you that the expected file, `crowsnest.py`, does not exist. Learning to read the test output is a skill in itself—it takes quite a bit of practice, so try not to feel overwhelmed. In my terminal (iTerm on a Mac), the output from pytest shows colors and bold print to highlight key failures. The text in bold, red letters is usually where I start, but your terminal may behave differently.

Let's take a gander at the output. It does look a bit daunting at first, but you'll get used to reading the messages and finding your errors.

The “E” at the beginning of this line is the “Error” you should read. It’s very difficult to understand what the test is trying to tell you, but essentially the `./crowsnest.py` file does not exist.

```
===== FAILURES =====
test_exists
```

```
def test_exists(): <-- This is the actual code inside
    """exists"""  

>     assert os.path.isfile(prg) <-- The ">" at the beginning of this line
E     AssertionError: assert False indicates this is the line where the
E         + where False = <function isfile at 0x1086f1310>('./crowsnest.py')
E         + where <function isfile at 0x1086f1310> = <module 'posixpath'
from '/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8 posixpath.
py'>.isfile
E         + where <module 'posixpath' from
'./crowsnest.py'> = os.path
```

If you haven't created it, this will fail as expected.

```
test.py:22: AssertionError
!!!!!!!!!!!!!! stopping after 1 failures !!!!!!!
===== 1 failed in 0.05s =====
```

This warns that no more tests will run after the one failure. This is because we ran it with the flag to stop testing at the first failure.

The first test for every program in the book checks that the expected file exists, so let's create it!

2.1.2 Creating programs with new.py

In order to pass the first test, you need to create a file called `crowsnest.py` inside the `02_crowsnest` directory where `test.py` is located. While it's perfectly fine to start writing from scratch, I suggest you use the `new.py` program to print some useful boilerplate code that you'll need in every exercise.

From the top level of the repository, you can run the following command to create the new program.

```
$ bin/new.py 02_crowsnest/crowsnest.py
Done, see new script "02_crowsnest/crowsnest.py."
```

If you don't want to use new.py, you can copy the template/template.py program:

```
$ cp template/template.py 02_crowsnest/crowsnest.py
```

You should now have the outline of a working program that accepts command-line arguments. If you run your new crowsnest.py with no arguments, it will print a short usage statement like the following (notice how “usage” is the first word of the output):

```
$ ./crowsnest.py
usage: crowsnest.py [-h] [-a str] [-i int] [-f FILE] [-o] str
crowsnest.py: error: the following arguments are required: str
```

Run it with `./crowsnest.py --help`. It will print a longer help message too.

NOTE Those are not the correct parameters for our program, just the default examples supplied by new.py. You will need to modify them to suit this program.

2.1.3 Write, test, repeat

You just created the program, so you ought to be able to pass the first test. The cycle I hope you'll develop is to write a very small amount of code—literally one or two lines at most—and then run the program or the tests to see how you're doing.

Let's run the tests again:

The expected file exists, so this test passes.

```
$ make test
pytest -xv test.py
=====
test session starts =====
...
collected 6 items
→ test.py::test_exists PASSED [ 16%]
→ test.py::test_usage PASSED [ 33%] ←
→ test.py::test_consonant FAILED [ 50%] ←
```

The test `_consonant()` test is failing. That's OK! We haven't even started writing the actual program, but at least we have a place to start.

As you can see, creating a new program with new.py will make you pass the first *two* tests:

- 1 Does the program exist? Yes, you just created it.
- 2 Does the program print a help message when you ask for help? Yes, you ran it above with no arguments and the `--help` flag, and you saw that it will produce help messages.

Now you have a working program that accepts some arguments (but not the right ones). Next you need to make your program accept the “narwhal” or “octopus” value that needs to be announced. We’ll use command-line arguments to do that.

2.1.4 Defining your arguments

Figure 2.1 is sure to shiver your timbers, showing the inputs (or *parameters*) and output of the program. We’ll use these diagrams throughout the book to imagine how code and data work together. In this program, the input is a word, and a phrase incorporating that word with the correct article is the output.

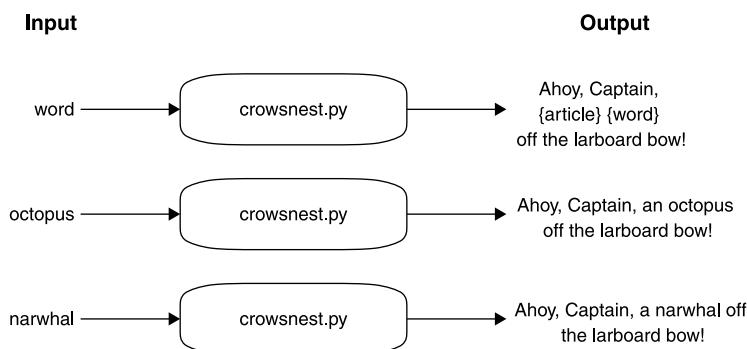


Figure 2.1 The input to the program is a word, and the output is that word plus its proper article (and some other stuff).

We need to modify the part of the program that gets the arguments—the aptly named `get_args()` function. This function uses the `argparse` module to parse the command-line arguments, and our program needs to accept a single, positional argument. If you’re unsure what a “positional” argument is, be sure to read the appendix, especially section A.4.1.

The `get_args()` function created by the template names the first argument `positional`. Remember that positional arguments are defined by their positions and don’t have names that start with dashes. You can delete all the arguments except for the positional word. Modify the `get_args()` part of your program until it will print this usage:

```
$ ./crowsnest.py
usage: crowsnest.py [-h] word
crowsnest.py: error: the following arguments are required: word
```

Likewise, it should print longer usage documentation for the `-h` or `--help` flag:

```
$ ./crowsnest.py -h
usage: crowsnest.py [-h] word
```

```
Crow's Nest -- choose the correct article
```

```
positional arguments:
  word      A word
```

You need to define a word parameter. Notice that it is listed as a positional argument.

```
optional arguments:
```

```
-h, --help show this help message and exit
```

The `-h` and `--help` flags are created automatically by argparse. You are not allowed to use these as options. They are used to create the documentation for your program.

Do not proceed until your usage matches the preceding!

When your program prints the correct usage, you can get the word argument inside the `main` function. Modify your program so that it will print the word:

```
def main():
    args = get_args()
    word = args.word
    print(word)
```

Then test that it works:

```
$ ./crowsnest.py narwhal
narwhal
```

And now run your tests again. You should still be passing two and failing the third. Let's read the test failure:

It's not terribly important right now to understand this line, but the `getoutput()` function is running the program with a word. We're going to talk about the f-string in this chapter. The output from running the program will go into the `out` variable, which will be used to see if the program created the correct output for a given word. None of the code in this function is anything you should worry about being able to write yet.

```
===== FAILURES =====
test_consonant
```

```
def test_consonant():
    """brigantine -> a brigantine"""

    for word in consonant_words:
```

The line starting with ">" shows the code that produced an error. The output of the program is compared to an expected string. Since it didn't match, the assert produces an exception.

```
>     out = getoutput(f'{prg} {word}')
E         assert out.strip() == template.format('a', word)
E             AssertionError: assert 'brigantine' == 'Ahoy, Captain...larboard bow!'
E                 - brigantine
E                 + Ahoy, Captain, a brigantine off the larboard bow!
```

This line starts with "E" to indicate the error.

The line starting with the plus sign (+) is what the test expected: "Ahoy, Captain, a brigantine off the larboard bow!"

The line starting with a hyphen (-) is what the test got when it ran with the argument "brigantine"—it got back the word "brigantine."

So we need to get the word into the "Ahoy" phrase. How can we do that?

2.1.5 Concatenating strings

Putting strings together is called *concatenating* or *joining* strings. To demonstrate, I'll enter some code directly into the Python interpreter. I want you to type along. No, really! Type everything you see, and try it for yourself.

Open a terminal and type `python3` or `ipython` to start a REPL. A REPL is a Read-Evaluate-Print-Loop—Python will *read* each line of input, *evaluate* it, and *print* the results in a *loop*. Here's what it looks like on my system:

```
$ python3
Python 3.8.1 (v3.8.1:1b293b6006, Dec 18 2019, 14:08:53)
[Clang 6.0 (clang-600.0.57)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>>
```

The “`>>>`” is a prompt where you can type code. Remember *not* to type that part! To exit the REPL, either type `quit()` or press Ctrl-D (the Control key plus the letter *D*).

NOTE You may prefer to use Python's IDLE (integrated development and learning environment) program, IPython, or Jupyter Notebooks to interact with the language. I'll stick to the `python3` REPL throughout the book.

Let's start off by assigning the variable `word` to the value “narwhal.” In the REPL, type `word = 'narwhal'` and press Enter:

```
>>> word = 'narwhal'
```

Note that you can put as many (or no) spaces around the `=` as you like, but convention and readability (and tools like Pylint and Flake8 that help you find errors in your code) ask you to use exactly one space on either side.

If you type `word` and press Enter, Python will print the current value of `word`:

```
>>> word
'narwhal'
```

Now type `werd` and press Enter:

```
>>> werd
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'werd' is not defined
```

WARNING There is no `werd` variable because we haven't set `werd` to be anything. Using an undefined variable causes an *exception* that will crash your program. Python will happily create `werd` for you when you assign it a value.

We need to insert the `word` between two other strings. The `+` operator can be used to join strings together:

```
>>> 'Ahoy, Captain, a ' + word + ' off the larboard bow!'
'Ahoy, Captain, a narwhal off the larboard bow!'
```

If you change your program to print() that string instead of just printing the word, you should be able to pass four tests:

```
test.py::test_exists PASSED [ 16%]
test.py::test_usage PASSED [ 33%]
test.py::test_consonant PASSED [ 50%]
test.py::test_consonant_upper PASSED [ 66%]
test.py::test_vowel FAILED [ 83%]
```

If you look closely at the failure, you'll see this:

```
E      - Ahoy, Captain, a aviso off the larboard bow!
E      + Ahoy, Captain, an aviso off the larboard bow!
E          ?           +
```

We hardcoded the “a” before the word, but we really need to figure out whether to use “a” or “an” depending on whether the word starts with a vowel. How can we do that?

2.1.6 Variable types

Before we go much further, I need to take a small step back and point out that our word variable is a *string*. Every variable in Python has a *type* that describes the kind of data it holds. Because we put the value for word in quotes ('narwhal'), word holds a *string*, which Python represents with a class called *str*. (A *class* is a collection of code and functions that we can use.)

The type() function will tell you what kind of data Python thinks something is:

```
>>> type(word)
<class 'str'>
```

Whenever you put a value in single quotes (' ') or double quotes (" "), Python will interpret it as a *str*:

```
>>> type("submarine")
<class 'str'>
```

WARNING If you forget the quotes, Python will look for some variable or function by that name. If there is no variable or function by that name, it will cause an exception:

```
>>> word = narwhal
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'narwhal' is not defined
```

Exceptions are bad, and we will try to write code that avoids them, or at least knows how to handle them gracefully.

2.1.7 Getting just part of a string

Back to our problem. We need to put either “a” or “an” in front of the word we’re given, based on whether the first character of word is a vowel or a consonant.

In Python, we can use square brackets and an *index* to get an individual character from a string. The index is the numeric position of an element in a sequence, and we must remember that indexing starts at 0.

```
>>> word = 'narwhal'
>>> word[0]
'n'
```

n a r w h a l
0 1 2 3 4 5 6



You can index into a literal string value too:

```
>>> 'narwhal'[0]
'n'
```

n a r w h a l
0 1 2 3 4 5 6



Because the index values start with 0, that means the last index is *one less than the length* of the string, which is often confusing. The length of “narwhal” is 7, but the last character is found at index 6:

```
>>> word[6]
'l'
```

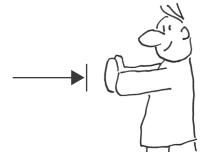
n a r w h a l
-6 -5 -4 -3 -2 -1



You can also use negative index numbers to count backwards from the end, so the last index is also -1:

```
>>> word[-1]
'l'
```

n a r w h a l
0 1 2 3 4 5 6



You can use *slice notation* [start:stop] to get a range of characters. Both start and stop are optional. The default value for start is 0 (the beginning of the string), and the stop value is *not inclusive*:

```
>>> word[:3]
'nar'
```

The default value for `stop` is the end of the string:

```
>>> word[3:]
'whal'
```

In the next chapter, you’ll see that this is the same as the syntax for slicing lists. A string is (sort of) a list of characters, so this isn’t too strange.

2.1.8 **Finding help in the REPL**

The `str` class has a ton of functions we can use to handle strings, but what are they? A large part of programming is knowing how to ask questions and where to look for answers. A common refrain you may hear is “RTFM”—Read the Fine Manual. The Python community has created reams of documentation, which are all available at <https://docs.python.org/3/>. You will need to refer to the documentation constantly to remind yourself (and discover) how to use certain functions. The docs for the string class are here: <https://docs.python.org/3/library/string.html>.

I prefer to read the docs directly inside the REPL, in this case by typing `help(str)`:

```
>>> help(str)
```

Inside the `help`, you move up and down in the text using the up and down cursor arrows on your keyboard. You can also press the spacebar or the letter `F` (or sometimes `Ctrl-F`) to jump forward to the next page, and the letter `B` (or sometimes `Ctrl-B`) to jump backward. You can search through the documentation by pressing `/` and then the text you want to find. If you press `N` (for “next”) after a search, you will jump to the next place that string is found. To leave the help, press `Q` (for “quit”).



2.1.9 **String methods**

Now that we know `word` is a string (`str`), we have all these incredibly useful methods we can call on the variable. (A *method* is a function that belongs to a variable like `word`.)

For instance, if I wanted to shout about the fact that we have a narwhal, I could print it in UPPERCASE LETTERS. If I search through the help, I will see that there is a function called `str.upper()`. Here is how you can *call* or *execute* that function:

```
>>> word.upper()
'NARWHAL'
```

You must include the parentheses, `()`, or else you’re talking about the *function itself*:

```
>>> word.upper
<built-in method upper of str object at 0x10559e500>
```



That will actually come in handy later, when we use functions like `map()` and `filter()`, but for now we want Python to execute the `str.upper()` function on the variable `word`, so we add the parentheses. Note that the function returns an uppercase version of the word but *does not* change the value of `word`:

```
>>> word
'narwhal'
```

There is another `str` function with “upper” in the name called `str.isupper()`. The name helps you know that this will return a true/false type answer. Let’s try it:

```
>>> word.isupper()
False
```

We can chain methods together like so:

```
>>> word.upper().isupper()
True
```

That makes sense. If I convert `word` to uppercase, then `word.isupper()` returns `True`.

I find it odd that the `str` class does not include a method to get the length of a string. For that, we must use a separate function called `len()`, short for “length”:

```
>>> len('narwhal')
7
>>> len(word)
7
```



Are you typing all this into Python yourself? I recommend you do! Find other methods in the `str` help, and try them out.

2.1.10 String comparisons

You now know how to get the first letter of `word` by using `word[0]`. Let’s assign it to the variable `char`:

```
>>> word = 'octopus'
>>> char = word[0]
>>> char
'o'
```

If you check the `type()` of your new `char` variable, it is a `str`. Even a single character is still considered by Python to be a string:

```
>>> type(char)
<class 'str'>
```

Now we need to figure out if `char` is a vowel or a consonant. We’ll say that the letters “a,” “e,” “i,” “o,” and “u” make up our set of “vowels.” You can use `==` to compare strings:

```
>>> char == 'a'
False
>>> char == 'o'
True
```

NOTE Be careful to always use one equal sign (=) when *assigning a value* to a variable, like `word = 'narwhal'` and two equal signs (==, which, in my head, I pronounce “equal-equal”) when you *compare two values* like `word == 'narwhal'`. The first is a statement that changes the value of `word`, and the second is an *expression* that returns True or False (see figure 2.2).

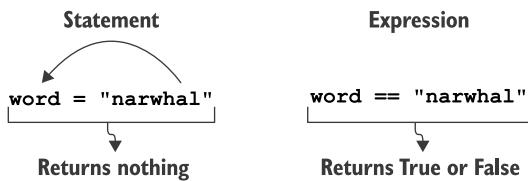


Figure 2.2 An expression returns a value. A statement does not.

We need to compare our `char` to *all* the vowels. You can use `and` and `or` in such comparisons, and they will be combined according to standard Boolean algebra:

```
>>> char == 'a' or char == 'e' or char == 'i' or char == 'o' or char == 'u'
True
```

What if the word is “Octopus” or “OCTOPUS”?

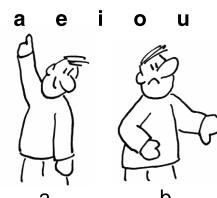
```
>>> word = 'OCTOPUS'
>>> char = word[0]
>>> char == 'a' or char == 'e' or char == 'i' or char == 'o' or char == 'u'
False
```

Do we have to make 10 comparisons in order to check the uppercase versions, too? What if we were to lowercase `word[0]`? Remember that `word[0]` returns a `str`, so we can chain other `str` methods onto that:

```
>>> word = 'OCTOPUS'
>>> char = word[0].lower()
>>> char == 'a' or char == 'e' or char == 'i' or char == 'o' or char == 'u'
True
```

An easier way to determine if `char` is a vowel would be to use Python’s `x in y` construct, which will tell us if the value `x` is in the collection `y`. We can ask whether the letter ‘a’ is in the longer string ‘aeiou’:

```
>>> 'a' in 'aeiou'
True
```



But the letter 'b' is not:

```
>>> 'b' in 'aeiou'
False
```

Let's use that to test the first character of the lowercased word (which is 'o'):

```
>>> word = 'OCTOPUS'
>>> word[0].lower() in 'aeiou'
True
```

2.1.11 Conditional branching

Once you have figured out if the first letter is a vowel, you will need to select an article. We'll use a very simple rule: if the word starts with a vowel, choose "an"; otherwise, choose "a." This misses exceptions like when the initial "h" in a word is silent. For instance, we say "a hat" but "an honor." Nor will we consider the case where an initial vowel has a consonant sound, as in "union," where the "u" sounds like a "y."

We can create a new variable called `article` that we will set to the empty string, and we'll use an `if/else` statement to figure out what to put in it:

```
article = ''           | Initialize article to
if word[0].lower() in 'aeiou':    | the empty string.
    article = 'an'      | Check if the first, lowercased
else:                      | character of word is a vowel.
    article = 'a'       | Set article to "an" if the
                        | first character is a vowel.
                        | Set article to "a" if the
                        | first character is not a vowel.
```

Here is a much shorter way to write that with an `if expression` (expressions return values; statements do not). The `if expression` is written a little backwards. First comes the value if the test (or "predicate") is True, then the predicate, and then the value if the predicate is False (figure 2.3).

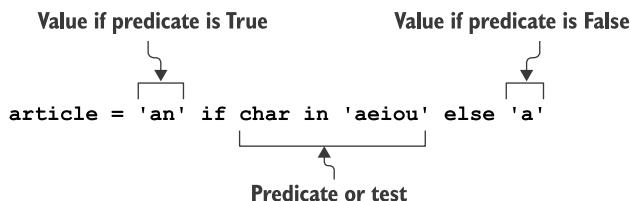


Figure 2.3 The `if` expression will return the first value if the predicate is True and the second value otherwise.

This approach is also safer because the `if expression` is *required* to have the `else`. There's no chance that we could forget to handle both cases:

```
>>> char = 'o'
>>> article = 'an' if char in 'aeiou' else 'a'
```

Let's verify that we have the correct article:

```
>>> article
'an'
```

2.1.12 String formatting

Now we have two variables, `article` and `word`, that need to be incorporated into our “Ahoy!” phrase. You saw earlier that we can use the plus sign (+) to concatenate strings. Another method for creating new strings from other strings is to use the `str.format()` method.

To do so, you create a string template with curly brackets {}, which indicate placeholders for values. The values that will be substituted are arguments to the `str.format()` method, and they are substituted in the same order that the {} appear (figure 2.4).

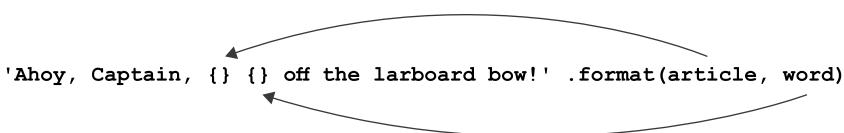


Figure 2.4 The `str.format()` method is used to expand the values of variables inside strings.

Here it is in code:

```
>>> 'Ahoy, Captain, {} {} off the larboard bow!'.format(article, word)
'Ahoy, Captain, an octopus off the larboard bow!'
```

Another method for combining strings uses the special “f-string” where you can put the variables directly into the curly brackets {}. It’s a matter of taste which approach you choose; I tend to prefer this style because I don’t have to think about which variable goes with which set of brackets:

```
>>> f'Ahoy, Captain, {article} {word} off the larboard bow!'
'Ahoy, Captain, an octopus off the larboard bow!'
```

NOTE In some programming languages, you have to declare the variable’s name and what *type* of data it will hold. If a variable is declared to be a number, it can never hold a different type of value, like a string. This is called *static typing* because the type of the variable can never change.

Python is a *dynamically typed* language, which means you do not have to declare a variable or what kind of data the variable will hold. You can change the value and type of data at any time. This could be either great or terrible news. As Hamlet says, “There is nothing either good or bad, but thinking makes it so.”



2.1.13 Time to write

Here are a few hints for writing your solution:

- Start your program with new.py and fill in `get_args()` with a single positional argument called `word`.
- You can get the first character of the word by indexing it like a list, `word[0]`.
- Unless you want to check both upper- and lowercase letters, you can use either the `str.lower()` or `str.upper()` method to force the input to one case for checking whether the first character is a vowel or consonant.
- There are fewer vowels (five, if you recall) than consonants, so it's probably easier to check whether the first character is one of those.
- You can use the `x in y` syntax to see if the element `x` is in the collection `y`, with the collection here being a list.
- Use `str.format()` or f-strings to insert the correct article for the given word into the longer phrase.
- Run `make test` (or `pytest -xv test.py`) after *every* change to your program to ensure that your program compiles and is on the right track.

Now go write the program before you turn the page and study my solution. Look alive, you ill-tempered shabaroon!

2.2 Solution

Following is one way you could write a program that satisfies the test suite:

```
#!/usr/bin/env python3
"""Crow's Nest"""

# -----
# Define the function get_args () to handle the
# command-line arguments. I like to put this first so
# I can see it right away when I'm reading the code.
# -----
def get_args():
    """Get command-line arguments"""
    # -----
    # ----- parser = argparse.ArgumentParser(
    # ----- description="Crow's Nest -- choose the correct article",
    # ----- formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    #
    # ----- parser.add_argument('word', metavar='word', help='A word')
    #
    # ----- return parser.parse_args()
    # -----
# -----
# Define the main() function
# where the program will start.
def main():
    """Make a jazz noise here"""
    # -----
    args = get_args()
    word = args.word
    # -----
```

The annotations provide context for specific parts of the code:

- The parser will parse the arguments.** Points to the `argparse.ArgumentParser` constructor.
- Show the default values for each parameter in the usage.** Points to the `formatter_class` parameter of the `ArgumentParser` constructor.
- args contains the return value from the get_args() function.** Points to the `args` variable assignment.
- Define the function get_args () to handle the command-line arguments. I like to put this first so I can see it right away when I'm reading the code.** Points to the `get_args` function definition.
- The description shows in the usage to describe what the program does.** Points to the `description` parameter of the `ArgumentParser` constructor.
- The result of parsing the arguments will be returned to main().** Points to the `parse_args` method call.
- Define a positional argument called word.** Points to the `'word'` argument in the `add_argument` call.
- Put the args.word value from the arguments into the word variable.** Points to the `word = args.word` assignment.

```
→ article = 'an' if word[0].lower() in 'aeiou' else 'a'

    print(f'Ahoy, Captain, {article} {word} off the larboard bow!') ←
```

```
# -----
if __name__ == '__main__':
    main()
```

Choose the correct article, using an if expression to see if the lowercased, first character of word is or is not in the set of vowels.

Check if we are in the “main” namespace, which means the program is running.

If we are in the “main” namespace, call the main() function to make the program start.

Print the output string using an f-string to interpolate the article and word variables inside the string.

2.3 Discussion

I’d like to stress that the preceding listing is *a* solution, not *the* solution. There are many ways to express the same idea in Python. As long as your code passes the test suite, it is correct.

That said, I created my program with new.py, which automatically gives me two functions:

- `get_args()`, where I define the arguments to the program
- `main()`, where the program starts

Let’s talk about these two functions.

2.3.1 Defining the arguments with `get_args()`

I prefer to put the `get_args()` function first so that I can see right away what the program expects as input. You don’t have to define this as a separate function—you could put all this code inside `main()`, if you prefer. Eventually our programs are going to get longer, though, and I think it’s nice to keep this as a separate idea. Every program I present will have a `get_args()` function that will define and validate the input.

Our program specifications (the “specs”) say that the program should accept one positional argument. I changed the ‘positional’ argument name to ‘word’ because I’m expecting a single word:

```
parser.add_argument('word', metavar='word', help='Word')
```

I recommend you never leave the positional argument named ‘positional’ because it is an entirely nondescriptive term. Naming your variables according to *what they are* will make your code more readable.

The program doesn’t need any of the other options created by `new.py`, so you can delete the rest of the `parser.add_argument()` calls.

The `get_args()` function will return the result of parsing the command-line arguments that I put into the variable `args`:

```
return parser.parse_args()
```

If argparse is not able to parse the arguments—for example, if there are none—it will never return from `get_args()` but will instead print the “usage” for the user and exit with an error code to let the operating system know that the program exited without success. (On the command line, an exit value of 0 means there were 0 errors. Anything other than 0 is considered an error.)

2.3.2 *The main() thing*

Many programming languages will automatically start from the `main()` function, so I always define a `main()` function and start my programs there. This is not a requirement, but it’s an extremely common idiom in Python. Every program I present will start with a `main()` function that will first call `get_args()` to get the program’s inputs:

```
def main():
    args = get_args()
```

I can now access the word by calling `args.word`. Note the lack of parentheses. It’s not `args.word()` because it is not a function call. Think of `args.word` as being like a slot where the value of the word lives:

```
word = args.word
```

I like to work through my ideas using the REPL, so I’m going to pretend that `word` has been set to “octopus”:

```
>>> word = 'octopus'
```

2.3.3 *Classifying the first character of a word*

To figure out whether the article I choose should be a or an, I need to look at the first character of the word. In the introduction, we used this:

```
>>> word[0]
'o'
```

I can check if the first character is in the string of vowels, both lower- and uppercase:

```
>>> word[0] in 'aeiouAEIOU'
True
```

I can make this shorter, however, if I use the `word.lower()` function. Then I’d only have to check the lowercase vowels:

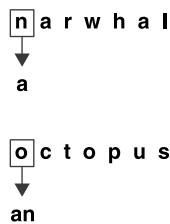
```
>>> word[0].lower() in 'aeiou'
True
```

Remember that the `x in y` form is a way to ask if element `x` is in the collection `y`. You can use it for letters in a longer string (like the list of vowels):

```
>>> 'a' in 'aeiou'
True
```

You can use membership in the list of vowels as a condition to choose “an”; otherwise, we choose “a.” As mentioned in the introduction, the `if` expression is the shortest and safest way to make a *binary choice* (where there are only two possibilities):

```
>>> article = 'an' if word[0].lower() in 'aeiou' else 'a'
>>> article
'an'
```



The safety of the `if` expression comes from the fact that Python will not even run this program if you forget the `else`. Try it and see what error you get.

Let’s change the value of `word` to “galleon” and check that it still works:

```
>>> word = 'galleon'
>>> article = 'an' if word[0].lower() in 'aeiou' else 'a'
>>> article
'a'
```

2.3.4 Printing the results

Finally we need to print out the phrase with our article and word. As noted in the introduction, you can use the `str.format()` function to incorporate the variables into a string:

```
>>> article = 'a'
>>> word = 'ketch'
>>> print('Ahoy, Captain, {} {} off the larboard bow!'.format(article, word))
Ahoy, Captain, a ketch off the larboard bow!
```

Python’s f-strings will *interpolate* any code inside the `{}` placeholders, so variables get turned into their contents:

```
>>> print(f'Ahoy, Captain, {article} {word} off the larboard bow!')
Ahoy, Captain, a ketch off the larboard bow!
```

However you choose to print out the article and word is fine, as long as it passes the tests. While it’s a matter of personal taste which you choose, I find f-strings a bit easier to read, as my eyes don’t have to jump back and forth from the `{}` placeholders to the variables that will go inside them.

2.3.5 Running the test suite

“A computer is like a mischievous genie. It will give you exactly what you ask for, but not always what you want.”

—Joe Sondow

Computers are a bit like bad genies. They will do exactly what you tell them, but not necessarily what you *want*. In an episode of *The X-Files*, the character Mulder wishes for peace on Earth, and a genie removes all humans but him.

Tests are what we can use to verify that our programs are doing what we *actually* want them to do. Tests can never prove that our program is truly free from errors, only that the bugs we imagined or found while writing the program no longer exist. Still, we write and run tests because they are really quite effective and much better than not doing so.

This is the idea behind *test-driven development*:

- Write tests *before* we write the software.
- Run the tests to verify that our as-yet-unwritten software fails to deliver on some task.
- Write the software to fulfill the request.
- Run the tests to check that it now *does* work.
- Keep running all the tests to ensure that when we add some new code we do not break existing code.

We won't be discussing how to *write* tests just yet. That will come later. For now, I've written all the tests for you. I hope that by the end of this book, you will see the value of testing and will always start off by writing *tests first and code second!*

2.4 Going further

- Have your program match the case of the incoming word (for example, "an octopus" and "An Octopus"). Copy an existing `test_` function in `test.py` to verify that your program works correctly while still passing all the other tests. Try writing the test first, and then make your program pass the test. That's *test-driven development*!
- Accept a new parameter that changes "larboard" (the left side of the boat) to "starboard" (the right side¹). You could either make an option called `--side` that defaults to "larboard," or you could make a `--starboard` flag that, if present, changes the side to "starboard."
- The provided tests only give you words that start with an actual alphabetic character. Expand your code to handle words that start with numbers or punctuation. Should your program reject these? Add more tests to ensure that your program does what you intend.



¹ "Starboard" has nothing to do with stars but with the "steering board" or rudder, which typically was on the right side of the boat for right-handed sailors.

Summary

- All Python's documentation is available at <https://docs.python.org/3/> and via the `help` command in the REPL.
- Variables in Python are dynamically typed according to whatever value you assign them, and they come into existence when you assign a value to them.
- Strings have methods like `str.upper()` and `str.isupper()` that you can call to alter them or get information.
- You can get parts of a string by using square brackets and indexes like `[0]` for the first letter or `[-1]` for the last.
- You can concatenate strings with the `+` operator.
- The `str.format()` method allows you to create a template with `{}` placeholders that get filled in with arguments.
- F-strings like `f'{article} {word}'` allow variables and code to go directly inside the brackets.
- The `x in y` expression will report whether the value `x` is present in the collection `y`.
- Statements like `if/else` do not return a value, whereas expressions like `x if y else z` do return a value.
- Test-driven development is a way to ensure programs meet some minimum criteria of correctness. Every feature of a program should have tests, and writing and running test suites should be an integral part of writing programs.

Going on a picnic: Working with lists

Writing code makes me hungry! Let's write a program to list some tasty foods we'd like to eat.

So far we've worked with single variables, like a name to say "hello" to or a nautical-themed object to point out. In this program, we want to keep track of one or more foods that we will store in a `list`, a variable that can hold any number of items. We use lists all the time in real life. Maybe it's your top-five favorite songs, your birthday wish list, or a bucket list of the best types of buckets.

In this chapter, we're going on a picnic, and we want to print a list of items to bring along. You will learn to

- Write a program that accepts multiple positional arguments
- Use `if`, `elif`, and `else` to handle conditional branching with three or more options
- Find and alter items in a list
- Sort and reverse lists
- Format a list into a new string

The items for the list will be passed as positional arguments. When there is only one item, you'll print that:

```
$ ./picnic.py salad  
You are bringing salad.
```



What? Who just brings salad on a picnic? When there are two items, you'll print "and" between them:

```
$ ./picnic.py salad chips
You are bringing salad and chips.
```



Hmm, chips. That's an improvement. When there are three or more items, you'll separate the items with commas:

```
$ ./picnic.py salad chips cupcakes
You are bringing salad, chips, and cupcakes.
```



There's one other twist. The program will also need to accept a --sorted argument that will require you to sort the items before you print them. We'll deal with that in a bit.

So, your Python program must do the following:

- Store one or more positional arguments in a list
- Count the number of arguments
- Possibly sort the items
- Use the list to print a new a string that formats the arguments according to how many items there are

How should we begin?

3.1 *Starting the program*

I will always recommend you start programming by running new.py or by copying template/template.py to the program name. This time the program should be called picnic.py, and you need to create it in the 03_picnic directory.

You can do this using the new.py program from the top level of your repository:

```
$ bin/new.py 03_picnic/picnic.py
Done, see new script "03_picnic/picnic.py."
```

Now go into the 03_picnic directory and run make test or pytest -xv test.py. You should pass the first two tests (program exists, program creates usage) and fail the third:

test.py::test_exists PASSED	[14%]
test.py::test_usage PASSED	[28%]
test.py::test_one FAILED	[42%]

The rest of the output complains that the test expected “You are bringing chips” but got something else:

```
=====
 FAILURES =====
=====
----- test_one -----
-----
```

```
def test_one():
    """one item"""

    out = getoutput(f'{prg} chips')
    assert out.strip() == 'You are bringing chips.'
    assert 'str_arg = ""...nal = "chips"' == 'You are bringing chips.'
    + You are bringing chips.
    - str_arg = ""
    - int_arg = "0"
    - file_arg = ""
    - flag_arg = "False"
    - positional = "chips"

test.py:31: AssertionError
===== 1 failed, 2 passed in 0.56 seconds =====
```

The program is being run with the argument “chips.”

The line starting with a + sign shows what was expected.

The lines starting with the - sign show what was returned by the program.

This line is causing the error. The output is tested to see if it is equal (==) to the string “You are bringing chips.”

Let’s run the program with the argument “chips” and see what it gets:

```
$ ./picnic.py chips
str_arg = ""
int_arg = "0"
file_arg = ""
flag_arg = "False"
positional = "chips"
```

Right, that’s not correct at all! Remember, the template doesn’t yet have the *correct* arguments, just some examples, so the first thing we need to do is fix the `get_args()` function. Your program should print a usage statement like the following if given *no arguments*:

```
$ ./picnic.py
usage: picnic.py [-h] [-s] str [str ...]
picnic.py: error: the following arguments are required: str
```

And here is the usage for the `-h` or `--help` flags:

```
$ ./picnic.py -h
usage: picnic.py [-h] [-s] str [str ...]

Picnic game

positional arguments:
  str           Item(s) to bring
```

optional arguments:

```
-h, --help      show this help message and exit
-s, --sorted   Sort the items (default: False)
```

We need one or more positional arguments and an optional flag called `--sorted`. Modify your `get_args()` until it produces the preceding output.

Note that there should be one or more of the `item` parameter, so you should define it with `nargs='+'`. Refer to section A.4.5 in the appendix for details.

3.2 Writing `picnic.py`

Figure 3.1 shows a tasty diagram of the inputs and outputs for the `picnic.py` program we'll write.

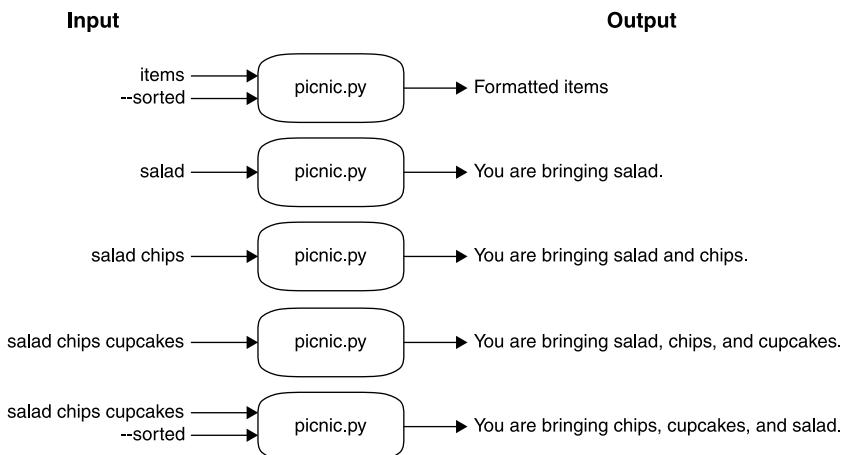


Figure 3.1 A string diagram of the picnic program showing the various inputs and outputs the program will handle

The program should accept one or more positional arguments for the items to bring on a picnic as well as an `-s` or `--sorted` flag to indicate whether or not to sort the items. The output will be “You are bringing” followed by the list of items formatted according to the following rules:

- If there’s one item, state the item:

```
$ ./picnic.py chips
You are bringing chips.
```

- If there are two items, put “and” in between the items. Note that “potato chips” is just *one string* that happens to contain *two words*. If you leave out the quotes,

there would be three arguments to the program. It doesn't matter here whether you use single or double quotes:

```
$ ./picnic.py "potato chips" salad
You are bringing potato chips and salad.
```

- If there are three or more items, place a comma and space between the items and the word “and” before the final element. Don’t forget the comma before the “and” (sometimes called the “Oxford comma”) because your author was an English lit major and, while I may have finally stopped using two spaces after the end of a sentence, you can pry the Oxford comma from my cold, dead hands:

```
$ ./picnic.py "potato chips" salad soda cupcakes
You are bringing potato chips, salad, soda, and cupcakes.
```

Be sure to sort the items if the `-s` or `--sorted` flag is specified:

```
$ ./picnic.py --sorted salad soda cupcakes
You are bringing cupcakes, salad, and soda.
```

To figure out how many items we have, how to sort and slice them, and how to format the output string, we need to talk about the `list` type in Python.

3.3 Introducing lists

It's time to learn how to define positional arguments so that they are available as a `list`. That is, if we run the program like this,

```
$ ./picnic.py salad chips cupcakes
```

the arguments `salad` `chips` `cupcakes` will be available as a `list` of strings inside the program. If you `print()` a `list` in Python, you'll see something like this:

```
['salad', 'chips', 'cupcakes']
```

The square brackets tell us this is a `list`, and the quotes around the elements tell us they are strings. Note that the items are shown in the same order as they were provided on the command line. Lists always keep their order!

['salad', 'chips', 'cupcakes']

Square brackets
mean a list.

Let's go into the REPL and create a variable called `items` to hold some scrumptious victuals to bring on our picnic. I really want you to type these commands yourself, whether in the `python3` REPL or IPython or a Jupyter Notebook. It's very important to interact in real time with the language.

To create a new, empty `list`, you can use the `list()` function:

```
>>> items = list()
```

Or you can use empty square brackets:

```
>>> items = []
```

Check what Python says for the `type()`. Yep, it's a `list`:

```
>>> type(items)
<class 'list'>
```

One of the first things we need to know is how many `items` we have for our picnic. Just as with a `str`, we can use `len()` (length) to get the number of elements in `items`:

```
>>> len(items)
0
```

The length of an empty list is 0.

3.3.1 Adding one element to a list

An empty list is not very useful. Let's see how we can add new items. We used `help(str)` in the last chapter to read documentation about the string methods—the functions that belong to every `str` in Python. Here I want you to use `help(list)` to learn about the `list` methods:

```
>>> help(list)
```

Remember that pressing the spacebar or *F* key (or Ctrl-F) will take you forward, and pressing *B* (or Ctrl-B) will take you back. Pressing the / key will let you search for a string.

You'll see lots of “double-under” methods, like `__len__`. Skip over those, and the first method is `list.append()`, which we can use to add items to the end of a `list`.

If we evaluate `items`, the empty brackets will tell us that it's empty:

```
>>> items
[]
```

Let's add “sammiches” to the end:

```
>>> items.append('sammiches')
```

Nothing happened, so how do we know if it worked? Let's check the length. It should be 1:

```
>>> len(items)
1
```

Hooray! That worked. In the spirit of testing, we'll use the `assert` statement to verify that the length is 1:

```
>>> assert len(items) == 1
```

The fact that nothing happens is good. When an assertion fails, it triggers an exception that results in a lot of messages.

If you type `items` and press Enter in the REPL, Python will show you the contents:

```
>>> items
['sammiches']
```

Cool, we added one element.

3.3.2 Adding many elements to a list

Let's try to add "chips" and "ice cream" to `items`:

```
>>> items.append('chips', 'ice cream')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: append() takes exactly one argument (2 given)
```

Here is one of those pesky exceptions, and these will cause your programs to *crash*, something we want to avoid at all costs. As you can see, `append()` takes exactly one argument, and we gave it two. If you look at `items`, you'll see that nothing was added:

```
>>> items
['sammiches']
```

OK, so maybe we were supposed to give it a list of items to add? Let's try that:

```
>>> items.append(['chips', 'ice cream'])
```

Well, that didn't cause an exception, so maybe it worked? We would expect there to be three `items`, so let's use an assertion to check that:

```
>>> assert len(items) == 3
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AssertionError
```

We get another exception, because `len(items)` is not 3. What is the length?

```
>>> len(items)
2
```

Only 2? Let's look at `items`:

```
>>> items
['sammiches', ['chips', 'ice cream']]
```

Check that out! Lists can hold any type of data, like strings and numbers and even other lists (see figure 3.2). We asked `items.append()` to add `['chips', 'ice cream']`, which is a list, and that's just what it did. Of course, it's not what we wanted.

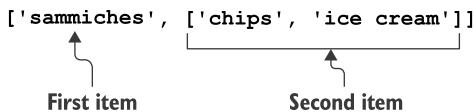


Figure 3.2 A list can hold any mix of values, such as a string and another list of strings.

Let's reset items so we can fix this:

```
>>> items = ['sammiches']
```

If you read further into the help, you will find the `list.extend()` method:

```
| extend(self, iterable, /)
| Extend list by appending elements from the iterable.
```

Let's try that:

```
>>> items.extend('chips', 'ice cream')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: extend() takes exactly one argument (2 given)
```



Well that's frustrating! Now Python is telling us that `extend()` takes exactly one argument, which, if you refer to the help, should be an iterable. A list is something you can iterate (travel over from beginning to end), so that will work:

```
>>> items.extend(['chips', 'ice cream'])
```

Nothing happened. No exception, so maybe that worked? Let's check the length. It *should* be 3:

```
>>> assert len(items) == 3
```

Yes! Let's look at the items we've added:

```
>>> items
['sammiches', 'chips', 'ice cream']
```

Great! This is sounding like a pretty delicious outing.

If you know everything that will go into the list, you can create it like so:

```
>>> items = ['sammiches', 'chips', 'ice cream']
```

The `list.append()` and `list.extend()` methods add new elements to the *end* of a given list. The `list.insert()` method allows you to place new items at any position

by specifying the index. I can use the index 0 to put a new element at the beginning of `items`:

```
>>> items.insert(0, 'soda')
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
```

I recommend you read through all the `list` functions so you get an idea of just how powerful this data structure is. In addition to `help(list)`, you can also find lots of great documentation here: <https://docs.python.org/3/tutorial/datastructures.html>.

3.3.3 Indexing lists

We now have a `list` of items. We know how to use `len()` to find how many items there are in the `items` list, and now we need to know how to get parts of the `list` to format.

Indexing a `list` in Python looks exactly the same as indexing a `str` (figure 3.3). (This actually makes me a bit uncomfortable, so I tend to imagine a `str` as a `list` of characters, and then I feel somewhat better.)

0	1	2	3
['soda', 'sammiches', 'chips', 'ice cream']			
-4	-3	-2	-1

Figure 3.3 Indexing lists and strings is the same. For both, you start counting at 0, and you can also use negative numbers to index from the end.

All indexing in Python is zero-offset, so the first element of `items` is at index `items[0]`:

```
>>> items[0]
'soda'
```

If the index is negative, Python starts counting backwards from the end of the list. The index -1 is the last element of the list:

```
>>> items[-1]
'ice cream'
```

You should be very careful when using indexes to reference elements in a `list`. This is unsafe code:

```
>>> items[10]
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
IndexError: list index out of range
```

WARNING Referencing an index that is not present will cause an exception.

You'll soon learn how to safely *iterate*, or travel through, a list so that you don't have to use indexes to get at elements.

3.3.4 Slicing lists

You can extract “slices” (sub-lists) of a list by using `list[start:stop]`. To get the first two elements, you use `[0:2]`. Remember that the 2 is actually the index of the *third* element, but it's not inclusive, as shown in figure 3.4.

```
>>> items[0:2]
['soda', 'sammiches']
```

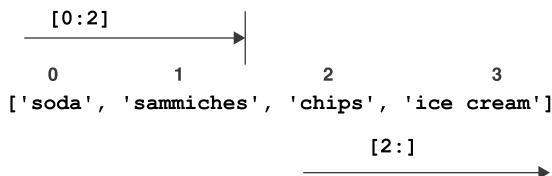


Figure 3.4 The stop value for a list slice is not included. If the stop value is omitted, the slice goes to the end of the list.

If you leave out start, it will default to a value of 0, so the following line does the same thing:

```
>>> items[:2]
['soda', 'sammiches']
```

If you leave out stop, it will go to the end of the list:

```
>>> items[2:]
['chips', 'ice cream']
```

Oddly, it is completely *safe* for slices to use list indexes that don't exist. For example, we can ask for all the elements from index 10 to the end, even though there is nothing at index 10. Instead of an exception, we get an empty list:

```
>>> items[10:]
[]
```

For this chapter's exercise, you're going to need to insert the word “and” into the list if there are three or more elements. Could you use a list index to do that?

3.3.5 Finding elements in a list

Did we remember to pack the chips?

Often you'll want to know if some item is in a list. The `index` method will return the location of an element in a list:

```
>>> items.index('chips')
2
```

Note that `list.index()` is unsafe code, because it will cause an exception if the argument is not present in the list. See what happens if we check for a fog machine:

```
>>> items.index('fog machine')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: 'fog machine' is not in list
```

You should never use `list.index()` unless you have first verified that an element is present. The `x in y` approach that we used in chapter 2 to see if a letter was in a string of vowels can also be used for lists. We get back a `True` value if `x` is in the collection of `y`:

```
>>> 'chips' in items
True
```

I hope they're salt and vinegar chips.

The same code returns `False` if the element is not present:

```
>>> 'fog machine' in items
False
```

We're going to need to talk to the planning committee. What's a picnic without a fog machine?



3.3.6 Removing elements from a list

The `list.pop()` method will remove *and return* the element at the index, as shown in figure 3.5. By default it will remove the *last* item (-1).

```
>>> items.pop()
'ice cream'
```

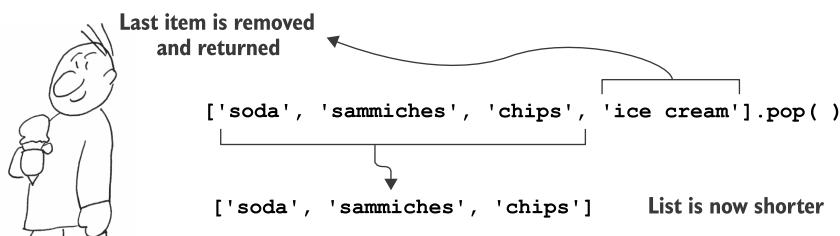


Figure 3.5 The `list.pop()` method will remove an element from the list.

If we look at `items`, we will see it's now shorter by one:

```
>>> items
['soda', 'sammiches', 'chips']
```

We can use an index value to remove an element at a particular location. For instance, we can use 0 to remove the first element (see figure 3.6):

```
>>> items.pop(0)
'soda'
```

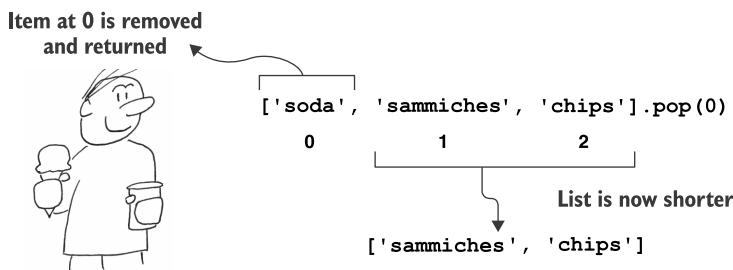


Figure 3.6 You can specify an index value to `list.pop()` to remove a particular element.

Now `items` is shorter still:

```
>>> items
['sammiches', 'chips']
```

You can also use the `list.remove()` method to remove the first occurrence of a given item (see figure 3.7):

```
>>> items.remove('chips')
>>> items
['sammiches']
```

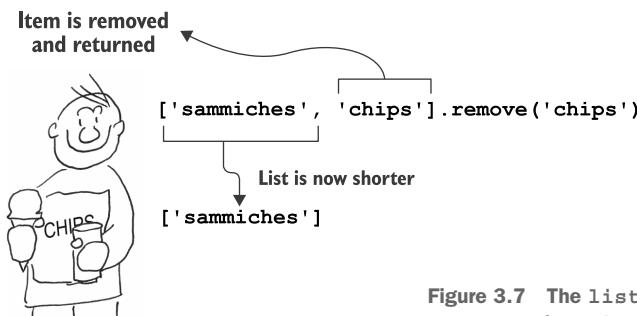


Figure 3.7 The `list.remove()` method will remove an element matching a given value.

WARNING The `list.remove()` method will cause an exception if the element is not present.

If we try to use `items.remove()` to remove the chips again, we'll get an exception:

```
>>> items.remove('chips')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: list.remove(x): x not in list
```



So don't use this code unless you've verified that a given element is in the list:

```
item = 'chips'
if item in items:
    items.remove(item)
```

3.3.7 Sorting and reversing a list

If the `--sorted` flag is used to call our program, we're going to need to sort the items. You might notice in the help documentation that two methods, `list.reverse()` and `list.sort()`, stress that they work *in place*. That means that the list itself will be either reversed or sorted, and nothing will be returned. So, given this list,

```
>>> items = ['soda', 'sammiches', 'chips', 'ice cream']
```

the `items.sort()` method will return nothing:

```
>>> items.sort()
```

If you inspect `items`, you will see that the items have been sorted alphabetically:

```
>>> items
['chips', 'ice cream', 'sammiches', 'soda']
```

`None` ← `items.sort()`

↓
Items are sorted,
and nothing is returned.

As with `list.sort()`, nothing is returned from the `list.reverse()` call:

```
>>> items.reverse()
```

But the items are now in the opposite order:

```
>>> items
['soda', 'sammiches', 'ice cream', 'chips']
```

The `list.sort()` and `list.reverse()` *methods* are easily confused with the `sorted()` and `reversed()` *functions*. The `sorted()` function accepts a list as an argument and *returns* a new list:

```
>>> items = ['soda', 'sammiches', 'chips', 'ice cream']
>>> sorted(items)
['chips', 'ice cream', 'sammiches', 'soda']
```

It's crucial to note that the `sorted()` function *does not alter* the given list:

```
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
```

Note that Python will sort a list of numbers *numerically*, so we've got that going for us, which is nice:

```
>>> sorted([4, 2, 10, 3, 1])
[1, 2, 3, 4, 10]
```

WARNING Sorting a list that mixes strings and numbers will cause an exception!

```
>>> sorted([1, 'two', 3, 'four'])
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: '<' not supported between instances of 'str' and 'int'
```

The `list.sort()` method is a function that belongs to the `list`. It can take arguments that affect the way the sorting happens. Let's look at `help(list.sort)`:

```
sort(self, /, *, key=None, reverse=False)
    Stable sort *IN PLACE*.
```

That means we can also `sort()` items in reverse, like so:

```
>>> items.sort(reverse=True)
```

Now they look like this:

```
>>> items
['soda', 'sammiches', 'ice cream', 'chips']
```



The `reversed()` function works a bit differently:

```
>>> reversed(items)
<list_reverseiterator object at 0x10e012ef0>
```

I bet you were expecting to see a new list with the items in reverse. This is an example of a *lazy* function in Python. The process of reversing a list might take a while, so Python is showing that it has generated an *iterator object* that will provide the reversed list when we actually need the elements.

We can see the values of our `reversed()` list in the REPL by using the `list()` function to evaluate the iterator:

```
>>> list(reversed(items))
['ice cream', 'chips', 'sammiches', 'soda']
```

As with the `sorted()` function, the original `items` remains unchanged:

```
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
```

If you use the `list.sort()` method instead of the `sorted()` function, you might end up deleting your data. Imagine you wanted to set `items` equal to the sorted list of items, like so:

```
>>> items = items.sort()
```

What is in `items` now? If you print `items` in the REPL, you won't see anything useful, so inspect the `type()`:

```
>>> type(items)
<class 'NoneType'>
```

It's no longer a list. We set it equal to the result of calling the `items.sort()` method, which changes `items` *in place* and returns `None`.

If the `--sorted` flag is given to your program, you will need to sort your items in order to pass the test. Will you use `list.sort()` or the `sorted()` function?

3.3.8 Lists are mutable

As you've seen, we can change a list quite easily. The `list.sort()` and `list.reverse()` methods change the whole list, but you can also change any single element by referencing it by index. Maybe we should make our picnic slightly healthier by swapping out the chips for apples:

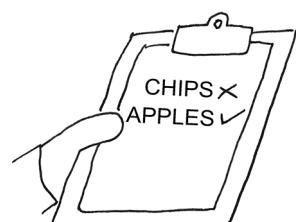
```
>>> items
['soda', 'sammiches', 'chips', 'ice cream']
>>> if 'chips' in items:
...     idx = items.index('chips')           See if the string 'chips' is
...     items[idx] = 'apples'                in the list of items.
...                                         Assign the index of 'chips'
...                                         to the variable idx.
...
...                                         Use the index idx to change
...                                         the element to 'apples'.
```

Let's look at `items` to verify the result:

```
>>> items
['soda', 'sammiches', 'apples', 'ice cream']
```

We can also write a couple of tests:

```
Make sure "chips" are no
longer on the menu.
>>> assert 'chips' not in items
>>> assert 'apples' in items
Check that we now have
some "apples."
```



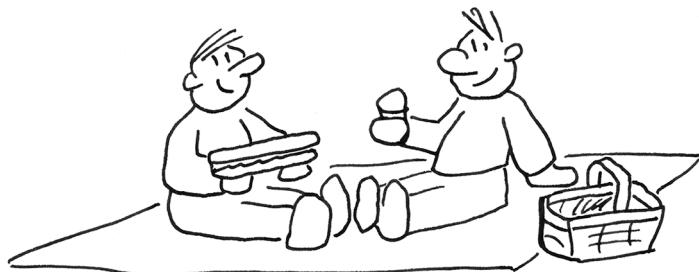
You will need to get the word “and” into your list just before the last element when there are three or more items. Could you use this idea?

3.3.9 ***Joining a list***

In this chapter’s exercise, you’ll need to print a string based on the number of elements in the given list. The string will intersperse other strings like a comma and a space (‘, ’) between the elements of the list.

The following syntax will join a list with a string made of a comma and a space:

```
>>> ', '.join(items)
'soda, sammiches, chips, ice cream'
```



The preceding code uses the `str.join()` method and passes the list as an argument. It always feels backwards to me, but that’s the way it goes.

The result of `str.join()` is a new string:

```
>>> type(', '.join(items))
<class 'str'>
```

The original list remains unchanged:

```
>>> items
['soda', 'sammiches', 'chips', 'apples']
```

We can do quite a bit more with Python’s list, but that should be enough for you to solve this chapter’s problem.

3.4 ***Conditional branching with if/elif/else***

You need to use conditional branching, based on the number of items, to correctly format the output. In chapter 2’s exercise, there were two conditions—either a vowel or not—so we used `if/else` statements. Here we have three options to consider, so you will have to use `elif` (`else-if`) as well.

For instance, suppose we want to classify someone by their age using three options:

- If their age is greater than 0, it is valid.
- If their age is less than 18, they are a minor.
- Otherwise, they are 18 years or older, which means they can vote.

Here is how we could write that code:

```
>>> age = 15
>>> if age < 0:
...     print('You are impossible.')
... elif age < 18:
...     print('You are a minor.')
... else:
...     print('You can vote.')
...
You are a minor.
```

See if you can use that example to figure out how to write the three options for picnic.py. First write the branch that handles one item. Then write the branch that handles two items. Then write the last branch for three or more items. Run the tests *after every change to your program*.

3.4.1 Time to write

Now go write the program yourself before you look at my solution. Here are a few hints:

- Go into your 03_picnic directory and run new.py picnic.py to create your program. Then run make test (or pytest -xv test.py). You should pass the first two tests.
- Next work on getting your --help usage looking like the example shown earlier in the chapter. It's very important to define your arguments correctly. For the items argument, look at nargs in argparse, as discussed in section A.4.5 of the appendix.
- If you use new.py to start your program, be sure to keep the Boolean flag and modify it for your sorted flag.
- Solve the tests in order! First handle one item, then handle two items, and then handle three. Then handle the sorted items.

You'll get the best benefit from this book if you try writing the programs and passing the tests before reading the solutions!

3.5 Solution

Here is one way to satisfy the tests. If you wrote something different that passed, that's great!

```
#!/usr/bin/env python3
"""Picnic game"""

import argparse
# -----
def get_args():
    """Get command-line arguments"""
    The get_args() function is placed first so we can
    easily see what the program accepts when we
    read it. Note that the function order here is not
    important to Python, only to us readers.
```

```

parser = argparse.ArgumentParser(
    description='Picnic game',
    formatter_class=argparse.ArgumentDefaultsHelpFormatter)

parser.add_argument('item',
                    metavar='str',
                    nargs='+',
                    help='Item(s) to bring')

parser.add_argument('-s',
                    '--sorted',
                    action='store_true',
                    help='Sort the items')

return parser.parse_args() ←

```

The item argument uses nargs='+' so that it will accept one or more positional arguments, which will be strings.

The dashes in the short (-s) and long (--sorted) names make this an option. There is no value associated with this argument. It's either present (in which case it will be True) or absent (False).

The main() function is where the program will start.

```

# -----
def main():
    """Make a jazz noise here"""

    args = get_args() ←
    items = args.item ←
    num = len(items) ←

```

Process the command-line arguments and return them to the caller.

Call the get_args() function and put the returned value into the variable args. If there is a problem parsing the arguments, the program will fail before the values are returned.

Copy the item list from args into the new variable items.

The args.sorted value will be either True or False.

Join the items on a string of comma and space.

Print the output string, using the str.format() method to interpolate the bringing variable.

```

    if args.sorted:
        items.sort() ←

```

If we are supposed to sort the items, call the items.sort() method to sort them in place.

```

    bringing = '' ←
    if num == 1:
        bringing = items[0] ←
    elif num == 2:
        bringing = ' and '.join(items) ←
    else:
        items[-1] = ' and ' + items[-1] ←
        bringing = ', '.join(items) ←

```

Use an empty string to initialize a variable to hold the items we are bringing.

If the number of items is 1, we will assign the one item to bringing.

If the number of items is 2, put the string ' and ' in between the items.

Otherwise, alter the last element in items to append the string 'and' before whatever is already there.

```

    print('You are bringing {}'.format(bringing)) ←

```

When Python runs the program, it will read all the lines to this point but will not run anything. Here we look to see if we are in the "main" namespace. If we are, we call the main() function to make the program begin.

```

# -----
if __name__ == '__main__':
    main()

```

3.6 Discussion

How did it go? Did it take you long to write your version? How different was it from mine? Let's talk about my solution. It's fine if yours is different from mine, just as long as you pass the tests!

3.6.1 Defining the arguments

This program can accept a variable number of arguments that are all the same thing (strings). In my `get_args()` method I define an item like so:

```
parser.add_argument('item',      ← A positional parameter
                  metavar='str',   ← called item
                  nargs='+',       ← An indication to the
                  help='Item(s) to bring') ← user in the usage that
                                         this should be a string
                                         ← The number of
                                         arguments, where '+' means one or more
                                         ← A longer help description that appears for the -h or --help options
```

This program also accepts `-s` and `--sorted` arguments. They are “flags,” which typically means that they are `True` if they are present and `False` if absent. Remember that the leading dashes makes them optional.

```
parser.add_argument('-s',        ← The short flag name
                  '--sorted',    ← The long flag name
                  action='store_true',
                  help='Sort the items') ← If the flag is present, store a True value. The default value will be False.
                                         ← The longer help description
```

3.6.2 Assigning and sorting the items

In `main()` I call `get_args()` to get the arguments, and I assign them to the `args` variable. Then I create the `items` variable to hold the `args.item` value(s):

```
def main():
    args = get_args()
    items = args.item
```

If `args.sorted` is `True`, I need to sort `items`. I chose the in-place `sort` method here:

```
if args.sorted:
    items.sort()
```

Now I have the items, sorted if needed, and I need to format them for output.

3.6.3 Formatting the items

I suggested you solve the tests in order. There are four conditions we need to solve:

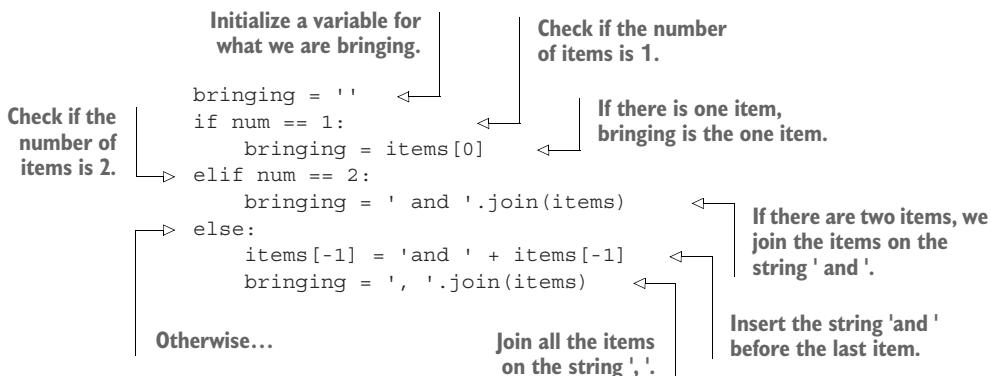
- Zero items
- One item

- Two items
- Three or more items

The first test is actually handled by argparse—if the user fails to provide any arguments, they get a usage message:

```
$ ./picnic.py
usage: picnic.py [-h] [-s] str [str ...]
picnic.py: error: the following arguments are required: str
```

Since argparse handles the case of no arguments, we have to handle the other three conditions. Here's one way to do that:



Can you come up with any other ways to do this?

3.6.4 Printing the items

Finally, to print() the output, I used a format string where the {} indicate a placeholder for a value, like so:

```
>>> print('You are bringing {}'.format(bringing))
You are bringing salad, soda, and cupcakes.
```

If you prefer, you could use an f'''-string:

```
>>> print(f'You are bringing {bringing}.')
You are bringing salad, soda, and cupcakes.
```

They both get the job done.

3.7 Going further

- Add an option so the user can choose not to print with the Oxford comma (even though that is a morally indefensible option).
- Add an option to separate items with a character passed in by the user (like a semicolon if the list of items needs to contain commas).

Be sure to add tests to the test.py program to ensure your new features are correct!

Summary

- Python lists are ordered sequences of other Python data types, such as strings and numbers.
- There are methods like `list.append()` and `list.extend()` to add elements to a list. Use `list.pop()` and `list.remove()` to remove elements.
- You can use `x in y` to ask if element `x` is in the list `y`. You can also use `list.index()` to find the index of an element, but this will cause an exception if the element is not present.
- Lists can be sorted and reversed, and elements within lists can be modified. Lists are useful when the order of the elements is important.
- Strings and lists share many features, such as using `len()` to find their lengths, using zero-based indexing where `0` is the first element and `-1` is the last, and using slices to extract smaller pieces from the whole.
- The `str.join()` method can be used to make a new `str` from a list.
- `if/elif/else` can be used to branch code depending on conditions.

Jump the Five: Working with dictionaries

“When I get up, nothing gets me down.”

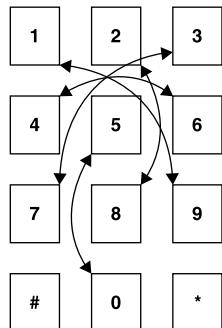
—D. L. Roth

In an episode of the television show *The Wire*, drug dealers assume the police are intercepting their text messages. Whenever a phone number needs to be texted in the course of a criminal conspiracy, the dealers will obfuscate the number. They use an algorithm we’ll call “Jump the Five” because each number is changed to its mate on the opposite of a US telephone pad if you jump over the 5. In this exercise, we’ll discuss how to encrypt messages using this algorithm, and then we’ll see how to use it to decrypt the encrypted messages, you feel me?

If we start with the 1 button and jump across the 5, we get to 9. The 6 jumps the 5 to become 4, and so forth. The numbers 5 and 0 will swap with each other.

In this exercise, we’re going to write a Python program called `jump.py` that will take in some text as a positional argument. Each number in the text will be encoded using this algorithm. All non-number text will pass through unchanged. Here are a couple of examples:

```
$ ./jump.py 867-5309  
243-0751  
$ ./jump.py 'Call 1-800-329-8044 today!'  
Call 9-255-781-2566 today!
```



You will need some way to inspect each character in the input text to identify the numbers—you will learn how to use a `for` loop for this. Then you’ll see how a `for` loop can be rewritten as a “list comprehension.” You’ll need some way to associate a number like 1 with the number 9, and so on for all the numbers—you’ll learn about a data structure in Python called a *dictionary* that will allow you to do exactly that.

In this chapter, you will learn to

- Create a dictionary
- Use a `for` loop and a list comprehension to process text, character by character
- Check if items exist in a dictionary
- Retrieve values from a dictionary
- Print a new string with the numbers substituted with their encoded values

Before we start writing, you need to learn about Python’s dictionaries.

4.1 Dictionaries

A Python dictionary allows us to relate some *thing* (a “key”) to some other *thing* (a “value”). An actual dictionary does this. If we look up a word like “quirky” in a dictionary (www.merriam-webster.com/dictionary/quirky), we can find a definition, as in figure 4.1. We can think of the word itself as the “key” and the definition as the “value.”



quirky ⇨ unusual, esp. in an interesting or appealing way

Figure 4.1 You can find the definition of a word by looking it up in a dictionary.

Dictionaries actually provide quite a bit more information about words, such as pronunciation, part of speech, derived words, history, synonyms, alternate spellings, etymology, first known use, and so on. (I really love dictionaries.) Each of those attributes has a value, so we could also think of the dictionary entry for a word as itself being another “dictionary” (see figure 4.2).

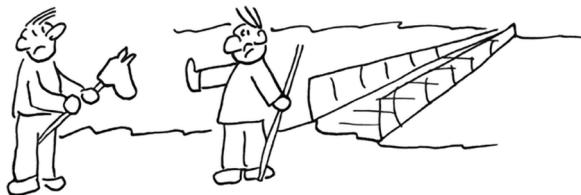
definition ⇨ unusual, esp. in an interesting or appealing way
pronunciation ⇨ 'kwar-kē
part of speech ⇨ adjective

Figure 4.2 The entry for “quirky” can contain much more than a single definition.

Let’s see how we can use Python’s dictionaries to go beyond word definitions.

4.1.1 Creating a dictionary

In the film *Monty Python and the Holy Grail*, King Arthur and his knights must cross The Bridge of Death. Anyone who wishes to cross must correctly answer three questions from the Keeper. Those who fail are cast into the Gorge of Eternal Peril.



Let us ride to CAMELOT.... No, sorry, let us create and use a dictionary to keep track of the questions and answers as key/value pairs. Once again, I want you to fire up your python3 or IPython REPL or Jupyter Notebook and type these out for yourself.

Lancelot goes first. We can use the `dict()` function to create an empty dictionary for his answers.

```
>>> answers = dict()
```

Or we can use empty curly brackets (both methods are equivalent):

```
>>> answers = {}
```

The Keeper's first question is, "What is your name?" Lancelot answers, "My name is Sir Lancelot of Camelot." We can add the key "name" to the answers dictionary by using square brackets ([])—not curlies! and the literal string 'name':

```
>>> answers['name'] = 'Sir Lancelot'
```

If you type `answers` and press Enter in the REPL, Python will show you a structure in curlies (see figure 4.3) to indicate that this is a dict:

```
>>> answers
{'name': 'Sir Lancelot'}
```

You can verify this with the `type()` function:

```
>>> type(answers)
<class 'dict'>
```

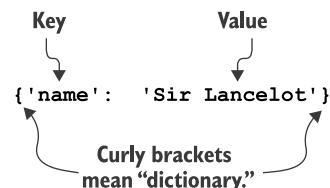
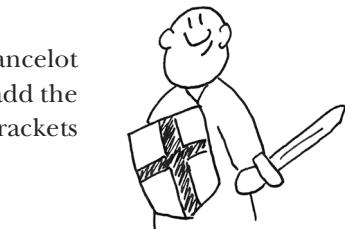


Figure 4.3 A dictionary is printed inside curly braces. The keys are separated from the values by a colon.

Next the Keeper asks, "What is your quest?" to which Lancelot answers "To seek the Holy Grail." Let's add "quest" to `answers`:

```
>>> answers['quest'] = 'To seek the Holy Grail'
```

There's no return value to let us know something happened, so type answers to inspect the variable again to ensure the new key/value was added:

```
>>> answers
{'name': 'Sir Lancelot', 'quest': 'To seek the
    Holy Grail'}
```

Finally the Keeper asks, “What is your favorite color?” and Lancelot answers, “Blue.”

```
>>> answers['favorite_color'] = 'blue'
>>> answers
{'name': 'Sir Lancelot', 'quest': 'To seek the Holy Grail', 'favorite_color':
    'blue'}
```

NOTE I'm using “favorite_color” (with an underscore) as the key, but I could use “favorite color” (with a space) or “FavoriteColor” or “Favorite color,” but each one of those would be a separate and distinct string, or key. I prefer to use the PEP 8 naming conventions for dictionary keys and variable and functions names. PEP 8, the “Style Guide for Python Code” (www.python.org/dev/peps/pep-0008/), suggests using lowercase names with words separated by underscores.

If you knew all the answers beforehand, you could create answers using the `dict()` function with the following syntax, where you do *not* have to quote the keys, and the keys are separated from the values with equal signs:

```
>>> answers = dict(name='Sir Lancelot', quest='To seek the Holy Grail',
    favorite_color='blue')
```

Or you could use the following syntax using curly braces {}, where the keys must be quoted and they are followed by a colon (:):

```
>>> answers = {'name': 'Sir Lancelot', 'quest': 'To seek the Holy Grail',
    'favorite_color': 'blue'}
```

It might be helpful to think of the answers dictionary as a box holding key/value pairs that describe Lancelot's answers (see figure 4.4), just the way the “quirky” dictionary holds all the information about that word.

answers	
name	⇒ Sir Lancelot
quest	⇒ To seek the Holy Grail
favorite color	⇒ blue



Figure 4.4 Just like the “quirky” dictionary entry, a Python dictionary can contain many key/value pairs.

4.1.2 Accessing dictionary values

To retrieve the values, you use the key name inside square brackets ([]). For instance, you can get the name like so:

```
>>> answers['name']
'Sir Lancelot'
```

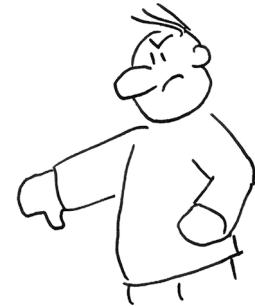
Let's request his "age":

```
>>> answers['age']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'age'
```

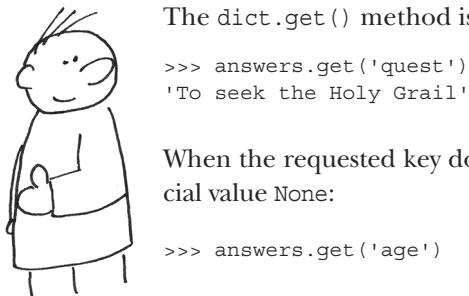
As you can see, you will cause an exception if you ask for a dictionary key that doesn't exist!

Just as with strings and lists, you can use `x in y` to first see if a key exists in the dict:

```
>>> 'quest' in answers
True
>>> 'age' in answers
False
```



The `dict.get()` method is a *safe* way to ask for a value:



```
>>> answers.get('quest')
'To seek the Holy Grail'
```

When the requested key does not exist in the dict, it will return the special value `None`:

```
>>> answers.get('age')
```

That doesn't print anything because the REPL won't print a `None`, but we can check the `type()`. Note that the type of `None` is the `NoneType`:

```
>>> type(answers.get('age'))
<class 'NoneType'>
```



There is an optional second argument you can pass to `dict.get()`, which is the value to return *if the key does not exist*:

```
>>> answers.get('age', 'NA')
'NA'
```

That's going to be important for the solution because we will only need to represent the characters 0–9.

4.1.3 Other dictionary methods

If you want to know how “big” a dictionary is, the `len()` (length) function on a dict will tell you how many key/value pairs are present:

```
>>> len(answers)
3
```

The `dict.keys()` method will give you just the keys:

```
>>> answers.keys()
dict_keys(['name', 'quest', 'favorite_color'])
```

And `dict.values()` will give you just the values:

```
>>> answers.values()
dict_values(['Sir Lancelot', 'To seek the Holy Grail', 'blue'])
```

Often we want both together, so you might see code like this:

```
>>> for key in answers.keys():
...     print(key, answers[key])
...
name Sir Lancelot
quest To seek the Holy Grail
favorite_color blue
```

An easier way to write this would be to use the `dict.items()` method, which will return the contents of the dictionary as a new list containing each key/value pair:

```
>>> answers.items()
dict_items([('name', 'Sir Lancelot'), ('quest', 'To seek the Holy Grail'),
('favorite_color', 'blue')])
```

The preceding `for` loop could also be written using the `dict.items()` method:

```
>>> for key, value in answers.items():
...     print(f'{key:15} {value}')
...
name           Sir Lancelot
quest          To seek the Holy Grail
favorite_color blue
```

Unpack each key/value pair into the variables `key` and `value` (see figure 4.5). Note that you don't have to call them `key` and `value`. You could use `k` and `v` or `question` and `answer`.

Print the key in a left-justified field 15 characters wide. The value is printed normally.

```
for key, value in [('name', 'Sir Lancelot'), ...]:
```

Figure 4.5 We can unpack the key/value pairs returned by `dict.items()` into variables.

In the REPL you can execute `help(dict)` to see all the methods available to you, like `dict.pop()`, which removes a key/value, or `dict.update()`, which merges one dictionary with another.

TIP Each key in the dict is unique.

That means if you set a value for a given key twice,

```
>>> answers = {}
>>> answers['favorite_color'] = 'blue'
>>> answers
{'favorite_color': 'blue'}
```

you will not have two entries but one entry with the *second* value:

```
>>> answers['favorite_color'] = 'red'
>>> answers
{'favorite_color': 'red'}
```

Keys don't have to be strings—you can also use numbers like the `int` and `float` types. Whatever value you use must be immutable. For instance, lists could not be used because they are mutable, as you saw in the previous chapter. You'll learn which types are immutable as we go further.

4.2 Writing jump.py

Now let's get started with writing our program. You'll need to create a program called `jump.py` in the `04_jump_the_five` directory so you can use the `test.py` that is there. Figure 4.6 shows a diagram of the inputs and outputs. Note that your program will only affect the numbers in the text. Anything that is *not* a number will remain unchanged.

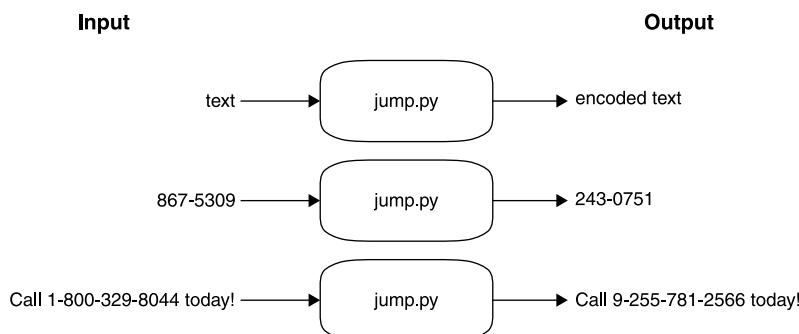


Figure 4.6 A string diagram for the `jump.py` program. Any number in the input text will be changed to a corresponding number in the output text.

When your program is run with no arguments, -h, or --help, it should print a usage message:

```
$ ./jump.py -h
usage: jump.py [-h] str

Jump the Five

positional arguments:
  str           Input text

optional arguments:
  -h, --help    show this help message and exit
```

Note that we will be processing *text* representations of the “numbers,” so the string '1' will be converted to the string '9'. We won’t be changing the actual integer value 1 to the integer value 9. Keep that in mind as you figure out a way to represent the substitutions in table 4.1.

Table 4.1 The encoding table for the numeric characters in the text

<pre>1 => 9 2 => 8 3 => 7 4 => 6 5 => 0 6 => 4 7 => 3 8 => 2 9 => 1 0 => 5</pre>	
--	--

How would you represent this using a dict? Try creating a dict called `jumper` in the REPL with the preceding key/value pairs, and then see if the following assert statements will execute with exceptions. Remember that assert will return nothing if the statement is True.

```
>>> assert jumper['1'] == '9'
>>> assert jumper['5'] == '0'
```

Next, you will need a way to visit each character. I suggest you use a for loop, like so:

```
>>> for char in 'ABC123':
...     print(char)
...
A
B
C
1
2
3
```

Rather than printing the char, print the value of char in the jumper table, or print the char itself. Look at the `dict.get()` method! Also, if you read `help(print)`, you'll see there is an `end` option to replace the newline that gets stuck onto the end with something else.

Here are some other hints:

- The numbers can occur anywhere in the text, so I recommend you process the input character by character with a `for` loop.
- Given any one character, how can you look it up in your table?
- If the character is in your table, how can you get the value (the translation)?
- How can you `print()` the translation or the value without printing a newline? Look at `help(print)` in the REPL to read about the options for `print()`.
- If you read `help(str)` on Python's `str` class, you'll see that there is a `str.replace()` method. Could you use that?

Now spend the time to write the program on your own before you look at the solutions. Use the tests to guide you.

4.3 Solution

Here is one solution that satisfies the tests. I will show some variations after we discuss this first version.

```
#!/usr/bin/env python3
"""Jump the Five"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Jump the Five',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='str', help='Input text')

    return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
```

The code is annotated with several callout boxes and arrows pointing to specific parts of the code:

- A box on the right side of the `get_args()` definition contains the text: "Define the `get_args()` function first, so it's easy to find when I read the program." An arrow points from the start of the `get_args()` definition to this text.
- A box on the right side of the `parser.add_argument()` call contains the text: "Define one positional argument called 'text.'" An arrow points from the `'text'` argument to this text.
- A box on the right side of the `main()` definition contains the text: "Define a `main()` function where the program starts." An arrow points from the `main()` definition to this text.
- A box on the right side of the `jumper` dictionary assignment contains the text: "Create a dictionary for the lookup table." An arrow points from the `jumper` assignment to this text.
- A box at the bottom center contains the text: "Get the command-line arguments from `get_args()`." An arrow points from the `args` variable in the `main()` function to this text.

```

for char in args.text:
    print(jumper.get(char, char), end='')
print()

# -----
if __name__ == '__main__':
    main()

```

Print a newline after I am done processing the characters.

Call the main() function if the program is in the “main” namespace.

Process each character in the input text.

Print either the value of the character from the “jumper” table or the character itself. Change the “end” value to print() so as to avoid adding a newline.

4.4 Discussion

Let’s break this program down into the big ideas, like how we define the parameters, define and use a dictionary, process the input text, and print the output.

4.4.1 Defining the parameters

As usual, the `get_args()` function is defined first. The program needs to define one positional argument. Since I’m expecting some “text,” I call the argument `'text'` and then assign that to a variable called `text`:

```
parser.add_argument('text', metavar='str', help='Input text')
```

While that seems rather obvious, I think it’s very important to name things for *what they are*. That is, please don’t leave the name of the argument as `'positional'`—that does not describe what it *is*.

It may seem like overkill to use `argparse` for such a simple program, but it handles the validation of the correct *number* and *type* of the arguments as well as generating help documentation, so it’s well worth the effort.

4.4.2 Using a dict for encoding

I suggested you could represent the substitution table as a `dict`, where each number key has its substitute as the value in the `dict`. For instance, I know that if I jump from 1 over the 5, I should land on 9:

```
>>> jumper = { '1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
...           '6': '4', '7': '3', '8': '2', '9': '1', '0': '5' }
>>> jumper['1']
'9'
```

Since there are only 10 numbers to encode, this is probably the easiest way to write this. Note that the numbers are written with quotes around them, so they are actually of the type `str` and not `int` (integers). I do this because I will be reading characters from a `str`. If I stored them as actual numbers, I would have to coerce the `str` types using the `int()` function:

```
>>> type('4')
<class 'str'>
```

```
>>> type(4)
<class 'int'>
>>> type(int('4'))
<class 'int'>
```

4.4.3 Various ways to process items in a series

As you've seen before, strings and lists in Python are similar in how you can index them. Both strings and lists are essentially sequences of elements—strings are sequences of characters, and lists can be sequences of anything at all.

There are several different ways to process any sequence of items, which here will be characters in a string.

METHOD 1: USING A FOR LOOP TO PRINT() EACH CHARACTER

As I suggested in the introduction, we can process each character of the text using a for loop. To start, I might first see if each character of the text is in the jumper table using the `x in y` construct:

```
>>> text = 'ABC123'
>>> for char in text:
...     print(char, char in jumper)
...
A False
B False
C False
1 True
2 True
3 True
```

NOTE When `print()` is given more than one argument, it will put a space between each bit of text. You can change that with the `sep` argument. Read `help(print)` to learn more.

Now let's try to translate the numbers. I could use an `if` expression, where I print the value from the jumper table if `char` is present, and, otherwise, print the `char`:

```
>>> for char in text:
...     print(char, jumper[char] if char in jumper else char)
...
A A
B B
C C
1 9
2 8
3 7
```

It's a bit laborious to check for every character, but it's necessary because, for instance, the letter "A" is not in `jumper`. If I try to retrieve that value, I'll get an exception:

```
>>> jumper['A']
Traceback (most recent call last):
```

```
File "<stdin>", line 1, in <module>
KeyError: 'A'
```

The `dict.get()` method allows me to safely ask for a value if it is present. Asking for “A” will not produce an exception, but it will also not show anything in the REPL because it returns the `None` value:

```
>>> jumper.get('A')
```

It’s a bit easier to see if we try to `print()` the values:

```
>>> for char in text:
...     print(char, jumper.get(char))
...
A None
B None
C None
1 9
2 8
3 7
```

I can provide a second, optional argument to `dict.get()`, which is the default value to return when the key does not exist. In this program, I want to print the character itself when it does not exist in `jumper`. For instance, if I had “A,” I’d want to print “A”:

```
>>> jumper.get('A', 'A')
'A'
```

But if I have “5,” I want to print “0”:

```
>>> jumper.get('5', '5')
'0'
```

I can use that to process all the characters:

```
>>> for char in text:
...     print(jumper.get(char, char))
...
A
B
C
9
8
7
```

I don’t want that newline printing after every character, so I can use `end=' '` to tell Python to put the empty string at the end instead of a newline.

When I run this in the REPL, the output is going to look funny because I have to press Enter after the `for` loop to run it. Then I’ll be left with ABC987 with no newline, and then the `>>>` prompt:

```
>>> for char in text:
...     print(jumper.get(char, char), end='')
...
ABC987>>>
```

In your code, you'll have to add another `print()`.

It's useful that you can change what is added at the end, and that you can `print()` with no arguments to print a newline. There are several other really cool things `print()` can do, so I encourage you to read `help(print)` and try them out.

METHOD 2: USING A FOR LOOP TO BUILD A NEW STRING

There are several other ways you could solve this. While it was fun to explore all the things we can do with `print()`, that code is a bit ugly. I think it's cleaner to create a `new_text` variable and call `print()` once with that:

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    new_text = ''                                ← Create an empty
    for char in args.text:                         new_text variable.
        new_text += jumper.get(char, char)          ← Append either the encoded
                                                       number or the original
                                                       char to the new_text.
    print(new_text)                                ← Use the same
                                                       for loop.
                                                       Print the new_text.
```

In this version, I start by setting `new_text` equal to the empty string:

```
>>> new_text = ''
```

I use the same for loop to process each character in the `text`. Each time through the loop, I use `+=` to append the right side of the equation to the left side. The `+=` adds the value on the right to the variable on the left:

```
>>> new_text += 'a'
>>> assert new_text == 'a'
>>> new_text += 'b'
>>> assert new_text == 'ab'
```

On the right, I'm using the `jumper.get()` method. Each character will be appended to the `new_text`, as shown in figure 4.7.

```
>>> new_text = ''
>>> for char in text:
...     new_text += jumper.get(char, char)
...
```

`new_text += jumper.get(char, char)`
 ↓
 The result of `jumper.get()`
 is appended to `new_text`.

Figure 4.7 The `+=` operator will append the string on the right to the variable on the left.

Now I can call `print()` once with the new value:

```
>>> print(new_text)
ABC987
```

METHOD 3: USING A FOR LOOP TO BUILD A NEW LIST

This method is the same as the preceding one, but rather than `new_text` being a `str`, it's a `list`:

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    new_text = []
    for char in args.text:
        new_text.append(jumper.get(char, char))
    print(''.join(new_text))

Join the new_text on the
empty string to create a
new string to print.

Initialize new_text
as an empty list.

Iterate through each
character of the text.

Append the results of the
jumper.get() call to the
new_text variable.
```

As we go through the book, I'll keep reminding you how Python treats strings and lists similarly. Here I'm using `new_text` exactly the same as I did before, starting with an empty structure and then making it longer for each character. I could actually use the exact same `+=` syntax instead of the `list.append()` method:

```
for char in args.text:
    new_text += jumper.get(char, char)
```

After the `for` loop is done, I have all the new characters that need to be put back together using `str.join()` into a new string that I can `print()`.

METHOD 4: TURNING A FOR LOOP INTO A LIST COMPREHENSION

A shorter solution uses a *list comprehension*, which is basically a one-line `for` loop inside square brackets `([])` that results in a new list (see figure 4.8).

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    print(''.join([jumper.get(char, char) for char in args.text]))
```

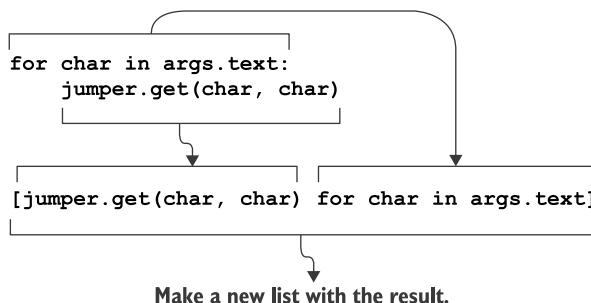


Figure 4.8 A list comprehension will generate a new list with the results of iterating with a `for` statement.

A list comprehension is read backwards from a `for` loop, but it's all there. It's one line of code instead four!

```
>>> text = '867-5309'
>>> [jumper.get(char, char) for char in text]
['2', '4', '3', '-', '0', '7', '5', '1']
```

You can use `str.join()` on the empty string to turn that list into a new string you can `print()`:

```
>>> print(''.join([jumper.get(char, char) for char in text]))
243-0751
```

The purpose of a list comprehension is to create a new list, which is what we were trying to do with the `for` loop code before. A list comprehension makes much more sense and uses far fewer lines of code.

METHOD 5: USING THE `STR.TRANSLATE()` FUNCTION

This last approach uses a really powerful method from the `str` class to change all the characters in one step:

```
def main():
    args = get_args()
    jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
              '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
    print(args.text.translate(str.maketrans(jumper)))
```

The argument to `str.translate()` is a translation table that describes how each character should be translated. That's exactly what `jumper` does.

```
>>> text = 'Jenny = 867-5309'
>>> text.translate(str.maketrans(jumper))
'Jenny = 243-0751'
```

I'll explain this in much greater detail in chapter 8.

4.4.4 (*Not*) using `str.replace()`

I asked earlier whether you could use `str.replace()` to change all the numbers. It turns out you cannot, because you'll end up changing some of the values twice so that they end up at their original values.

Watch how we start off with this string:

```
>>> text = '1234567890'
```

When you change “1” to “9,” now you have two 9’s:

```
>>> text = text.replace('1', '9')
>>> text
'9234567890'
```

This means that when you try to change all the 9's to 1's, you end up with two 1's. The 1 in the first position is changed to 9 and then back to 1 again:

```
>>> text = text.replace('9', '1')
>>> text
'1234567810'
```

So if you go through each number in “1234567890” and try to change them using `str.replace()`, you’ll end up with the value “1234543215”:

```
>>> text = '1234567890'
>>> for n in jumper.keys():
...     text = text.replace(n, jumper[n])
...
>>> text
'1234543215'
```

But the correctly encoded string is “9876043215.” The `str.translate()` function exists to change all the values in one move, all while leaving the unchanging characters alone.

4.5 Going further

- Try creating a similar program that encodes the numbers with strings (for example, “5” becomes “five,” “7” becomes “seven”). Be sure to write the necessary tests in `test.py` to check your work!
- What happens if you feed the output of the program back into itself? For example, if you run `./jump.py 12345`, you should get 98760. If you run `./jump.py 98760`, do you recover the original numbers? This is called *round-tripping*, and it’s a common operation with algorithms that encode and decode text.

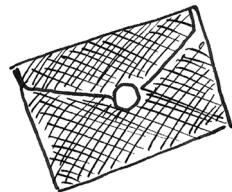
Summary

- You can create a new dictionary using the `dict()` function or with empty curly brackets (`{}`).
- Dictionary values are retrieved using their keys inside square brackets or by using the `dict.get()` method.
- For a dict called `x`, you can use `'key' in x` to determine if a key exists.
- You can use a for loop to iterate through the characters of a str just like you can iterate through the elements of a list. You can think of strings as lists of characters.
- The `print()` function takes optional keyword arguments like `end=' '`, which you can use to print a value to the screen without a newline.

5

Howler: Working with files and STDOUT

In the Harry Potter stories, a “Howler” is a nasty-gram that arrives by owl at Hogwarts. It will tear itself open, shout a blistering message at the recipient, and then combust. In this exercise, we’re going to write a program that will transform text into a rather mild-mannered version of a Howler by **MAKING ALL THE LETTERS UPPERCASE**. The text that we’ll process will be given as a single positional argument.

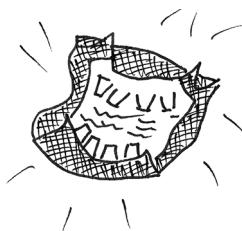


For instance, if our program is given the input, “How dare you steal that car!” it should scream back “HOW DARE YOU STEAL THAT CAR!” Remember spaces on the command line delimit arguments, so multiple words need to be enclosed in quotes to be considered one argument:

```
$ ./howler.py 'How dare you steal that car!'  
HOW DARE YOU STEAL THAT CAR!
```

The argument to the program may also name a file, in which case we need to read the file for the input:

```
$ ./howler.py ../inputs/fox.txt  
THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
```

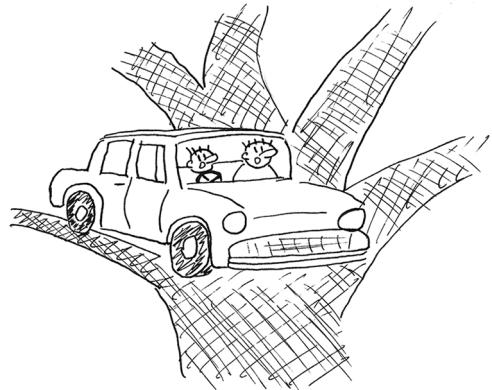


Our program will also accept an `-o` or `--outfile` option that names an output file into which the output text should be written. In that case, *nothing* will be printed on the command line:

```
$ ./howler.py -o out.txt 'How dare
you steal that car!'
```

There should now be a file called `out.txt` that has the output:

```
$ cat out.txt
HOW DARE YOU STEAL THAT CAR!
```



In this exercise, you will learn to

- Accept text input from the command line or from a file
- Change strings to uppercase
- Print output either to the command line or to a file that needs to be created
- Make plain text behave like a file handle

5.1 **Reading files**

This is our first exercise that will involve reading files. The argument to the program will be some text that might name an input file, in which case you will open and read the file. If the text is not the name of a file, you'll use the text itself.

The built-in `os` (operating system) module has a method for detecting whether a string is the name of a file. To use it, you must import the `os` module. For instance, there's probably not a file called "blargh" on your system:

```
>>> import os
>>> os.path.isfile('blargh')
False
```



The `os` module contains loads of useful submodules and functions. Consult the documentation at <https://docs.python.org/3/library/os.html> or use `help(os)` in the REPL.

For instance, `os.path.basename()` and `os.path.dirname()` can return a file's name or directory from a path, respectively (see figure 5.1):

```
>>> file = '/var/lib/db.txt'
>>> os.path.dirname(file)
'/var/lib'
>>> os.path.basename(file)
'db.txt'
```

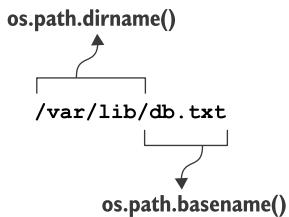


Figure 5.1 The `os` module has handy functions like `os.path.dirname()` and `os.path.basename()` for getting parts of file paths.

In the top level of the GitHub source repository, there is a directory called “inputs” that contains several files we’ll use for many of the exercises. Here I’ll use a file called `inputs/fox.txt`. Note that you will need to be in the main directory of the repo for this to work.

```
>>> file = 'inputs/fox.txt'
>>> os.path.isfile(file)
True
```

Once you’ve determined that the argument is the name of a file, you must `open()` it to `read()` it. The return from `open()` is a *file handle*. I usually call this variable `fh` to remind me that it’s a file handle. If I have more than one open file handle, like both `input` and `output` handles, I may call them `in_fh` and `out_fh`.

```
>>> fh = open(file)
```

NOTE Per PEP 8 (www.python.org/dev/peps/pep-0008/#function-and-variable-names), function and variable “names should be lowercase, with words separated by underscores as necessary to improve readability.”

If you try to `open()` a file that does not exist, you’ll get an exception. This is unsafe code:

```
>>> file = 'blargh'
>>> open(file)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
FileNotFoundException: [Errno 2] No such file or directory: 'blargh'
```

Always check that the file exists!

```
>>> file = 'inputs/fox.txt'
>>> if os.path.isfile(file):
...     fh = open(file)
```

We will use the `fh.read()` method to get the contents of the file. It might be helpful to think of a file as a can of tomatoes. The file’s name, like “`inputs/fox.txt`,” is the label on the can, which is not the same as the *contents*. To get at the text inside (or the “tomatoes”), we need to *open* the can.

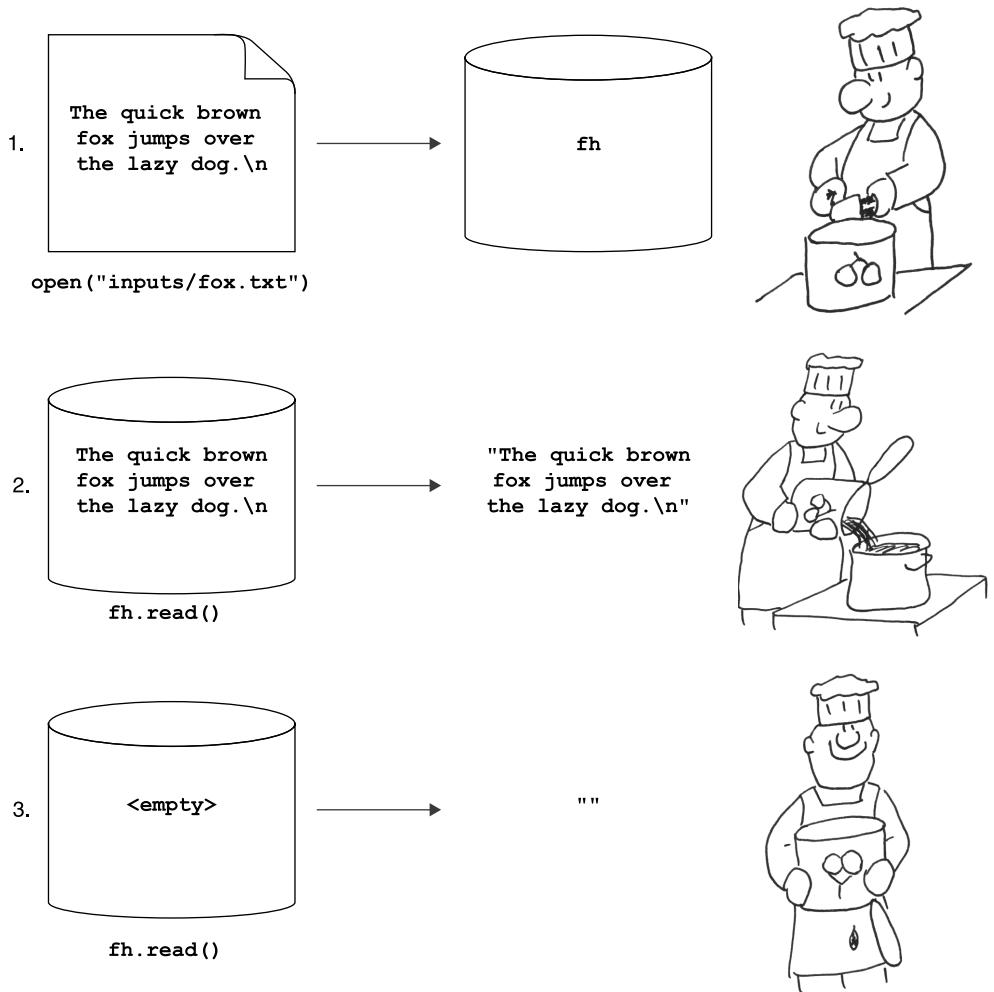


Figure 5.2 A file is a bit like a can of tomatoes. We have to open it first so that we can read it, after which the file handle is exhausted.

Take a look at figure 5.2:

- 1 The file handle (`fh`) is a mechanism we can use to get at the contents of the file. To get at the tomatoes, we need to `open()` the can.
- 2 The `fh.read()` method returns what is inside the file. With the can opened, we can get at the contents.
- 3 Once the file handle has been read, there's nothing left.

NOTE You can use `fh.seek(0)` to reset the file handle to the beginning if you really want to read it again.

Let's see what type() the fh is:

```
>>> type(fh)
<class '_io.TextIOWrapper'>
```

In computer lingo, “io” means “input/output.” The fh object is something that handles I/O operations. You can use help(fh) (using the name of the variable itself) to read the docs on the class TextIOWrapper.

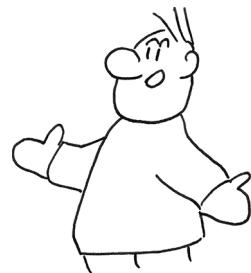
The two methods you’ll use quite often are `read()` and `write()`. Right now, we care about `read()`. Let’s see what that gives us:

```
>>> fh.read()
'The quick brown fox jumps over the lazy dog.\n'
```

Do me a favor and execute that line one more time. What do you see?

```
>>> fh.read()
''
```

A file handle is different from something like a `str`. Once you read a file handle, it’s empty. It’s like pouring the tomatoes out of the can. Now that the can is empty, you can’t empty it again.



We can actually compress `open()` and `fh.read()` into one line of code by *chaining* those methods together. The `open()` method returns a file handle that can be used for the call to `fh.read()` (see figure 5.3). Run this:

```
>>> open(file).read()
'The quick brown fox jumps over the lazy dog.\n'
```

```
open(file).read()
  ↓
  fh.read()
```

Figure 5.3 The `open()` function returns a file handle, so we can chain it to a call to `read()`.

And now run it again:

```
>>> open(file).read()
'The quick brown fox jumps over the lazy dog.\n'
```

Each time you `open()` the file, you get a fresh file handle to `read()`.

If we want to preserve the contents, we’ll need to copy them into a variable.

```
>>> text = open(file).read()
>>> text
'The quick brown fox jumps over the lazy dog.\n'
```

The `type()` of the result is a `str`:

```
>>> type(text)
<class 'str'>
```

If you want, you can chain any `str` method onto the end of that. For instance, maybe you want to remove the trailing newline. The `str.rstrip()` method will remove any whitespace (which includes newlines) from the *right* end of a string (see figure 5.4).

```
>>> text = open(file).read().rstrip()
>>> text
'The quick brown fox jumps over the lazy dog.'
```

```
open(file).read().rstrip()
  ↓
  fh.read().rstrip()
    ↓
  str.rstrip()
```

Figure 5.4 The `open()` method returns a file handle, to which we chain `read()`, which returns a string, to which we chain the `str.rstrip()` call.



Once you have your input text—whether it is from the command line or from a file—you need to UPPERCASE it. The `str.upper()` method is probably what you want.

5.2 Writing files

The output of the program should either appear on the command line or be written to a file. Command-line output is also called *standard out* or `STDOUT`. (It's the *standard* or normal place for *output* to occur.) Now let's look at how to write the output to a file.

We still need to `open()` a file handle, but we have to use an optional second argument, the string '`w`', to instruct Python to open it for *writing*. Other modes include '`r`' for *reading* (the default) and '`a`' for *appending*, as listed in table 5.1.

Table 5.1 File-writing modes

Mode	Meaning
<code>w</code>	Write
<code>r</code>	Read
<code>a</code>	Append

You can additionally describe the kind of content, whether '`t`' for *text* (the default) or '`b`' for *binary*, as listed in table 5.2.

Table 5.2 File-content modes

Mode	Meaning
t	Text
b	Bytes

You can combine the values in these two tables, like '`'rb'`' to *read* a *binary* file or '`'at'`' to *append* to a *text* file. Here we will use '`'wt'`' to *write* a *text* file.

I'll call my variable `out_fh` to remind me that this is the output file handle:

```
>>> out_fh = open('out.txt', 'wt')
```

If the file does not exist, it will be created. If the file does exist, it will be *overwritten*, which means that all the previous data will be lost! If you don't want an existing file to be lost, you can use the `os.path.isfile()` function you saw earlier to first check if the file exists, and perhaps use `open()` in the "append" mode instead. For this exercise, we'll use the '`'wt'`' mode to write text.

You can use the `write()` method of the file handle to put text into the file. Whereas the `print()` function will append a newline (`\n`) unless you instruct it not to, the `write()` method will *not* add a newline, so you have to explicitly add one.

If you use the `out_fh.write()` method in the REPL, you will see that it returns the number of bytes written. Here each character, including the newline (`\n`), is a byte:

```
>>> out_fh.write('this is some text\n')
18
```

You can check that this is correct:

```
>>> len('this is some text\n')
18
```

Most code tends to ignore this return value; that is, we don't usually bother to capture the results in a variable or check that we got a nonzero return. If `write()` fails, there's usually some much bigger problem with your system.

You can also use the `print()` function with the optional `file` argument. Notice that I don't include a newline with `print()` because it will add one. This method returns `None`:

```
>>> print('this is some more text', file=out_fh)
```

When you are done writing to a file handle, you should `out_fh.close()` it so that Python can clean up the file and release the memory associated with it. This method also returns `None`:

```
>>> out_fh.close()
```



Let's check if the lines of text we printed to our out.txt file made it by opening the file and reading it. Note that the newline appears here as \n. We need to print() the string for it to create an actual newline:

```
>>> open('out.txt').read()
'this is some text\nthis is some more text\n'
```

When we print() on an open file handle, the text will be appended to any previously written data. Look at this code, though:

```
>>> print("I am what I am an' I'm not ashamed.", file=open('hagrid.txt', 'wt'))
```

If you run that line twice, will the file called hagrid.txt have the line once or twice? Let's find out:

```
>>> open('hagrid.txt').read()
"I am what I am an' I'm not ashamed\n"
```

Just once! Why is that? Remember, each call to open() gives us a new file handle, so calling open() twice results in new file handles. Each time you run that code, the file is opened anew in *write* mode and the existing data is *overwritten*. To avoid confusion, I recommend you write code more along these lines:

```
fh = open('hagrid.txt', 'wt')
fh.write("I am what I am an' I'm not ashamed.\n")
fh.close()
```

5.3 Writing howler.py

You'll need to create a program called howler.py in the 05_howler directory. You can use the new.py program for this, copy template.py, or start however you prefer. Figure 5.5 is a string diagram showing an overview of the program and some example inputs and outputs.

When run with no arguments, it should print a short usage message:

```
$ ./howler.py
usage: howler.py [-h] [-o str] text
howler.py: error: the following arguments are required: text
```

When run with -h or --help, the program should print a longer usage statement:

```
$ ./howler.py -h
usage: howler.py [-h] [-o str] text

Howler (upper-cases input)

positional arguments:
  text                  Input string or file

optional arguments:
  -h, --help            show this help message and exit
  -o str, --outfile str
                        Output filename (default: )
```

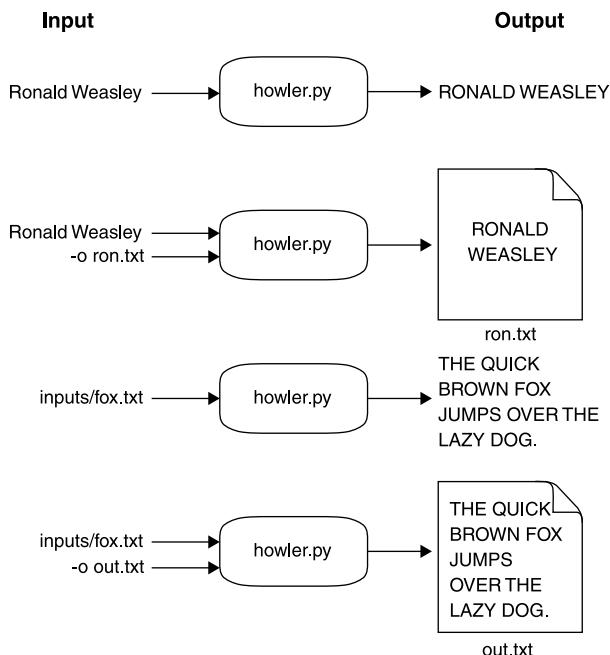


Figure 5.5 A string diagram showing that our `howler.py` program will accept strings or files as inputs and possibly an output filename.

If the argument is a regular string, it should uppercase that:

```
$ ./howler.py 'How dare you steal that car!'
HOW DARE YOU STEAL THAT CAR!
```

If the argument is the name of a file, it should uppercase the *contents of the file*:

```
$ ./howler.py ../inputs/fox.txt
THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
```

If given an `--outfile` filename, the uppercased text should be written to the indicated file and nothing should be printed to STDOUT:

```
$ ./howler.py -o out.txt ../inputs/fox.txt
$ cat out.txt
THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
```

Here are a few hints:

- Start with `new.py` and alter the `get_args()` section until your usage statements match the ones above.
- Run the test suite and try to pass just the first test that handles text on the command line and prints the uppercased output to STDOUT.
- The next test is to see if you can write the output to a given file. Figure out how to do that.

- The next test is for reading input from a file. Don't try to pass all the tests at once!
- There is a special file handle that always exists called "standard out" (often STDOUT). If you `print()` without a file argument, it defaults to `sys.stdout`. You will need to import `sys` in order to use it.

Be sure you really try to write the program and pass all the tests before moving on to read the solution. If you get stuck, maybe whip up a batch of Polyjuice Potion and freak out your friends.

5.4 Solution

Here is a solution that will pass the tests. It's rather short because Python allows us to express some really powerful ideas very concisely.

```
#!/usr/bin/env python3
"""Howler"""

import argparse
import os
import sys

# -----
def get_args():
    """get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Howler (upper-case input)',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text',
                        metavar='text',
                        type=str,
                        help='Input string or file')

    parser.add_argument('-o',
                        '--outfile',
                        help='Output filename',
                        metavar='str',
                        type=str,
                        default='')

    args = parser.parse_args()

    if os.path.isfile(args.text):
        args.text = open(args.text).read().rstrip()

    return args

    Return the arguments
    to the caller.

    Check if args.text is the
    name of an existing file.

    If so, overwrite the
    value of args.text with
    the results of reading
    the file.

    Parse the command-line arguments
    into the variable args so that we can
    manually check the text argument.

    The --outfile option is also
    a string that names a file.

    The text argument is a string
    that may be the name of a file.
```

```

# -----
def main():
    """Make a jazz noise here"""
    Use the
    opened file
    handle to
    write the
    output
    converted to
    uppercase.
    →
        args = get_args()           ← Call get_args () to get the
        out_fh = open(args.outfile, 'wt') if args.outfile else sys.stdout
        out_fh.write(args.text.upper() + '\n')
        out_fh.close()             ← Close the file
                                    handle.

# -----
if __name__ == '__main__':
    main()

```

Use an if expression to choose either sys.stdout or a newly opened file handle to write the output.

5.5 Discussion

How did it go for you this time? I hope you didn't sneak into Professor Snape's office again. You really don't want more Saturday detentions.

5.5.1 Defining the arguments

The `get_args()` function, as always, comes first. Here I define two arguments. The first is a positional text argument. Since it may or may not name a file, all I can know is that it will be a string.

```
parser.add_argument('text',
                    metavar='text',
                    type=str,
                    help='Input string or file')
```

NOTE If you define multiple positional parameters, their order *relative to each other* is important. The first positional parameter you define will handle the first positional argument provided. It's not important, however, to define positional parameters before or after options and flags. You can declare those in any order you like.

The other argument is an option, so I give it a short name of `-o` and a long name of `--outfile`. Even though the default type for all arguments is `str`, I like to state this explicitly. The default value is the empty string. I could just as easily use the special `None` type, which is also the default value, but I prefer to use a defined argument like the empty string.

```
parser.add_argument('-o',
                    '--outfile',
                    help='Output filename',
                    metavar='str',
                    type=str,
                    default='')
```

5.5.2 Reading input from a file or the command line

This is a deceptively simple program that demonstrates a couple of very important elements of file input and output. The text input might be a plain string, or it might be the name of a file. This pattern will come up repeatedly in this book:

```
if os.path.isfile(args.text):
    args.text = open(args.text).read().rstrip()
```

The `os.path.isfile()` function will tell me if there is a file with the specified name in `text`. If that returns True, I can safely `open(file)` to get a file handle, which has a method called `read` and which will return *all* the contents of the file.

WARNING You should be aware that `fh.read()` will return the *entire file* as a single string. Your computer must have more memory available than the size of the file. For all the programs in this book, you will be safe as the files are small. In my day job, I regularly deal with gigabyte-sized files, so calling `fh.read()` would likely crash my program if not my whole system, because I would exceed my available memory.

The result of `open(file).read()` is a `str`, which has a method called `str.rstrip()` that will return a copy of the string *stripped* of any whitespace on the *right* side (see figure 5.6). I call this so that the input text will look the same whether it comes from a file or directly from the command line. When you provide the input text directly on the command line, you have to press Enter to terminate the command. That Enter is a newline, and the operating system automatically removes it before passing it to the program.

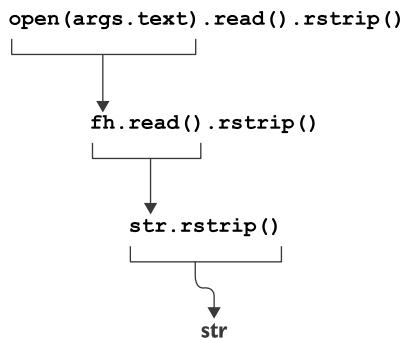


Figure 5.6 The `open()` function returns a file handle (`fh`). The `fh.read()` function returns a `str`. The `str.rstrip()` function returns a new `str` with the whitespace removed from the right side. All these functions can be chained together.

The longer way to write the preceding statement would be

```
if os.path.isfile(text):
    fh = open(text)
    text = fh.read()
    text = text.rstrip()
    fh.close()
```

In my version, I chose to handle this inside the `get_args()` function. This is the first time I've shown you that you can intercept and alter arguments before passing them on to `main()`. We'll use this idea quite a bit in later exercises.

I like to do all the work to validate the user's arguments inside `get_args()`. I could just as easily do this in `main()` after the call to `get_args()`, so this is entirely a style issue.

5.5.3 Choosing the output file handle

The following line decides where to put the output of the program:

```
out_fh = open(args.outfile, 'wt') if args.outfile else sys.stdout
```

The `if` expression will open `args.outfile` for writing text (`'wt'`) if the user provided that argument; otherwise, it will use `sys.stdout`, which is a file handle to `STDOUT`. Note that I don't have to call `open()` on `sys.stdout` because it is always available and open for business (figure 5.7).

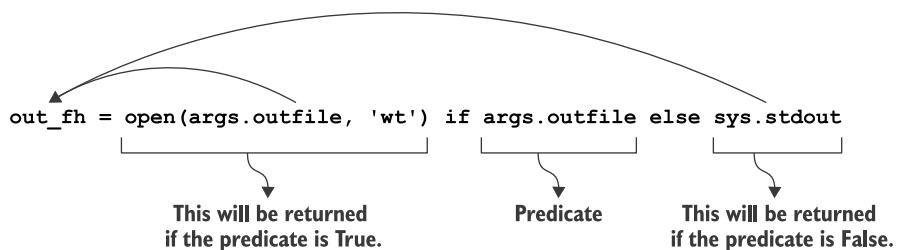


Figure 5.7 An `if` expression succinctly handles a binary choice. Here we want the output file handle to be the result of opening the `outfile` argument if present; otherwise, it should be `sys.stdout`.

5.5.4 Printing the output

To get the uppercase text, I can use the `text.upper()` method. Then I need to find a way to print it to the output file handle. I chose to do this:

```
out_fh.write(text.upper())
```

Alternatively, you could do this:

```
print(text.upper(), file=out_fh)
```

Finally, I need to close the file handle with `out_fh.close()`.

5.5.5 A low-memory version

There is a potentially serious problem waiting to bite us in this program. In `get_args()`, we're reading the entire file into memory with this line:

```
if os.path.isfile(args.text):
    args.text = open(args.text).read().rstrip()
```

We could, instead, only open() the file:

```
if os.path.isfile(args.text):
    args.text = open(args.text)
```

Later we could read it line by line:

```
for line in args.text:
    out_fh.write(line.upper())
```

The problem, though, is how to handle the times when the `text` argument is actually text and not the name of a file. The `io` (input-output) module in Python has a way to represent text as a *stream*:

```
>>> import io           ← Import the io module.
>>> text = io.StringIO('foo\nbar\nbaz\n')   ← Use the io.StringIO() function to turn
                                                the given str value into something we
                                                can treat like an open file handle.
>>> for line in text:
...     print(line, end='')           ← Use a for loop to iterate
...                                     through the “lines” of text
...                                     separated by newlines.
foo                         ← Print the line using the
bar                          end=' option to avoid
baz                         having two newlines.
```

This is the first time you’re seeing that you can treat a regular string value as if it were a generator of values similar to a file handle. This is a particularly useful technique for testing any code that needs to read an input file. You can use the return from `io.StringIO()` as a “mock” file handle so that your code doesn’t have to read an *actual* file, just a given value that can produce “lines” of text.

To make this work, we can change how we handle `args.text`, like so:

```
#!/usr/bin/env python3
"""Low-memory Howler"""

import argparse
import os
import io
import sys

# -----
def get_args():
    """get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Howler (upper-cases input)',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text',
                       metavar='text',
                       type=str,
                       help='Input string or file')

    parser.add_argument('-o',
                       '--outfile',
                       help='Output filename',
```

```

metavar='str',
type=str,
default='')

args = parser.parse_args()

if os.path.isfile(args.text):
    args.text = open(args.text)
else:
    args.text = io.StringIO(args.text + '\n')

return args

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    out_fh = open(args.outfile, 'wt') if args.outfile else sys.stdout
    for line in args.text:
        out_fh.write(line.upper())
    out_fh.close()

# -----
if __name__ == '__main__':
    main()

```

Check if args.text is a file.

If it is, replace args.text with the file handle created by opening the file.

Otherwise, replace args.text with an io.StringIO() value that will act like an open file handle. Note that we need to add a newline to the text so that it will look like the lines of input coming from an actual file.

Read the input (whether io.StringIO() or a file handle) line by line.

Process the line as before.

5.6 Going further

- Add a flag that will lowercase the input instead. Maybe call it `--ee` for the poet e e cummings, who liked to write poetry devoid of uppercase letters.
- Alter the program to handle multiple input files. Change `--outfile` to `--outdir`, and write each input file to the same filename in the output directory.

Summary

- To read or write files, you must first `open()` them.
- The default mode for `open()` is for reading a file.
- To write a text file, you must use `'wt'` as the second argument to `open()`.
- Text is the default type of data that you `write()` to a file handle. You must use the `'b'` flag to indicate that you want to write binary data.
- The `os.path` module contains many useful functions, such as `os.path.isfile()`, that will tell you if a file exists with a given name.
- `STDOUT` (standard output) is always available via the special `sys.stdout` file handle, which is always open.
- The `print()` function takes an optional `file` argument specifying where to put the output. That argument must be an open file handle, such as `sys.stdout` (the default) or the result of `open()`.



Words count: Reading files and STDIN, iterating lists, formatting strings

“I love to count!”

—Count von Count

Counting things is a surprisingly important programming skill. Maybe you’re trying to find out how many pizzas were sold each quarter or how many times you see certain words in a set of documents. Usually the data we deal with in computing comes to us in files, so in this chapter, we’re going to push a little further into reading files and manipulating strings.

We’re going to write a Python version of the venerable `wc` (“word count”) program. Ours will be called `wc.py`, and it will count the lines, words, and bytes found in each input supplied as one or more positional arguments. The counts will appear in columns eight characters wide, and they will be followed by the name of the file. For instance, here is what `wc.py` should print for one file:

```
$ ./wc.py ../inputs/scarlet.txt
    7035   68061  396320 ../inputs/scarlet.txt
```



When counting multiple files, there will be an additional “total” line summing each column:

```
$ ./wc.py ../inputs/const.txt ../inputs/sonnet-29.txt
    865      7620    44841 ../inputs/const.txt
      17       118      661 ../inputs/sonnet-29.txt
    882     7738   45502 total
```

There may also be *no* arguments, in which case we'll read from *standard in*, which is often written as `STDIN`. We started talking about `STDOUT` in chapter 5 when we used `sys.stdout` as a file handle. `STDIN` is the complement to `STDOUT`—it's the “standard” place to read input on the command line. When our program is given *no* positional arguments, it will read from `sys.stdin`.

`STDIN` and `STDOUT` are common file handles that many command-line programs recognize. We can chain the `STDOUT` from one program to the `STDIN` of another to create ad hoc programs. For instance, the `cat` program will print the contents of a file to `STDOUT`. We can use the pipe operator (`|`) to funnel that output as input into our program via `STDIN`:

```
$ cat ../inputs/fox.txt | ./wc.py
      1         9        45 <stdin>
```

Another option is to use the `<` operator to redirect input from a file:

```
$ ./wc.py < ../inputs/fox.txt
      1         9        45 <stdin>
```

One of the handiest command-line tools is `grep`, which can find patterns of text in files. If, for instance, we wanted to find all the lines of text that contain the word “scarlet” in all the files in the `inputs` directory, we could use this command:

```
$ grep scarlet ../inputs/*.txt
```

On the command line, the asterisk (*) is a wildcard that will match anything, so `*.txt` will match any file ending with “.txt.” If you run the preceding command, you'll see quite a bit of output.

To count the lines found by `grep`, we can pipe that output into our `wc.py` program like so:

```
$ grep scarlet ../inputs/*.txt | ./wc.py
    108     1192     9201 <stdin>
```

We can verify that this matches what `wc` finds:

```
$ grep scarlet ../inputs/*.txt | wc
    108     1192     9201
```

In this chapter, you will

- Learn how to process zero or more positional arguments
- Validate input files
- Read from files or from standard input
- Use multiple levels of `for` loops
- Break files into lines, words, and bytes
- Use counter variables
- Format string output

6.1 Writing *wc.py*

Let's get started! Create a program called *wc.py* in the `06_wc` directory, and modify the arguments until it will print the following usage if run with the `-h` or `--help` flags:

```
$ ./wc.py -h
usage: wc.py [-h] [FILE [FILE ...]]

Emulate wc (word count)

positional arguments:
  FILE      Input file(s) (default: [<_io.TextIOWrapper name='<stdin>' mode='r' encoding='UTF-8']])

optional arguments:
  -h, --help show this help message and exit
```

Given a nonexistent file, your program should print an error message and exit with a nonzero exit value:

```
$ ./wc.py blargh
usage: wc.py [-h] [FILE [FILE ...]]
wc.py: error: argument FILE: can't open 'blargh': \
[Errno 2] No such file or directory: 'blargh'
```

Figure 6.1 is a string diagram that will help you think about how the program should work.

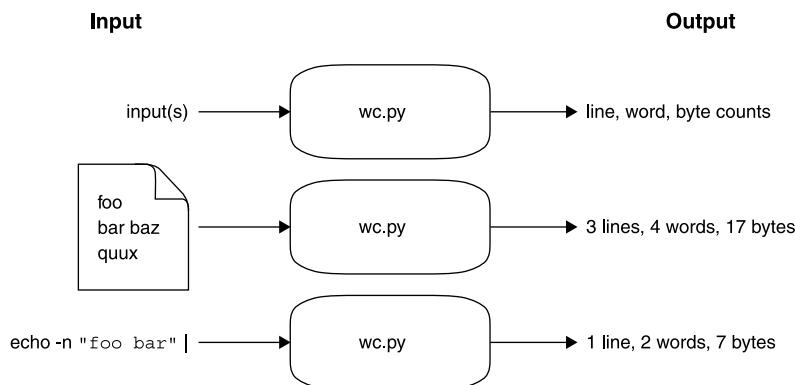


Figure 6.1 A string diagram showing that `wc.py` will read one or more file inputs or possibly `STDIN` and will produce a summary of the words, lines, and bytes contained in each input.

6.1.1 Defining file inputs

Let's talk about how we can define the program's parameters using argparse. This program takes *zero or more* positional arguments and nothing else. Remember that you never have to define the -h or --help arguments, as argparse handles those automatically.

In chapter 3 we used nargs='+' to indicate one or more items for our picnic. Here we want to use nargs='*' to indicate *zero or more*. When there are no arguments, the default value will be None. For this program, we'll read STDIN when there are no arguments.

All of the possible values for nargs are listed in table 6.1.

Table 6.1 Possible values for nargs

Symbol	Meaning
?	Zero or one
*	Zero or more
+	One or more

Any arguments that are provided to our program *must be readable files*. In chapter 5 you learned how to test whether the input argument was a file by using os.path.isfile(). The input was allowed to be either plain text or a filename, so you had to check this yourself.

In this program, the input arguments are required to be readable text files, so we can define our arguments using type=argparse.FileType('rt'). This means that argparse takes on all the work of validating the inputs from the user and producing useful error messages. If the user provides valid input, argparse will provide a list of *open file handles*. All in all, this will save us quite a bit of time. (Be sure to review section A.4.6 on file arguments in the appendix.)

In chapter 5 we used sys.stdout to write to STDOUT. To read from STDIN here, we'll use Python's sys.stdin file handle. Like sys.stdout, the sys.stdin file handle does not need an open()—it's always present and available for reading.

Because we are using nargs='*' to define our argument, the result will always be a list. To set sys.stdin as the default value, we should place it in a list like so:

```
parser.add_argument('file',
                    metavar='FILE',
                    nargs='*',
                    type=argparse.FileType('rt'),
                    default=[sys.stdin],
                    help='Input file(s)')

The default will be a list containing sys.stdin, which is like an
open file handle to STDIN. We do not need to open it.

Zero or more of
this argument

If arguments are provided,
they must be readable text
files. The files will be opened
by argparse and will be
provided as file handles.
```

6.1.2 Iterating lists

Your program will end up with a list of file handles that will need to be processed. In chapter 4 we used a for loop to iterate through the characters in the input text. Here we can use a for loop over the args.file inputs, which will be open file handles:

```
for fh in args.file:  
    # read each file
```

You can give whatever name you like to the variable you use in your for loop, but I think it's very important to give it a semantically meaningful name. Here the variable name `fh` reminds me that this is an open file handle. You saw in chapter 5 how to manually `open()` and `read()` a file. Here `fh` is already open, so we can use it directly to read the contents.

There are many ways to read a file. The `fh.read()` method will give you the *entire contents* of the file in one go. If the file is large—if it exceeds the available memory on your machine—your program will crash. I would recommend, instead, that you use another for loop on the `fh`. Python will understand this to mean that you wish to read each line of the file handle, one at a time.

```
for fh in args.file: # ONE LOOP!  
    for line in fh: # TWO LOOPS!  
        # process the line
```

That's two levels of for loops, one for each file handle and then another for each line in each file handle. ONE LOOP! TWO LOOPS! I LOVE TO COUNT!

6.1.3 What you're counting

The output for each file will be the number of lines, words, and bytes (like characters and whitespace), each of which is printed in a field eight characters wide, followed by a space and then the name of the file, which will be available to you via `fh.name`.

Let's take a look at the output from the standard `wc` program on my system. Notice that when it's run with just one argument, it produces counts only for that file:

```
$ wc fox.txt  
     1      9     45 fox.txt
```

The `fox.txt` file is short enough that you could manually verify that it does in fact contain 1 line, 9 words, and 45 bytes, which includes all the characters, spaces, and the trailing newline (see figure 6.2).

When run with multiple files, the standard `wc` program also shows a “total” line:

```
$ wc fox.txt sonnet-29.txt  
     1      9     45 fox.txt  
    17    118    669 sonnet-29.txt  
    18    127    714 total
```

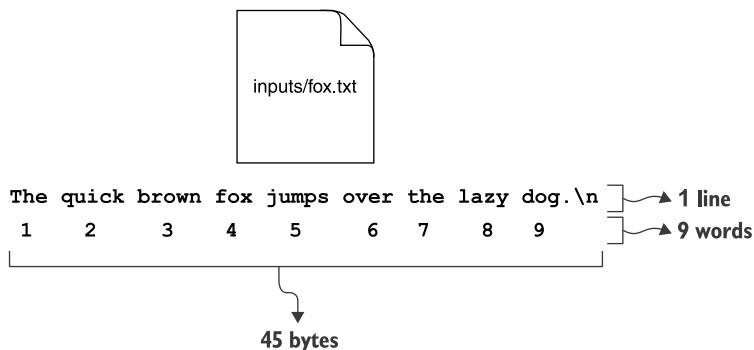


Figure 6.2 The fox.txt file contains 1 line of text, 9 words, and a total of 45 bytes.

We are going to emulate the behavior of this program. For each file, you will need to create variables to hold the numbers of lines, words, and bytes. For instance, if you use the `for line in fh` loop that I suggest, you will need to have a variable like `num_lines` to increment on each iteration.

That is, somewhere in your code you will need to set a variable to 0 and then, inside the `for` loop, make it go up by 1. The idiom in Python is to use the `+=` operator to add some value on the right side to the variable on the left side (as shown in figure 6.3):

```
num_lines = 0
for line in fh:
    num_lines += 1
```

You will also need to count the number of words and bytes, so you'll need similar `num_words` and `num_bytes` variables.

To get the words, we'll use the `str.split()` method to break each line on spaces. You can then use the length of the resulting list as the number of words. For the number of bytes, you can use the `len()` (length) function on the `line` and add that to a `num_bytes` variable.

NOTE Splitting the text on spaces doesn't actually produce "words" because it won't separate the punctuation, like commas and periods, from the letters, but it's close enough for this program. In chapter 15, we'll look at how to use a regular expression to separate strings that look like words from others that do not.

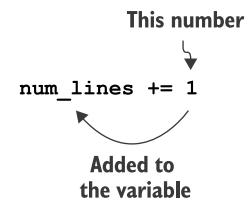


Figure 6.3 The `+=` operator will add the value on the right to the variable on the left.

6.1.4 Formatting your results

This is the first exercise where the output needs to be formatted in a particular way. Don't try to handle this part manually—that way lies madness. Instead, you need to learn the magic of the `str.format()` method. The help doesn't have much in the way

of documentation, so I recommend you read PEP 3101 on advanced string formatting (www.python.org/dev/peps/pep-3101/).

The `str.format()` method uses a template that contains curly brackets (`{}`) to create placeholders for the values passed as arguments. For example, we can print the raw value of `math.pi` like so:

```
>>> import math  
>>> 'Pi is {}'.format(math.pi)  
'Pi is 3.141592653589793'
```

You can add formatting instructions after a colon (`:`) to specify how you want the value displayed. If you are familiar with `printf()` from C-type languages, this is the same idea. For instance, I can print `math.pi` with two numbers after the decimal by specifying `0.02f`:

```
>>> 'Pi is {:.02f}'.format(math.pi)  
'Pi is 3.14'
```

In the preceding example, the colon (`:`) introduces the formatting options, and the `0.02f` describes two decimal points of precision.

You can also use the f-string method, where the variable comes *before* the colon:

```
>>> f'Pi is {math.pi:0.02f}'  
'Pi is 3.14'
```

In this chapter's exercise, you need to use the formatting option `{:8}` to align each of the lines, words, and characters into columns. The `8` describes the width of the field. The text is usually left-justified, like so:

```
>>> '{:8}'.format('hello')  
'hello '
```

But the text will be right-justified when you are formatting numeric values:

```
>>> '{:8}'.format(123)  
'    123'
```

You will need to place a single space between the last column and the name of the file, which you can find in `fh.name`.

Here are a few hints:

- Start with `new.py` and delete all the nonpositional arguments.
- Use `nargs='*' to indicate zero or more positional arguments for your file argument.`
- Try to pass one test at a time. Create the program, get the help right, and then worry about the first test, then the next, and so on.
- Compare the results of your version to the `wc` installed on your system. Note that not every system has the same version of `wc`, so results may vary.

It's time to write this yourself before you read the solution. Fear is the mind killer. You can do this.

6.2 Solution

Here is one way to satisfy the tests. Remember, it's fine if you wrote it differently, as long as it's correct and you understand your code!

```
#!/usr/bin/env python3
"""Emulate wc (word count)"""

import argparse
import sys

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Emulate wc (word count)',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file',
                        metavar='FILE',
                        nargs='*',
                        default=[sys.stdin],
                        type=argparse.FileType('rt'),
                        help='Input file(s)')

    return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    total_lines, total_bytes, total_words = 0, 0, 0
    for fh in args.file:
        num_lines, num_words, num_bytes = 0, 0, 0
        for line in fh:
            num_lines += 1
            num_bytes += len(line)
            num_words += len(line.split())
        total_lines += num_lines
        total_bytes += num_bytes
        total_words += num_words

    print(f'{total_lines} {total_bytes} {total_words}')


Initialize variables to
count the lines, words,
and bytes in just this file.

To get the number of words, we can
call line.split() to break the line on
whitespace. The length of that list is
added to the count of words.

These are the variables
for the "total" line, if I
need them.

The number of bytes is
incremented by the
length of the line.

Add all the counts for lines, words, and bytes for
this file to the variables for counting the totals.

If you set the default to a
list with sys.stdin, you have
handled the STDIN option.

If the user supplies any
arguments, argparse will
check if they are valid file
inputs. If there is a problem,
argparse will halt execution
of the program and show the
user an error message.

Iterate through the list of
arg.file inputs. I use the
variable fh to remind me
that these are open file
handles, even STDIN.

Iterate through each
line of the file handle.

For each line, increment
the number of lines by 1.
```

```

    print(f'{num_lines:8}{num_words:8}{num_bytes:8} {fh.name}')

    if len(args.file) > 1:
        print(f'{total_lines:8}{total_words:8}{total_bytes:8} total')

# -----
if __name__ == '__main__':
    main()

Print the counts for this file using the {:8} option
to print in a field 8 characters wide followed by a
single space and then the name of the file.

Print the
"total" line.

Check if we had
more than 1 input.

```

6.3 Discussion

This program is rather short and seems rather simple, but it's not exactly easy. Let's break down the main ideas in the program.

6.3.1 Defining the arguments

One point of this exercise is to get familiar with argparse and the trouble it can save you. The key is in defining the `file` parameter. We use `type=argparse.FileType('rt')` to indicate that any arguments provided must be readable text files. We use `nargs='*'` to indicate zero or more arguments, and we set the default to be a list containing `sys.stdin`. This means we know that `argparse` will always give us a list of one or more open file handles.

That's really quite a bit of logic packed into a small space, and most of the work validating the inputs, generating error messages, and handling the defaults is all done for us!

6.3.2 Reading a file using a for loop

The values that `argparse` returns for `args.file` will be a list of *open file handles*. We can create such a list in the REPL to mimic what we'd get from `args.file`:

```
>>> files = [open('../inputs/fox.txt')]
```

Before we use a `for` loop to iterate through them, we need to set up three variables to track the *total* number of lines, words, and characters. We could define them on three separate lines:

```
>>> total_lines = 0
>>> total_words = 0
>>> total_bytes = 0
```

Or we can declare them on a single line like the following:

```
>>> total_lines, total_words, total_bytes = 0, 0, 0
```

Technically we're creating a tuple on the right side by placing commas between the three zeros and then "unpacking" them into three variables on the left side. I'll have more to say about tuples much later.

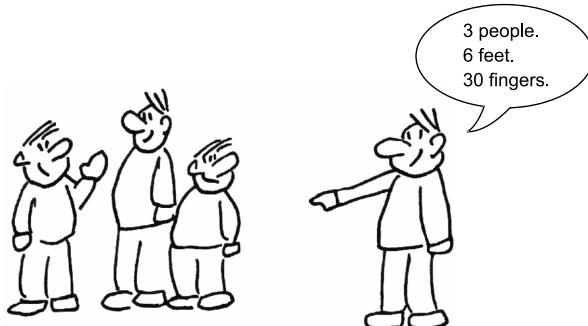
Inside the `for` loop for each file handle, we initialize three more variables to hold the count of lines, characters, and words *for this particular file*. We can then use another `for` loop to iterate over each line in the file handle (`fh`). For `lines`, we can add 1 on each pass through the `for` loop. For `bytes`, we can add the length of the line (`len(line)`) to track the number of "characters" (which may be printable characters or whitespace, so it's easiest to call them "bytes"). Lastly, for `words`, we can use `line.split()` to break the line on whitespace to create a list of "words." It's not a perfect way to count actual words, but it's close enough. We can use the `len()` function on the list to add to the `words` variable.

The `for` loop ends when the end of the file is reached. Next we can `print()` out the counts and the filename, using `{:8}` placeholders in the print template to indicate a text field 8 characters wide:

```
>>> for fh in files:
...     lines, words, bytes = 0, 0, 0
...     for line in fh:
...         lines += 1
...         bytes += len(line)
...         words += len(line.split())
...     print(f'{lines:8}{words:8}{bytes:8} {fh.name}')
...     total_lines += lines
...     total_bytes += bytes
...     total_words += words
...
1      9      45 ./inputs/fox.txt
```

Notice that the preceding call to `print()` lines up with the *second* `for` loop, so that it will run after we're done iterating over the lines in `fh`. I chose to use the f-string method to print each of `lines`, `words`, and `bytes` in a space eight characters wide, followed by one space and then the `fh.name` of the file.

After printing, we can add the counts to the "total" variables to keep a running total.



Lastly, if the number of file arguments is greater than 1, we need to print the totals:

```
if len(args.file) > 1:  
    print(f'{total_lines:8}{total_words:8}{total_bytes:8} total')
```

6.4 Going further

- By default, wc will print all the columns like our program does, but it will also accept flags to print -c for number of characters, -l for number of lines, and -w for number of words. When any of these flags are present, only columns for the specified flags are shown, so wc.py -wc would show just the columns for words and characters. Add short and long flags for these options to your program so that it behaves exactly like wc.
- Write your own implementation of other system tools like cat (to print the contents of a file to STDOUT), head (to print just the first n lines of a file), tail (to print the last n lines of a file), and tac (to print the lines of a file in reverse order).

Summary

- The nargs (number of arguments) option to argparse allows you to validate the number of arguments from the user. The asterisk ('*') means zero or more, whereas '+' means one or more.
- If you define an argument using type=argparse.FileType('rt'), argparse will validate that the user has provided a readable text file and will make the value available in your code as an open file handle.
- You can read and write from the standard in/out file handles by using sys.stdin and sys.stdout.
- You can nest for loops to handle multiple levels of processing.
- The str.split() method will split a string on spaces.
- The len() function can be used on both strings and lists. For lists, it will tell you the number of elements the list contains.
- Both str.format() and Python's f-strings recognize printf-style formatting options to allow you to control how a value is displayed.



Gashlycrumb: Looking items up in a dictionary

In this chapter, we’re going to look up lines of text from an input file that start with the letters provided by the user. The text will come from an input file that will default to Edward Gorey’s “The Gashlycrumb Tinies,” an abecedarian book that describes various and ghastly ways in which children expire. For instance, figure 7.1 shows “N is for Neville who died of ennui.”

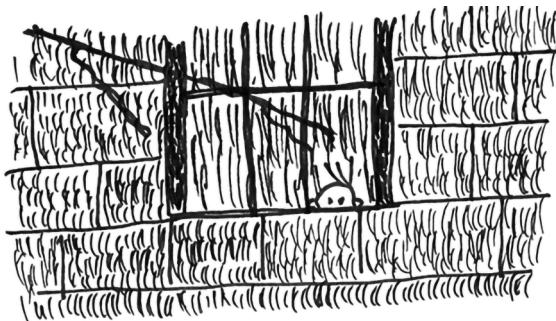


Figure 7.1 N is for Neville who died of ennui.

Our gashlycrumb.py program will take one or more letters as positional arguments and will look up the lines of text that start with that letter from an *optional* input file. We will look up the letters in a *case-insensitive* fashion.

The input file will have the value for each letter on a separate line:

```
$ head -2 gashlycrumb.txt
A is for Amy who fell down the stairs.
B is for Basil assaulted by bears.
```

When our unfortunate user runs this program, here is what they will see:

```
$ ./gashlycrumb.py e f
E is for Ernest who choked on a peach.
F is for Fanny sucked dry by a leech.
```

In this exercise, you will

- Accept one or more positional arguments that we'll call letter.
- Accept an optional --file argument, which must be a readable text file. The default value will be 'gashlycrumb.txt' (provided).
- Read the file, find the first letter of each line, and build a data structure that associates the letter to the line of text. (We'll only be using files where each line starts with a single, unique letter. This program would fail with any other format of text.)
- For each letter provided by the user, either print the line of text for the letter if present, or print a message if it isn't.
- Learn how to "pretty-print" a data structure.

You can draw from several previous programs:

- From chapter 2 you know how to get the first letter of a bit of text.
- From chapter 4 you know how to build a dictionary and look up a value.
- From chapter 6 you know how to accept a file input argument and read it line by line.



Now you'll put all those skills together to recite morbid poetry!

7.1 Writing *gashlycrumb.py*

Before you begin writing, I encourage you to run the tests with `make test` or `pytest -xv test.py` in the `07_gashlycrumb` directory. The first test should fail:

```
test.py::test_exists FAILED
```

This is just a reminder that the first thing you need to do is create the file called `gashlycrumb.py`. You can do this however you like, such as by running `new.py gashlycrumb.py` in the `07_gashlycrumb` directory, by copying the template/`template.py` file, or by just starting a new file from scratch. Run your tests again, and you should pass the first test and possibly the second if your program produces a usage statement.

Next, let's get the arguments straight. Modify your program's parameters in the `get_args()` function so that it will produce the following usage statement when the program is run with no arguments or with the `-h` or `--help` flags:

```
$ ./gashlycrumb.py -h
usage: gashlycrumb.py [-h] [-f FILE] letter [letter ...]
```

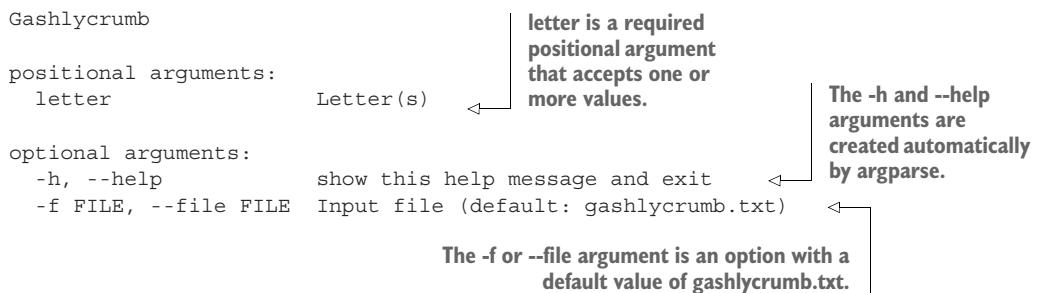


Figure 7.2 shows a string diagram of how the program will work.

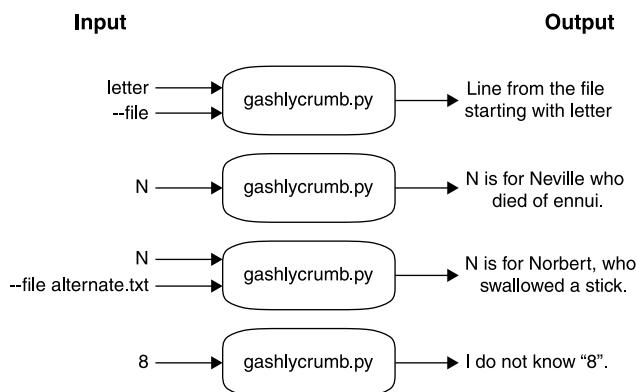


Figure 7.2 Our program will accept some letter(s) and possibly a file. It will then look up the line(s) of the file starting with the given letter(s).

In the `main()` function, start off by echoing each of the `letter` arguments:

```
def main():
    args = get_args()
    for letter in args.letter:
        print(letter)
```

Try running it to make sure it works:

```
$ ./gashlycrumb.py a b
a
b
```

Next, read the file line by line using a `for` loop:

```
def main():
    args = get_args()
```

```

for letter in args.letter:
    print(letter)

for line in args.file:
    print(line, end='')

```

Note that I'm using `end=''` with `print()` so that it won't print the newline that's already attached to each line of the file:

Try running it to ensure you can read the input file:

```
$ ./gashlycrumb.py a b | head -4
a
b
A is for Amy who fell down the stairs.
B is for Basil assaulted by bears.
```

Use the alternate.txt file too:

```
$ ./gashlycrumb.py a b --file alternate.txt | head -4
a
b
A is for Alfred, poisoned to death.
B is for Bertrand, consumed by meth.
```

If your program is provided a `--file` argument that does not exist, it should exit with an error and message. Note that if you declare the parameter in `get_args()` using `type=argparse.FileType('rt')` as we did in the previous chapter, this error should be produced automatically by `argparse`:

```
$ ./gashlycrumb.py -f blargh b
usage: gashlycrumb.py [-h] [-f FILE] letter [letter ...]
gashlycrumb.py: error: argument -f/--file: can't open 'blargh': \
[Errno 2] No such file or directory: 'blargh'
```

Now think about how you can use the first letter of each line to create an entry in a dict. Use `print()` to look at your dictionary. Figure out how to check if the given letter is *in* (wink, wink, nudge, nudge) your dictionary.

If your program is given a value that does not exist in the list of first characters on the lines from the input file (when searched without regard to case), you should print a message:

```
$ ./gashlycrumb.py 3
I do not know "3".
$ ./gashlycrumb.py CH
I do not know "CH".
```



If the given letter is in the dictionary, print the value for it (see figure 7.3):

```
$ ./gashlycrumb.py a
A is for Amy who fell down the stairs.
```

```
$ ./gashlycrumb.py z
Z is for Zillah who drank too much gin.
```

```
{ "A": "A is for Amy who fell down the stairs" }
```

Figure 7.3 We need to create a dictionary where the first letter of each line is the key and the line is the value.

Run the test suite to ensure your program meets all the requirements. Read the errors closely and fix your program.

Here are some hints:

- Start with new.py and remove everything but the positional letter and optional --file parameters.
- Use type=argparse.FileType('rt') to validate the --file argument.
- Use nargs='+' to define the positional argument letter so it will require one or more values.
- A dictionary is a natural data structure for associating a value like the letter "A" to a phrase like "A is for Amy who fell down the stairs." Create a new, empty dict.
- Once you have an open file handle, you can read the file line by line with a for loop.
- Each line of text is a string. How can you get the first character of a string?
- Create an entry in your dictionary using the first character as the key and the line itself as the value.
- Iterate through each letter argument. How can you check that a given value is in the dictionary?

No skipping ahead to the solution until you have written your own version! If you peek, you will die a horrible death: stamped by kittens.

7.2 **Solution**

I really hope you looked at Gorey's artwork for his book. Now let's talk about how to build a dictionary from a file input:

```
#!/usr/bin/env python3
"""Lookup tables"""

import argparse
```

```

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Gashlycrumb',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('letter',
                        help='Letter(s) ',
                        metavar='letter',
                        nargs='+',
                        type=str)                                ← A positional argument called letter
                                                               uses nargs='+' to indicate that
                                                               one or more values are required.

    parser.add_argument('-f',
                        '--file',
                        help='Input file',
                        metavar='FILE',
                        type=argparse.FileType('rt'),           ← The optional --file
                                                               argument must be a
                                                               readable file because of
                                                               type=argparse.FileType('rt').
                                                               The default value is
                                                               gashlycrumb.txt, which I
                                                               know exists.

    return parser.parse_args()

```



```

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    lookup = {}                                ← Create an empty dictionary
                                                to hold the lookup table.

    for line in args.file:                      ← Iterate through each line of
                                                the args.file, which will be
                                                an open file handle.

        lookup[line[0].upper()] = line.rstrip()   ← Uppercase the first
                                                character of the line to
                                                use as the key into the
                                                lookup table and set
                                                the value to be the line
                                                stripped of whitespace
                                                on the right side.

    for letter in args.letter:
        if letter.upper() in lookup:             ← Check if the letter is in the
                                                lookup dictionary, using
                                                letter.upper() to disregard case.

            print(lookup[letter.upper()])          ← If so, print the line of
                                                text from the
                                                lookup for
                                                the letter.

        else:
            print(f'I do not know "{letter}".')  ← Otherwise, print a
                                                message saying the
                                                letter is unknown.

    if __name__ == '__main__':
        main()

```

7.3 Discussion

Did the frightful paws of the kittens hurt much? Let's talk about how I solved this problem. Remember, mine is just one of many possible solutions.

7.3.1 Handling the arguments

I prefer to have all the logic for parsing and validating the command-line arguments in the `get_args()` function. In particular, `argparse` can do a fine job of verifying tedious things, such as an argument being an existing, readable text file, which is why

I use `type=argparse.FileType('rt')` for that argument. If the user doesn't supply a valid argument, argparse will throw an error, printing a helpful message along with the short usage statement, and will exit with an error code.

By the time I get to the line `args = get_args()`, I know that I have one or more "letter" arguments and a valid, open file handle in the `args.file` slot. In the REPL, I can use `open` to get a file handle, which I usually like to call `fh`. For copyright purposes, I'll use my alternate text:

```
>>> fh = open('alternate.txt')
```

7.3.2 *Reading the input file*

We want to use a dictionary where the keys are the first letters of each line and the values are the lines themselves. That means we need to start by creating a new, empty dictionary, either by using the `dict()` function or by setting a variable equal to an empty set of curly brackets (`{}`). Let's call the variable `lookup`:

```
>>> lookup = {}
```

We can use a `for` loop to read each line of text. From the Crow's Nest program in chapter 2, you know we can use `line[0].upper()` to get the first letter of `line` and uppercase it. We can use that as the key into `lookup`.

Each line of text ends with a newline that I'd like to remove. The `str.rstrip()` method will strip whitespace from the right side of the line ("rstrip" = *right strip*). The result of that will be the value for my `lookup`:

```
for line in fh:
    lookup[line[0].upper()] = line.rstrip()
```

Let's look at the resulting `lookup` dictionary. We can `print()` it from the program or type `lookup` in the REPL, but it's going to be hard to read. I encourage you to try it.

Luckily there is a lovely module called `pprint` to "pretty-print" data structures. Here is how you can import the `pprint()` function from the `pprint` module with the alias `pp`:

```
>>> from pprint import pprint as pp
```

Figure 7.4 illustrates how this works.

```
from pprint import pprint as pp
          ↓      ↓      ↓
Module   Function  Alias
```

Figure 7.4 We can specify exactly which functions to import from a module and even give the function an alias.

Now let's take a peek at the lookup table:

```
>>> pp(lookup)
{'A': 'A is for Alfred, poisoned to death.',
 'B': 'B is for Bertrand, consumed by meth.',
 'C': 'C is for Cornell, who ate some glass.',
 'D': 'D is for Donald, who died from gas.',
 'E': 'E is for Edward, hanged by the neck.',
 'F': 'F is for Freddy, crushed in a wreck.',
 'G': 'G is for Geoffrey, who slit his wrist.',
 'H': "H is for Henry, who's neck got a twist.",
 'I': 'I is for Ingrid, who tripped down a stair.',
 'J': 'J is for Jered, who fell off a chair.',
 'K': 'K is for Kevin, bit by a snake.',
 'L': 'L is for Lauryl, impaled on a stake.',
 'M': 'M is for Moira, hit by a brick.',
 'N': 'N is for Norbert, who swallowed a stick.',
 'O': 'O is for Orville, who fell in a canyon.',
 'P': 'P is for Paul, strangled by his banyan.',
 'Q': 'Q is for Quintanna, flayed in the night.',
 'R': 'R is for Robert, who died of spite.',
 'S': 'S is for Susan, stung by a jelly.',
 'T': 'T is for Terrange, kicked in the belly.',
 'U': "U is for Uma, who's life was vanquished.",
 'V': 'V is for Victor, consumed by anguish.',
 'W': "W is for Walter, who's socks were too long.",
 'X': 'X is for Xavier, stuck through with a prong.',
 'Y': 'Y is for Yoeman, too fat by a piece.',
 'Z': 'Z is for Zora, smothered by a fleece.'}
```

Hey, that looks like a handy data structure. Hooray for us! Please don't discount the value of using lots of `print()` calls when you are trying to write and understand a program, and of using the `pprint()` function whenever you need to see a complex data structure.

7.3.3 Using a dictionary comprehension

In chapter 4 you saw that you can use a list comprehension to build a list by putting a `for` loop inside `[]`. If we change the brackets to curlies (`{}`), we create a dictionary comprehension:

```
>>> fh = open('gashlycrumb.txt')
>>> lookup = { line[0].upper(): line.rstrip() for line in fh }
```

See in figure 7.5 how we can rearrange three lines of our `for` loop into a single line of code.

If you print the `lookup` table again, you should see the same output as before. It may seem like showing off to write one line of code instead of three, but it really does make a good deal of sense to write compact, idiomatic code. More code always means more chances for bugs, so I usually try to write code that is as simple as possible (but no simpler).

```

lookup = {}
for line in fh:
    lookup[line[0].upper()] = line.rstrip()

```

→ `lookup = { line[0].upper(): line.rstrip() for line in fh }`

Figure 7.5 The `for` loop we used to build a dictionary can be written using a dictionary comprehension.

7.3.4 Dictionary lookups

Now that I have a lookup table, I can ask whether some value is in the keys. I know the letters are in uppercase, and since the user could give me a lowercase letter, I use `letter.upper()` to only compare that case:

```

>>> letter = 'a'
>>> letter.upper() in lookup
True
>>> lookup[letter.upper()]
'A is for Amy who fell down the stairs.'

```



If the letter is found, I can print the line of text for that letter; otherwise, I can print a message saying that I don't know that letter:

```

>>> letter = '4'
>>> if letter.upper() in lookup:
...     print(lookup[letter.upper()])
... else:
...     print('I do not know "{}".format(letter)')
...
I do not know "4".

```

An even shorter way to write that would be to use the `dict.get()` method:

```

def main():
    args = get_args()
    lookup = {line[0].upper(): line.rstrip() for line in args.file}

    for letter in args.letter:
        print(lookup.get(letter.upper(), f'I do not know "{letter}".')) ←

```

lookup.get() will return the value for `letter.upper()` or the warning about a value not being found in our lookup.

7.4 Going further

- Write a phonebook that reads a file and creates a dictionary from the names of your friends and their email or phone numbers.
- Create a program that uses a dictionary to count the number of times you see each word in a document.

- Write an interactive version of the program that takes input directly from the user. Use while True to set up an infinite loop and keep using the input() function to get the user's next letter:

```
$ ./gashlycrumb_interactive.py
Please provide a letter [! to quit]: t
T is for Titus who flew into bits.
Please provide a letter [! to quit]: 7
I do not know "7".
Please provide a letter [! to quit]: !
Bye
```

- Interactive programs are fun to write, but how would you go about testing them? In chapter 17 I'll show you one way to do this.

Summary

- A dictionary comprehension is a way to build a dictionary in a one-line for loop.
- Defining file input arguments using argparse.FileType saves you time and code.
- Python's pprint module is used to pretty-print complex data structures.

Apples and Bananas: Find and replace

Have you ever misspelled a word? I haven't, but I've heard that many other people often do. We can use computers to find and replace all instances of a misspelled word with the correction. Or maybe you'd like to replace all mentions of your ex's name in your poetry with your new love's name? Find and replace is your friend.

To get us started, let's consider the children's song "Apples and Bananas," wherein we intone our favorite fruits to consume:

I like to eat, eat, eat apples and bananas

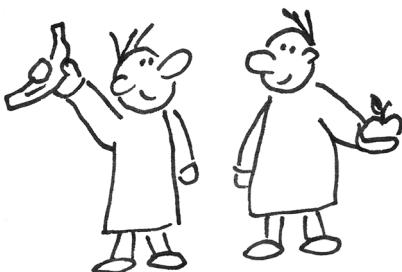
Subsequent verses substitute the main vowel sound in the fruits for various other vowel sounds, such as the long "a" sound (as in "hay"):

I like to ate, ate, ate ay-ples and ba-nay-nays

Or the ever-popular long "e" (as in "knee"):

I like to eat, eat, eat ee-ples and bee-nee-nees

And so forth. In this exercise, we'll write a Python program called `apples.py` that takes some text, given as a single positional argument, and replaces all the vowels in the text with the given `-v` or `--vowel` options (with the default being `a`).



The program should be written in the 08_apples_and_bananas directory and should handle text on the command line:

```
$ ./apples.py foo
faa
```

And accept the `-v` or `--vowel` option:

```
$ ./apples.py foo -v i
fii
```

Your program should *preserve the case* of the input vowels:

```
$ ./apples.py -v i "APPLES AND BANANAS"
IPPLIS IND BININIS
```

As with the Howler program in chapter 5, the text argument may name a file, in which case your program should read the contents of the file:

```
$ ./apples.py ../inputs/fox.txt
Tha qaack brawn fax jamps avar tha lazy dag.
$ ./apples.py --vowel e ../inputs/fox.txt
The qeeck brewn fex jemps ever the lezy deg.
```

Figure 8.1 shows a diagram of the program's inputs and output.

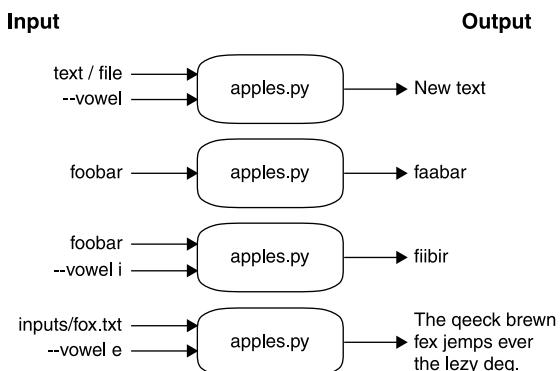


Figure 8.1 Our program will accept some text and possibly a vowel. All the vowels in the given text will be changed to the same vowel, resulting in hilarity.

Here is the usage statement that should print when there are no arguments:

```
$ ./apples.py
usage: apples.py [-h] [-v vowel] text
apples.py: error: the following arguments are required: text
```

And the program should always print usage for the `-h` and `--help` flags:

```
$ ./apples.py -h
usage: apples.py [-h] [-v vowel] text
```

```
Apples and bananas
```

positional arguments:

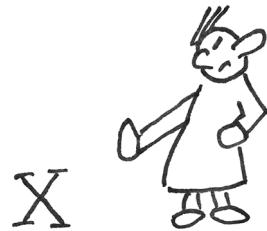
text	Input text or file
------	--------------------

optional arguments:

-h, --help	show this help message and exit
-v vowel, --vowel vowel	The vowel to substitute (default: a)

The program should complain if the --vowel argument is not a single, lowercase vowel:

```
$ ./apples.py -v x foo
usage: apples.py [-h] [-v str] str
apples.py: error: argument -v/--vowel: \
invalid choice: 'x' (choose from 'a', 'e', 'i', 'o', 'u')
```



Your program is going to need to do the following:

- Take a positional argument that might be some plain text or may name a file
- If the argument is a file, use the contents as the input text
- Take an optional -v or --vowel argument that should default to the letter “a”
- Verify that the --vowel option is in the set of vowels “a,” “e,” “i,” “o,” and “u”
- Replace all instances of vowels in the input text with the specified (or default) --vowel argument
- Print the new text to STDOUT

8.1 Altering strings

So far in our discussions of Python strings, numbers, lists, and dictionaries, we’ve seen how easily we can change or *mutate* variables. There is a problem, however, in that *strings are immutable*. Suppose we have a text variable that holds our input text:

```
>>> text = 'The quick brown fox jumps over the lazy dog.'
```

If we wanted to turn the first “e” (at index 2) into an “i,” we cannot do this:

```
>>> text[2] = 'i'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'str' object does not support item assignment
```

To change text, we need to set it equal to an entirely new value. In chapter 4 you saw that you can use a for loop to iterate over the characters in a string. For instance, I could laboriously uppercase text like so:

```

new = ''           ← Initialize a variable equal
for char in text: ← to the empty string.
    new += char.upper() ← Iterate through each
                        ← character in the text.
                        ← Append the uppercase version of
                        ← the character to the variable.

```

We can inspect the value of new to verify that it is all uppercase:

```
>>> new
'THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.'
```

Using this idea, you could iterate through the characters of text and build up a new string. Whenever the character is a vowel, you could change it for the given vowel; otherwise, you could use the character itself. We had to identify vowels in chapter 2, so you can refer back to how you did that.

8.1.1 Using the `str.replace()` method

In chapter 4 we talked about using the `str.replace()` method to replace all the numbers in a string with a different number. Maybe that would be a good way to solve this problem? Let's look at the documentation for that using `help(str.replace)` in the REPL:

```
>>> help(str.replace)
replace(self, old, new, count=-1, /)
    Return a copy with all occurrences of substring old replaced by new.

count
    Maximum number of occurrences to replace.
    -1 (the default value) means replace all occurrences.
```

If the optional argument count is given, only the first count occurrences are replaced.

Let's give that a try. We could replace "T" with "X":

```
>>> text.replace('T', 'X')
'Xhe quick brown fox jumps over the lazy dog.'
```

This seems promising! Can you see a way to replace all the vowels using this idea? Remember that this method never mutates the given string but instead returns a new string that you will need to assign to a variable.

8.1.2 Using `str.translate()`

We also looked at the `str.translate()` method in chapter 4. There we created a dictionary that described how to turn one character, like "1," into another string like "9." Any character not mentioned in the dictionary was left alone.

The documentation for this method is a bit more cryptic:

```
>>> help(str.translate)
translate(self, table, /)
    Replace each character in the string using the given translation table.

    table
        Translation table, which must be a mapping of Unicode ordinals to
        Unicode ordinals, strings, or None.

The table must implement lookup/indexing via __getitem__, for instance a
dictionary or list. If this operation raises LookupError, the character is
left untouched. Characters mapped to None are deleted.
```

In my solution, I created the following dictionary:

```
jumper = {'1': '9', '2': '8', '3': '7', '4': '6', '5': '0',
          '6': '4', '7': '3', '8': '2', '9': '1', '0': '5'}
```

That is the argument to the `str.maketrans()` function, which creates a translation table that is then used with `str.translate()` to change all the characters present as keys in the dictionary to their corresponding values:

```
>>> '876-5309'.translate(str.maketrans(jumper))
'234-0751'
```

What keys and values should you have in a dictionary if you want to change all the vowels, both lower- and uppercase, to some other value?

8.1.3 Other ways to mutate strings

If you know about regular expressions, that's a strong solution. If you haven't heard of them, don't worry—I'll introduce them in the discussion.

The point is for you to *play* with this and come up with a solution. I found eight ways to change all the vowels to a new character, so there are many ways you could approach this. How many *different* methods can you find on your own before you look at my solution?

Here are a few hints:

- Consider using the `choices` option in the `argparse` documentation to constrain the `--vowel` options. Be sure to read section A.4.3 in the appendix for an example.
- Be sure to change both lower- and uppercase versions of the vowels, preserving the case of the input characters.

Now is the time to dig in and see what you can do before you look at my solution.

8.2 Solution

Here is the first solution I wanted to share. After this, we'll explore several more.

```
#!/usr/bin/env python3
"""Apples and Bananas"""

import argparse
import os

# -----
def get_args():
    """get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Apples and bananas',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='text', help='Input text or file') ←

    parser.add_argument('-v',
                        '--vowel',
                        help='The vowel(s) allowed',
                        metavar='vowel',
                        type=str,
                        default='a',
                        choices=list('aeiou')) ←
    The input might be
    text or a filename, so I
    defined it as a string.

    args = parser.parse_args()

    if os.path.isfile(args.text):
        args.text = open(args.text).read().rstrip() ←
    If it is, read the file
    using str.rstrip() to
    remove any trailing
    whitespace.

    return args

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    text = args.text
    vowel = args.vowel
    new_text = [] ←
    Create a new
    list to hold the
    characters for the
    transformed text.

    for char in text:
        if char in 'aeiou':
            new_text.append(vowel)
        elif char in 'AEIOU':
            new_text.append(vowel.upper())
        else:
            new_text.append(char) ←
    Iterate through each
    character of the text.

    Check if the current
    character is in the list
    of lowercase vowels.

    If it is, use the vowel value
    instead of the character.

    Check if the current character is
    in the list of uppercase vowels.

    If it is, use the value of vowel.upper()
    instead of the character.

Otherwise,
use the
character
itself. ←
```

```
print(''.join(new_text))      ← Print a new string made by joining the
# -----  
if __name__ == '__main__':  
    main()
```

8.3 Discussion

I came up with eight ways to write my solution. All of them start with the same `get_args()` function, so let's look at that first.

8.3.1 Defining the parameters

This is one of those problems that has many valid and interesting solutions. The first problem to solve is, of course, getting and validating the user's input. As always, I will use `argparse`.

I usually define all my required parameters first. The `text` parameter is a positional string that *might* be a filename:

```
parser.add_argument('text', metavar='str', help='Input text or file')
```

The `--vowel` option is also a string, and I decided to use the `choices` option to have `argparse` validate that the user's input is in the `list('aeiou')`:

```
parser.add_argument('-v',
                    '--vowel',
                    help='The vowel to substitute',
                    metavar='str',
                    type=str,
                    default='a',
                    choices=list('aeiou'))
```

That is, `choices` wants a list of options. I could pass in `['a', 'e', 'i', 'o', 'u']`, but that's a lot of typing on my part. It's much easier to type `list('aeiou')` and have Python turn the str "aeiou" into a list of the characters. Both approaches produce the same results, because `list(str)` creates a list of the individual characters in a given string. And remember, the use of single or double quotes doesn't matter. Any value enclosed in either type of quotes is a `str`, even if it's just one character:

```
>>> ['a', 'e', 'i', 'o', 'u']
['a', 'e', 'i', 'o', 'u']
>>> list('aeiou')
['a', 'e', 'i', 'o', 'u']
```

We can even write a test for this. The absence of any error means that it's OK:

```
>>> assert ['a', 'e', 'i', 'o', 'u'] == list('aeiou')
```

The next task is detecting whether `text` is the name of a file that should be read for the text, or if it is the text itself. This is the same code I used in chapter 5, and again I

chose to handle the `text` argument inside the `get_args()` function so that, by the time I get `text` inside `main()`, it's all been handled. Figure 8.2 illustrates how we can chain the `open()` function to the `read()` method of a file handle to the `rstrip()` method of a string.

```
if os.path.isfile(args.text):
    args.text = open(args.text).read().rstrip()
```

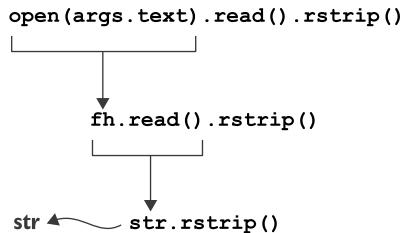


Figure 8.2 We can chain methods together to create pipelines of operations. The `open()` returns a file handle that we can read. The `read()` operation returns a string that we strip of whitespace.

At this point, the user's arguments to the program have been fully vetted. We've got `text` either from the command line or from a file, and we've verified that the `--vowel` value is one of the allowed characters. To me, this code is a single “unit” where I've handled the arguments. Processing can now go forward by returning the arguments:

```
return args
```

8.3.2 Eight ways to replace the vowels

How many ways did you find to replace the vowels? You only needed one, of course, to pass the tests, but I hope you probed the edges of the language to see how many different techniques there are. I know that the Zen of Python says

There should be one—and preferably only one—obvious way to do it.

www.python.org/dev/peps/pep-0020/

But I really come from the Perl mentality, where “There Is More Than One Way To Do It” (TIMTOWTDI or “Tim Toady”).

METHOD 1: ITERATING THROUGH EVERY CHARACTER

The first method is similar to what we did in chapter 4, where we used a `for` loop on a string to access each character. Here is some code you can copy and paste into the ipython REPL:

<p>Set the <code>new_text</code> variable to an empty list.</p> <pre>>>> new_text = []</pre>	<p>Set <code>text</code> to the string “Apples and Bananas!”</p> <pre>>>> text = "Apples and Bananas!"</pre>	<p>Set the <code>vowel</code> variable to the string “o”. That is, we'll replace all the vowels with this one.</p> <pre>>>> vowel = 'o'</pre>
<pre>>>> for char in text:</pre>		<p>Use a <code>for</code> to iterate <code>text</code>, putting each character into the <code>char</code> variable.</p>

```

...     if char in 'aeiou':
...         new_text.append(vowel)
...     elif char in 'AEIOU':
...         new_text.append(vowel.upper())
...     else:
...         new_text.append(char)
...
>>> text = ''.join(new_text)    ←
>>> text
'Opplos ond Bononos!'

```

If the character is in the set of lowercase vowels, add the vowel “o” to the new text.

If the character is in the set of uppercase vowels, substitute the `vowel.upper()` version “O” into the new text.

Otherwise, add the current character to the new text.

Turn the `new_text` list into a new str by joining it on the empty string ("").

Note that it would be just fine to start off making `new_text` an empty string and then concatenating the new characters. With that approach, you wouldn’t have to `str.join()` them at the end. Whatever you prefer:

```
new_text += vowel
```

Next I’m going to show you several alternate solutions. They’re all functionally equivalent because they all pass the tests—the point here is to explore the Python language and understand it. For the alternate solutions, I’ll just show the `main()` function.

METHOD 2: USING THE STR.REPLACE() METHOD

Here is a way to solve the problem using the `str.replace()` method:

```

def main():
    args = get_args()
    text = args.text
    vowel = args.vowel
    for v in 'aeiou':    ←
        text = text.replace(v, vowel).replace(v.upper(), vowel.upper()) ←
    print(text)

```

Iterate through the list of vowels. We don’t have to say `list('aeiou')` here—Python will automatically treat the string ‘aeiou’ like a list because we are using it in a list context with the `for` loop.

Use the `str.replace()` method twice to replace both the lower- and uppercase versions of the vowel in the text.

Earlier in the chapter, I mentioned the `str.replace()` method, which will return a new string with all instances of one string replaced by another:

```

>>> s = 'foo'
>>> s.replace('o', 'a')
'faa'
>>> s.replace('oo', 'x')
'fx'

```

Note that the original string remains unchanged:

```
>>> s
'foo'
```

You don’t have to chain the two `str.replace()` methods. It could be written as two separate statements, as illustrated in figure 8.3.

```

text = text.replace(v, vowel).replace(v.upper(), vowel.upper())
      [           ] [           ]
      ↓   ↓
text = text.replace(v, vowel)
      ↓
text = text.replace(v.upper(), vowel.upper())

```

Figure 8.3 The chained calls to `str.replace()` can be written as two separate statements if you prefer.

METHOD 3: USING THE STR.TRANSLATE() METHOD

Can we use the `str.translate()` method to solve this? I showed in chapter 4 how you could use a dictionary called `jumper` to change a character like “1” to the character “9.” In this problem, we need to change all the lower- and uppercase vowels (10 total) to some given vowel. For instance, to change all the vowels into the letter “o,” we could create a translation table `t` like so:

```
t = {'a': 'o',
     'e': 'o',
     'i': 'o',
     'o': 'o',
     'u': 'o',
     'A': 'O',
     'E': 'O',
     'I': 'O',
     'O': 'O',
     'U': 'O'}
```

We could use `t` with the `str.translate()` method:

```
>>> 'Apples and Bananas'.translate(str.maketrans(t))
'Opplos ond Bononos'
```

If you read the documentation for `str.maketrans()`, you will find that another way to specify the translation table is to supply two strings of equal lengths:

```
maketrans(x, y=None, z=None, /)
    Return a translation table usable for str.translate().
```

If there is only one argument, it must be a dictionary mapping Unicode ordinals (integers) or characters to Unicode ordinals, strings or None. Character keys will be then converted to ordinals.

If there are two arguments, they must be strings of equal length, and in the resulting dictionary, each character in `x` will be mapped to the character at the same position in `y`. If there is a third argument, it must be a string, whose characters will be mapped to None in the result.

The first string should contain the letters you want to replace, which are the lower- and uppercase vowels ‘aeiouAEIOU’. The second string is composed of the letters to use for substitution. We want to use ‘ooooo’ for ‘aeiou’ and ‘OOOOO’ for ‘AEIOU’.

We can repeat `vowel` five times using the `*` operator that you'll normally associate with numeric multiplication. This is (sort of) “multiplying” a string, so, OK, I guess:

```
>>> vowel * 5
'oooooo'
```

Next we handle the uppercase version:

```
>>> vowel * 5 + vowel.upper() * 5
'ooooooOOOOO'
```

And now we can make the translation table in one line of code like this:

```
>>> trans = str.maketrans('aeiouAEIOU', vowel * 5 + vowel.upper() * 5)
```

Let's inspect the `trans` table. We'll use the `pprint.pprint()` (pretty-print) function so we can read it easily:

```
>>> from pprint import pprint as pp
>>> pp(trans)
{65: 79,
 69: 79,
 73: 79,
 79: 79,
 85: 79,
 97: 111,
 101: 111,
 105: 111,
 111: 111,
 117: 111}
```

The enclosing curly braces {} tell us that `trans` is a dict. Each character is represented by its *ordinal value*, which is the character's position in the ASCII table (www.asciitable.com).

You can go back and forth between characters and their ordinal values by using the `chr()` and `ord()` functions. We will explore and use these functions later in chapter 18. Here are the `ord()` values for the vowels:

```
>>> for char in 'aeiou':
...     print(char, ord(char))
...
a 97
e 101
i 105
o 111
u 117
```

You can create the same output by starting with the `ord()` values to get the `chr()` values:

```
>>> for num in [97, 101, 105, 111, 117]:
...     print(chr(num), num)
...
a 97
e 101
```

```
i 105
o 111
u 117
>>>
```

If you'd like to inspect all the ordinal values for all the printable characters, you can run this:

```
>>> import string
>>> for char in string.printable:
...     print(char, ord(char))
```

I haven't included the output because there are 100 printable characters:

```
>>> print(len(string.printable))
100
```

So the `trans` table is a mapping from one character to another, just like in the "Jump the Five" exercise in chapter 4. The lowercase vowels ("aeiou") all map to the ordinal value 111, which is "o." The uppercase vowels ("AEIOU") map to 79, which is "O." You can use the `dict.items()` method to iterate over the key/value pairs of `trans` to verify that this is the case:

```
>>> for x, y in trans.items():
...     print(f'{chr(x)} => {chr(y)}')
...
a => o
e => o
i => o
o => o
u => o
A => O
E => O
I => O
O => O
U => O
```

The original text will be unchanged by the `str.translate()` method, so we can overwrite text with the new version. Here's how I wrote that idea in my solution:

```
def main():
    args = get_args()
    vowel = args.vowel
    trans = str.maketrans('aeiouAEIOU', vowel * 5 + vowel.upper() * 5)
    text = args.text.translate(trans)
    print(text)
```

Create a translation table from each of the vowels, both lower- and uppercase, to their respective characters. The lowercase vowels will be matched to the lowercase vowel argument, and the uppercase vowels will be matched to the uppercase vowel argument.

Call the `str.translate()` method on the `text` variable, passing the translation table as an argument.

That was a lot of explanation about `ord()` and `chr()` and dictionaries and such, but look how simple and elegant that solution is. This is much shorter than method 1. Fewer lines of code (LOC) means fewer opportunities for bugs!

METHOD 4: USING A LIST COMPREHENSION

Following up on method 1, we can use a *list comprehension* to significantly shorten the `for` loop. In chapter 7 we looked at a dictionary comprehension as a one-line method to create a new dictionary using a `for` loop. Here we can do the same, creating a new list:

```
def main():
    args = get_args()
    vowel = args.vowel
    text = [
        vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU' else c
        for c in args.text
    ]
    print(''.join(text))
```

The code is annotated with three callouts:

- Use a list comprehension to process all the characters in args.text to create a new list called text.** Points to the first line of the list comprehension: `text = [`.
- Print the translated string by joining the text list on the empty string.** Points to the final line: `print(''.join(text))`.
- Use a compound if expression to handle three cases: lowercase vowel, uppercase vowel, and the default.** Points to the middle part of the list comprehension: `vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU' else c`.

Let's talk just a bit more about list comprehensions. As an example, we can generate a list of the squared values of the numbers 1 through 4 by using the `range()` function to get the numbers from a starting number to an ending number (not inclusive). In the REPL, we must use the `list()` function to force the production of the values, but usually your code won't need to do this:

```
>>> list(range(1, 5))
[1, 2, 3, 4]
```

NOTE `range()` is another example of a *lazy* function in Python, which means it won't actually produce values until your program needs them—a lazy function is a promise to do something. If your program branches in such a way that you never need to produce the values, the work is never done, meaning your code is more efficient.

We can write a `for` loop to `print()` the squares:

```
>>> for num in range(1, 5):
...     print(num ** 2)
...
1
4
9
16
```

Instead of printing the values, imagine that we wanted to create a new list that contains those values. One way to do this would be to create an empty list and then use `list.append()` to add each value in a `for` loop:

```
>>> squares = []
>>> for num in range(1, 5):
...     squares.append(num ** 2)
```

Now we can verify that we have our squares:

```
>>> assert len(squares) == 4
>>> assert squares == [1, 4, 9, 16]
```

We can achieve the same result in fewer lines of code using a list comprehension to generate our new list, as shown in figure 8.4.

```
>>> [num ** 2 for num in range(1, 5)]
[1, 4, 9, 16]
```

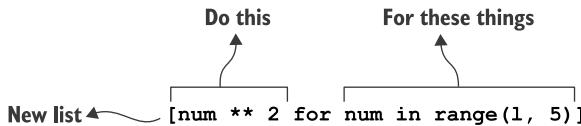


Figure 8.4 A list comprehension creates a new list using a `for` loop to iterate over the source values.

We can assign this list to the variable `squares` and verify that we still have what we expected. Ask yourself which version of the code you'd rather maintain: the longer one with the `for` loop, or the shorter one with the list comprehension?

```
>>> squares = [num ** 2 for num in range(1, 5)]
>>> assert len(squares) == 4
>>> assert squares == [1, 4, 9, 16]
```

For this version of the program, we'll condense the `if/elif/else` logic from method 1 into a compound `if` expression. First let's see how we could shorten the `for` loop version:

```
>>> text = 'Apples and Bananas!'
>>> new = []
>>> for c in text:
...     new.append(vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU'
...             else c)
...
>>> ''.join(new)
'Opplos ond Bononos!'
```

Figure 8.5 shows how the parts of the expression match up to the original `if/elif/else`:

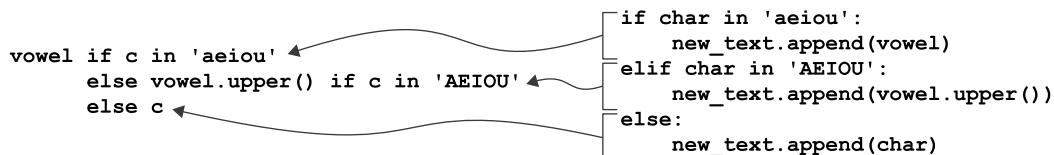


Figure 8.5 The three conditional branches can be written using two `if` expressions.

Now let's turn that into a list comprehension:

```
>>> text = 'Apples and Bananas!'
>>> new_text = [
...     vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU' else c
...     for c in text ]
...
>>> ''.join(new_text)
'Opplos ond Bononos!'
```

Select the character using the compound if expression.

Perform this action for each character in the text.

The code is denser than the previous for loop, but it has advantages in that

- The list comprehension is shorter and generates our list rather than using the side effects of `list.append()`.
- The compound if expression will not compile if we forget one of the conditional branches.

METHOD 5: USING A LIST COMPREHENSION WITH A FUNCTION

The compound if expression inside the list comprehension is complicated enough that it probably should be a function. We can *define* a new function with the `def` statement and call it `new_char()`. It accepts a character we'll call `c`. After that, we can use the same compound if expression as before:

```
def main():
    args = get_args()
    vowel = args.vowel

    def new_char(c):
        return vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU' else c

    text = ''.join([new_char(c) for c in args.text])
    print(text)
```

Define a function to choose a new character. Note that it uses the `vowel` variable because the function has been declared in the same scope. This is called a closure, because `new_char()` closes over the variable.

Use a list comprehension to process all the characters in text.

Use the compound if expression to select the correct character.

You can play with the `new_char()` function by putting this into your REPL:

```
vowel = 'o'
def new_char(c):
    return vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU' else c
```

It should always return the letter “o” if the argument is a lowercase vowel:

```
>>> new_char('a')
'o'
```

It should return “O” if the argument is an uppercase vowel:

```
>>> new_char('A')
'O'
```

Otherwise, it should return the given character:

```
>>> new_char('b')
'b'
```

We can use the `new_char()` function to process all the characters in `text`, using a list comprehension:

```
>>> text = 'Apples and Bananas!'
>>> text = ''.join([new_char(c) for c in text])
>>> text
'Opplos ond Bononos!'
```

Note that the `new_char()` function is declared *inside* the `main()` function. Yes, you can do that! The function is then only “visible” inside the `main()` function. I’ve done this because we want to reference the `vowel` variable inside the function without passing it as an argument.

As an example, let’s define a `foo()` function that has a `bar()` function inside it. We can call `foo()`, and it will call `bar()`. But from outside of `foo()`, the `bar()` function does not exist (it “is not visible” or “is not in scope”).

```
>>> def foo():
...     def bar():
...         print('This is bar')
...     bar()
...
>>> foo()
This is bar
>>> bar()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'bar' is not defined
```

I declared the `new_char()` function inside `main()` because I wanted to reference the `vowel` variable inside the function, as shown in figure 8.6. Because `new_char()` “closes” around the `vowel`, it is a special type of function called a *closure*.

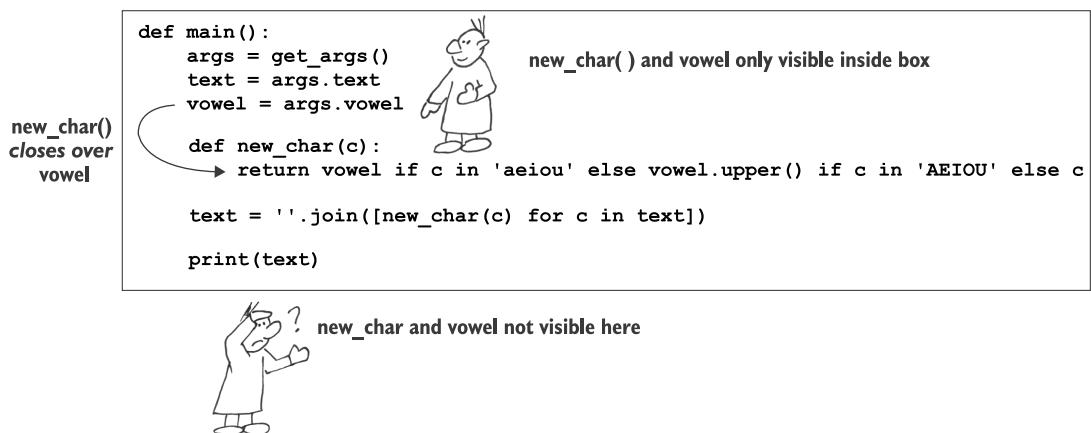


Figure 8.6 The `new_char()` function can only be seen within the `main()` function. It creates a closure because it references the `vowel` variable. Code outside of `main()` cannot see or call `new_char()`.

If we don't write this as a closure, we will have to pass the vowel as an argument:

```
def main():
    args = get_args()
    print(''.join([new_char(c, args.vowel) for c in args.text]))
```

```
→ def new_char(char, vowel):
    return vowel if char in 'aeiou' else \
        vowel.upper() if char in 'AEIOU' else char
```

We need to pass
args.vowel as an
argument to the
new_char()
function.

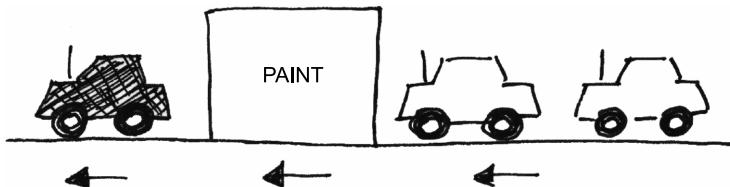
The vowel is only visible inside the main() function. Since new_char() is no longer declared in the same scope, we need to accept vowel as an argument.

While the closure method is interesting, this version is arguably easier to understand. It would also be easier to write a unit test for it, which is something we'll start doing soon.

METHOD 6: USING THE MAP() FUNCTION

For this method, I'll introduce the `map()` function, as it's quite similar to a list comprehension. The `map()` function accepts two arguments:

- A function
- An iterable like a list, a lazy function, or a generator



I like to think of `map()` like a paint booth—you load up the booth with, say, blue paint. Unpainted cars go in, blue paint is applied, and blue cars come out.

We can create a function to “paint” cars by adding the string “blue” to the beginning:

```
>>> list(map(lambda car: 'blue ' + car, ['BMW', 'Alfa Romeo', 'Chrysler']))
['blue BMW', 'blue Alfa Romeo', 'blue Chrysler']
```

The first argument you see here starts with the keyword `lambda`, which is used to create an *anonymous* function. With the regular `def` keyword, the function name follows. With `lambda`, there is no name, only the list of parameters and the function body.

For example, an `add1()` function that adds 1 to a value is a regular named function:

```
def add1(n):
    return n + 1
```



It works as expected:

```
>>> assert add1(10) == 11
>>> assert add1(add1(10)) == 12
```

Compare the preceding definition to one created using `lambda`, which we assign to the variable `add1`:

```
>>> add1 = lambda n: n + 1
```

This definition of `add1` is functionally equivalent to the first version. We call it just like the `add1()` function:

```
>>> assert add1(10) == 11
>>> assert add1(add1(10)) == 12
```

The body for a `lambda` is a brief (usually one-line) expression. There is no `return` statement because the final evaluation of the expression is returned automatically. In figure 8.7, you can see that the `lambda` will return the result of $n + 1$.

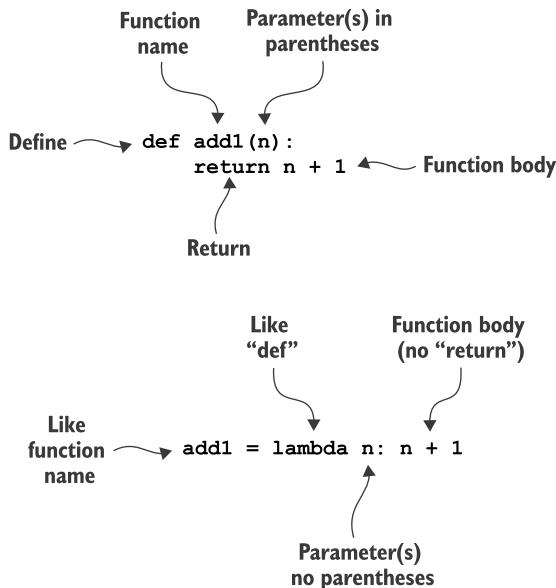


Figure 8.7 Both `def` and `lambda` are used to create functions.

In both versions of the `add1` definition, using `def` and `lambda`, the argument to the function is `n`. In the usual named function, `def add(n)`, the argument is defined in the parentheses just after the function name. In the `lambda n` version, there is no function name and no parentheses around the function's parameter, `n`.

There is no difference in how you can use the two types of functions. They are both functions:

```
>>> type(lambda x: x)
<class 'function'>
```

If you are comfortable with using `add1()` in a list comprehension, like this,

```
>>> [add1(n) for n in [1, 2, 3]]
[2, 3, 4]
```

it's a short step to using the `map()` function.

The `map()` function is a lazy function, like the `range()` function we looked at earlier. It won't create the values until you actually need them, as compared to a list comprehension, which will produce the resulting list immediately. I don't personally tend to worry about the performance of the code as much as I do the readability. When I write code for myself, I prefer to use `map()`, but you should write code that makes the most sense for you and your teammates.

To force the results from `map()` in the REPL, we need to use the `list()` function:

```
>>> list(map(add1, [1, 2, 3]))
[2, 3, 4]
```

We can write the list comprehension with the `add1()` code in line:

```
>>> [n + 1 for n in [1, 2, 3]]
[2, 3, 4]
```

That looks very similar to the `lambda` code (as illustrated in figure 8.8):

```
>>> list(map(lambda n: n + 1, [1, 2, 3]))
[2, 3, 4]
```

Here is how we could use `map()`:

```
def main():
    args = get_args()
    vowel = args.vowel
    text = map(
        lambda c: vowel if c in 'aeiou' else vowel.upper() ←
        if c in 'AEIOU' else c, args.text) ←
    print(''.join(text)) ←
```

The `map()` function wants a function for the first argument and an iterable for the second.

Use lambda to create an anonymous function that accepts a character, `c`.

`args.text` is the second argument to `map()`. Technically, `args.text` is a string, but, because `map()` expects this argument to be a list, the string will be coerced to a list.

`map()` returns a new list to the `text` variable. We join it on the empty string to print it.

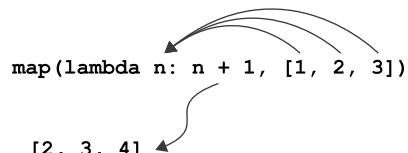


Figure 8.8 The `map()` function will create a new list from processing each element of an iterable through a given function.

Higher-order functions

The `map()` function is called a *higher-order function* (HOF) because it takes another function as an argument, which is wicked cool. Later we'll use another HOF called `filter()`.

METHOD 7: USING MAP() WITH A NAMED FUNCTION

We are not required to use `map()` with a lambda expression. Any function at all will work, so let's go back to using our `new_char()` function:

```
def main():
    args = get_args()
    vowel = args.vowel
    def new_char(c):
        return vowel if c in 'aeiou' else vowel.upper() if c in 'AEIOU' else c
    print(''.join(map(new_char, args.text)))
```

Define a function that will return the proper character. Note that I'm using the closure version so as to reference the "vowel" argument.

Use map() to apply new_char() to all the characters in args.text. The result is a list of characters, and we can use str.join() to turn them into a new string for print().

Notice that `map()` uses `new_char` *without parentheses* as the first argument. If you added the parentheses, you'd be *calling* the function and would see this error:

```
>>> text = ''.join(map(new_char(), text))
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: new_char() missing 1 required positional argument: 'c'
```

As shown in figure 8.9, `map()` takes each character from `text` and passes it as the argument to the `new_char()` function, which decides whether to return a vowel or the original character. The result of mapping these characters is a new list of characters that we `str.join()` on the empty string to create a new version of `text`.

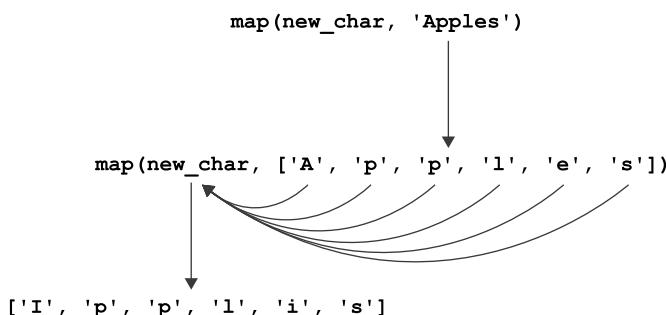


Figure 8.9 `map()` will apply a given function to each element of an iterable. A string will be processed as a list of characters.

METHOD 8: USING REGULAR EXPRESSIONS

A *regular expression* is a way to describe patterns of text. Regular expressions (also called “*regexes*”) are a separate domain-specific language (DSL). They really have nothing whatsoever to do with Python. They have their own syntax and rules, and they are used in many places, from command-line tools to databases. Regexes are incredibly powerful and well worth the effort to learn.

To use regular expressions, you must `import re` in your code to import the regular expression module:

```
>>> import re
```

In this example, we’re trying to find characters that are vowels, which we can define as the letters “a,” “e,” “i,” “o,” and “u.” To describe this idea using a regular expression, we put those characters inside square brackets:

```
>>> pattern = '[aeiou]'
```

We can use the “*substitute*” function, `re.sub()`, to find all the vowels and replace them with the given vowel. The square brackets around the vowels '`[aeiou]`' create a *character class*, meaning anything matching one of the characters listed inside the brackets.

The second argument is the string that will replace the found strings—here it is the vowel provided by the user. The third argument is the string we want to change, which is the `text` from the user:

```
>>> vowel = 'o'
>>> re.sub(pattern, vowel, 'Apples and bananas!')
'Applos ond bononos!'
```

That misses the capital “A,” so we’ll have to handle both lower- and uppercase. Here is how we could write that:

```
def main():
    args = get_args()
    text = args.text
    vowel = args.vowel
    text = re.sub('[aeiou]', vowel, text)           ← Substitute any of the lowercase vowels with the given vowel (which is lowercase because of the restrictions in get_args()).
    text = re.sub('[AEIOU]', vowel.upper(), text)   ← Substitute any of the uppercase vowels with the uppercase vowel.
    print(text)
```

If you prefer, we could squash the two calls to `re.sub()` into one, just as we did with the `str.replace()` method shown earlier:

```
>>> text = 'Apples and Bananas!'
>>> text = re.sub('[AEIOU]', vowel.upper(), re.sub('[aeiou]', vowel, text))
>>> text
'Opplos ond Bononos!'
```

One of the biggest differences between this solution and all the others is that we use regular expressions to describe what we are looking for. We didn’t have to write the

code to identify the vowels. This is more along the lines of *declarative* programming. We declare what we want, and the computer does the grunt work!

8.4 Refactoring with tests

There are many ways to solve this problem. The most important step is to get your program to work properly. Tests let you know when you've reached that point. From there, you can explore other ways to solve the problem and keep using the tests to ensure you still have a correct program.

Tests provide you with great freedom to be creative. Always be thinking about tests you can write for your own programs, so that when you change them later, they will always keep working.

I showed many ways to solve this seemingly trivial problem. Some of the techniques using higher-order functions and regular expression are quite advanced techniques. It might seem like driving a finishing nail with a sledgehammer, but I want to start introducing you to programming ideas that I'll visit again and again in later chapters.

If you only really understood the first few solutions, that's fine! Just stick with me. The more times you see these ideas applied in different contexts, the more they will begin to make sense.

8.5 Going further

Write a version of the program that collapses multiple adjacent vowels into a single substituted value. For example, “quick” should become “qack” and not “qaack.”

Summary

- You can use argparse to limit an argument's values to a list of choices that you define.
- Strings cannot be directly modified, but the str.replace() and str.translate() methods can create a *new, modified string* from an existing string.
- A for loop on a string will iterate the characters of the string.
- A list comprehension is a shorthand way to write a for loop inside [] to create a new list.
- Functions can be defined inside other functions. Their visibility is then limited to the enclosing function.
- Functions can reference variables declared within the same scope, creating a closure.
- The map() function is similar to a list comprehension. It will create a new, modified list by applying some function to every member of a given list. The original list will not be changed.
- Regular expressions provide a syntax for describing patterns of text with the re module. The re.sub() method will substitute found patterns with new text. The original text will be unchanged.

Dial-a-Curse: Generating random insults from lists of words

“He or she is a slimy-sided, frog-mouthed, silt-eating slug with the brains of a turtle.”

—Dial-A-Curse

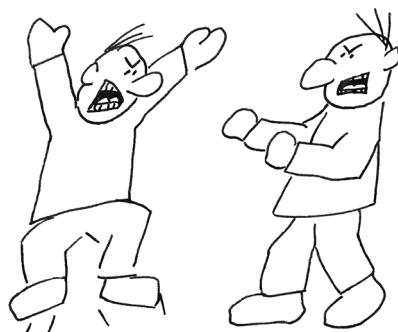
Random events are at the heart of interesting games and puzzles. Humans quickly grow bored of things that are always the same. I think one reason people may choose to have pets and children is to inject some randomness into their lives. Let’s learn how to make our programs more interesting by having them behave differently each time they are run.

This exercise will show you how to randomly select one or more elements from lists of options. To explore randomness, we’ll create a program called `abuse.py` that will insult the user by randomly selecting adjectives and nouns to create slanderous epithets.

In order to test randomness, though, we need to control it. It turns out that “random” events on computers are rarely truly random but only *pseudo-random*, which means we can control them by using a “seed.”¹ Each time you use the same seed, you will get the same “random” choices!

Shakespeare had some of the best insults, so we’ll draw from the vocabulary of his works. Here is the list of adjectives you should use:

bankrupt base caterwauling corrupt cullionly detestable dishonest false filthsome filthy foolish foul gross heedless indistinguishable infected insatiate irksome lascivious



¹ “The generation of random numbers is too important to be left to chance.”—Robert R. Coveyou

*lecherous loathsome lubbery old peevish rascaly rotten ruinous scurilous scurvy slanderous
sodden-witted thin-faced toad-spotted unmannered vile wall-eyed*

And these are the nouns:

*Judas Satan ape ass barbermanger beggar block boy braggart butt carbuncle coward
coxcomb cur dandy degenerate fiend fishmonger fool gull harpy jack jolthead knave liar
lunatic maw milksoop minion ratcatcher recreant rogue scold slave swine traitor varlet
villain worm*

For instance, it might produce the following:

```
$ ./abuse.py
You slanderous, rotten block!
You lubbery, scurilous ratcatcher!
You rotten, foul liar!
```

In this exercise, you will learn to

- Use `parser.error()` from `argparse` to throw errors
- Control randomness with random seeds
- Take random choices and samples from Python lists
- Iterate an algorithm a specified number of times with a `for` loop
- Format output strings

9.1 Writing abuse.py

You should go into the `09_abuse` directory to create your new program. Let's start by looking at the usage statement it should produce:

```
$ ./abuse.py -h
usage: abuse.py [-h] [-a adjectives] [-n insults] [-s seed]

Heap abuse

optional arguments:
  -h, --help            show this help message and exit
  -a adjectives, --adjectives adjectives
                        Number of adjectives (default: 2)
  -n insults, --number insults
                        Number of insults (default: 3)
  -s seed, --seed seed  Random seed (default: None)
```

All parameters are options that have default values, so our program will be able to run with no arguments at all.

For instance, the `-n` or `--number` option will have a default of 3 and will control the number of insults:

```
$ ./abuse.py --number 2
You filthsome, cullionly fiend!
You false, thin-faced minion!
```

The `-a` or `--adjectives` option should default to 2 and will determine how many adjectives are used in each insult:

```
$ ./abuse.py --adjectives 3
You caterwauling, heedless, gross coxcomb!
You sodden-witted, rascaly, lascivious varlet!
You dishonest, lecherous, foolish varlet!
```

Lastly, the `-s` or `--seed` option will control the random choices in the program by setting an initial value. The default should be the special `None` value, which is like an undefined value.

Because the program will use a random seed, the following output should be exactly reproducible by any user on any machine at any time:

```
$ ./abuse.py --seed 1
You filthsome, cullionly fiend!
You false, thin-faced minion!
You sodden-witted, rascaly cur!
```

When run with no arguments, the program should generate insults using the defaults:

```
$ ./abuse.py
You foul, false varlet!
You filthy, insatiate fool!
You lascivious, corrupt recreant!
```



I recommend you start by copying the template/template.py file to abuse/abuse.py or by using new.py to create the abuse.py program in the 09_abuse directory of your repository.

Figure 9.1 is a string diagram that illustrates the parameters for the program.

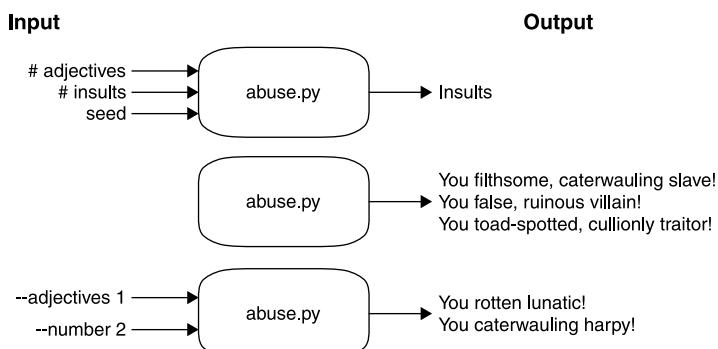


Figure 9.1 The abuse.py program will accept options for the number of insults to create, the number of adjectives per insult, and a random seed value.

9.1.1 Validating arguments

The options for the number of insults and adjectives and the random seed should all be `int` values. If you define each using `type=int` (remember, there are no quotes around the `int`), `argparse` will handle the validation and conversion of the arguments to `int` values for you. That is, just by defining `type=int`, the following error will be generated for you if a string is entered:

```
$ ./abuse.py -n foo
usage: abuse.py [-h] [-a adjectives] [-n insults] [-s seed]
abuse.py: error: argument -n/--number: invalid int value: 'foo'
```

Not only must the value be a number, but it must be an *integer* which means it must be a whole number, so `argparse` will complain if you give it something that looks like a float. Note that you can use `type=float` when you actually want a floating-point value:

```
$ ./abuse.py -a 2.1
usage: abuse.py [-h] [-a adjectives] [-n insults] [-s seed]
abuse.py: error: argument -a/--adjectives: invalid int value: '2.1'
```

Additionally, if either `--number` or `--adjectives` is less than 1, your program should exit with an error code and message:

```
$ ./abuse.py -a -4
usage: abuse.py [-h] [-a adjectives] [-n insults] [-s seed]
abuse.py: error: --adjectives "-4" must be > 0
$ ./abuse.py -n -4
usage: abuse.py [-h] [-a adjectives] [-n insults] [-s seed]
abuse.py: error: --number "-4" must be > 0
```

As you start to write your own programs and tests, I recommend you steal from the tests I've written.² Let's take a look at one of the tests in `test.py` to see how the program is tested:

Run the program using `getstatusoutput()` from the `subprocess`³ module using a bad `-a` value. This function returns the exit value (which I put into `rv` for "return value") and standard out (`out`).

```
def test_bad_adjective_num():
    """bad_adjectives"""

    n = random.choice(range(-10, 0))
    rv, out = getstatusoutput(f'{prog} -a {n}')
    assert rv != 0
    assert re.search(f'--adjectives "\{n\}" must be > 0', out)
```

The name of the function must start with "test_" in order for Pytest to find and run it.

Use the `random.choice()` function to randomly select a value from the `range()` of numbers from -10 to 0. We will use this same function in our program, so note here how it is called.

Assert that the return value (`rv`) is not 0, where "0" would indicate success (or "zero errors").

Assert that the output somewhere contains the statement that the `--adjectives` argument must be greater than 0.

² "Good composers borrow. Great ones steal." — Igor Stravinsky

³ The `subprocess` module allows you to run a command from inside your program. The `subprocess.getoutput()` function will capture the output from the command, while the `subprocess.getstatusoutput()` will capture both the exit value and the output from the command.

There's no simple way to tell argparse that the numbers for adjectives and insults must be greater than zero, so we'll have to check those values ourselves. We'll use the verification ideas from section A.4.7 in the appendix. There I introduce the `parser.error()` function, which you can call inside the `get_args()` function to do the following:

- 1 Print the short usage statement
- 2 Print an error message to the user
- 3 Stop execution of the program
- 4 Exit with a nonzero exit value to indicate an error

That is, `get_args()` normally finishes with this:

```
return args.parse_args()
```

Instead, we'll put the `args` into a variable and check the `args.adjectives` value to see if it's less than 1. If it is, we'll call `parser.error()` with an error message to report to the user:

```
args = parser.parse_args()

if args.adjectives < 1:
    parser.error(f'--adjectives "{args.adjectives}" must be > 0')
```

We'll also do this for `args.number`. If they are both fine, you can return the arguments to the calling function:

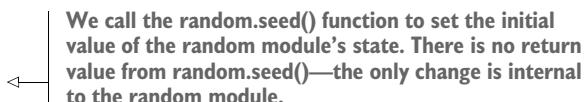
```
return args
```

9.1.2 Importing and seeding the random module

Once you have defined and validated all the program's arguments, you are ready to heap scorn upon the user. First, we need to add `import random` to our program so we can use functions from that module to select adjectives and nouns. It's best practice to list all the `import` statements, one module at a time, at the top of a program.

In `main()`, the first thing we need to do is call `get_args()` to get our arguments. The next step is to pass the `args.seed` value to the `random.seed()` function:

```
def main():
    args = get_args()
    random.seed(args.seed)
```



We call the `random.seed()` function to set the initial value of the `random` module's state. There is no return value from `random.seed()`—the only change is internal to the `random` module.

You can read about the `random.seed()` function in the REPL:

```
>>> import random
>>> help(random.seed)
```

There you'll learn that the function will “initialize internal state [of the `random` module] from hashable object.” That is, we set an initial value from some *hashable* Python type. Both the `int` and `str` types are hashable, but the tests are written with

the expectation that you will define the `seed` argument as an `int`. (Remember that the character '`1`' is different from the *integer value 1*!)

The default value for `args.seed` should be `None`. If the user has not indicated any seed, then setting `random.seed(None)` is the same as not setting it at all.

If you look at the `test.py` program, you will notice that all the tests that expect a particular output will pass an `-s` or `--seed` argument. Here is the first test for output:

```
Run the program using getoutput() from the subprocess module
using a seed value of 1 and requesting 1 insult. This function
returns only the output from the program.

def test_01():
    out = getoutput(f'{prg} -s 1 -n 1')
    assert out.strip() == 'You filthsome, cullionly fiend!'
```

Verify that the entire output is the one expected insult.

This means `test.py` will run your program and capture the output into the `out` variable:

```
$ ./abuse.py -s 1 -n 1
You filthsome, cullionly fiend!
```

It will then verify that the program did in fact produce the expected number of insults with the expected selection of words.

9.1.3 Defining the adjectives and nouns

Earlier in the chapter, I gave you a long list of adjectives and nouns that you should use in your program. You could create a list by individually quoting each word:

```
>>> adjectives = ['bankrupt', 'base', 'caterwauling']
```

Or you could save yourself a good bit of typing by using the `str.split()` method to create a new list from a `str` by splitting on spaces:

```
>>> adjectives = 'bankrupt base caterwauling'.split()
>>> adjectives
['bankrupt', 'base', 'caterwauling']
```

If you try to make one giant string of all the adjectives, it will be very long and so will wrap around in your code editor and look ugly. I recommend you use triple quotes (either single or double quotes), which will allow you to include newlines:

```
>>> """
... bankrupt base
... caterwauling
... """ .split()
['bankrupt', 'base', 'caterwauling']
```

Once you have variables for adjectives and nouns, you should check that you have the right number of each:

```
>>> assert len(adjectives) == 36
>>> assert len(nouns) == 39
```

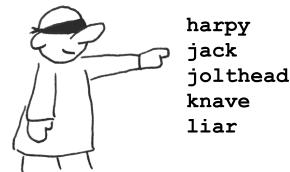
NOTE In order to pass the tests, your adjectives and nouns must be in alphabetical order as they were provided.

9.1.4 Taking random samples and choices

In addition to the `random.seed()` function, we will also use the `random.choice()` and `random.sample()` functions. In the `test_bad_adjective_num` function in section 9.1.1, you saw one example of using `random.choice()`. We can use it similarly to select a noun from the list of nouns.

Notice that this function returns a single item, so, given a list of `str` values, it will return a single `str`:

```
>>> random.choice(nouns)
'braggart'
>>> random.choice(nouns)
'milksop'
```



For the adjectives, you should use `random.sample()`. If you read the `help(random.sample)` output, you will see that this function takes some list of items and a `k` parameter for how many items to return:

```
sample(population, k) method of random.Random instance
    Chooses k unique random elements from a population sequence or set.
```

Note that this function returns a new list:

```
>>> random.sample(adjectives, 2)
['detestable', 'peevish']
>>> random.sample(adjectives, 3)
['slanderous', 'detestable', 'base']
```

There is also a `random.choices()` function that works similarly but which might select the same items twice because it samples “with replacement.” We will not use that.

9.1.5 Formatting the output

The output of the program is a --number of insults, which you could generate using a `for` loop and the `range()` function. It doesn’t matter here that `range()` starts at zero. What’s important is that it generates three values:

```
>>> for n in range(3):
...     print(n)
...
0
1
2
```

You can loop the --number of times needed, select your sample of adjectives and your noun, and then format the output. Each insult should start with the string “You”, then

have the adjectives joined on a comma and a space, then the noun, and finish with an exclamation point (figure 9.2). You could use either an f-string or the `str.format()` function to print() the output to STDOUT.

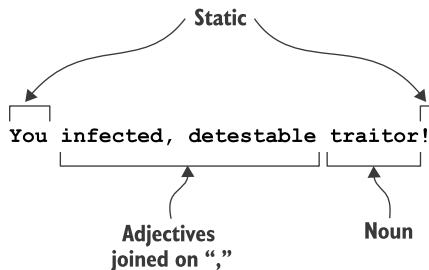


Figure 9.2 Each insult will combine the chosen adjectives joined on commas with the selected noun and some static bits of text.

Here are a few hints:

- Perform the check for positive values for --adjectives and --number *inside* the `get_args()` function, and use `parser.error()` to throw the error while printing a message and the usage.
- If you set the default of `args.seed` to `None` and use `type=int`, you can directly pass the value to `random.seed()`. When the value is `None`, it will be like not setting the value at all.
- Use a `for` loop with the `range()` function to create a loop that will execute the --number of times to generate each insult.
- Look at the `random.sample()` and `random.choice()` functions for help in selecting some adjectives and a noun.
- You can use three single quotes ('''') or double quotes ("""") to create a multi-line string and then use `str.split()` to get a list of strings. This is easier than individually quoting a long list of shorter strings (such as the list of adjectives and nouns).
- To construct an insult to print, you can use the `+` operator to concatenate strings, use the `str.join()` method, or use format strings.

Now give this your best shot before reading ahead to the solution, you snotty-faced heap of parrot droppings!

9.2 Solution

This is the first solution where I use `parser.error()` to augment the validation of the arguments. I also incorporate triple-quoted strings and introduce the `random` module, which is quite fun unless you're a vacuous, coffee-nosed, malodorous git.

```
#!/usr/bin/env python3
"""Heap abuse"""
import argparse
import random
```

Bring in the random module so we can call functions.

```

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Heap abuse',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('-a',
                        '--adjectives',
                        help='Number of adjectives',
                        metavar='adjectives',
                        type=int,
                        default=2)

    parser.add_argument('-n',
                        '--number',
                        help='Number of adjectives',
                        metavar='adjectives',
                        type=int,
                        default=3)

    parser.add_argument('-s',
                        '--seed',
                        help='Random seed',
                        metavar='seed',
                        type=int,
                        default=None)

    args = parser.parse_args()

    if args.adjectives < 1:
        parser.error('--adjectives "{}" must be > 0'.format(args.adjectives))

    if args.number < 1:
        parser.error('--number "{}" must be > 0'.format(args.number))

    return args

```

The random seed default should be None.

Similarly check args.number.

Define the parameter for the number of adjectives, setting type=int and the default value.

Similarly define the parameter for the number of insults as an integer with a default.

Get the result of parsing the command-line arguments. The argparse module will handle errors such as non-integer values.

Check that args.adjectives is greater than 0. If there is a problem, call parser.error() with the error message.

At this point, all the user's arguments have been validated, so return the arguments to the caller.

```

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    random.seed(args.seed)

    adjectives = """
bankrupt base caterwauling corrupt cullionly detestable dishonest false
filthsome filthy foolish foul gross heedless indistinguishable infected

```

Set random.seed() using whatever value was passed by the user. Any integer value is valid, and I know that argparse has handled the validation and conversion of the argument to an integer.

Create a list of adjectives by splitting the very long string contained in the triple quotes.

```

insatiate irksome lascivious lecherous loathsome lubbery old peeves
rascally rotten ruinous scurilous scurvy slanderous sodden-witted
thin-faced toad-spotted unmannered vile wall-eyed
""".strip().split()

nouns = """
Judas Satan ape ass barbarmonger beggar block boy braggart butt
carbuncle coward coxcomb cur dandy degenerate fiend fishmonger fool
gull harpy jack jolthead knave liar lunatic maw milksop minion
ratcatcher recreant rogue scold slave swine traitor varlet villain worm
""".strip().split()

for _ in range(args.number):
    adjs = ', '.join(random.sample(adjectives, k=args.adjectives))
    print(f'You {adjs} {random.choice(nouns)}!')

# -----
if __name__ == '__main__':
    main()

```

Do the same for the list of nouns.

Use an f-string to format the output to print().

Use the random.sample() function to select the correct number of adjectives and join them on the comma-space string.

Use a for loop over the range() of the args.number. Since I don't actually need the value from range(), I can use the _ to disregard it.

9.3 Discussion

I trust you did not peek at the solution before you passed all the tests or else you are a rascally, filthsome swine.

9.3.1 Defining the arguments

More than half of my solution is defining the program's arguments to argparse. The effort is well worth the result. Because I set type=int, argparse will ensure that each argument is a valid integer value. Notice that there are no quotes around the int—it's not the string 'int' but a reference to the class in Python:

```

parser.add_argument('-a',      ← The short flag
                    '--adjectives',
                    help='Number of adjectives',
                    metavar='adjectives',
                    type=int,
                    default=2)           ← A description of the parameter

```

The actual Python type for converting the input; note that this is the bare word int for the integer class

The default value for the number of adjectives per insult

The long flag

The help message

I set reasonable defaults for all the program's options so that no input is required from the user. The --seed option should default to None so that the default behavior is to generate pseudo-random insults. This value is only important for testing purposes.

9.3.2 Using `parser.error()`

I really love the argparse module for all the work it saves me. In particular, I often use `parser.error()` when I find there is a problem with an argument. This function will do four things:

- 1 Print the short usage of the program to the user
- 2 Print a specific message about the problem
- 3 Halt execution of the program
- 4 Return an error code to the operating system

I'm using `parser.error()` here because, while I can ask argparse to verify that a given value is an int, I can't as easily say that it must be a *positive* value. I can, however, inspect the value myself and halt the program if there is a problem. I do all this inside `get_args()` so that, by the time I get the args in my `main()` function, I know they have been validated.

I highly recommend you tuck this tip into your back pocket. It can prove quite handy, saving you loads of time validating user input and generating useful error messages. (And it's quite likely that the future user of your program will be *you*, so you will really appreciate your efforts.)

9.3.3 Program exit values and `STDERR`

I would like to highlight the exit value of the program. Under normal circumstances, programs should exit with a value of 0. In computer science, we often think of 0 as a False value, but here it's quite positive. In this instance we should think of it as "zero errors."

If you use `sys.exit()` in your code to exit a program prematurely, the default exit value is 0. If you want to indicate to the operating system or some calling program that your program exited with an error, you should return *any value other than 0*. You can also call the function with a string, which will be printed as an error message, and Python will exit with the value 1. If you run this in the REPL, you will be returned to the command line:

```
>>> import sys  
>>> sys.exit('You gross, thin-faced worm!')  
You gross, thin-faced worm!
```

Additionally, it's common for all error messages to be printed not to STDOUT (standard out) but to STDERR (standard error). Many command shells (like Bash) can segregate these two output channels using 1 for STDOUT and 2 for STDERR. When using the Bash shell, note how I can use 2> to redirect STDERR to the file called err so that nothing appears on STDOUT:

```
$ ./abuse.py -a -1 2>err
```

I can verify that the expected error messages are in the err file:

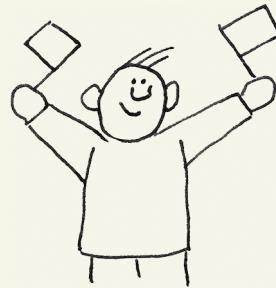
```
$ cat err
usage: abuse.py [-h] [-a adjectives] [-n insults] [-s seed]
abuse.py: error: --adjectives "-1" must be > 0
```

If you were to handle all of this yourself, you would need to write something like this:

```
if args.adjectives < 1:
    parser.print_usage()           ← Print the short usage. You can also
                                    use parser.print_help() to print the
                                    more verbose output for -h.
    print(f"--adjectives '{args.adjectives}' must be > 0", file=sys.stderr) ← Print the error message to the sys.stderr file
    sys.exit(1)                   ← Exit the program with a value that is not 0 to indicate an error.
```

Writing pipelines

As you write more and more programs, you may eventually start chaining them together. We often call these *pipelines*, as the output of one program is “piped” to become the input for the next program. If there is an error in any part of the pipeline, you’ll generally want the entire operation to stop so that the problems can be fixed. A nonzero return value from any program is a warning flag to halt operations.



9.3.4 Controlling randomness with random.seed()

The pseudo-random events in the `random` module follow from a given starting point. That is, each time you start from a given state, the events will happen in the same way. We can use the `random.seed()` function to set that starting point.

The seed value must be *hashable*. According to the Python documentation (<https://docs.python.org/3.1/glossary.html>), “all of Python’s immutable built-in objects are hashable, while no mutable containers (such as lists or dictionaries) are.” In this program, we have to use an integer value because the tests were written using integer seeds. When you write your own programs, you may choose to use a string or other hashable type.

The default for our seed is the special `None` value, which is a bit like an undefined state. Calling `random.seed(None)` is essentially the same as not setting the seed at all, so it makes it safe to write this:

```
random.seed(args.seed)
```

9.3.5 Iterating with range() and using throwaway variables

To generate some --number of insults, we can use the `range()` function. Because we don't need the numbers returned by `range()`, we use the underscore (`_`) as the variable name to indicate this is throwaway value:

```
>>> num_insults = 2
>>> for _ in range(num_insults):
...     print('An insult!')
...
An insult!
An insult!
```

The underscore is a valid variable name in Python. You can assign to it and use it:

```
>>> _ = 'You indistinguishable, filthsome carbuncle!'
>>> _
'You indistinguishable, filthsome carbuncle!'
```

The use of the underscore as a variable name is a convention to indicate that we don't intend to use the value. That is, if we had said `for num in range(...)`, some tools like Pylint will see that the `num` variable is not used and will report this as a possible error (and well it could be). The `_` indicates that you're throwing this value away, which is good information for your future self, some other user, or external tools to know.

Note that you can use multiple `_` variables in the same statement. For instance, I can unpack a 3-tuple so as to get the middle value:

```
>>> x = 'Jesus', 'Mary', 'Joseph'
>>> _, name, _ = x
>>> name
'Mary'
```

9.3.6 Constructing the insults

To create my list of adjectives, I used the `str.split()` method on a long, multiline string enclosed in triple quotes. I think this is probably the easiest way to get all these strings into my program. The triple quotes allow us to enter line breaks, which single quotes would not allow:

```
>>> adjectives = """
... bankrupt base caterwauling corrupt cullionly detestable dishonest
... false filthsome filthy foolish foul gross heedless indistinguishable
... infected insatiate irksome lascivious lecherous loathsome lubbery old
... peevish rascaly rotten ruinous scurilous scurvy slanderous
... sodden-witted thin-faced toad-spotted unmannered vile wall-eyed
... """".strip().split()
>>> nouns = """
... Judas Satan ape ass barbermonger beggar block boy braggart butt
... carbuncle coward coxcomb cur dandy degenerate fiend fishmonger fool
... gull harpy jack jolthead knave liar lunatic maw milksop minion
... ratcatcher recreant rogue scold slave swine traitor varlet villain worm
... """".strip().split()
```

```
>>> len(adjectives)
36
>>> len(nouns)
39
```

Because we need one or more adjectives, the `random.sample()` function is a good choice. It will return a list of items randomly selected from a given list:

```
>>> import random
>>> random.sample(adjectives, k=3)
['filthsome', 'cullionly', 'insatiate']
```

The `random.choice()` function is appropriate for selecting just one item from a list, such as the noun for our invective:

```
>>> random.choice(nouns)
'boy'
```

Next we need to concatenate the epithets using ', ' (a comma and a space) similar to what we did in chapter 3 for our picnic items. The `str.join()` function is perfect for this:

```
>>> adjs = random.sample(adjectives, k=3)
>>> adjs
['thin-faced', 'scurvy', 'sodden-witted']
>>> ', '.join(adjs)
'thin-faced, scurvy, sodden-witted'
```

To create the insult, we can combine the adjectives and nouns inside our template using an f-string:

```
>>> adjs = ', '.join(random.sample(adjectives, k=3))
>>> print(f'You {adjs} {random.choice(nouns)}!')
You heedless, thin-faced, gross recreant!
```

And now I have a handy way to make enemies and influence people.

9.4 Going further

- Read your adjectives and nouns from files that are passed as arguments.
- Add tests to verify that the files are processed correctly and new insults are still stinging.

Summary

- Use the `parser.error()` function to print a short usage statement, report the problem, and exit the program with an error value.
- Triple-quoted strings may contain line breaks, unlike regular single- or double-quoted strings.



- The `str.split()` method is a useful way to create a list of string values from a long string.
- The `random.seed()` function can be used to make reproducible pseudo-random selections each time a program is run.
- The `random.choice()` and `random.sample()` functions are useful for randomly selecting one or several items from a list of choices, respectively.

10

Telephone: Randomly mutating strings

“What we have here is a failure to communicate.”

—Captain

Now that we’ve played with randomness, let’s apply the idea to randomly mutating a string. This is interesting, because strings are actually *immutable* in Python. We’ll have to figure out a way around that.

To explore these ideas, we’ll write a version of the game of Telephone where a secret message is whispered through a line or circle of people. Each time the message is transmitted, it’s usually changed in some unpredictable way. The last person to receive the message will say it out loud to compare it to the original message. Often the results are nonsensical and possibly comical.

We will write a program called `telephone.py` that will mimic this game. It will print “You said: ” and the original text, followed by “I heard: ” with a modified version of the message. As in chapter 5, the input text may come from the command line:

```
$ ./telephone.py 'The quick brown fox jumps over the lazy dog.'  
You said: "The quick brown fox jumps over the lazy dog."  
I heard : "TheMquick brown fox jumps ovMr t:e lamy dog."
```

Or it may come from a file:

```
$ ./telephone.py ../inputs/fox.txt  
You said: "The quick brown fox jumps over the lazy dog."  
I heard : "The quick]b'own fox jumps ovek the la[y dog."
```



The program should accept an `-m` or `--mutations` option, which should be a floating-point number between 0 and 1 with a default value of 0.1 (10%). This will be the percentage of the number of letters that should be altered. For instance, `.5` means that 50% of the letters should be changed:

```
$ ./telephone.py ../inputs/fox.txt -m .5
You said: "The quick brown fox jumps over the lazy dog."
I heard : "F#eYquJsY ZrHnna"o. Muz/$ Nver t/Relazy dA!."
```

Because we are using the `random` module, we'll accept an `int` value for the `-s` or `--seed` option, so that we can reproduce our pseudo-random selections:

```
$ ./telephone.py ../inputs/fox.txt -s 1
You said: "The quick brown fox jumps over the lazy dog."
I heard : "The 'uicq brown *ox jumps over the l-zy dog."
```

Figure 10.1 shows a string diagram of the program.

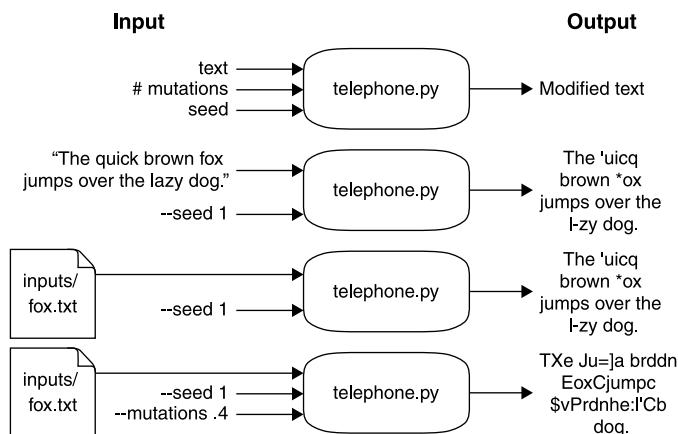


Figure 10.1 The telephone program will accept text and possibly some percentage of mutations, along with a random seed. The output will be a randomly mutated version of the input text.

In this exercise, you will learn to

- Round numbers
- Use the `string` module
- Modify strings and lists to introduce random mutations

10.1 Writing telephone.py

I recommend you use the new.py program to create a new program called telephone.py in the 10_telephone directory. You could do this from the top level of the repository like so:

```
$ ./bin/new.py 10_telephone/telephone.py
```

You could also copy template/template.py to 10_telephone/telephone.py. Modify the `get_args()` function until your `-h` output matches the following. I would recommend you use `type=float` for the `mutations` parameter:

```
$ ./telephone.py -h
usage: telephone.py [-h] [-s seed] [-m mutations] text

Telephone

positional arguments:
  text                  Input text or file

optional arguments:
  -h, --help            show this help message and exit
  -s seed, --seed seed  Random seed (default: None)
  -m mutations, --mutations mutations
                        Percent mutations (default: 0.1)
```

Now run the test suite. You should pass at least the first two tests (the telephone.py program exists and prints a usage statement when run with `-h` or `--help`).

The next two tests check that your `--seed` and `--mutations` options both reject non-numeric values. This should happen automatically if you define these parameters using the `int` and `float` types, respectively. That is, your program should behave like this:

```
$ ./telephone.py -s blargh foo
usage: telephone.py [-h] [-s seed] [-m mutations] text
telephone.py: error: argument -s/--seed: invalid int value: 'blargh'
$ ./telephone.py -m blargh foo
usage: telephone.py [-h] [-s seed] [-m mutations] text
telephone.py: error: argument -m/--mutations: invalid float value: 'blargh'
```

The next test checks if the program rejects values for `--mutations` outside the range 0–1 (where both bounds are inclusive). This is not a check that you can easily describe to argparse, so I suggest you look at how we handled the validation of the arguments in abuse.py in chapter 9. In the `get_args()` function of that program, we manually checked the value of the arguments and used the `parser.error()` function to throw an error. Note that a `--mutations` value of 0 is acceptable, in which case we will print out the input text without modifications. Your program should do this:

```
$ ./telephone.py -m -1 foobar
usage: telephone.py [-h] [-s seed] [-m mutations] text
telephone.py: error: --mutations "-1.0" must be between 0 and 1
```

This is another program that accepts input text either from the command line or from a file, and I suggest you look at the solution in chapter 5. Inside the `get_args()` function, you can use `os.path.isfile()` to detect whether the `text` argument is a file. If it is a file, read the contents of the file for the `text` value.

Once you have taken care of all the program parameters, start your `main()` function with setting the `random.seed()` and echoing back the given text:

```
def main():
    args = get_args()
    random.seed(args.seed)
    print(f'You said: "{args.text}"')
    print(f'I heard : "{args.text}"')
```

Your program should handle command-line text:

```
$ ./telephone.py 'The quick brown fox jumps over the lazy dog.'
You said: "The quick brown fox jumps over the lazy dog."
I heard : "The quick brown fox jumps over the lazy dog."
```

And it should handle an input file:

```
$ ./telephone.py ../inputs/fox.txt
You said: "The quick brown fox jumps over the lazy dog."
I heard : "The quick brown fox jumps over the lazy dog."
```

At this point, your code should pass up to `test_for_echo()`. The next tests start asking you to mutate the input, so let's discuss how to do that.

10.1.1 Calculating the number of mutations

The number of letters that need to be changed can be calculated by multiplying the length of the input text by the `args.mutations` value. If we want to change 20% of the characters in “The quick brown fox...” string, we’ll find that it is not a whole number:

```
>>> text = 'The quick brown fox jumps over the lazy dog.'
>>> mutations = .20
>>> len(text) * mutations
8.8
```

We can use the `round()` function to give us the nearest integer value. Read `help(round)` to learn how to round floating-point numbers to a specific number of digits:

```
>>> round(len(text) * mutations)
9
```

Note that you could also convert a `float` to an `int` by using the `int` function, but this truncates the fractional part of the number rather than rounding it:

```
>>> int(len(text) * mutations)
8
```

You will need this value for later, so let's save it in a variable:

```
>>> num_mutations = round(len(text) * mutations)
>>> assert num_mutations == 9
```

10.1.2 The mutation space

When we change a character, what will we change it to? For this, we'll use the `string` module. I encourage you to take a look at the documentation by importing the module and reading `help(string)`:

```
>>> import string
>>> help(string)
```

We can, for instance, get all the lowercase ASCII letters as follows. Note that this is not a method call as there are no parentheses () at the end:

```
>>> string.ascii_lowercase
'abcdefghijklmnopqrstuvwxyz'
```

This returns a `str`:

```
>>> type(string.ascii_lowercase)
<class 'str'>
```

For our program, we can use `string.ascii_letters` and `string.punctuation` to get strings of all the letters and punctuation. To concatenate the two strings together, we can use the + operator. We'll draw from this string to randomly select a character to replace another:

```
>>> alpha = string.ascii_letters + string.punctuation
>>> alpha
'abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ!"#$%&\'()*+, - ./:;<=>?@[\\\]^_`{|}~'
```

Note that even if we both use the same random seed, you and I will get different results if our letters are in a different order. To ensure our results match, we'll both need to sort the alpha characters so they are in a consistent order.

10.1.3 Selecting the characters to mutate

There are at least two approaches we could take to choosing which characters to change: a *deterministic approach* where the results are always guaranteed to be the same and a *non-deterministic approach* where we employ chance to get close to a target. Let's examine the latter one first.



NON-DETERMINISTIC SELECTION

One way to choose the characters to change would be to mimic method 1 in chapter 8. We could iterate through each of the characters in our text and select a random number to decide whether to keep the original character or change it to some randomly selected value. If our random number is less than or equal to our mutations setting, we should change the character:

```

new_text = ''           ← Initialize new_text as
for char in args.text:   ← an empty string.
    new_text += random.choice(alpha) if random.random() <= args.mutations
        else char
print(new_text)          ← Iterate through each
                        ← character in the text.
                        ← Use random.random() to generate a floating-point value from a
                        ← uniform distribution between 0 and 1. If that value is less than
                        ← or equal to the args.mutation value, we randomly choose from
                        ← alpha; otherwise, we use the original character.
                        ← Print the resulting
                        ← new_text.

```

We used the `random.choice()` function in `abuse.py` in chapter 9 to randomly select *one* value from a list of choices. We can use it here to select a character from `alpha` if the `random.random()` value falls within the range of the `args.mutation` value (which we know is also a float).

The problem with this approach is that, by the end of the `for` loop, we are not guaranteed to have made exactly the correct number of changes. That is, we calculated that we should change 9 characters out of 44 when the mutation rate is 20%. We would expect to end up changing about 20% of the characters with this code, because a random value from a uniform distribution of values between 0 and 1 should be less than or equal to 0.2 about 20% of the time. Sometimes we might end up only changing 8 characters or other times we might change 10. Because of this uncertainty, this approach would be considered *non-deterministic*.

Still, this is a really useful technique that you should note. Imagine you have an input file with millions or potentially billions of lines of text, and you want to randomly sample approximately 10% of the lines. The preceding approach would be reasonably fast and accurate. A larger sample size will help you get closer to the desired number of mutations.

RANDOMLY SAMPLING CHARACTERS

A deterministic approach to the million-line file would require first reading the entire input to count the number of lines, choosing which lines to take, and then going back through the file a second time to take those lines. This approach would take *much* longer than the method described above. Depending on how large the input file is, how the program is written, and how much memory your computer has, the program could possibly even crash your computer!

Our input is rather small, however, so we will use this algorithm because it has the advantages of being exact and testable. Rather than focusing on lines of text, though, we'll consider indexes of characters. You've seen the `str.replace()` method (in chapter 8), which allows us to change all instances of one string to another:

```
>>> 'foo'.replace('o', 'a')
'faa'
```



We can't use `str.replace()` because it will change every occurrence of some character, and we only want to change individual characters. Instead we can use the `random.sample()` function to select some indexes of the characters in the text. The first argument to `random.sample()` needs to be something like a `list`. We can give it a `range()` of numbers up to the length of our text.

Suppose our text is 44 characters long:

```
>>> text
'The quick brown fox jumps over the lazy dog.'
>>> len(text)
44
```

We can use the `range()` function to make a list of numbers up to 44:

```
>>> range(len(text))
range(0, 44)
```

Note that `range()` is a lazy function. It won't actually produce the 44 values until we force it, which we can do in the REPL using the `list()` function:

```
>>> list(range(len(text)))
```

We calculated earlier that the `num_mutations` value for altering 20% of `text` is 9. Here is one selection of indexes that could be changed:

```
>>> indexes = random.sample(range(len(text)), num_mutations)
>>> indexes
[13, 6, 31, 1, 24, 27, 0, 28, 17]
```

I suggest you use a `for` loop to iterate through each of these index values:

```
>>> for i in indexes:
...     print(f'{i:2} {text[i]}')
...
13 w
 6 i
31 t
 1 h
24 s
```

```
27 v
0 T
28 e
17 o
```

You should replace the character at each index position with a randomly selected character from alpha:

```
>>> for i in indexes:
...     print(f'{i:2} {text[i]} changes to {random.choice(alpha)}')
...
13 w changes to b
 6 i changes to W
31 t changes to B
 1 h changes to #
24 s changes to d
27 v changes to :
 0 T changes to C
28 e changes to %
17 o changes to ,
```

I will introduce one other twist—we don't want the replacement value to ever be the same as the character it is replacing. Can you figure out how to get a subset of alpha that *does not* include the character at the position?

10.1.4 Mutating a string

Python str variables are *immutable*, meaning we cannot directly modify them. For instance, suppose we want to change the character 'w' at position 13 to a 'b'. It would be handy to directly modify `text[13]`, but that will create an exception:

```
>>> text[13] = 'b'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'str' object does not support item assignment
```

The only way to modify the str value `text` is to overwrite it with a new str. We need to create a new str with the following, as shown in figure 10.2:

- The part of text before a given index
- The randomly selected value from alpha
- The part of text after a given index

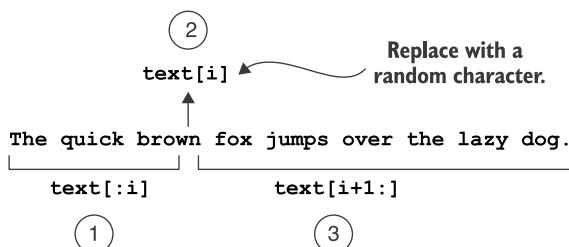


Figure 10.2 Create a new string by selecting the portion of the string up to the index, a new character, and the portion of the string after the index.

For 1 and 3, you can use string *slices*. For example, if the index *i* is 13, the slice before it is

```
>>> text[:13]
'The quick bro'
```

The part after it is

```
>>> text[14:]
'n fox jumps over the lazy dog.'
```

Using the three parts listed earlier, your for loop should be

```
for i in index:
    text = 1 + 2 + 3
```

Can you figure that out?

10.1.5 Time to write

OK, the lesson is over. You have to go write this now. Use the tests. Solve them one at a time. You can do this.

10.2 Solution

How different was your solution from mine? Let's look at one way to write a program that satisfies the tests:

```
#!/usr/bin/env python3
"""Telephone"""

import argparse
import os
import random
import string
# -----
def get_args():
    """Get command-line arguments"""
    parser = argparse.ArgumentParser(
        description='Telephone',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    parser.add_argument('text', metavar='text', help='Input text or file')
    parser.add_argument('-s', '--seed',
                        help='Random seed',
                        metavar='seed',
                        type=int,
                        default=None)

    import the string module
    we'll need to select a
    random character.
```

Define a positional argument for the text. This could be either a string of text or a file that needs to be read.

The --seed parameter is an integer value with a default of None.

```

If args.mutations is not in the acceptable range of 0–1, use
parser.error() to halt the program and print the given message. Note
the use of feedback to echo the bad args.mutation value to the user.

    parser.add_argument('-m',
                        '--mutations',
                        help='Percent mutations',
                        metavar='mutations',
                        type=float,
                        default=0.1)

    args = parser.parse_args()

    if not 0 <= args.mutations <= 1:
        parser.error(f'--mutations "{args.mutations}" must be between 0 and 1')

    if os.path.isfile(args.text):
        args.text = open(args.text).read().rstrip()

    return args
}

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    text = args.text
    random.seed(args.seed)
    alpha = ''.join(sorted(string.ascii_letters + string.punctuation))
    len_text = len(text)
    num_mutations = round(args.mutations * len_text)
    new_text = text

    for i in random.sample(range(len_text), num_mutations):
        new_char = random.choice(alpha.replace(new_text[i], ''))
        new_text = new_text[:i] + new_char + new_text[i + 1:]

    print(f'You said: "{text}"\nI heard : "{new_text}"')

#
if __name__ == '__main__':
    main()

```

If args.mutations is not in the acceptable range of 0–1, use parser.error() to halt the program and print the given message. Note the use of feedback to echo the bad args.mutation value to the user.

The --mutations parameter is a floating-point value with a default of 0.1.

Process the arguments from the command line. If argparse detects problems, such as non-numeric values for the seed or mutations, the program dies here and the user sees an error message. If this call succeeds, argparse has validated the arguments and converted the values.

If args.text names an existing file, read that file for the contents and overwrite the original value of args.text.

Return the processed arguments to the caller.

Set the random.seed() to the value provided by the user. Remember that the default value for args.seed is None, which is the same as not setting the seed.

Make a copy of text.

Print the text.

Overwrite the text by concatenating the slice before the current index with the new_char and then the slice after the current index.

Use random.choice () to select a new_char from a string created by replacing the current character (text[i]) in the alpha variable with nothing. This ensures that the new character cannot be the same as the one we are replacing.

Use random.sample () to get num_mutations indexes to change. This function returns a list that we can iterate using the for loop.

Figure the num_mutations by multiplying the mutation rate by the length of the text.

Since we use len(text) more than once, we put it into a variable.

Set alpha to be the characters we'll use for replacements. The sorted() function will return a new list of the characters in the right order, and then we can use the str.join() function to turn that back into a str value.

10.3 Discussion

There's nothing in `get_args()` that you haven't seen before. The `--seed` argument is an `int` that we will pass to the `random.seed()` function so as to control the randomness for testing. The default seed value is `None` so that we can call `random.seed(args.seed)` where `None` is the same as not setting it. The `--mutations` parameter is a `float` with a reasonable default, and we use `parser.error()` to create an error message if the value is not in the proper range. As in other programs, we test whether the `text` argument is a file and read the contents if it is.

10.3.1 Mutating a string

You saw earlier that we can't just change the `text` string:

```
>>> text = 'The quick brown fox jumps over the lazy dog.'
>>> text[13] = 'b'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'str' object does not support item assignment
```

We have to create a *new* string using the text before and after `i`, which we can get with string slices using `text[start:stop]`. If you leave out `start`, Python starts at 0 (the beginning of the string), and if you leave out `stop`, it goes to the end, so `text[:]` is a copy of the entire string.

If `i` is 13, the bit before `i` is

```
>>> i = 13
>>> text[:i]
'The quick bro'
```

The bit after `i + 1` is

```
>>> text[i+1:]
'n fox jumps over the lazy dog.'
```

Now for what to put in the middle. I noted that we should use `random.choice()` to select a character from `alpha`, which is the combination of all the ASCII letters and punctuation *without* the current character. I use the `str.replace()` method to get rid of the current letter:

```
>>> alpha = ''.join(sorted(string.ascii_letters + string.punctuation))
>>> alpha.replace(text[i], '')
'!#"#$%&\'()*+,-
./:;<=>?@ABCDEFGHIJKLMNPQRSTUVWXYZ[\]\\]^_`abcdefghijklmnopqrstuvwxyz{|}~'
```

Then I use that to get a new letter that won't include what it's replacing:

```
>>> new_char = random.choice(alpha.replace(text[i], ''))
>>> new_char
'Q'
```

There are many ways to join strings together into new strings. The `+` operator is perhaps the simplest:

```
>>> text = text[:i] + new_char + text[i+1:]
>>> text
'The quick broQn fox jumps over the lazy dog.'
```

I do this for each index in the `random.sample()` of indexes, each time overwriting `text`. After the `for` loop is done, I have mutated all the positions of the input string, and I can `print()` it.

10.3.2 Using a list instead of a str

Strings are immutable, but lists are not. You've seen that a move like `text[13] = 'b'` creates an exception, but we can change `text` into a list and directly modify it with the same syntax:

```
>>> text = list(text)
>>> text[13] = 'b'
```

We can then turn that list back into a `str` by joining it on the empty string:

```
>>> ''.join(text)
'The quick brobn fox jumps over the lazy dog.'
```

Here is a version of `main()` that uses this approach:

```
def main():
    args = get_args()
    text = args.text
    random.seed(args.seed)
    alpha = ''.join(sorted(string.ascii_letters + string.punctuation))
    len_text = len(text)
    num_mutations = round(args.mutations * len_text)
    new_text = list(text)                                ← Initialize new text as a list
                                                        of the original text value.

    for i in random.sample(range(len_text), num_mutations):
        new_text[i] = random.choice(alpha.replace(new_text[i], ''))

    print('You said: "{}"\nI heard : "{}"'.format(text, ''.join(new_text))) ← Join new_list on the empty
                                                                string to make a new str.

Now we can directly modify
a value in new_text.
```

There's no particular advantage of one approach over the other, but I would personally choose the second method because I don't like messing around with slicing strings. To me, modifying a list in place makes much more sense than repeatedly chopping up and piecing together a `str`.

Mutations in DNA

For what it's worth, this program mimics (kind of, sort of) how DNA changes over time. The machinery to copy DNA makes mistakes, and mutations randomly occur. Often the change has no deleterious effect on the organism.

Our example only changes characters to other characters—what biologists call “point mutations,” “single nucleotide variations” (SNV), or “single nucleotide polymorphisms” (SNP). We could instead write a version that would also randomly delete or insert new characters, which are called “in-del’s” (insertion-deletions). Mutations (that don’t result in the demise of the organism) occur at a fairly standard rate, so counting the number of mutations between a conserved region of any two organisms can allow an estimate of how long ago they diverged from a common ancestor.

10.4 Going further

- Apply the mutations to randomly selected words instead of the whole string.
- Perform insertions and deletions in addition to mutations; maybe create arguments for the percentage of each, and choose to add or delete characters at the indicated frequency.
- Add an option for `-o` or `--output` that names a file to write the output to. The default should be to print to `STDOUT`.
- Add a flag to limit the replacements to character values only (no punctuation).
- Add tests to `test.py` for every new feature, and ensure your program works properly.

Summary

- A string cannot be directly modified, but the variable containing the string can be repeatedly overwritten with new values.
- Lists can be directly modified, so it can sometimes help to use `list` on a string to turn it into a list, modify that list, and then use `str.join()` to change it back to a `str`.
- The `string` module has handy definitions of various strings.



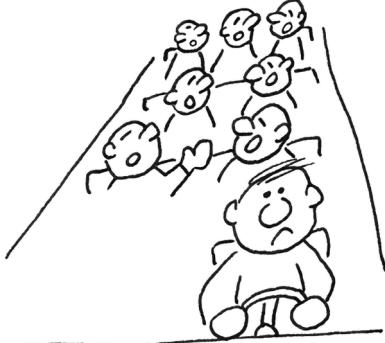
Bottles of Beer Song: Writing and testing functions

Few songs are as annoying as “99 Bottles of Beer on the Wall.” Hopefully you’ve never had to ride for hours in a van with middle school boys who like to sing this. I have. It’s a fairly simple song that we can write an algorithm to generate. This will give us an opportunity to play with counting up and down, formatting strings, and—new to this exercise—writing functions and tests for those functions!

Our program will be called `bottles.py` and will take one option, `-n` or `--num`, which must be a *positive int* (the default will be 10). The program should print all the verses from `--num` down to 1. There should be two newlines between each verse to visually separate them, but there must be only one newline after the last verse (for one bottle), which should print “No more bottles of beer on the wall” rather than “0 bottles”:

```
$ ./bottles.py -n 3
3 bottles of beer on the wall,
3 bottles of beer,
Take one down, pass it around,
2 bottles of beer on the wall!

2 bottles of beer on the wall,
2 bottles of beer,
Take one down, pass it around,
1 bottle of beer on the wall!
```



```
1 bottle of beer on the wall,
1 bottle of beer,
Take one down, pass it around,
No more bottles of beer on the wall!
```

In this exercise, you will

- Learn how to produce a list of numbers decreasing in value
- Write a function to create a verse of the song, using a test to verify when the verse is correct
- Explore how for loops can be written as list comprehensions, which in turn can be written with the `map()` function

11.1 Writing bottles.py

We'll be working in the `11_bottles_of_beer` directory. Start off by copying `template.py` or using `new.py` to create your `bottles.py` program there. Then modify the `get_args()` function until your usage matches the following usage statement. You need to define only the `--num` option with `type=int` and `default=10`:

```
$ ./bottles.py -h
usage: bottles.py [-h] [-n number]

Bottles of beer song

optional arguments:
  -h, --help            show this help message and exit
  -n number, --num number
                        How many bottles (default: 10)
```

If the `--num` argument is not an `int` value, your program should print an error message and exit with an error value. This should happen automatically if you define your parameter to `argparse` properly:

```
$ ./bottles.py -n foo
usage: bottles.py [-h] [-n number]
bottles.py: error: argument -n/--num: invalid int value: 'foo'
$ ./bottles.py -n 2.4
usage: bottles.py [-h] [-n number]
bottles.py: error: argument -n/--num: invalid int value: '2.4'
```

Since we can't sing zero or fewer verses, we'll need to check if `--num` is less than 1. To handle this, I suggest you use `parser.error()` inside the `get_args()` function, as in previous exercises:

```
$ ./bottles.py -n 0
usage: bottles.py [-h] [-n number]
bottles.py: error: --num "0" must be greater than 0
```

Figure 11.1 shows a string diagram of the inputs and outputs.

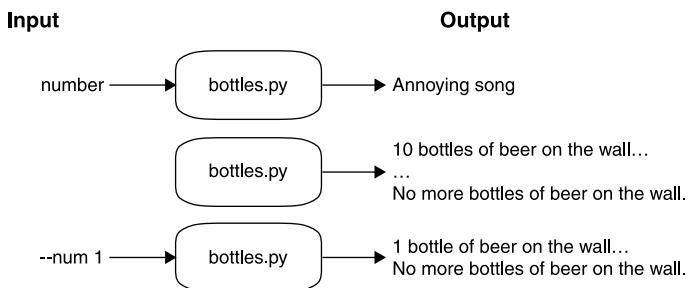
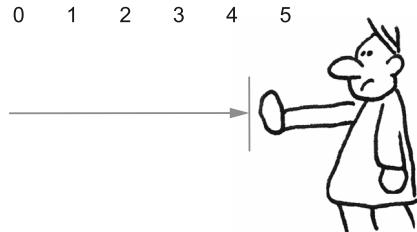


Figure 11.1 The bottles program may take a number for the verse to start, or it will sing the song starting at 10.

11.1.1 Counting down

The song starts at the given `--num` value, like 10, and needs to count down to 9, 8, 7, and so forth. How can we do that in Python? We've seen how to use `range(start, stop)` to get a list of integers that go *up* in value. If you give it just one number, that will be considered the `stop`, and it will assume 0 as the `start`:

```
>>> list(range(5))
[0, 1, 2, 3, 4]
```



Because this is a lazy function, we must use the `list()` function in the REPL to force it to produce the numbers. Remember that the `stop` value is never included in the output, so the preceding output stopped at 4, not 5.

If you give `range()` two numbers, they are considered to be `start` and `stop`:

```
>>> list(range(1, 5))
[1, 2, 3, 4]
```

To reverse this sequence, you might be tempted to swap the `start` and `stop` values. Unfortunately, if `start` is greater than `stop`, you get an empty list:

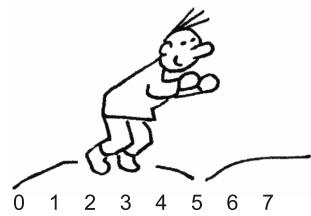
```
>>> list(range(5, 1))
[]
```

You saw in chapter 3 that we can use the `reversed()` function to reverse a list. This is another lazy function, so again I'll use the `list()` function to force the values in the REPL:

```
>>> list(reversed(range(1, 5)))
[4, 3, 2, 1]
```

The `range()` function can also take an optional third argument for a step value. For instance, you could use this to count by fives:

```
>>> list(range(0, 50, 5))
[0, 5, 10, 15, 20, 25, 30, 35, 40, 45]
```



Another way to count down is to swap the start and stop and use `-1` for the step:

```
>>> list(range(5, 0, -1))
[5, 4, 3, 2, 1]
```

So you have couple of ways to count in reverse.

11.1.2 Writing a function

Up to this point, I've suggested that all your code go into the `main()` function. This is the first exercise where I suggest you write a function. I would like you to consider how to write the code to sing *just one verse*. The function could take the number of the verse and return the text for that verse.

You can start off with something like the example in figure 11.2. The `def` keyword "defines" a function, and the name of the function follows. Function names should contain only letters, numbers, and underscores and cannot start with a number. After the name comes parentheses, which describe any parameters that the function accepts. Here our function will be called `verse()`, and it has the parameter `bottle` (or `number` or whatever you want to call it). After the parameters comes a colon to indicate the end of the `def` line. The function body comes next, with all lines being indented at least four spaces.

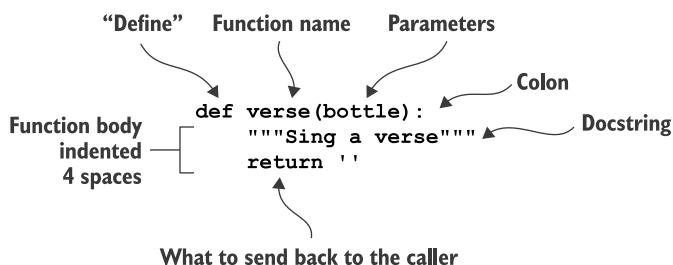


Figure 11.2 The elements of a function definition in Python

The docstring in figure 11.2 is a string just after the function definition. It will show up in the help for your function.

You can enter this function into the REPL:

```
>>> def verse(bottle):
...     """Sing a verse"""
...     return ''
...
>>> help(verse)
```

When you do, you will see this:

```
Help on function verse in module __main__:

verse(bottle)
    Sing a verse
```

The return statement tells Python what to send back from the function. It's not very interesting right now because it will just send back the empty string:

```
>>> verse(10)
''
```

It's also common practice to use the `pass` statement for the body of a dummy function. The `pass` will do nothing, and the function will return `None` instead of the empty string, as we have done here. When you start writing your own functions and tests, you might like to use `pass` when you stub out a new function, until you decide what the function will do.

11.1.3 Writing a test for `verse()`

In the spirit of *test-driven development*, let's write a test for `verse()` before we go any further. The following listing shows a test you can use. Add this code into your `bottles.py` program just after your `main()` function :

```
def verse(bottle):
    """Sing a verse"""

    return ''


def test_verse():
    """Test verse"""

    last_verse = verse(1)
    assert last_verse == '\n'.join([
        '1 bottle of beer on the wall,', '1 bottle of beer,',

        'Take one down, pass it around,',

        'No more bottles of beer on the wall!'
    ])

    two_bottles = verse(2)
    assert two_bottles == '\n'.join([
        '2 bottles of beer on the wall,', '2 bottles of beer,',

        'Take one down, pass it around,',

        '1 bottle of beer on the wall!'
    ])
```

```
'Take one down, pass it around,', '1 bottle of beer on the wall!'
])
```

There are many, many ways you could write this program. I have in mind that my `verse()` function will produce a single verse of the song, returning a new `str` value that is the lines of the verse joined on newlines.

You don't have to write your program this way, but I'd like you to consider what it means to write a function and a *unit test*. If you read about software testing, you'll find that there are different definitions of what a "unit" of code is. In this book, I consider a *function* to be a *unit*, so my unit tests are tests of individual functions.



Even though the song has potentially hundreds of verses, these two tests should cover everything you need to check. It may help to look at the musical notation in figure 11.3 for the song, as this does a nice job of graphically showing the structure of the song and, hence, our program.

99 Bottles of Beer

Anonymous

Figure 11.3 The musical notation for the song shows there are two cases to handle: one for all the verses up to the last, and then the last one.

I've taken a few liberties with the notation by mixing in some programming ideas. If you don't know how to read music, let me briefly explain the important parts. The `N` is the current *number*, like "99" so that `(N - 1)` would be "98." The endings are noted `1 - (N - 1)`, which is a bit confusing because we're using the hyphen to indicate both a range and subtraction in the same "equation." Still, the first ending is used for the first time through the penultimate repeat. The colon before the bar lines in the first ending means to repeat the song from the beginning. Then the `N` ending is taken on the last repeat, and the double bar indicates the end of the song/program.

What we can see from the music is that there are only two cases we need to handle: the last verse, and all the other verses. So first we check the last verse. We're looking for "1 bottle" (singular) and not "1 bottles" (plural). We also need to check that the last line says "No more bottles" instead of "0 bottles." The second test, for "2 bottles of beer," is making sure that the numbers are "2 bottles" and then "1 bottle." If we managed to pass these two tests, our program ought to be able to handle all the verses.

I wrote `test_verse()` to test just the `verse()` function. The name of the function matters because I am using the `pytest` module to find all the functions in my code that start with `test_` and run them. If your `bottles.py` program has the preceding functions for `verse()` and `test_verse()`, you can run `pytest bottles.py`.

Try it, and you should see something like this:

```
$ pytest bottles.py
=====
test session starts =====
...
collected 1 item

bottles.py F [100%]

=====
FAILURES =====
test_verse

def test_verse():
    """Test verse"""
    last_verse = verse(1)      ←
    assert last_verse == '\n'.join([
        '1 bottle of beer on the wall,', '1 bottle of beer,',

        'Take one down, pass it around,',

        'No more bottles of beer on the wall!'
    ])
E     AssertionError: assert '' == '1 bottle of beer on the wal...ottles of
beer on the wall!'                                ←
E         + 1 bottle of beer on the wall,
E         + 1 bottle of beer,
E         + Take one down, pass it around,
E         + No more bottles of beer on the wall!

bottles.py:49: AssertionError
=====
1 failed in 0.10 seconds =====
```

The `>` at the beginning of this line indicates this is the source of the error. The test checks if the value of `last_verse` is equal to an expected str value. Since it's not, this line throws an exception, causing the assertion to fail.

Call the `verse()` function with the argument 1 to get the last verse of the song.

The "E" lines show the difference between what was received and what was expected. The value of `last_verse` is the empty string (""), which does not match the expected string "1 bottle of beer..." and so on.

To pass the first test, you could copy the code for the expected value of `last_verse` directly from the test. Change your `verse()` function to match this:

```
def verse(bottle):
    """Sing a verse"""

    return '\n'.join([
        '1 bottle of beer on the wall,', '1 bottle of beer,', 
        'Take one down, pass it around,', 
        'No more bottles of beer on the wall!'
    ])
```

Now run your test again. The first test should pass, and the second one should fail. Here are the relevant error lines:

```
===== FAILURES =====
test_verse _____
```

```
def test_verse() -> None:
    """Test verse"""

    last_verse = verse(1)
    assert last_verse == '\n'.join([
        '1 bottle of beer on the wall,', '1 bottle of beer,', 
        'Take one down, pass it around,', 
        'No more bottles of beer on the wall!'
    ]) This test now passes.

    two_bottles = verse(2) Call verse() with the value of 2 to get the “Two bottles...” verse.
    assert two_bottles == '\n'.join([
        '2 bottles of beer on the wall,', '2 bottles of beer,', 
        'Take one down, pass it around,', '1 bottle of beer on the wall!'
    ])

E     AssertionError: assert '1 bottle of ... on the wall!' == '2 bottles of ...
... on the wall!' These E lines are showing you the problem. The verse() function returned '1 bottle' but the test expected '2 bottles', etc.
E         - 1 bottle of beer on the wall,
E         ? ^
E         + 2 bottles of beer on the wall,
E         ? ^      +
E         - 1 bottle of beer,
E         ? ^
E         + 2 bottles of beer, ...
E
E     ...Full output truncated (7 lines hidden), use '-vv' to show
```

Assert that this verse is equal to the expected string.

Go back and look at your `verse()` definition. Look at figure 11.4 and think about which parts need to change—the first, second, and fourth lines. The third line is always the same. You’re given a value for `bottle` that needs to be used in the first two lines, along with either “bottle” or “bottles,” depending on the value of `bottle`. (Hint: It’s only singular for the value 1; otherwise, it’s plural.) The fourth line needs the value of `bottle - 1` and, again, the proper singular or plural depending on that value. Can you figure out how to write this?

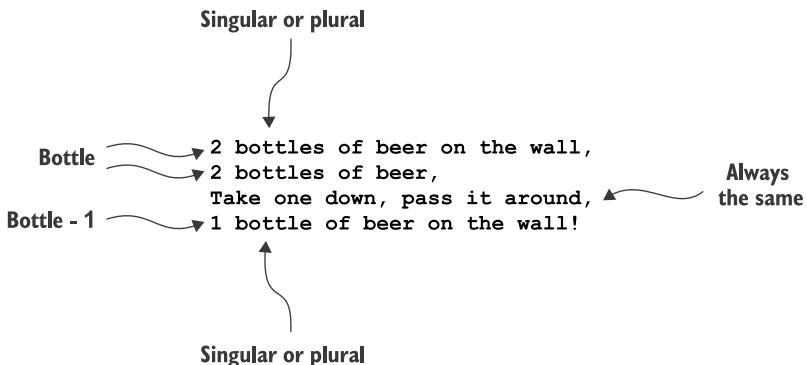


Figure 11.4 Each verse has four lines, where the first two and last are very similar. The third line is always the same. Find the parts that vary.

Focus on passing those two tests before you move to the next stage of printing the whole song. That is, do not attempt anything until you see this:

```
$ pytest bottles.py
=====
test session starts =====
...
collected 1 item

bottles.py . [100%]

===== 1 passed in 0.05 seconds =====
```

11.1.4 Using the `verse()` function

At this point, you know

- That the `--num` value is a valid integer value greater than 0
- How to count from that `--num` value backwards down to 0
- That the `verse()` function will print any one verse properly

Now you need to put them together. I suggest you start by using a `for` loop with the `range()` function to count down. Use each value from that to produce a `verse()`. There should be two newlines after every verse except for the last.

You will use `pytest -xv test.py` (or `make test`) to test the program at this point. In the parlance of testing, `test.py` is an *integration test* because it checks that the program *as a whole* is working. From this point on, we'll focus on how to write *unit* tests to check individual functions in addition to integration tests to ensure that all the functions work together.

Once you can pass the test suite using a `for` loop, try to rewrite it using either a list comprehension or a `map()`. Rather than starting again from scratch, I suggest

you comment out your working code by adding # to the beginnings of the lines, and then try other ways to write the algorithm. Use the tests to verify that your code still passes. If it is at all motivating, my solution is one line long. Can you write a single line of code that combines the `range()` and `verse()` functions to produce the expected output?

Here are a few hints:

- Define the `--num` argument as an `int` with a default value of 10.
- Use `parser.error()` to get argparse to print an error message for a negative `--num` value.
- Write the `verse()` function. Use the `test_verse()` function and Pytest to make that work properly.
- Combine the `verse()` function with `range()` to create all the verses.

Do try your best to write the program before reading the solution. Also feel free to solve the problem in a completely different way, even writing your own unit tests.

11.2 Solution

I've decided to show you a slightly fancy-pants version that uses `map()`. Later I'll show you how to write it using a `for` loop and a list comprehension.

```
#!/usr/bin/env python3
"""Bottles of beer song"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Bottles of beer song',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('-n',
                        '--num',
                        metavar='number',
                        type=int,
                        default=10,
                        help='How many bottles')

    args = parser.parse_args() →
    if args.num < 1: ←
        parser.error(f"--num '{args.num}' must be greater than 0")

    return args

Parse the command-line argument into the variable args. →
```

Define the `--num` argument as an `int` with a default value of 10.

If `args.num` is less than 1, use `parser.error()` to display an error message and exit the program with an error value.

The `map()` function expects a function as the first argument and some iterable as the second argument. Here I feed the descending numbers from the `range()` function to my `verse()` function. The result from `map()` is a new list of verses that can be joined on two newlines.

```
# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    print('\n\n'.join(map(verse, range(args.num, 0, -1))))
```

Define a function that can create a single verse().

Do the same for s2 (the second “s”), depending on the value of next_bottle.

```
# -----
def verse(bottle):
    """Sing a verse"""

    next_bottle = bottle - 1
    s1 = '' if bottle == 1 else 's'
    s2 = '' if next_bottle == 1 else 's'
    num_next = 'No more' if next_bottle == 0 else next_bottle
    return '\n'.join([
        f'{bottle} bottle{s1} of beer on the wall,',
        f'{bottle} bottle{s1} of beer,',
        f'Take one down, pass it around,',
        f'{num_next} bottle{s2} of beer on the wall!',
    ])
```

Define a next_bottle that is one less than the current bottle.

Define an s1 (the first “s”) that is either the character ‘s’ or the empty string, depending on the value of current bottle.

Define a value for next_num depending on whether the next value is 0 or not.

Create a return string by joining the four lines of text on the newline. Substitute in the variables to create the correct verse.

Test the last verse() with the value 1.

```
# -----
def test_verse():
    """Test verse"""

    last_verse = verse(1)
    assert last_verse == '\n'.join([
        '1 bottle of beer on the wall,',
        '1 bottle of beer,',
        'Take one down, pass it around,',
        'No more bottles of beer on the wall!'
    ])
```

Define a unit test called `test_verse()` for the `verse()` function. The `test_` prefix means that the `pytest` module will find this function and execute it.

Test a verse() with the value 2.

```
# -----
two_bottles = verse(2)
assert two_bottles == '\n'.join([
    '2 bottles of beer on the wall,',
    '2 bottles of beer,',
    'Take one down, pass it around,',
    '1 bottle of beer on the wall!'
])
```

```
# -----
if __name__ == '__main__':
    main()
```

11.3 Discussion

There isn't anything new in the `get_args()` function in this program. By this point, you have had several opportunities to define an optional integer parameter with a default argument and to use `parser.error()` to halt your program if the user provides a bad argument. By relying on `argparse` to handle so much busy work, you are saving yourself loads of time as well as ensuring that you have good data to work with. Let's move on to the new stuff!

11.3.1 Counting down

You know how to count down from the given `--num`, and you know you can use a `for` loop to iterate:

```
>>> for n in range(3, 0, -1):
...     print(f'{n} bottles of beer')
...
3 bottles of beer
2 bottles of beer
1 bottles of beer
```

Instead of directly creating each verse inside the `for` loop, I suggested that you could create a function called `verse()` to create any given verse and use that with the `range()` of numbers. Up to this point, we've been doing all our work in the `main()` function. As you grow as a programmer, though, your programs will become longer—hundreds to even thousands of lines of code (LOC). Long programs and functions can get very difficult to test and maintain, so you should try to break ideas into small, functional units that you can understand and test. Ideally, functions should do *one* thing. If you understand and trust your small, simple *functions*, then you know you can safely compose them into longer, more complicated *programs*.

11.3.2 Test-driven development

I wanted you to add a `test_verse()` function to your program to use with Pytest to create a working `verse()` function. This idea follows the principles described by Kent Beck in his book, *Test-Driven Development* (Addison-Wesley Professional, 2002):

- 1 Add a new test for an unimplemented unit of functionality.
- 2 Run all previously written tests and see the newly added test fails.
- 3 Write code that implements the new functionality.
- 4 Run all tests and see them succeed.
- 5 Refactor (rewrite to improve readability or structure).
- 6 Start at the beginning (repeat).

For instance, suppose we want a function that adds 1 to any given number. We'll call it `add1()` and define the function body as `pass` to tell Python "nothing to see here":

```
def add1(n):
    pass
```

Now write a `test_add1()` function where you pass some arguments to the function, and use `assert` to verify that you get back the value that you expect:

```
def test_add1():
    assert add1(0) = 1
    assert add1(1) = 2
    assert add1(-1) = 0
```

Run pytest (or whatever testing framework you like) and verify that the function *does not work* (of course it won't, because it just executes `pass`). Then go fill in some function code that *does work* (return `n + 1` instead of `pass`). Pass all manner of arguments you can imagine, including nothing, one thing, and many things.¹

11.3.3 The `verse()` function

I provided you with a `test_verse()` function that shows you exactly what is expected for the arguments of 1 and 2. What I like about writing my tests first is that it gives me an opportunity to think about how I'd like to use the code, what I'd like to give as arguments, and what I expect to get back in return. For instance, what *should* the function `add1()` return if given



- No arguments
- More than one argument
- The value `None`
- Anything other than a numeric type (`int`, `float`, or `complex`) like a `str` value or a `dict`

You can write tests to pass both good and bad values and decide how you want your code to behave under both favorable and adverse conditions.

Here's the `verse()` function I wrote, which passes the `test_verse()` function:

```
def verse(bottle):
    """Sing a verse"""

    next_bottle = bottle - 1
    s1 = '' if bottle == 1 else 's'
    s2 = '' if next_bottle == 1 else 's'
    num_next = 'No more' if next_bottle == 0 else next_bottle
    return '\n'.join([
        f'{bottle} bottle{s1} of beer on the wall,',
        f'{bottle} bottle{s1} of beer,',
        f'Take one down, pass it around,',
        f'{num_next} bottle{s2} of beer on the wall!',
    ])
```

¹ A CS professor once told me in office hours to handle the cases of 0, 1, and n (infinity), and that has always stuck with me.

This code is annotated in section 11.2, but I essentially isolate all the parts of the return string that change, and I create variables to substitute into those places. I use `bottle` and `next_bottle` to decide if there should be an “s” or not after the “bottle” strings. I also need to figure out whether to print the next bottle as a number, or if I should print the string “No more” (when `next_bottle` is 0). Choosing the values for `s1`, `s2`, and `num_next` all involve *binary* decisions, meaning they are a choice between *two* values, so I find it best to use an `if` expression.

This function passes `test_verse()`, so I can move on to using it to generate the song.

11.3.4 Iterating through the verses

I could use a `for` loop to count down and `print()` each `verse()`:

```
>>> for n in range(3, 0, -1):
...     print(verse(n))
...
3 bottles of beer on the wall,
3 bottles of beer,
Take one down, pass it around,
2 bottles of beer on the wall!
2 bottles of beer on the wall,
2 bottles of beer,
Take one down, pass it around,
1 bottle of beer on the wall!
1 bottle of beer on the wall,
1 bottle of beer,
Take one down, pass it around,
No more bottles of beer on the wall!
```

That’s *almost* correct, but we need two newlines in between all the verses. I could use the `end` option to print to include two newlines for all values greater than 1:

```
>>> for n in range(3, 0, -1):
...     print(verse(n), end='\n' * (2 if n > 1 else 1))
...
3 bottles of beer on the wall,
3 bottles of beer,
Take one down, pass it around,
2 bottles of beer on the wall!

2 bottles of beer on the wall,
2 bottles of beer,
Take one down, pass it around,
1 bottle of beer on the wall!

1 bottle of beer on the wall,
1 bottle of beer,
Take one down, pass it around,
No more bottles of beer on the wall!
```

I would rather use the `str.join()` method to put two newlines in between items in a list. My items are the verses, and I can turn a for loop into a list comprehension as shown in figure 11.5.

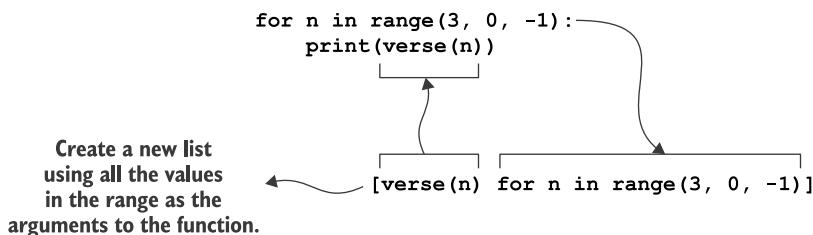


Figure 11.5 A for loop compared to a list comprehension

```
>>> verses = [verse(n) for n in range(3, 0, -1)]
>>> print('\n\n'.join(verses))
3 bottles of beer on the wall,
3 bottles of beer,
Take one down, pass it around,
2 bottles of beer on the wall!

2 bottles of beer on the wall,
2 bottles of beer,
Take one down, pass it around,
1 bottle of beer on the wall!

1 bottle of beer on the wall,
1 bottle of beer,
Take one down, pass it around,
No more bottles of beer on the wall!
```

That is a fine solution, but I would like you to start noticing a pattern we will see repeatedly: applying a function to every element of a sequence, which is exactly what `map()` does! As shown in figure 11.6, our list comprehension can be rewritten very concisely using `map()`.

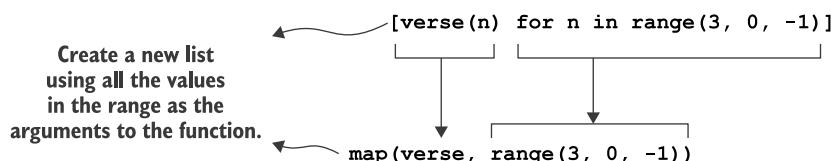


Figure 11.6 A list comprehension can be replaced with `map()`. They both return a new list.

In our case, our sequence is a descending `range()` of numbers, and we want to apply our `verse()` function to each number and collect the resulting verses. It's like the paint booth idea in chapter 8, where the function "painted" the cars "blue" by adding the word "blue" to the start of the string. When we want to apply a function to every element in a sequence, we might consider refactoring the code using `map()`:

```
>>> verses = map(verse, range(3, 0, -1))
>>> print('\n\n'.join(verses))
3 bottles of beer on the wall,
3 bottles of beer,
Take one down, pass it around,
2 bottles of beer on the wall!

2 bottles of beer on the wall,
2 bottles of beer,
Take one down, pass it around,
1 bottle of beer on the wall!

1 bottle of beer on the wall,
1 bottle of beer,
Take one down, pass it around,
No more bottles of beer on the wall!
```

Whenever I need to transform some sequence of items with some function, I like to start off by thinking about how I'll handle just *one* of the items. I find it's much easier to write and test one function with one input rather than some possibly huge list of operations. List comprehensions are often considered more "Pythonic," but I tend to favor `map()` because it usually involves shorter code. If you search the internet for "python list comprehension map," you'll find that some people think list comprehensions are easier to read than `map()`, but `map()` might possibly be somewhat faster. I wouldn't say either approach is better than the other. It really comes down to taste or perhaps a discussion with your teammates.

If you want to use `map()`, remember that it wants a *function* as the first argument and then a sequence of elements that will become arguments to the function. The `verse()` function (which you've tested!) is the first argument, and the `range()` provides the list. The `map()` function will pass each element of the `range()` as an argument to the `verse()` function, as shown in figure 11.7. The result is a new list with the return values from all those function calls. Many are the `for` loops that can be better written as mapping a function over a list of arguments!

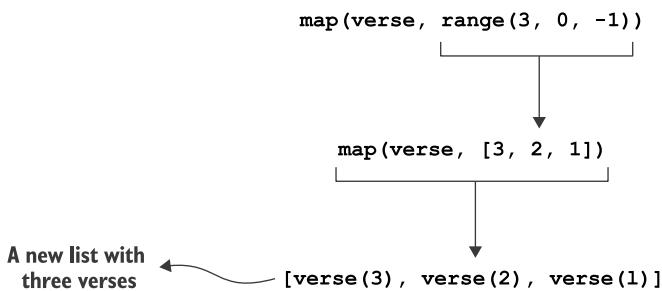


Figure 11.7 The `map()` function will call the `verse()` function with each element produced by the `range()` function. It's functions all the way down.

11.3.5 1,500 other solutions

There are literally hundreds of ways to solve this problem. The “99 Bottles of Beer” website (www.99-bottles-of-beer.net) claims to have 1,500 variations in various languages. Compare your solution to others there. Trivial as the program may be, it has allowed us to explore some really interesting ideas in Python, testing, and algorithms.



11.4 Going further

- Replace the Arabic numbers (1, 2, 3) with text (one, two, three).
- Add a --step option (positive int, default 1) that allows the user to skip numbers, like by twos or fives.
- Add a --reverse flag to reverse the order of the verses, counting up instead of down.

Summary

- Test-driven development (TDD) is central to developing dependable, reproducible code. Tests also give you the freedom to refactor your code (reorganize and improve it for speed or clarity), knowing that you can always verify your new version still works properly. As you write your code, always write tests!
- The `range()` function will count backwards if you swap `start` and `stop` and supply the optional third `step` value of `-1`.
- A `for` loop can often be replaced with a list comprehension or a `map()` for shorter, more concise code.

Ransom: Randomly capitalizing text

All this hard work writing code is getting on my nerves. I'm ready to turn to a life of crime! I've kidnapped (cat-napped?) the neighbor's cat, and I want to send them a ransom note. In the good old days, I'd cut letters from magazines and paste them onto a piece of paper to spell out my demands. That sounds like too much work. Instead, I'm going to write a Python program called `ransom.py` that will encode text into randomly capitalized letters:

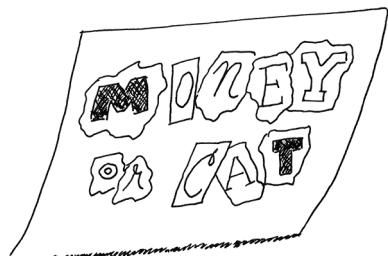
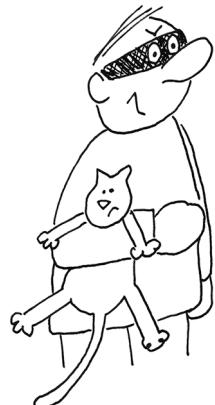
```
$ ./ransom.py 'give us 2 million dollars or the cat  
    gets it!'  
gIVE uS 2 MILLioN doLLaRS OR tHe cAt GETs It!
```

As you can see, my diabolical program accepts the heinous input text as a positional argument. Since this program uses the `random` module, I want to accept an `-s` or `--seed` option so I can replicate the vile output:

```
$ ./ransom.py --seed 3 'give us 2 million dollars or the cat gets it!'  
gIVE uS 2 MILLioN doLLaRS OR tHe cAt GETs It!
```

The dastardly positional argument might name a vicious file, in which case that should be read for the demoniac input text:

```
$ ./ransom.py --seed 2 ../inputs/fox.txt  
the qUICK BROWN FOX JUmPs oVer tHe LAZY DOG.
```



If the unlawful program is run with no arguments, it should print a short, infernal usage statement:

```
$ ./ransom.py
usage: ransom.py [-h] [-s int] text
ransom.py: error: the following arguments are required: text
```

If the nefarious program is run with -h or --help flags, it should print a longer, fiendish usage:

```
$ ./ransom.py -h
usage: ransom.py [-h] [-s int] text

Ransom Note

positional arguments:
  text           Input text or file

optional arguments:
  -h, --help      show this help message and exit
  -s int, --seed int  Random seed (default: None)
```

Figure 12.1 shows a noxious string diagram to visualize the inputs and outputs.

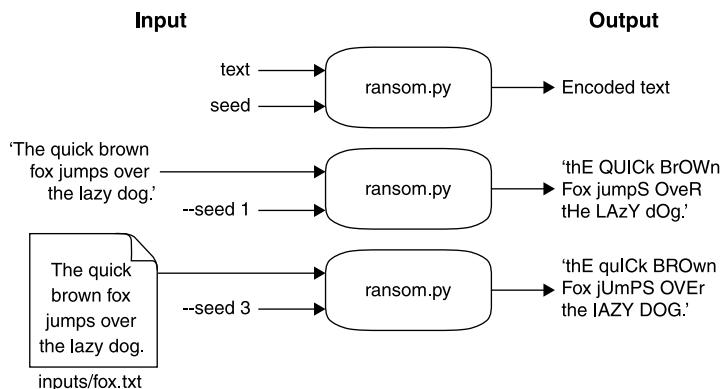


Figure 12.1 The awful program will transform input text into a ransom note by randomly capitalizing letters.

In this chapter, you will

- Learn how to use the `random` module to figuratively “flip a coin” to decide between two choices
- Explore ways to generate new strings from an existing one, incorporating random decisions
- Study the similarities of `for` loops, list comprehensions, and the `map()` function

12.1 Writing ransom.py

I suggest starting with new.py or copying the template/template.py file to create ransom.py in the 12_ransom directory. This program, like several before it, accepts a required, positional string for the text and an optional integer (default None) for the --seed. Also, as in previous exercises, the text argument may name a file that should be read for the text value.

To start out, use this for your main() code:

```
def main():
    args = get_args()
    random.seed(args.seed)
    print(args.text)
```

Get the processed command-line arguments.

Set the random.seed() with the value from the user. The default is None, which is the same as not setting it.

Start off by echoing back the input.

If you run this program, it should echo the input from the command line:

```
$ ./ransom.py 'your money or your life!'
your money or your life!
```

Or the text from an input file:

```
$ ./ransom.py ../inputs/fox.txt
The quick brown fox jumps over the lazy dog.
```

The important thing when writing a program is to take baby steps. You should run your program *after every change*, checking manually and with the tests to see if you are progressing.

Once you have this working, it's time to think about how to randomly capitalize this awful message.

12.1.1 Mutating the text

You've seen before that you can't directly modify a str value:

```
>>> text = 'your money or your life!'
>>> text[0] = 'Y'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'str' object does not support item assignment
```

So how can we randomly change the case of some of the letters?

I suggest that instead of thinking about how to change many letters, you should think about how to change *one* letter. That is, given a single letter, how can you randomly return the upper- or lowercase version of the letter? Let's create a dummy choose() function that accepts a single character. For now, we'll have the function return the character unchanged:

```
def choose(char):
    return char
```

Here's a test for it:

```
The state of the random module is global to the program.  
Any change we make here could affect unknown parts of  
the program, so we save our current state.
```

```
def test_choose():
    state = random.getstate()
    random.seed(1)
    assert choose('a') == 'a'
    assert choose('b') == 'b'
    assert choose('c') == 'C'
    assert choose('d') == 'd'
    random.setstate(state)
```

```
Set the random seed to a known  
value. This is a global change to our  
program. Any other calls to  
functions from the random module  
will be affected!
```

```
The choose() function is given a  
series of letters, and we use the  
assert statement to test if the  
value returned by the function  
is the expected letter.
```

```
Reset the global state to the original value.
```

Random seeds

Have you wondered how I knew what would be the result of `choose()` for a given random seed? Well, I confess that I wrote the function, then set the seed, and ran it with the given inputs. I recorded the results as the assertions you see. In the future, these results should still be the same. If they are not, I've changed something and probably broken my program.

12.1.2 Flipping a coin

We need to `choose()` between returning the upper- or lowercase version of the character you are given. It's a *binary* choice, meaning we have two options, so we can use the analogy of flipping a coin. Heads or tails? Or, for our purposes, 0 or 1:

```
>>> import random
>>> random.choice([0, 1])
1
```

Or `True` or `False` if you prefer:

```
>>> random.choice([False, True])
True
```



Think about using an `if` expression where you return the uppercase answer when the 0 or `False` option is selected and the lowercase version otherwise. My entire `choose()` function is this one line.

12.1.3 Creating a new string

Now we need to apply our `choose()` function to each character in the input string. I hope this is starting to feel like a familiar tactic. I encourage you to start by mimicking the first approach from chapter 8 where we used a `for` loop to iterate through each

character of the input text and replace all the vowels with a single vowel. In this program, we can iterate through the characters of text and use them as the argument to the `choose()` function. The result will be a new list (or `str`) of the transformed characters. Once you can pass the test with a `for` loop, try to rewrite it as a list comprehension, and then a `map()`.

Now off you go! Write the program, pass the tests.

12.2 Solution

We're going to explore many ways to process all the characters in the input text. We'll start off with a `for` loop that builds up a new list, and I hope to convince you that a list comprehension is a better way to do this. Finally, I'll show you how to use `map()` to create a very terse (perhaps even elegant) solution.

```
#!/usr/bin/env python3
"""Ransom note"""

import argparse
import os
import random

# -----
def get_args():
    """get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Ransom Note',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='text', help='Input text or file') ←

    parser.add_argument('-s', '--seed',
                        help='Random seed',
                        metavar='int',
                        type=int,
                        default=None) ←
    The --seed option
    is an integer that
    defaults to None.

    args = parser.parse_args() ←
    Process the command-line
    arguments into the args
    variable.

    if os.path.isfile(args.text): ←
        args.text = open(args.text).read().rstrip()

    return args ←
    Return the arguments
    to the caller.

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    text = args.text
```

The text argument is a positional string value.

If the args.text is a file, use the contents of that as the new args.text value.

```

Set the random.seed() to the given args.seed value. The
default is None, which is the same as not setting it. That
means the program will appear random when no seed is
given but will be testable when we do provide a seed value.

Create an empty list to
hold the new ransom
message.

Use a for loop to iterate
through each character
of args.text.

Join the ransom list on the
empty string to create a
new string to print.

Define a function to
randomly return the
uppercase or lowercase
version of a given
character.

Append the
chosen
letter to
the ransom
list.

Save the
current
state of the
random
module.

Set the
random.seed()
to a known
value for the
purposes of
the test.

random.seed(args.seed)
ransom = []
for char in args.text:
    ransom.append(choose(char))
print('.join(ransom))

# -----
def choose(char):
    """Randomly choose an upper or lowercase letter to return"""

    return char.upper() if random.choice([0, 1]) else char.lower()

# -----
def test_choose():
    """Test choose"""

    state = random.getstate()
    random.seed(1)
    assert choose('a') == 'a'
    assert choose('b') == 'b'
    assert choose('c') == 'C'
    assert choose('d') == 'd'
    random.setstate(state)

# -----
if __name__ == '__main__':
    main()

```

Append the chosen letter to the ransom list.

Save the current state of the random module.

Set the random.seed() to a known value for the purposes of the test.

Create an empty list to hold the new ransom message.

Use a for loop to iterate through each character of args.text.

Join the ransom list on the empty string to create a new string to print.

Define a function to randomly return the uppercase or lowercase version of a given character.

Use random.choice() to select either 0 or 1, which, in the Boolean context of the if expression, evaluates to False or True, respectively.

Use the assert statement to verify that we get the expected result from the choose() for a known argument.

Reset the random module's state so that our changes won't affect any other part of the program.

12.3 Discussion

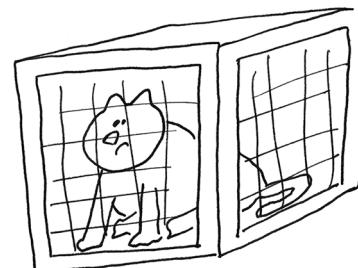
I like this problem because there are so many interesting ways to solve it. I know, I know, Python likes there to be “one obvious way” to solve it, but let’s explore, shall we? There’s nothing in `get_args()` that we haven’t seen several times by now, so let’s skip that.

12.3.1 Iterating through elements in a sequence

Assume that we have the following cruel message:

```
>>> text = '2 million dollars or the cat sleeps with the fishes!'
```

I want to randomly upper- and lowercase the letters. As suggested in the earlier description of the problem, we can use a for loop to iterate over each character.



One way to print an uppercase version of the text is to print an uppercase version of each letter:

```
for char in text:
    print(char.upper(), end='')
```

That would give me “2 MILLION DOLLARS OR THE CAT SLEEPS WITH THE FISHES!” Now, instead of always printing `char.upper()`, I can randomly choose between `char.upper()` and `char.lower()`. For that, I’ll use `random.choice()` to choose between two values like `True` and `False` or 0 and 1:

```
>>> import random
>>> random.choice([True, False])
False
>>> random.choice([0, 1])
0
>>> random.choice(['blue', 'green'])
'blue'
```



Following the first solution from chapter 8, I created a new list to hold the ransom message and added these random choices:

```
ransom = []
for char in text:
    if random.choice([False, True]):
        ransom.append(char.upper())
    else:
        ransom.append(char.lower())
```

Then I joined the new characters on the empty string to print a new string:

```
print(''.join(ransom))
```

It’s far less code to write this with an `if` expression to select whether to take the uppercase or lowercase character, as shown in figure 12.2:

```
ransom = []
for char in text:
    ransom.append(char.upper() if random.choice([False, True]) else char.lower())
```

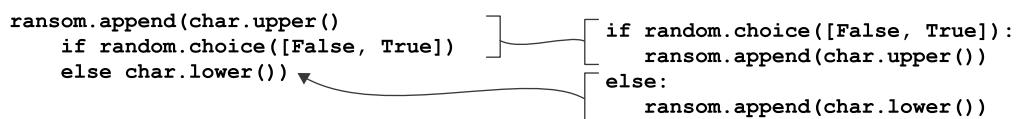


Figure 12.2 A binary if/else branch is more succinctly written using an if expression.

You don't have to use actual Boolean values (`False` and `True`). You could use 0 and 1 instead:

```
ransom = []
for char in text:
    ransom.append(char.upper() if random.choice([0, 1]) else char.lower())
```

When numbers are evaluated *in a Boolean context* (that is, in a place where Python expects to see a Boolean value), 0 is considered `False`, and every other number is `True`.

12.3.2 Writing a function to choose the letter

The `if` expression is a bit of code that could be put into a function. I find it hard to read inside the `ransom.append()`.

By putting it into a function, I can give it a descriptive name and write a test for it:

```
def choose(char):
    """Randomly choose an upper or lowercase letter to return"""

    return char.upper() if random.choice([0, 1]) else char.lower()
```

Now I can run the `test_choose()` function to test that my function does what I think. This code is much easier to read:

```
ransom = []
for char in text:
    ransom.append(choose(char))
```

12.3.3 Another way to write `list.append()`

The solution in section 12.2 creates an empty list, to which I `list.append()` the return from `choose()`. Another way to write `list.append()` is to use the `+=` operator to add the right-hand value (the element to add) to the left-hand side (the list), as in figure 12.3.

```
def main():
    args = get_args()
    random.seed(args.seed)

    ransom = []
    for char in args.text:
        ransom += choose(char)

    print(''.join(ransom))
```

`ransom.append(choose(char))`



`ransom += choose(char)`

Add the result
of the function
to `ransom`.

Figure 12.3 The `+=` operator is another way to write `list.append()`.

This is the same syntax for concatenating a character to a string or adding a number to another number.

12.3.4 Using a str instead of a list

The two previous solutions require that the lists be joined on the empty string to make a new string to print. We could, instead, start off with an empty string and build that up, one character at a time, using the `+=` operator:

```
def main():
    args = get_args()
    random.seed(args.seed)

    ransom = ''
    for char in args.text:
        ransom += choose(char)

    print(ransom)
```

As we just noted, the `+=` operator is another way to append an element to a list. Python often treats strings and lists interchangeably, often implicitly, for better or worse.

12.3.5 Using a list comprehension

The previous patterns all initialize an empty `str` or `list` and then build it up with a `for` loop. I'd like to convince you that it's almost always better to express this using a list comprehension, because its entire *raison d'être* is to return a new list. We can condense our three lines of code to just one:

```
def main():
    args = get_args()
    random.seed(args.seed)
    ransom = [choose(char) for char in args.text]
    print(''.join(ransom))
```

Or you can skip creating the `ransom` variable altogether. As a general rule, I only assign a value to a variable if I use it more than once or if I feel it makes my code more readable:

```
def main():
    args = get_args()
    random.seed(args.seed)
    print(''.join([choose(char) for char in args.text]))
```

A `for` loop is really for iterating through some sequence and producing *side effects*, like printing values or handling lines in a file. If your goal is to create a new list, a list comprehension is probably the best tool. Any code that would go into the body of the `for` loop to process an element is better placed in a function with a test.

12.3.6 Using a map() function

I've mentioned before that `map()` is just like a list comprehension, though usually with less typing. Both approaches generate a new list from some iterable, as shown in figure 12.4. In this case, the resulting list from `map()` is created by applying the `choose()` function to each character of `args.text`:

```
def main():
    args = get_args()
    random.seed(args.seed)
    ransom = map(choose, args.text)
    print(''.join(ransom))
```

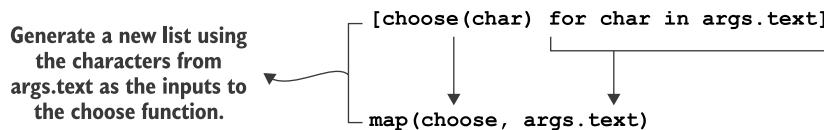


Figure 12.4 The ideas of the list comprehension can be expressed more succinctly with `map()`.

Or, again, you could leave out the `ransom` assignment and use the list that comes back from `map()` directly:

```
def main():
    args = get_args()
    random.seed(args.seed)
    print(''.join(map(choose, args.text)))
```

12.4 Comparing methods

It may seem silly to spend so much time working through so many ways to solve what is essentially a trivial problem, but one of the goals of this book is to explore the various ideas available in Python. The first solution in section 12.2 is a very imperative solution that a C or Java programmer would probably write. The version using a list comprehension is very idiomatic to Python—it is “Pythonic,” as Pythonistas would say. The `map()` solution would look very familiar to someone coming from a purely functional language like Haskell.

All these approaches accomplish the same goal, but they embody different aesthetics and programming paradigms. My preferred solution would be the last one, using `map()`, but you should choose an approach that makes the most sense to you.

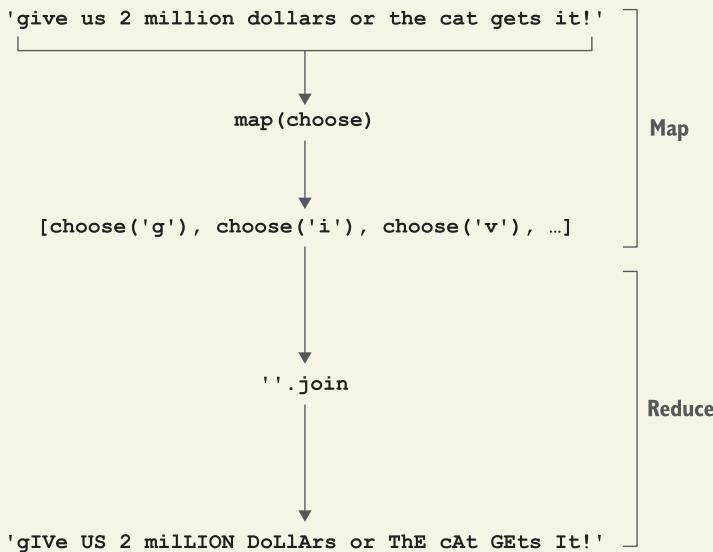
MapReduce

In 2004, Google released a paper on their “MapReduce” algorithm. The “map” phase applies some transformation to all the elements in a collection, such as all the pages of the internet that need to be indexed for searching. These operations can happen in *parallel*, meaning you can use many machines to process the pages separately from each other and in any order. The “reduce” phase then brings all the processed elements back together, maybe to put the results into a unified database.

In our `ransom.py` program, the “map” part selected a randomized case for the given letter, and the “reduce” part was putting all those bits back together into a new string. Conceivably, `map()` could make use of multiple processors to run the functions *in parallel* as opposed to *sequentially* (like with a `for` loop), possibly cutting the time to produce the results.

The ideas of map/reduce can be found in many places, from indexing the internet to our ransom program.

Learning about MapReduce was, to me, a bit like learning the name of a new bird. I never even noticed that bird before, but, once I was told its name, I saw it everywhere. Once you understand this pattern, you’ll begin to see it in many places.



12.5 Going further

Write a version of `ransom.py` that represents letters in other ways by combining ASCII characters, such as the following. Feel free to make up your own substitutions. Be sure to update your tests.

A	4	K	<
B	3	L	_

C	(M	\ /
D)	N	\
E	3	P	^
F	=	S	5
G	(-	T	+
H	-	V	\ /
I	1	W	\ / \ /
J	_		

Summary

- Whenever you have lots of things to process, try to think about how you'd process just one of them.
- Write a test that helps you imagine how you'd like to use the function to process one item. What will you pass in, and what do you expect back?
- Write your function to pass your test. Be sure to think about what you'll do with both good and bad input.
- To apply your function to each element in your input, use a for loop, a list comprehension, or a map().



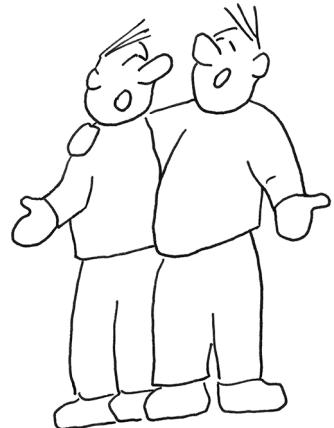
12 *Twelve Days of Christmas: Algorithm design*

Perhaps one of the worst songs of all time, and the one that is sure to ruin my Christmas spirit, is “The Twelve Days of Christmas.” WILL IT EVER STOP!? AND WHAT IS WITH ALL THE BIRDS?! Still, it’s pretty interesting to write an algorithm to generate the song starting from any given day because you have to count *up* as you add each verse (day) and then count *down* inside the verses (recapitulating the previous days’ gifts). You’ll be able to build on what you learned writing the program for “99 Bottles of Beer.”

Our program in this chapter will be called `twelve_days.py`, and it will generate the “Twelve Days of Christmas” song up to a given day, specified by the `-n` or `--num` argument (default 12). Note that there should be two newlines between verses but only one at the end:

```
$ ./twelve_days.py -n 3
On the first day of Christmas,
My true love gave to me,
A partridge in a pear tree.
```

```
On the second day of Christmas,
My true love gave to me,
Two turtle doves,
And a partridge in a pear tree.
```



On the third day of Christmas,
 My true love gave to me,
 Three French hens,
 Two turtle doves,
 And a partridge in a pear tree.

The text will be printed to STDOUT unless there is an -o or --outfile argument, in which case the text should be placed inside a file with the given name. Note that there should be 113 lines of text for the entire song:

```
$ ./twelve_days.py -o song.txt
$ wc -l song.txt
 113 song.txt
```

In this exercise, you will

- Create an algorithm to generate “The Twelve Days of Christmas” from any given day in the range 1–12
- Reverse a list
- Use the range() function
- Write text to a file or to STDOUT

13.1 Writing twelve_days.py

As always, I suggest you create your program by running new.py or by copying the template/template.py file. This one must be called twelve_days.py and live in the 13_twelve_days directory.

Your program should take two options:

- -n or --num—An int with a default of 12
- -o or --outfile—An optional filename for writing the output

For the second option, you can go back to chapter 5 to see how we handled this in the Howler solution. That program writes its blistering output to the given filename if one is supplied, and otherwise writes to sys.stdout. For this program, I suggest you declare the --outfile using type=argparse.FileType('wt') to indicate that argparse will require an argument to name a *writable text* file. If the user supplies a valid argument, args.outfile will be an *open, writable file handle*. If you also use a default of sys.stdout, you’ll have quickly handled both options of writing to a text file or STDOUT!

The only downside to this approach is that the usage statement for the program looks a little funny in describing the default for the --outfile parameter:

```
$ ./twelve_days.py -h
usage: twelve_days.py [-h] [-n days] [-o FILE]

Twelve Days of Christmas

optional arguments:
  -h, --help            show this help message and exit
```

```
-n days, --num days    Number of days to sing (default: 12)
-o FILE, --outfile FILE
                    Outfile (default: <_io.TextIOWrapper name='<stdout>' mode='w' encoding='utf-8'>)
```

Once you've completed the usage, your program should pass the first two tests.

Figure 13.1 shows a holly, jolly string diagram to get you in the mood for writing the rest of the program.

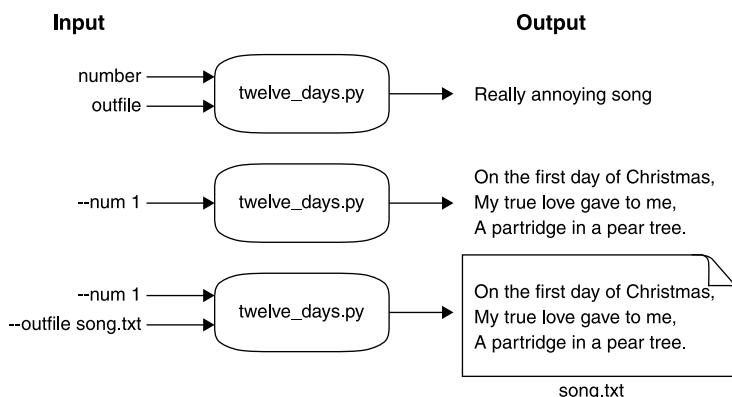


Figure 13.1 The twelve_days.py program takes options for which day to start on and an output file to write.

The program should complain if the `--num` value is not in the range 1–12. I suggest you check this inside the `get_args()` function and use `parser.error()` to halt with an error and usage message:

```
$ ./twelve_days.py -n 21
usage: twelve_days.py [-h] [-n days] [-o FILE]
twelve_days.py: error: --num "21" must be between 1 and 12
```

Once you've handled the bad `--num`, you should pass the first three tests.

13.1.1 Counting

In the “99 Bottles of Beer” song, we needed to count down from a given number. Here we need to count up to `--num` and then count back down through the gifts. The `range()` function will give us what we need, but we must remember to start at 1 because we don't start singing “On the zeroth day of Christmas.” Keep in mind that the upper bound is not included:

```
>>> num = 3
>>> list(range(1, num))
[1, 2]
```

You'll need to add 1 to whatever you're given for `--num`:

```
>>> list(range(1, num + 1))
[1, 2, 3]
```

Let's start by printing something like the first line of each verse:

```
>>> for day in range(1, num + 1):
...     print(f'On the {day} day of Christmas,')
...
On the 1 day of Christmas,
On the 2 day of Christmas,
On the 3 day of Christmas,
```

At this point, I'm starting to think about how we wrote “99 Bottles of Beer.” There we ended up creating a `verse()` function that would generate any *one* verse. Then we used `str.join()` to put them all together with two newlines. I suggest we try the same approach here, so I'll move the code inside the `for` loop into its own function:

```
def verse(day):
    """Create a verse"""
    return f'On the {day} day of Christmas,'
```

Notice that the function will not `print()` the string but will return the verse, so that we can test it:

```
>>> assert verse(1) == 'On the 1 day of Christmas,'
```

Let's see how we can use this `verse()` function:

```
>>> for day in range(1, num + 1):
...     print(verse(day))
...
On the 1 day of Christmas,
On the 2 day of Christmas,
On the 3 day of Christmas,
```

Here's a simple `test_verse()` function we could start off with:

```
def test_verse():
    """ Test verse """
    assert verse(1) == 'On the 1 day of Christmas,'
    assert verse(2) == 'On the 2 day of Christmas,'
```

This is incorrect, of course, because it should say “On the *first* day” or the “*second* day,” not “1 day” or “2 day.” Still, it's a place to start. Add the `verse()` and `test_verse()` functions to your `twelve_days.py` program, and then run `pytest twelve_days.py` to verify this much works.

13.1.2 Creating the ordinal value

Maybe the first thing to do is to change the numeric value to its ordinal position, that is “1” to “first,” “2” to “second.” You could use a dictionary like we used in “Jump The Five” to associate each `int` value 1–12 with its `str` value. That is, you might create a new dict called `ordinal`:

```
>>> ordinal = {} # what goes here?
```

Then you could do this:

```
>>> ordinal[1]  
'first'  
>>> ordinal[2]  
'second'
```

You could also use a list, if you think about how you could use each day in the `range()` to index into a list of ordinal strings.

```
>>> ordinal = [] # what goes here?
```

Your `verse()` function might look something like this now:

```
def verse(day):  
    """Create a verse"""  
    ordinal = [] # something here!  
    return f'On the {ordinal[day]} of Christmas,'
```

You can update your test with your expectations:

```
def test_verse():  
    """ Test verse """  
    assert verse(1) == 'On the first day of Christmas,'  
    assert verse(2) == 'On the second day of Christmas,'
```

Once you have this working, you should be able to replicate something like this:

```
>>> for day in range(1, num + 1):  
...     print(verse(day))  
...  
On the day first day of Christmas,  
On the day second day of Christmas,  
On the day third day of Christmas,
```

If you put the `test_verse()` function inside your `twelve_days.py` program, you can verify that your `verse()` function works by running `pytest twelve_days.py`. The `pytest` module will run any function that has a name starting with `test_`.

Shadowing

You might be tempted to use the variable name `ord`, and you would be allowed by Python to do this. The problem is that Python has a function called `ord()` that returns “the Unicode code point for a one-character string”:

```
>>> ord('a')
97
```

Python will not complain if you define a variable or another function with the name `ord`,

```
>>> ord = {}
```

such that you could do this:

```
>>> ord[1]
'first'
```

But that overwrites the actual `ord` function and so breaks a function call:

```
>>> ord('a')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'dict' object is not callable
```

This is called “shadowing,” and it’s quite dangerous. Any code in the scope of the shadowing would be affected by the change.

Tools like Pylint can help you find problems like this in your programs. Assume you have the following code:

```
$ cat shadow.py
#!/usr/bin/env python3

ord = {}
print(ord('a'))
```

Here is what Pylint has to say:

```
$ pylint shadow.py
*****
Module shadow
shadow.py:3:0: W0622: Redefining built-in 'ord' (redefined-builtin)
shadow.py:1:0: C0111: Missing module docstring (missing-docstring)
shadow.py:4:6: E1102: ord is not callable (not-callable)

-----
Your code has been rated at -25.00/10
```

It’s good to double-check your code with tools like Pylint and Flake8!

13.1.3 Making the verses

Now that we have the basic structure of the program, let's focus on creating the *correct* output. We'll update `test_verse()` with the actual values for the first two verses. You can, of course, add more tests, but presumably if we can manage the first two days, we can handle all the other days:

```
def test_verse():
    """Test verse"""

    assert verse(1) == '\n'.join([
        'On the first day of Christmas,', 'My true love gave to me,',
        'A partridge in a pear tree.'
    ])

    assert verse(2) == '\n'.join([
        'On the second day of Christmas,', 'My true love gave to me,',
        'Two turtle doves,', 'And a partridge in a pear tree.'
    ])
```

If you add this to your `twelve_days.py` program, you can run `pytest twelve_days.py` to see how your `verse()` function is failing:

```
===== FAILURES =====
test_verse
```

```
def test_verse():
    """Test verse"""

>     assert verse(1) == '\n'.join([
        'On the first day of Christmas,', 'My true love gave to me,',
        'A partridge in a pear tree.'
    ])
E     AssertionError: assert 'On the first...of Christmas,' == 'On the first
... a pear tree.'
E         - On the first day of Christmas,
E         + On the first day of Christmas,
E         ?
E         + My true love gave to me,
E         + A partridge in a pear tree.
```

The leading > shows that this is the code that is creating an exception. We are running `verse(1)` and asking if it's equal to the expected verse.

This is the text that `verse(1)` actually produced, which is only the first line of the verse.

The lines following are what was expected.

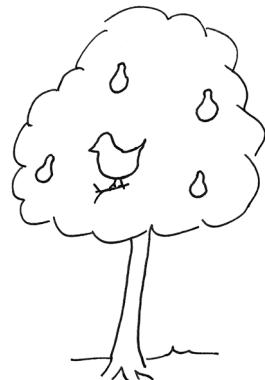
```
twelve_days.py:88: AssertionError
===== 1 failed in 0.11 seconds =====
```

Now we need to supply the rest of the lines for each verse. They all start off the same:

```
On the {ordinal}[day] } day of Christmas,
My true love gave to me,
```

Then we need to add these gifts for each day:

- 1** A partridge in a pear tree
- 2** Two turtle doves
- 3** Three French hens
- 4** Four calling birds
- 5** Five gold rings
- 6** Six geese a laying
- 7** Seven swans a swimming
- 8** Eight maids a milking
- 9** Nine ladies dancing
- 10** Ten lords a leaping
- 11** Eleven pipers piping
- 12** Twelve drummers drumming



Note that for every day greater than 1, the last line changes “A partridge...” to “And a partridge in a pear tree.”

Each verse needs to count backwards from the given day. For example, if the day is 3, then the verse lists

- 1** Three French hens
- 2** Two turtle doves
- 3** And a partridge in a pear tree

We talked in chapter 3 about how you can reverse a list, either with the `list.reverse()` method or the `reversed()` function. We also used these ideas in chapter 11 to get the bottles of beer off the wall, so this code should not be unfamiliar:

```
>>> day = 3
>>> for n in reversed(range(1, day + 1)):
...     print(n)
...
3
2
1
```

Try to make the function return the first two lines and then the countdown of the days:

```
>>> print(verse(3))
On the third day of Christmas,
My true love gave to me,
3
2
1
```

Then, instead of 3 2 1, add the actual gifts:

```
>>> print(verse(3))
On the third day of Christmas,
My true love gave to me,
```

```
Three French hens,
Two turtle doves,
And a partridge in a pear tree.
```

If you can get that to work, you ought to be able to pass the `test_verse()` test.

13.1.4 Using the `verse()` function

Once you have that working, think about a final structure that calls your `verse()`. It could be a `for` loop:

```
verses = []
for day in range(1, args.num + 1):
    verses.append(verse(day))
```

Since we're trying to create a list of the verses, a list comprehension is a better choice:

```
verses = [verse(day) for day in range(1, args.num + 1)]
```

Or it could be a `map()`:

```
verses = map(verse, range(1, args.num + 1))
```



13.1.5 Printing

Once you have all the verses, you can use the `str.join()` method to print the output. The default is to print this to “standard out” (`STDOUT`), but the program will also take an optional `--outfile` that names a file to write the output to. You can copy exactly what we did in chapter 5, but it's really worth your time to learn how to declare output files using `type=argparse.FileType('wt')`. You can even set the default to `sys.stdout` so that you'll never have to open() the output file yourself!

13.1.6 Time to write

It's not at all mandatory that you solve the problem the way that I describe. The “correct” solution is one that you write and understand and that passes the test suite. It's fine if you like the idea of creating a function for `verse()` and using the provided test. It's also fine if you want to go another way, but do try to think of writing small functions *and tests* to solve small parts of your problem, and then combining them to solve the larger problem.

If you need more than one sitting or even several days to pass the tests, take your time. Sometimes a good walk or a nap can do wonders for solving problems. Don't neglect your hammock¹ or a nice cup of tea.

¹ Search the internet for the talk “Hammock Driven Development” by Rich Hickey, the creator of the Clojure language.

13.2 Solution

A person would receive almost 200 birds in this song! Anyway, here is a solution that uses `map()`. After that you'll see versions that use `for` and list comprehensions.

```
#!/usr/bin/env python3
"""Twelve Days of Christmas"""

import argparse
import sys

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Twelve Days of Christmas',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('-n',           ←
                       '--num',
                       help='Number of days to sing',
                       metavar='days',
                       type=int,
                       default=12)           ← The --num option is an
                                                int with a default of 12.

    parser.add_argument('-o',           ←
                       '--outfile',
                       help='Outfile',
                       metavar='FILE',
                       type=argparse.FileType('wt'),
                       default=sys.stdout)   ← The --outfile option is a
                                                type=argparse.FileType('wt') with a
                                                default of sys.stdout. If the user supplies
                                                a value, it must be the name of a writable
                                                file, in which case argparse will open the
                                                file for writing.

    args = parser.parse_args()       ← Capture the results of parsing the command-
                                    ← line arguments into the args variable.

    if args.num not in range(1, 13):
        parser.error(f"--num '{args.num}' must be between 1 and 12")   ←

    return args

Check that
the given
args.num is
in the
allowed
range 1–12,
inclusive.   ←

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()           ←
    verses = map(verse, range(1, args.num + 1))
    print('\n\n'.join(verses), file=args.outfile)   ← Join the verses on two newlines
                                                    ← and print to args.outfile, which is
                                                    ← an open file handle, or sys.stdout.

    # -----
def verse(day):
    """Create a verse"""\n
                                                ← Define a function to create any one
                                                ← verse from a given number.

If args.num is invalid, use
parser.error() to print a
short usage statement and
the error message to STDERR
and exit the program with an
error value. Note that the
error message includes the
bad value for the user and
explicitly states that a good
value should be in the
range 1–12.
```

```

The ordinal values is a list of str values.    ordinal = [
    'first', 'second', 'third', 'fourth', 'fifth', 'sixth', 'seventh',
    'eighth', 'ninth', 'tenth', 'eleventh', 'twelfth'
]

The gifts for the days is a list of str values.  gifts = [
    'A partridge in a pear tree.', 'Two turtle doves,', 'Three French hens,', 'Four calling birds,', 'Five gold rings,', 'Six geese a laying,', 'Seven swans a swimming,', 'Eight maids a milking,', 'Nine ladies dancing,', 'Ten lords a leaping,', 'Eleven pipers piping,', 'Twelve drummers drumming,', ]
]

Check if this is for a day greater than 1.  lines = [
    f'On the {ordinal[day - 1]} day of Christmas,', 'My true love gave to me,'
]
lines.extend(reversed(gifts[:day]))  ← Use the list.extend() method to add the gifts, which are a slice from the given day and then reversed().

Return the lines joined on the newline.  if day > 1:
    lines[-1] = 'And ' + lines[-1].lower()  ← Change the last of the lines to add "And " to the beginning, appended to the lowercased version of the line.

# -----  return '\n'.join(lines)  ← The unit test for the verse() function

def test_verse():
    """Test verse"""
    assert verse(1) == '\n'.join([
        'On the first day of Christmas,', 'My true love gave to me,', 'A partridge in a pear tree.'
    ])

    assert verse(2) == '\n'.join([
        'On the second day of Christmas,', 'My true love gave to me,', 'Two turtle doves,', 'And a partridge in a pear tree.'
    ])

# -----  if __name__ == '__main__':
    main()  ←

```

13.3 Discussion

Not much in `get_args()` is new, so we'll throw it a sidelong, cursory glance. The `--num` option is an `int` value with a default value of 12, and we use `parser.error()` to halt the program if the user provides a bad value. The `--outfile` option is a bit different, though, as we're declaring it with `type=argparse.FileType('wt')` to indicate the value must be a writable file. This means the value we get from `argparse` will be an open, writable file. We set the default to `sys.stdout`, which is also an open, writable file, so we've handled the two output options entirely through `argparse`, which is a real time saver!

13.3.1 Making one verse

I chose to make a function called `verse()` to create any one verse given an `int` value of the day:

```
def verse(day):
    """Create a verse"""
```

I decided to use a list to represent the ordinal value of the day:

```
ordinal = [
    'first', 'second', 'third', 'fourth', 'fifth', 'sixth', 'seventh',
    'eighth', 'ninth', 'tenth', 'eleventh', 'twelfth'
]
```

Since the day is based on counting from 1, but Python lists start from 0 (see figure 13.2), I have to subtract 1:

```
>>> day = 3
>>> ordinal[day - 1]
'third'
```

	Index	Day
ordinal = [
'first'	0	1
'second'	1	2
'third'	2	3
'fourth'	3	4
'fifth'	4	5
'sixth'	5	6
'seventh'	6	7
'eighth'	7	8
'ninth'	8	9
'tenth'	9	10
'eleventh'	10	11
'twelfth'	11	12

Figure 13.2 Our days start counting from 1, but Python indexes from 0.

I could just as easily have used a dict:

```
ordinal = {
    1: 'first', 2: 'second', 3: 'third', 4: 'fourth',
    5: 'fifth', 6: 'sixth', 7: 'seventh', 8: 'eighth',
    9: 'ninth', 10: 'tenth', 11: 'eleventh', 12: 'twelfth',
}
```

In this case you don't have to subtract 1. Whatever works for you:

```
>>> ordinal[3]
'third'
```

I also used a list for the gifts:

```
gifts = [
    'A partridge in a pear tree.',
    'Two turtle doves,',
    'Three French hens,',
    'Four calling birds,',
    'Five gold rings,',
    'Six geese a laying,',
    'Seven swans a swimming,',
    'Eight maids a milking,',
    'Nine ladies dancing,',
    'Ten lords a leaping,',
    'Eleven pipers piping,',
    'Twelve drummers drumming,',
]
```

This makes a bit more sense, as I can use a list slice to get the gifts for a given day (see figure 13.3):

```
>>> gifts[:3]
['A partridge in a pear tree.',
 'Two turtle doves,',
 'Three French hens,']
```

```
gifts[:3]
[ 'A partridge in a pear tree.', 0
  'Two turtle doves,', 1
  'Three French hens,', 2
  'Four calling birds,', 3
  'Five gold rings,', 4
  'Six geese a laying,', 5
  'Seven swans a swimming,', 6
  'Eight maids a milking,', 7
  'Nine ladies dancing,', 8
  'Ten lords a leaping,', 9
  'Eleven pipers piping,', 10
  'Twelve drummers drumming,' ] 11
```



Figure 13.3 The gifts are listed by their days in ascending order.

But I want them in reverse order. The `reversed()` function is lazy, so I need to use the `list()` function in the REPL to coerce the values:

```
>>> list(reversed(gifts[:3]))
['Three French hens,',
 'Two turtle doves,',
 'A partridge in a pear tree.]
```

The first two lines of any verse are the same, substituting in the ordinal value for the day:

```
lines = [
    f'On the {ordinal[day - 1]} day of Christmas,',
    'My true love gave to me,'
]
```

I need to put these two lines together with the gifts. Since each verse is made of some number of lines, I think it will make sense to use a list to represent the entire verse.

I need to add the gifts to the lines, and I can use the `list.extend()` method to do that:

```
>>> lines.extend(reversed(gifts[:day]))
```

Now there are five lines:

```
>>> lines
['On the third day of Christmas,',
 'My true love gave to me,',
 'Three French hens,',
 'Two turtle doves,',
 'A partridge in a pear tree.']
>>> assert len(lines) == 5
```

Note that I cannot use the `list.append()` method. It's easy to confuse it with the `list.extend()` method, which takes another list as its argument, expands it, and adds all of the individual elements to the original list. The `list.append()` method is meant to add *just one* element to a list, so if you give it a list, it will tack that entire list onto the end of the original list!

Here the `reversed()` iterator will be added to the end of `lines`, such that it would have three elements rather than the desired five:

```
>>> lines.append(reversed(gifts[:day]))
>>> lines
['On the third day of Christmas,',
 'My true love gave to me,',
 <list_reverseiterator object at 0x105bc8588>]
```

Maybe you're thinking you could coerce `reversed()` with the `list()` function? Thinking you are, young Jedi, but, alas, that will still add a new list to the end:

```
>>> lines.append(list(reversed(gifts[:day])))
>>> lines
```

```
[ 'On the third day of Christmas,',
  'My true love gave to me,',
  ['Three French hens,', 'Two turtle doves,', 'A partridge in a pear tree.']]
```

And we still have three lines rather than five:

```
>>> len(lines)
3
```

If day is greater than 1, I need to change the last line to say “And a” instead of “A”:

```
if day > 1:
    lines[-1] = 'And ' + lines[-1].lower()
```

Note that this is another good reason to represent the lines as a list, because the elements of a list are *mutable*. I could have represented the lines as a str, but strings are *immutable*, so it would be much harder to change the last line.

I want to return a single str value from the function, so I join the lines on a newline:

```
>>> print('\n'.join(lines))
On the third day of Christmas,
My true love gave to me,
Three French hens,
Two turtle doves,
A partridge in a pear tree.
```

My function returns the joined lines and will pass the `test_verse()` function I provided.

13.3.2 Generating the verses

Given the `verse()` function, I can create all the needed verses by iterating from 1 to the given `--num`. I could collect them in a list of verses:

```
day = 3
verses = []
for n in range(1, day + 1):
    verses.append(verse(n))
```

I can test that I have the right number of verses:

```
>>> assert len(verses) == day
```

Whenever you see this pattern of creating an empty str or list and then using a for loop to add to it, consider instead using a list comprehension:

```
>>> verses = [verse(n) for n in range(1, day + 1)]
>>> assert len(verses) == day
```

I personally prefer using `map()` over list comprehensions. See figure 13.4 to review how the three methods fit together. I need to use the `list()` function to coerce the lazy `map()` function in the REPL, but it's not necessary in the program code:

```
>>> verses = list(map(verse, range(1, day + 1)))
>>> assert len(verses) == day
```

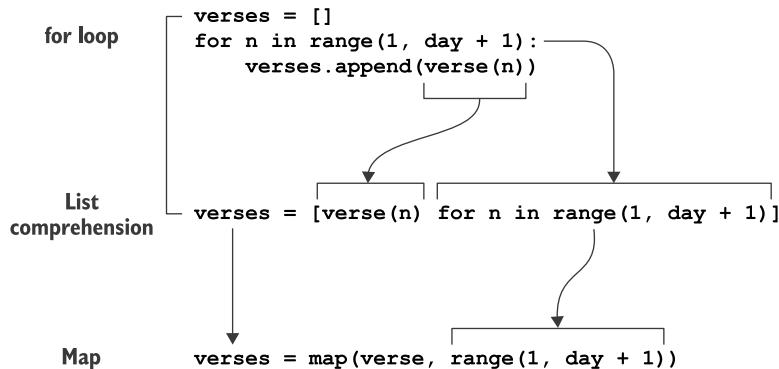


Figure 13.4 Building a list using a `for` loop, a list comprehension, and `map()`.

All of these methods will produce the correct number of verses. Choose whichever one makes the most sense to you.

13.3.3 Printing the verses

Just like with “99 Bottles of Beer” in chapter 11, I want to `print()` the verses with two newlines in between. The `str.join()` method is a good choice:

```
>>> print('\n\n'.join(verses))
On the first day of Christmas,
My true love gave to me,
A partridge in a pear tree.
```

```
On the second day of Christmas,
My true love gave to me,
Two turtle doves,
And a partridge in a pear tree.
```

```
On the third day of Christmas,
My true love gave to me,
Three French hens,
Two turtle doves,
And a partridge in a pear tree.
```

You can use the `print()` function with the optional `file` argument to put the text into an open file handle. The `args.outfile` value will be either the file indicated by the user or `sys.stdout`:

```
print('\n\n'.join(verses), file=args.outfile)
```

Or you can use the `fh.write()` method, but you need to remember to add the trailing newline that `print()` adds for you:

```
args.outfile.write('\n\n'.join(verses) + '\n')
```

There are dozens to hundreds of ways to write this algorithm, just as there are for “99 Bottles of Beer.” If you came up with an entirely different approach that passed the test, that’s terrific! Please share it with me. I wanted to stress the idea of how to write, test, and use a single `verse()` function, but I’d love to see other approaches!



13.4 Going further

Install the `emoji` module (<https://pypi.org/project/emoji/>) and print various emojis for the gifts rather than text. For instance, you could use `:bird:` to print `🐦` for every bird, like a hen or dove. I also used `:man:`, `:woman:`, and `:drum:`, but you can use whatever you like:

```
On the twelfth day of Christmas,  
My true love gave to me,  
Twelve 🎅s drumming,  
Eleven 🎅s piping,  
Ten 🎅s a leaping,  
Nine 🎅s dancing,  
Eight 🐄s a milking,  
Seven 🐓s a swimming,  
Six 🐓s a laying,  
Five gold 💰s,  
Four calling 🐦s,  
Three French 🐓s,  
Two turtle 🐢s,  
And a 🐦 in a pear tree.
```

Summary

- There are many ways to encode algorithms to perform repetitive tasks. In my version, I wrote and tested a function to handle one task and then mapped a range of input values over that.
- The `range()` function will return `int` values between given start and stop values, the latter of which is not included.
- You can use the `reversed()` function to reverse the values returned by `range()`.

- If you use `type=argparse.FileType('wt')` to define an argument with argparse, you get a file handle that is open for writing text.
- The `sys.stdout` file handle is always open and available for writing.
- Modeling gifts as a list allowed me to use a list slice to get all the gifts for a given day. I used the `reversed()` function to put them into the right order for the song.
- I modeled `lines` as a list because a list is mutable, which I needed in order to change the last line when the day is greater than 1.
- Shadowing a variable or function is reusing an existing variable or function name. If, for instance, you create a variable with the name of an existing function, that function is effectively hidden because of the shadow. Avoid shadowing by using tools like Pylint to find these and many other common coding problems.

Rhymer: Using regular expressions to create rhyming words

In the movie *The Princess Bride*, the characters Inigo and Fezzik have a rhyming game they like to play, especially when their cruel boss, Vizzini, yells at them:

Inigo: That Vizzini, he can fuss.

Fezzik: I think he likes to scream at us.

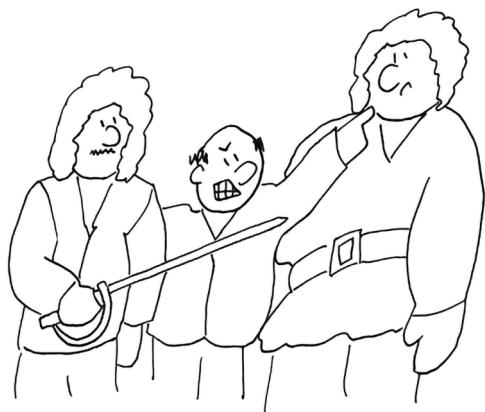
Inigo: Probably he means no harm.

Fezzik: He's really very short on charm.

When I was writing the alternate.txt for chapter 7, I would come up with a word like “cyanide” and wonder what I could rhyme with that. Mentally I started with the first consonant sound of the alphabet and substituted “b” for “byanide,” skipped “c” because that’s already the first character, then “d” for “dyanide,” and so forth. This is effective but tedious, so I decided to write a program to do this for me, as one does.

This is basically another find-and-replace type of program, like swapping all the numbers in a string in chapter 4 or all the vowels in a string in chapter 8. We wrote those programs using very manual, *imperative* methods, like iterating through all the characters of a string, comparing them to some wanted value, and possibly returning a new value.

In the final solution for chapter 8, we briefly touched on “regular expressions” (also called “regexes”—pronounced with a soft “g” like in “George”), which give us



a *declarative* way to describe patterns of text. The material here may seem a bit of a reach, but I really want to help you dig into regexes to see what they can do.

In this chapter, we’re going to take a given word and create “words” that rhyme. For instance, the word “bake” rhymes with words like “cake,” “make,” and “thrake,” the last of which isn’t actually a dictionary word but just a new string I created by replacing the “b” in “bake” with “thr.”

The algorithm we’ll use will split a word into any initial consonants and the rest of the word, so “bake” is split into “b” and “ake.” We’ll replace the “b” with all the other consonants from the alphabet plus these consonant clusters:

```
bl br ch cl cr dr fl fr gl gr pl pr sc sh sk sl sm sn sp st
sw th tr tw thw wh wr sch scr shr sph spl spr squ str thr
```



These are the first three words our program will produce for “cake”:

```
$ ./rhymers.py cake | head -3
bake
blake
brake
```

And these are the last three:

```
$ ./rhymers.py cake | tail -3
xake
yake
zake
```

Make sure your output is sorted alphabetically as this is important for the tests.

We’ll replace any leading consonants with a list of other consonant sounds to create a total of 56 words:

```
$ ./rhymers.py cake | wc -l
56
```

Note that we’ll replace *all* the leading consonants, not just the first one. For instance, with the word “chair” we need to replace “ch”:

```
$ ./rhymers.py chair | tail -3
xair
yair
zair
```

If a word like “apple” does not start with a consonant, we’ll append all the consonant sounds to the beginning to create words like “bapple” and “shrapple.”

```
$ ./rhymers.py apple | head -3
bapple
blapple
brapple
```

Because there is no consonant to *replace*, words that start with a vowel will produce 57 rhyming words:

```
$ ./rhymers.py apple | wc -l
57
```

To make this a bit easier, the output should always be all lowercase, even if the input has uppercase letters:

```
$ ./rhymers.py GUITAR | tail -3
xuitar
yuitar
zuitar
```

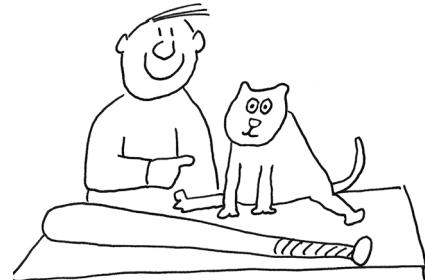
If a word contains nothing but consonants, we'll print a message stating that the word cannot be rhymed:

```
$ ./rhymers.py RDNZL
Cannot rhyme "RDNZL"
```

The task of finding the initial consonants is made significantly easier with regexes.

In this program, you will

- Learn to write and use regular expressions
- Use a guard with a list comprehension
- Explore the similarities of list comprehension with a guard to the `filter()` function
- Entertain ideas of “truthiness” when evaluating Python types in a Boolean context



14.1 Writing rhymers.py

The program takes a single, positional argument, which is the string to rhyme. Figure 14.1 shows a snazzy, jazzy, frazzy, thwazzy string diagram.

If given no arguments or the `-h` or `--help` flags, it should print a usage statement:

```
$ ./rhymers.py -h
usage: rhymers.py [-h] word

Make rhyming "words"

positional arguments:
  word        A word to rhyme

optional arguments:
  -h, --help  show this help message and exit
```

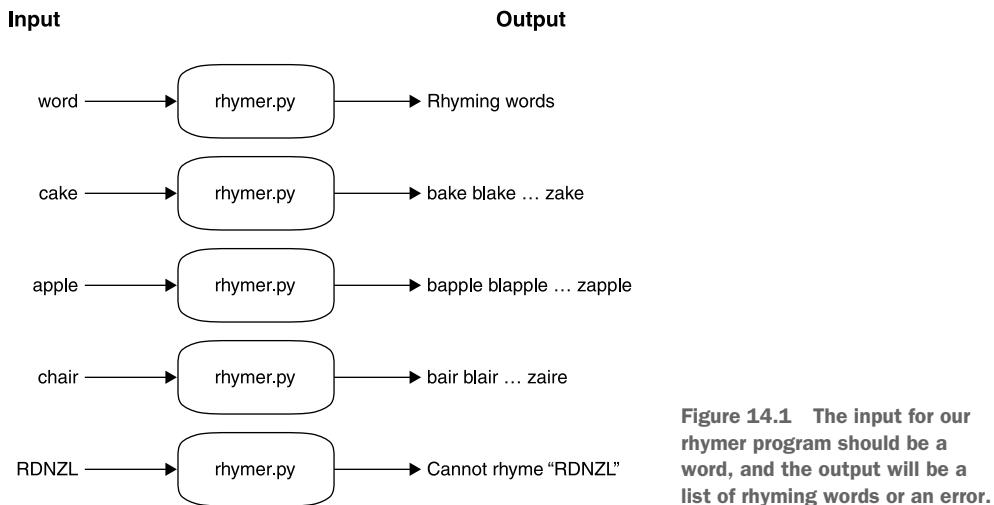


Figure 14.1 The input for our rhyming program should be a word, and the output will be a list of rhyming words or an error.

14.1.1 Breaking a word

To my mind, the main problem of the program is breaking the given word into the leading consonant sounds and the rest—something like the “stem” of the word.

To start out, we can define a placeholder for a function I call `stemmer()` that does nothing right now:

```
def stemmer():
    """Return leading consonants (if any), and 'stem' of word"""
    pass
```

The `pass` statement will do nothing at all. Since the function does not return a value, Python will return `None` by default.

Then we can define a `test_stemmer()` function to help us think about the values we might give the function and what we expect it to return. We want a test with good values like “cake” and “apple” that can be rhymed as well as values like the empty string or a number, which cannot:

```
def test_stemmer():
    """ Test stemmer """
    assert stemmer('') == ('', '')
    ② assert stemmer('cake') == ('c', 'ake')
    assert stemmer('chair') == ('ch', 'air') ③
    ④ assert stemmer('APPLE') == ('', 'apple')
    assert stemmer('RDNZL') == ('rdnzl', '') ⑤
    ⑥ assert stemmer('123') == ('123', '')
```

The tests cover the following good and bad inputs:

- ① The empty string
- ② A word with a single leading consonant

- ③ A word with a leading consonant cluster
- ④ A word with no initial consonants; also an uppercase word, so this checks that lowercase is returned
- ⑤ A word with no vowels
- ⑥ Something that isn't a word at all

I decided that my `stemmer()` function will always return a 2-tuple of the `(start, rest)` of the word. (You can write a function that does something different, but be sure to change the test to match.) It's the second part of that tuple—the `rest`—that we can use to create rhyming words. For instance, the word “cake” produces a tuple with `('c', 'ake')`, and “chair” is split into `('ch', 'air')`. The argument “APPLE” has no start and only the rest of the word, which is lowercase.

When I'm writing tests, I usually try to provide both good and bad data to my functions and programs. Three of the test values cannot be rhymed: the empty string (''), a string with no vowels ('RDNZL'), and a string with no letters ('123'). The `stemmer()` function will still return a tuple containing the lowercased word in the first position of the tuples and the empty string in the second position for the rest of the word. It is up to the calling code to deal with a word that has no part that can be used to rhyme.

14.1.2 Using regular expressions

It's certainly *possible* to write this program without regular expressions, but I hope you'll see how radically different using regexes can be from manually writing your own search-and-replace code.

To start off, we need to bring in the `re` module:

```
>>> import re
```

I encourage you to read `help(re)` to get a feel for all that you can do with regexes. They are a deep subject with many books and whole branches of academia devoted to them (*Mastering Regular Expressions* by Jeffrey Friedl (O'Reilly, 2006) is one book I would recommend). There are many helpful websites that can further explain regexes, and some can help you write them (such as <https://regexr.com/>). We will only scratch the surface of what you can do with regexes.

Our goal in this program is to write a regex that will find consonants at the beginning of a string. We can define consonants as the characters of the English alphabet that are not vowels (“a,” “e,” “i,” “o,” and “u”). Our `stemmer()` function will only return lowercase letters, so there are only 21 consonants we need to define. You could write them out, but I'd rather write a bit of code!

We can start with `string.ascii_lowercase`:

```
>>> import string
>>> string.ascii_lowercase
'abcdefghijklmnopqrstuvwxyz'
```

Next, we can use a list comprehension with a “guard” clause to filter out the vowels. As we want a str of consonants and not a list, we can use str.join() to make a new str value:

```
>>> import string as s
>>> s.ascii_lowercase
'abcdefghijklmnopqrstuvwxyz'
>>> consonants = ''.join([c for c in s.ascii_lowercase if c not in 'aeiou'])
>>> consonants
'bcdfghjklmnpqrstvwxyz'
```

The longer way to write this with a for loop and an if statement is as follows (see figure 14.2):

```
consonants = ''
for c in string.ascii_lowercase:
    if c not in 'aeiou':
        consonants += c
```

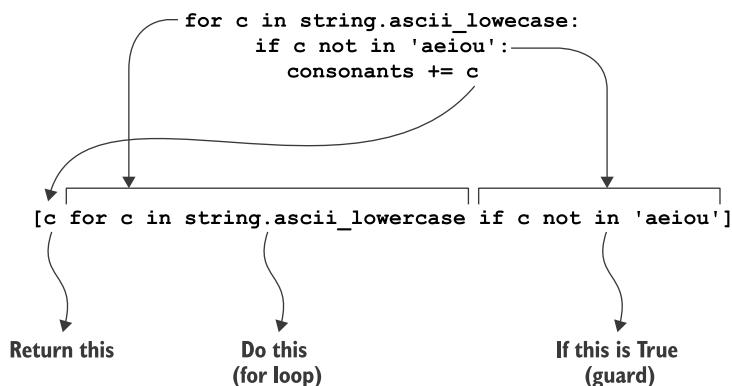


Figure 14.2 The for loop (top) can be written as a list comprehension (bottom). This list comprehension includes a guard so that only consonants are selected, which is like the if statement at the top.

In chapter 8 we created a “character class” for matching the vowels by listing them in square brackets, like '[aeiou]'. We can do the same here with our consonants, like so:

```
>>> pattern = '[' + consonants + ']'
>>> pattern
'[bcdfghjklmnpqrstvwxyz]'
```

The re module has two search-like functions called `re.match()` and `re.search()`, and I always get them confused. They both look for a pattern (the first argument) in some text, but the `re.match()` function starts *from the beginning* of the text, whereas the `re.search()` function will match starting *anywhere* in the text.

As it happens, `re.match()` is just fine because we are looking for consonants at the beginning of a string (see figure 14.3).

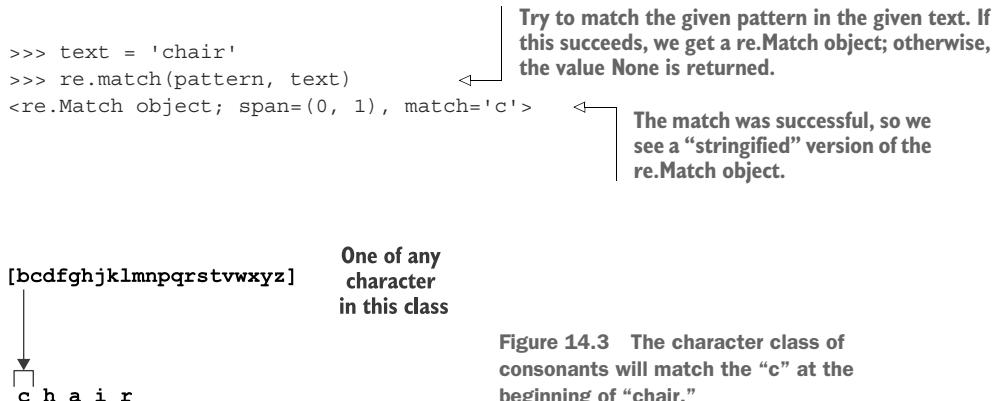


Figure 14.3 The character class of consonants will match the “c” at the beginning of “chair.”

The `match='c'` shows us that the regular expression found the string '`c`' at the beginning. Both the `re.match()` and `re.search()` functions will return a `re.Match` object on success. You can read `help(re.Match)` to learn more about all the cool things you can do with them:

```
>>> match = re.match(pattern, text)
>>> type(match)
<class 're.Match'>
```

How do we get our regex to match the letters '`ch`'? We can put a '+' sign after the character class to say we want *one or more* (see figure 14.4). (Does this sound a bit like `nargs='+'` to say one or more arguments?) I will use an f-string here to create the pattern:

```
>>> re.match(f'{consonants}+', 'chair')
<re.Match object; span=(0, 2), match='ch'>
```

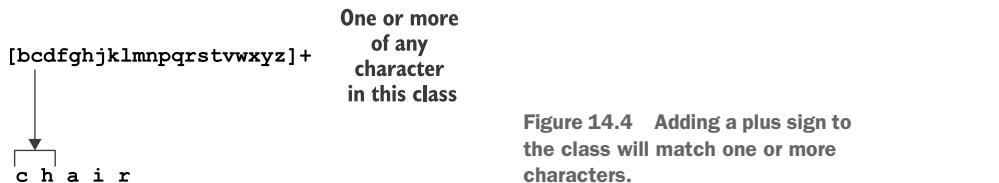


Figure 14.4 Adding a plus sign to the class will match one or more characters.

What does it give us for a string with no leading consonants like “apple,” as in figure 14.5?

```
>>> re.match(f'{consonants}+', 'apple')
```

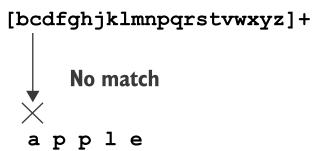


Figure 14.5 This regex fails to match a word that does not start with a consonant.

It seems we got nothing back from that. What is the `type()` of that return value?

```
>>> type(re.match(f'[{consonants}]+', 'apple'))
<class 'NoneType'>
```

Both the `re.match()` and `re.search()` functions return `None` to indicate a failure to match any text. We know that only some words will have a leading consonant sound, so this is not surprising. We'll see in a moment how to make this an optional match.

14.1.3 Using capture groups

It's all well and good to have found (or not) the leading consonants, but the goal here is to split the text into two parts: the consonants (if any) and the rest of the word.

We can wrap parts of the regex in parentheses to create “capture groups.” If the regex matches successfully, we can recover the parts using the `re.Match.groups()` method (see figure 14.6):

```
>>> match = re.match(f'([{consonants}]+)', 'chair')
>>> match.groups()
('ch',)
```

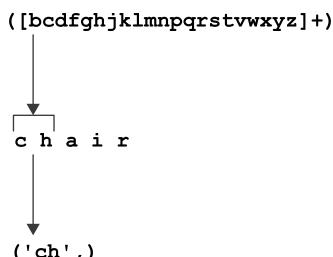


Figure 14.6 Adding parentheses around a pattern causes the matching text to be available as a capture group.

To capture everything that comes after the consonants, we can use a period (.) to match anything, and add a plus sign (+) to mean one or more. We can put that into parentheses to capture it (see figure 14.7):

```
>>> match = re.match(f'([{consonants}]+)(.+)', 'chair')
>>> match.groups()
('ch', 'air')
```

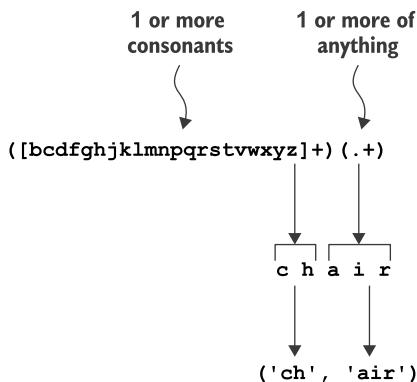


Figure 14.7 We define two capture groups to access the leading consonant sound and whatever follows.

What happens when we try to use this on “apple”? It fails to make the first match on the consonants, so *the whole match fails* and returns None (see figure 14.8):

```

>>> match = re.match(f'(\{consonants\}+) (.+)', 'apple')
>>> match.groups()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'NoneType' object has no attribute 'groups'
  
```

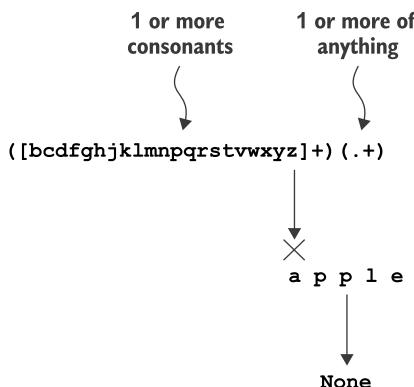


Figure 14.8 The pattern still fails when the text starts with a vowel.

Remember that `re.match()` returns `None` when it fails to find the pattern. We can add a question mark (?) at the end of the consonants pattern to make it optional (see figure 14.9):

```

>>> match = re.match(f'(\{consonants\}+)? (.+)', 'apple')
>>> match.groups()
(None, 'apple')
  
```

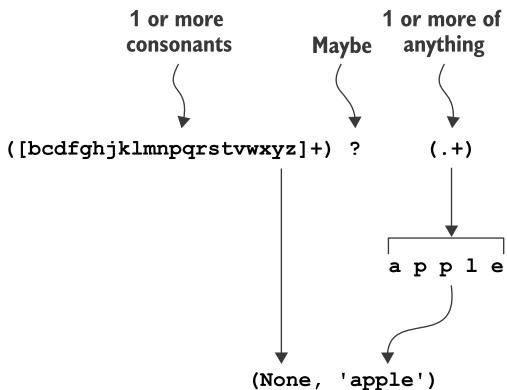


Figure 14.9 A question mark after a pattern makes it optional.

The `match.groups()` function returns a tuple containing the matches for each grouping created by the parentheses. You can also use `match.group()` (singular) with a group number to get a specific group. Note that these start numbering from 1:

```

>>> match.group(1)      There was no match for the first
                        group on "apple," so this is a None.
>>> match.group(2)      The second group captured
'apple'                  the entire word.

```

If you match on “chair,” there are values for both groups:

```

>>> match = re.match(f'(\{{consonants}\}+)?(.+)', 'chair')
>>> match.group(1)
'ch'
>>> match.group(2)
'air'

```

So far we’ve only dealt with lowercase text because our program will always emit lowercase values. Still, let’s explore what happens when we try to match uppercase text:

```

>>> match = re.match(f'(\{{consonants}\}+)?(.+)', 'CHAIR')
>>> match.groups()
(None, 'CHAIR')

```

Not surprisingly, that fails. Our pattern only defines lowercase characters. We could add all the uppercase consonants, but it’s a bit easier to use a third optional argument to `re.match()` to specify case-insensitive searching:

```

>>> match = re.match(f'(\{{consonants}\}+)?(.+)', 'CHAIR', re.IGNORECASE)
>>> match.groups()
('CH', 'AIR')

```

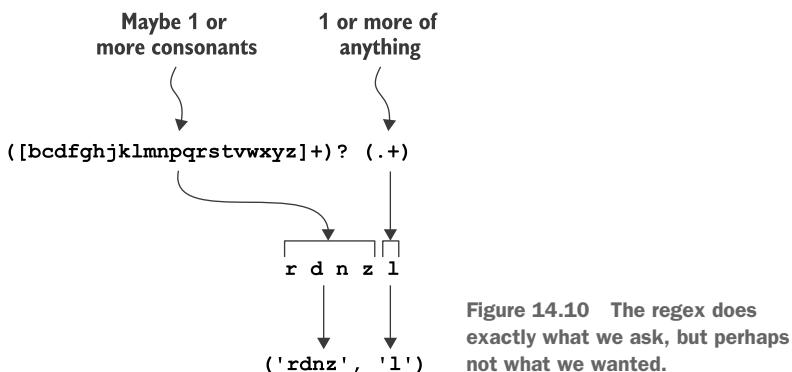
Or you can force the text you are searching to lowercase:

```
>>> match = re.match(f'(\{{consonants}\}+)?(.+)', 'CHAIR'.lower())
>>> match.groups()
('ch', 'air')
```

What do you get when you search on text that has nothing but consonants?

```
>>> match = re.match(f'(\{{consonants}\}+)?(.+)', 'rdnzl')
>>> match.groups()
('rdnz', 'l')
```

Were you expecting the first group to include *all* the consonants and the second group to have nothing? It might seem a bit odd that it decided to split off the “l” into the last group, as shown in figure 14.10, but we have to think *extremely literally* about how the regex engine is working. We described an optional group of one or more consonants that *must be followed* by one or more of anything else. The “l” counts as one or more of anything else, so the regex matched exactly what we requested.



If we change the `(.+)` to `(.*)` to make it *zero or more*, it works as expected:

```
>>> match = re.match(f'(\{{consonants}\}+)?(.*)', 'rdnzl')
>>> match.groups()
('rdnzl', '')
```

Our regex is not quite complete, as it doesn't handle matching on something like 123. That is, it matches too well because the period (.) will match the digits, which we don't want:

```
>>> re.match(f'(\{{consonants}\}+)?(.*)?', '123')
<re.Match object; span=(0, 3), match='123'>
```

We need to indicate that there should be *at least one vowel* after the consonants, which may be followed by anything else. We can use another character class to describe any

vowel. Since we need to capture this, we'll put it in parentheses, so `([aeiou])`. That may be followed by *zero or more* of anything, which also needs to be captured, so `(.*)`, as shown in figure 14.11.

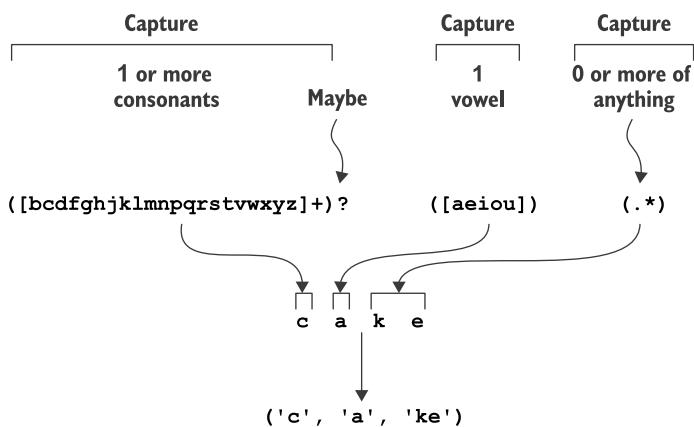


Figure 14.11 The regex now requires the presence of a vowel.

Let's go back and try this on values we expect to work:

```

>>> re.match(f'(\{consonants\}+)?([aeiou])(.*)', 'cake').groups()
('c', 'a', 'ke')
>>> re.match(f'(\{consonants\}+)?([aeiou])(.*)', 'chair').groups()
('ch', 'a', 'ir')
>>> re.match(f'(\{consonants\}+)?([aeiou])(.*)', 'apple').groups()
(None, 'a', 'pple')
  
```

As you can see, this fails to match when the string contains no vowels or letters:

```

>>> type(re.match(f'(\{consonants\}+)?([aeiou])(.*)', 'rdnzl'))
<class 'NoneType'>
>>> type(re.match(f'(\{consonants\}+)?([aeiou])(.*)', '123'))
<class 'NoneType'>
  
```

14.1.4 Truthiness

We know that our program will receive some inputs that cannot be rhymed, so what should the `stemmer()` function do with these? Some people like to use exceptions in cases like this. We've encountered exceptions like asking for a list index or a dictionary key that does not exist. If exceptions are not caught and handled, they cause our programs to crash!

I try to avoid writing code that creates exceptions. I decided that my `stemmer()` function would always return a 2-tuple of `(start, rest)`, and that I would always use

the empty string to denote a missing value rather than a `None`. Here is one way I could write the code for returning those tuples:

```

The match will be None if the regex
failed, which is "falsey." If it succeeds,
then it will be "truthy."
if match:
    p1 = match.group(1) or ''
    p2 = match.group(2) or ''
    p3 = match.group(3) or ''
    return (p1, p2 + p3)
else:
    return (word, '')

```

There are three capture groups that we can put into three variables. We want to ensure we don't return any `None` values, so we can use an "or" to evaluate the left side as "truthy" and take the empty string on the right if it's not.

Return a tuple that has the first part of the word (maybe consonants) and the "rest" of the word (the vowel plus anything else).

If the match was `None`, return a tuple of the word and an empty string to indicate there is no "rest" of the word to rhyme.

Let's take a moment to think about the `or` operator, which we're using to decide between something on the left *or* something on the right. The `or` will return the first "truthy" value, the one that—sort of, kind of—evaluates to `True` in a Boolean context:

```

It's easiest to see with literal
True and False values.
>>> True or False
True
>>> False or True
True
>>> 1 or 0
1
>>> 0 or 1
1
>>> 0.0 or 1.0
1.0
>>> '0' or ''
'0'
>>> 0 or False
False
>>> [] or ['foo']
['foo']
>>> {} or dict(foo=1)
{'foo': 1}

```

No matter the order, the True value will be taken.

In a Boolean context, the integer value 0 is "falsey," and any other value is "truthy."

The number values behave exactly like actual Boolean values.

Floating-point values also behave like integer values, where 0.0 is "falsey" and anything else is "truthy."

With string values, the empty string is "falsey" and anything else is "truthy." It may look odd because it returns '0', but that's not the numeric value zero but the string we use to represent the value of zero. Wow, so philosophical.

If no value is "truthy," the last value is returned.

The empty dict is "falsey," and any non-empty dict is "truthy."

The empty list is "falsey," so any non-empty list is "truthy."

You should be able to use these ideas to write a `stemmer()` function that will pass the `test_stemmer()` function. Remember, if both of these functions are in your `rhymer.py` program, you can run the `test_` functions like so:

```
$ pytest -xv rhymer.py
```

14.1.5 Creating the output

Let's review what the program should do:

- 1 Take a positional string argument.
- 2 Try to split it into two parts: any leading consonants and the rest of the word.
- 3 If the split is successful, combine the "rest" of the word (which might actually be the entire word if there are no leading consonants) with all the other consonant sounds. Be sure to *not* include the original consonant sound and to sort the rhyming strings.
- 4 If you are unable to split the word, print the message Cannot rhyme "<word>".

Now it's time to write the program. Have fun storming the castle!

14.2 Solution

"No more rhymes now, I mean it!"

"Anybody want a peanut?"

Let's take a look at one way to solve this problem. How different was your solution?

```
#!/usr/bin/env python3
"""Make rhyming words"""

import argparse
import re
import string
    ← The re module is for
    regular expressions.

# -----
def get_args():
    """get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Make rhyming "words"',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('word', metavar='word', help='A word to rhyme')

    return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    prefixes = list('bcdfghjklmnpqrstvwxyz') + (
        'bl br ch cl cr dr fl fr gl gr pl pr sc '
        'sh sk sl sm sn sp st sw th tr tw thw wh wr '
        'sch scr shr sph spl spr squ str thr').split()

    start, rest = stemmer(args.word)
    if rest:

Check if
there is a
part of the
word that
we can use
to create
rhyming
strings.    ← Get the command-line
arguments.

    Define all the prefixes that
will be added to create
rhyming words.

    Split the word argument into
two possible parts. Because the
stemmer() function always
returns a 2-tuple, we can unpack
the values into two variables.
```

If there is, use a list comprehension to iterate through all the prefixes and add them to the stem of the word. Use a guard to ensure that any given prefix is not the same as the beginning of the word. Sort all the values and print them, joined on newlines.

```
→     print('\n'.join(sorted([p + rest for p in prefixes if p != start])))
else:
    print(f'Cannot rhyme "{args.word}"')
```

If there is nothing for the “rest” of the word that can be used to create rhymes, let the user know.

```
# -----
def stemmer(word):
    """Return leading consonants (if any), and 'stem' of word"""

Lowercase the word.
```

```
→ word = word.lower()
vowels = 'aeiou'
consonants = ''.join(
    [c for c in string.ascii_lowercase if c not in vowels])
pattern = (
    '([' + consonants + ']+)?' # capture one or more, optional
    '([' + vowels + '])' # capture at least one vowel
    '(.*')' # capture zero or more of anything
)
```

Since we will use the vowels more than once, assign them to a variable.

The consonants are the letters that are not vowels. We will only match to lowercase letters.

The pattern is defined using consecutive literal strings that Python will join together into one string. By breaking up the pieces onto separate lines, we can comment on each part of the regular expression.

```
match = re.match(pattern, word)
if match:
    p1 = match.group(1) or ''
    p2 = match.group(2) or ''
    p3 = match.group(3) or ''
    return (p1, p2 + p3)
else:
    return (word, '')
```

Use the `re.match()` function to start matching at the beginning of the word.

The `re.match()` function will return `None` if the pattern failed to match, so check if the match is “truthy” (`not None`).

Put each group into a variable, always ensuring that we use the empty string rather than `None`.

The tests for the `stemmer()` function. I usually like to put my unit tests directly after the functions they test.

```
# -----
def test_stemmer():
    """test the stemmer"""

    assert stemmer('') == ('', '')
    assert stemmer('cake') == ('c', 'ake')
    assert stemmer('chair') == ('ch', 'air')
    assert stemmer('APPLE') == ('', 'apple')
    assert stemmer('RDNZL') == ('rdnzl', '')
    assert stemmer('123') == ('', '')
```

Return a new tuple that has the “first” part of the word (possible leading consonants) and the “rest” of the word (the vowel plus anything else).

```
# -----
if __name__ == '__main__':
    main()
```

If the match failed, return the word and an empty string for the “rest” of the word to indicate there is nothing to rhyme.

14.3 Discussion

There are many ways you could have written this, but, as always, I wanted to break the problem down into units I could write and test. For me, this came down to splitting the word into a possible leading consonant sound and the rest of the word. If I can manage that, I can create rhyming strings; if I cannot, then I need to alert the user.

14.3.1 Stemming a word

For the purposes of this program, the “stem” of a word is the part after any initial consonants, which I define using a list comprehension with a guard to take only the letters that are not vowels:

```
>>> vowels = 'aeiou'  
>>> consonants = ''.join([c for c in string.ascii_lowercase if c not in vowels])
```

Throughout the chapters, I have shown how a list comprehension is a concise way to generate a list and is preferable to using a for loop to append to an existing list. Here we have added an if statement to only include some characters if they are not vowels. This is called a *guard* statement, and only those elements that evaluate as “truthy” will be included in the resulting list.

We've looked at `map()` several times now and talked about how it is a *higher-order function* (HOF) because it takes *another function* as the first argument and will apply it to all the elements from some *iterable* (something that can be *iterated*, like a list). Here I'd like to introduce another HOF called `filter()`, which also takes a function and an iterable (see figure 14.12). As with the list comprehension with the guard, only those elements that return a "truthy" value from the function are allowed in the resulting list.

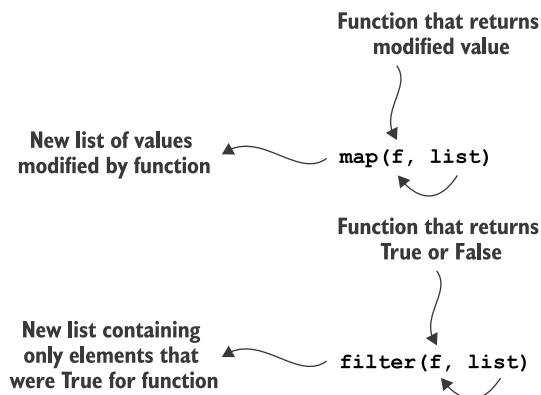


Figure 14.12 The `map()` and `filter()` functions both take a function and an iterable, and both produce a new list.

Here is another way to write the idea of the list comprehension using filter():

```
>>> consonants = ''.join(filter(lambda c: c not in vowels,  
    string.ascii_lowercase))
```

Just as with `map()`, I use the `lambda` keyword to create an *anonymous function*. The `c` is the variable that will hold the argument, which, in this case, will be each character from `string.ascii_lowercase`. The entire body of the function is the evaluation `c not in vowels`. Each of the vowels will return `False` for this:

```
>>> 'a' not in vowels  
False
```

And each of the consonants will return `True`:

```
>>> 'b' not in vowels  
True
```

Therefore, only the consonants will be allowed to pass through `filter()`. Think back to our “blue” cars; let’s write a `filter()` that only accepts cars that start with the string “blue”:

```
>>> cars = ['blue Honda', 'red Chevy', 'blue Ford']  
>>> list(filter(lambda car: car.startswith('blue '), cars))  
['blue Honda', 'blue Ford']
```

When the `car` variable has the value “red Chevy,” the `lambda` returns `False`, and that value is rejected:

```
>>> car = 'red Chevy'  
>>> car.startswith('blue ')  
False
```

Note that if none of the elements from the original iterable are accepted, `filter()` will produce an empty list (`[]`). For example, I could `filter()` for numbers greater than 10. Note that `filter()` is another *lazy* function that I must coerce using the `list` function in the REPL:

```
>>> list(filter(lambda n: n > 10, range(0, 5)))  
[]
```

A list comprehension would also return an empty list:

```
>>> [n for n in range(0, 5) if n > 10]  
[]
```

Figure 14.13 shows the relationship between creating a new list called `consonants` using an imperative `for`-loop approach, an idiomatic list comprehension with a guard, and a purely functional approach using `filter()`. All of these are perfectly acceptable, though the most Pythonic technique is probably the list comprehension. The `for` loop would be very familiar to a C or Java programmer, while the `filter()` approach would be immediately recognizable to the Haskeller or even someone from a Lisp-like language. The `filter()` might be slower than the list comprehension, especially if the iterable were large. Choose whichever way makes more sense for your style and application.

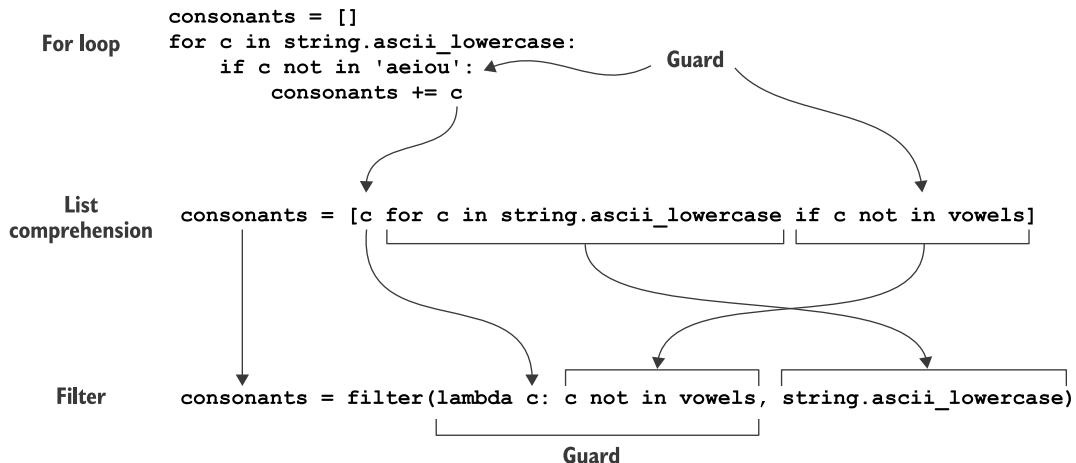


Figure 14.13 Three ways to create a list of consonants: using a `for` loop with an `if` statement, a list comprehension with a guard, and a `filter()`

14.3.2 Formatting and commenting the regular expression

We talked in the introduction about the individual parts of the regular expression I ended up using. I'd like to take a moment to mention the way I formatted the regex in the code. I used an interesting trick of the Python interpreter that will implicitly concatenate adjacent string literals. See how these four strings become one:

```
>>> this_is_just_to_say = ('I have eaten '
... 'the plums '
... 'that were in '
... 'the icebox')
>>> this_is_just_to_say
'I have eaten the plums that were in the icebox'
```

Note that there are no commas after each string, as that would create a tuple with four individual strings:

```
>>> this_is_just_to_say = ('I have eaten ',
... 'the plums ',
... 'that were in ',
... 'the icebox')
>>> this_is_just_to_say
('I have eaten ', 'the plums ', 'that were in ', 'the icebox')
```

The advantage of writing out the regular expression on separate lines is that you can add comments to help your reader understand each part:

```
pattern = (
    '[' + consonants + ']+)?' # capture one or more, optional
    '([' + vowels + '])' # capture at least one vowel
    '(.*)' # capture zero or more of anything
)
```

The individual strings will be concatenated by Python into a single string:

```
>>> pattern
'([bcdfghjklmnpqrstvwxyz]+)?([aeiou])(.*)'
```

I could have written the entire regex on one line, but ask yourself which version would you rather read and maintain, the preceding version or the following:¹

```
pattern = f'({{consonants}})+?({{vowels}})(.*)'
```

14.3.3 Using the stemmer() function outside your program

One of the very interesting things about Python code is that your `rhymers.py` program is also—kind of, sort of—a sharable *module* of code. That is, you haven’t explicitly written it to be a container of reusable (and tested!) functions, but it is. You can even run the functions from inside the REPL.

For this to work, be sure you run `python3` inside the same directory as the `rhymers.py` code:

```
>>> from rhymers import stemmer
```

Now you can run and test your `stemmer()` function manually:

```
>>> stemmer('apple')
(' ', 'apple')
>>> stemmer('banana')
('b', 'anana')
>>> import string
>>> stemmer(string.punctuation)
('!'#$%&\()'*+, -./:;=>?@[\\\]^_`{|}~, ' ')
```

The deeper meaning of `if __name__ == '__main__':`

Note that if you were to change the last two lines of `rhymers.py` from this,

```
if __name__ == '__main__':
    main()
```

to this,

```
main()
```

the `main()` function would be run when you try to import the module:

```
>>> from rhymers import stemmer
usage: [-h] str
: error: the following arguments are required: str
```

¹ “Looking at code you wrote more than two weeks ago is like looking at code you are seeing for the first time.”—Dan Hurvitz

(continued)

This is because `import rhymer` causes Python to execute the `rhymer.py` file to the end. If the last line of the module calls `main()`, then `main()` will run!

The `__name__` variable is set to '`__main__`' when `rhymer.py` is being *run as a program*. That is the only time `main()` is executed. When the module is being imported by another module, then `__name__` is equal to `rhymer`.

If you don't explicitly import a function, you can use the fully qualified function name by adding the module name to the front:

```
>>> import rhymer
>>> rhymer.stemmer('cake')
('c', 'ake')
>>> rhymer.stemmer('chair')
('ch', 'air')
```

There are many advantages to writing many small functions rather than long, sprawling programs. One is that small functions are much easier to write, understand, and test. Another is that you can put your tidy, tested functions into modules and share them across different programs you write.

As you write more and more programs, you will find yourself solving some of the same problems repeatedly. It's far better to create modules with reusable code than to copy pieces from one program to another. If you ever find a bug in a shared function, you can fix it once, and all the programs sharing the function get the fix. The alternative is to find the duplicated code in every program and change it (hoping that this doesn't introduce even more problems because the code is entangled with other code).



14.3.4 Creating rhyming strings

I decided that my `stemmer()` function would always return a 2-tuple of the `(start, rest)` for any given word. As such, I can unpack the two values into two variables:

```
>>> start, rest = stemmer('cat')
>>> start
'c'
>>> rest
'at'
```

If there is a value for `rest`, I can add all my prefixes to the beginning:

```
>>> prefixes = list('bcd...xyz') + (
...     'bl br ch cl cr dr fl fr gl gr pl pr sc '
```

```
...     'sh sk sl sm sn sp st sw th tr tw wh wr'
...     'sch scr shr sph spl spr squ str thr').split()
```

I decided to use another list comprehension with a guard to skip any prefix that is the same as the start of the word. The result will be a new list that I pass to the sorted() function to get the correctly ordered strings:

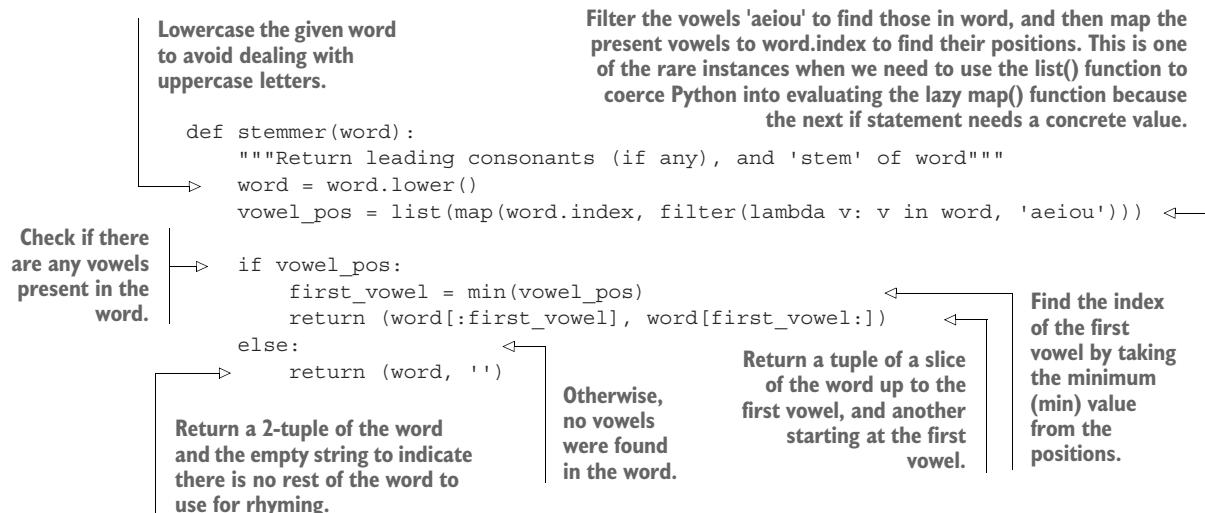
```
>>> sorted([p + rest for p in prefixes if p != start])
['bat', 'blat', 'brat', 'chat', 'clat', 'crat', 'dat', 'drat', 'fat',
 'flat', 'frat', 'gat', 'glat', 'grat', 'hat', 'jat', 'kat', 'lat',
 'mat', 'nat', 'pat', 'plat', 'prat', 'qat', 'rat', 'sat', 'scat',
 'schat', 'scrat', 'shat', 'shrat', 'skat', 'slat', 'smat', 'snat',
 'spat', 'sphat', 'splat', 'sprat', 'squat', 'stat', 'strat', 'swat',
 'tat', 'that', 'thrat', 'thwat', 'trat', 'twat', 'vat', 'wat',
 'what', 'wrat', 'xat', 'yat', 'zat']
```

I then print() that list, joined on newlines. If there is no rest of the given word, I print() a message that the word cannot be rhymed:

```
if rest:
    print('\n'.join(sorted([p + rest for p in prefixes if p != start])))
else:
    print(f'Cannot rhyme "{args.word}"')
```

14.3.5 Writing stemmer() without regular expressions

It is certainly possible to write a solution that does not use regular expressions. We could start by finding the first position of a vowel in the given string. If one is present, we could use a list slice to return the portion of the string up to that position and the portion starting at that position:



This function will also pass the `test_stemmer()` function. By writing a test just for the idea of this one function, and exercising it with all the different values I would expect, I'm free to *refactor* my code. In my mind, the `stemmer()` function is a black box. What goes on inside the function is of no concern to the code that calls it. As long as the function passes the tests, it is "correct" (for certain values of "correct").

Small functions and their *tests* will set you free to improve your programs. First make something work, and make it beautiful. Then try to make it better, using your tests to ensure it keeps working as expected.



14.4 Going further

- Add an `--output` option to write the words to a given file. The default should be to write to `STDOUT`.
- Read an input file and create rhyming words for all the words in the file. You can borrow from the program in chapter 6 to read a file and break it into words, then iterate each word, and create an output file for each word with the rhyming words.
- Write a new program that finds all unique consonant sounds in a dictionary of English words. (I have included `inputs/words.txt.zip`, which is a compressed version of the dictionary from my machine. Unzip the file to use `inputs/words.txt`.) Print the output in alphabetical order and use those to expand this program's consonants.
- Alter your program to only emit words that are found in the system dictionary (for example, `inputs/words.txt`).
- Write a program to create Pig Latin, where you move the initial consonant sound from the beginning of the word to the end and add "-ay," so that "cat" becomes "at-cay." If a word starts with a vowel, add "-yay" to the end so that "apple" becomes "apple-yay."
- Write a program to create spoonerisms, where the initial consonant sounds of adjacent words are switched, so you get "blushing crow" instead of "crushing blow."

Summary

- Regular expressions allow you to declare a pattern that you wish to find. The regex engine will sort out whether the pattern is found or not. This is a *declarative* approach to programming, in contrast to the *imperative* method of manually seeking out patterns by writing code ourselves.

- You can wrap parts of the pattern in parentheses to “capture” them into groups that you can fetch from the result of `re.match()` or `re.search()`.
- You can add a guard to a list comprehension to avoid taking some elements from an iterable.
- The `filter()` function is another way to write a list comprehension with a guard. Like `map()`, it is a lazy, higher-order function that takes a function that will be applied to every element of an iterable. Only those elements that are deemed “truthy” by the function are returned.
- Python can evaluate many types—including strings, numbers, lists, and dictionaries—in a Boolean context to arrive at a sense of “truthiness.” That is, you are not restricted to just `True` and `False` in `if` expressions. The empty string `''`, the `int 0`, the `float 0.0`, the empty `list []`, and the empty `dict {}` are all considered “falsey,” so any non-falsey value from those types, like the non-empty `str`, `list`, or `dict`, or any numeric value not zero-ish, will be considered “truthy.”
- You can break long string literals into shorter adjacent strings in your code to have Python join them into one long string. It’s advisable to break long regexes into shorter strings and add comments on each line to document the function of each pattern.
- Write small functions and tests, and share them in modules. Every `.py` file can be a module from which you can `import` functions. Sharing small, tested functions is better than writing long programs and copying/pasting code as needed.

The Kentucky Friar: More regular expressions

I grew up in the American Deep South where we tend to drop the final “g” of words ending in “ing,” like “cookin’” instead of “cooking.” We also tend to say “y’all” for the second-person plural pronoun, which makes sense because Standard English is missing a distinctive word for this. In this exercise, we’ll write a program called `friar.py` that will accept some input as a single positional argument and transform the text by replacing the final “g” with an apostrophe (‘) for two-syllable words ending in “ing” and changing “you” to “y’all.” Granted, we have no way to know if we’re changing the first- or second-person “you,” but it makes for a fun challenge nonetheless.



Figure 15.1 is a string diagram that will help you see the inputs and outputs. When run with no arguments or with the `-h` or `--help` flags, your program should present the following usage statement:

```
$ ./friar.py -h
usage: friar.py [-h] text

Southern fry text

positional arguments:
  text            Input text or file

optional arguments:
  -h, --help      show this help message and exit
```

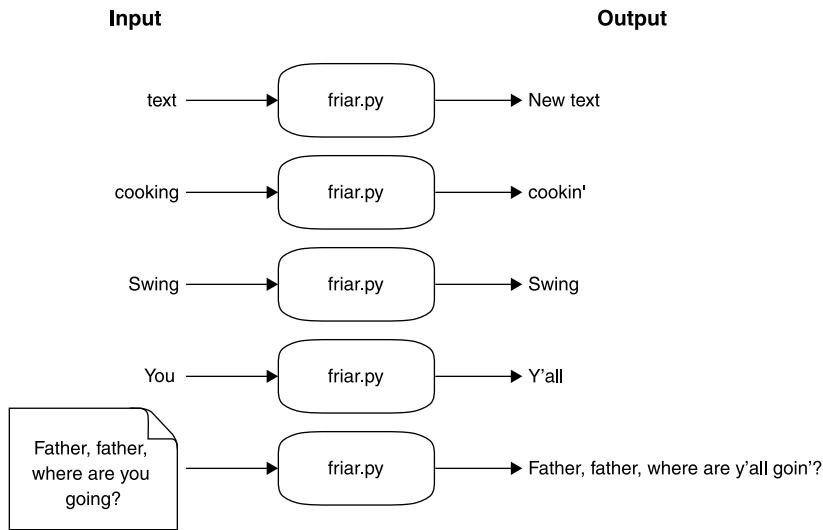


Figure 15.1 Our program will modify the input text to give it a Southern lilt.

We will only change “-ing” words with *two syllables*, so “cooking” becomes “cookin’” but “swing” will stay the same. Our heuristic for identifying two-syllable “-ing” words is to inspect the part of the word before the “-ing” ending to see if it contains a vowel, which in this case will include “y.” We can split “cooking” into “cook” and “ing,” and because there is an “o” in “cook,” we should drop the final “g”:

```
$ ./friar.py Cooking
Cookin'
```

When we remove “ing” from “swing,” though, we’re left with “sw,” which contains no vowel, so it will remain the same:

```
$ ./friar.py swing
swing
```

When changing “you” to “y’all,” be mindful to keep the case the same on the first letter. For example, “You” should become “Y’all”:

```
$ ./friar.py you
y'all
$ ./friar.py You
Y'all
```

As in several previous exercises, the input may name a file, in which case you should read the file for the input text. To pass the tests, you will need to preserve the line structure of the input, so I recommend you read the file line by line. Given this input,

```
$ head -2 inputs/banner.txt
O! Say, can you see, by the dawn's early light,
What so proudly we hailed at the twilight's last gleaming -
```

the output should have the same line breaks:

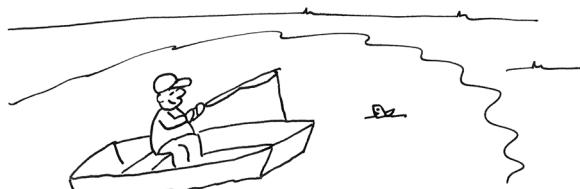
```
$ ./friar.py inputs/banner.txt | head -2
O! Say, can y'all see, by the dawn's early light,
What so proudly we hailed at the twilight's last gleamin' -
```

To me, it's quite amusing to transform texts this way, but maybe I'm just weird:

```
$ ./friar.py inputs/raven.txt
Presently my soul grew stronger; hesitatin' then no longer,
"Sir," said I, "or Madam, truly your forgiveness I implore;
But the fact is I was nappin', and so gently y'all came rappin',
And so faintly y'all came tappin', tappin' at my chamber door,
That I scarce was sure I heard y'all" - here I opened wide the door: -
Darkness there and nothin' more.
```

In this exercise you will

- Learn more about using regular expressions
- Use both `re.match()` and `re.search()` to find patterns anchored to the beginning of a string or anywhere in the string, respectively
- Learn how the `$` symbol in a regex anchors a pattern to the *end* of a string
- Learn how to use `re.split()` to split a string
- Explore how to write a manual solution for finding two-syllable “-ing” words or the word “you”



15.1 Writing `friar.py`

As usual, I recommend you start with `new.py friar.py` or copy the template/`template.py` file to `15_friar/friar.py`. I suggest you start with a simple version of the program that echoes back the input from the command line:

```
$ ./friar.py cooking
cooking
```

Or from a file:

```
$ ./friar.py inputs/blake.txt
Father, father, where are you going?
Oh do not walk so fast!
Speak, father, speak to your little boy,
Or else I shall be lost.
```

We need to process the input line by line, and then word by word. You can use the `str.splitlines()` method to get each line of the input, and then use the `str.split()` method to break the line on spaces into word-like units. This code,

```
for line in args.text.splitlines():
    print(line.split())
```

should create this output:

```
$ ./friar.py tests/blake.txt
['Father,', 'father,', 'where', 'are', 'you', 'going?']
['Oh', 'do', 'not', 'walk', 'so', 'fast!']
['Speak,', 'father,', 'speak', 'to', 'your', 'little', 'boy,']
['Or', 'else', 'I', 'shall', 'be', 'lost.']}
```

If you look closely, it's going to be difficult to handle some of these word-like units because the adjacent punctuation is still attached to the words, as in 'Father,' and 'going?' Splitting the text on spaces is not sufficient, so I'll show you how to split the text *using a regular expression*.

15.1.1 Splitting text using regular expressions

As in chapter 14, we need to import `re` to use regexes:

```
>>> import re
```

For demonstration purposes, I'm going to set `text` to the first line:

```
>>> text = 'Father, father, where are you going?'
```

By default, `str.split()` breaks text on spaces. Note that whatever text is used for splitting will be missing from the result, so here there are no spaces:

```
>>> text.split()
['Father,', 'father,', 'where', 'are', 'you', 'going?']
```

You can pass an optional value to `str.split()` to indicate the string you want to use for splitting. If we choose the comma, we'll end up with three strings instead of six. Note that there are no commas in the resulting list, as that is the argument to `str.split()`:

```
>>> text.split(',')
['Father', ' father', ' where are you going?']
```

The `re` module has a function called `re.split()` that works similarly. I recommend you read `help(re.split)`, as this is a very powerful and flexible function. Like `re.match()`, which we used in chapter 14, this function wants at least a pattern and a string. We can use `re.split()` with a comma to get the same output as `str.split()`, and, as before, the commas are missing from the result:

```
>>> re.split(',', text)
['Father', ' father', ' where are you going?']
```

15.1.2 Shorthand classes

We are after the things that look like “words,” in that they are composed of the characters that normally occur in words. The characters that *don’t* normally occur in words (things like punctuation) are what we want to use for splitting. You’ve seen before that we can create a *character class* by putting literal values inside square brackets, like `'[aeiou]'` for the vowels. What if we create a character class where we enumerate all the non-letter characters? We could do something like this:

```
>>> import string
>>> ''.join([c for c in string.printable if c not in string.ascii_letters])
'0123456789!#$%&\!()*+,-./:;=>?@[\\\]^_`{|}~ \t\n\r\x0b\x0c'
```

That won’t be necessary, because almost every implementation of regular expression engines define shorthand character classes. Table 15.1 lists some of the most common shorthand classes and how they can be written longhand.

Table 15.1 Regex shorthand classes

Character class	Shorthand	Other ways to write the class
Digits	<code>\d</code>	<code>[0123456789]</code> , <code>[0-9]</code>
Whitespace	<code>\s</code>	<code>[\t\n\r\x0b\x0c]</code> , same as <code>string.whitespace</code>
Word characters	<code>\w</code>	<code>[a-zA-Z0-9_-]</code>

NOTE There is a basic flavor of regular expression syntax that is recognized by everything from Unix command-line tools like `awk` to regex support inside of languages like Perl, Python, and Java. Some tools add extensions to their regexes that may not be understood by other tools. For example, there was a time when Perl’s regex engine added many new ideas that eventually became a dialect known as “PCRE” (Perl-Compatible Regular Expressions). Not every tool that understands regexes will understand every flavor of regex, but in all my years of writing and using regexes, I’ve rarely found this to be a problem.

The shorthand `\d` means any *digit* and is equivalent to `'[0123456789]'`. I can use the `re.search()` method to look anywhere in a string for any digit. In the following example, it will find the character ‘1’ in the string `'abc123!'` because this is the first digit in the string (see figure 15.2):

```
>>> re.search('\d', 'abc123!')
<re.Match object; span=(3, 4), match='1'>
```



Figure 15.2 The digit shorthand will match any single digit.

That is the same as using the longhand version (see figure 15.3):

```
>>> re.search('[0123456789]', 'abc123!')
<re.Match object; span=(3, 4), match='1'>
```

[0123456789]



abc123!

Figure 15.3 We can also create a character class enumerating all the digits.

It's also the same as the version that uses the range of characters '[0-9]' (see figure 15.4):

```
>>> re.search('[0-9]', 'abc123!')
<re.Match object; span=(3, 4), match='1'>
```

[0-9]



abc123!

Figure 15.4 Character classes can use a range of contiguous values, like 0-9.

To have it find *one or more digits in a row*, add the + (see figure 15.5):

```
>>> re.search('\d+', 'abc123!')
<re.Match object; span=(3, 6), match='123'>
```

\d+



abc123!

Figure 15.5 The plus signs means to match one or more of the preceding expression.

The \w shorthand means “any word-like character.” It includes all the Arabic numbers, the letters of the English alphabet, the dash ('-'), and the underscore ('_'). The first match in the string is 'a' (see figure 15.6):

```
>>> re.search('\w', 'abc123!')
<re.Match object; span=(0, 1), match='a'>
```

\w



abc123!

Figure 15.6 The shorthand for word characters is \w.

If you add the `+` as in figure 15.7, it matches one or more word characters in a row, which includes `abc123` but not the exclamation mark (!):

```
>>> re.search('\w+', 'abc123!')
<re.Match object; span=(0, 6), match='abc123'>
```

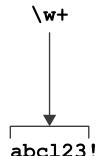


Figure 15.7 Add the plus sign to match one or more word characters.

15.1.3 Negated shorthand classes

You can complement or “negate” a character class by putting the caret (`^`) *immediately inside* the character class as in figure 15.8. One or more of any character *not* a digit is `[^0-9]+`. With it, `'abc'` is found:

```
>>> re.search('[^0-9]+', 'abc123!')
<re.Match object; span=(0, 3), match='abc'>
```



Figure 15.8 A caret just inside a character class will negate or complement the characters. This regex matches non-digits.

The shorthand class of non-digits `[^0-9]+` can also be written as `\D+` as in figure 15.9:

```
>>> re.search('\D+', 'abc123!')
<re.Match object; span=(0, 3), match='abc'>
```

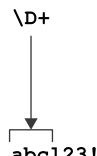


Figure 15.9 The shorthand `\D+` matches one or more non-digits.

The shorthand for non-word characters is `\W`, which will match the exclamation point (see figure 15.10):

```
>>> re.search('\W', 'abc123!')
<re.Match object; span=(6, 7), match='! '>
```



Figure 15.10 The `\W` will match anything that is not a letter, digit, underscore, or dash.

Table 15.2 summarizes these shorthand classes and how they can be expanded.

Table 15.2 Negated regex shorthand classes

Character class	Shorthand	Other ways to write the class
Not a digit	\D	[^0123456789], [^0-9]
Not whitespace	\S	[^ \t\n\r\x0b\x0c]
Not word characters	\W	[^a-zA-Z0-9_-]

15.1.4 Using `re.split()` with a captured regex

We can use \W as the argument to `re.split()`:

```
>>> re.split('\W', 'abc123!')
['abc123', '']
```

NOTE Pylint will complain if we use '\W' in a regular expression in our program, returning the message "Anomalous backslash in string: '\W'. String constant might be missing an r prefix." We can use the r prefix to create a "raw" string, one where Python does not try to interpret the \W as it will, for instance, interpret \n to mean a newline or \r to mean a carriage return. From this point on, I will use the r-string syntax to create a raw string.

There is a problem, though, because the result of `re.split()` omits those strings matching the pattern. Here we've lost the exclamation point! If we read `help(re.split)` closely, we can find the solution:

If capturing parentheses are used in [the] pattern, then the text of all groups in the pattern are also returned as part of the resulting list.

We used capturing parentheses in chapter 14 to tell the regex engine to "remember" certain patterns, like the consonant(s), vowel, and the rest of a word. When the regex matched, we were able to use `match.groups()` to retrieve strings that were found by the patterns. Here we will use the parentheses around the pattern to `re.split()` so that the strings matching the pattern will also be returned:

```
>>> re.split(r'(\W)', 'abc123!')
['abc123', '!', '']
```

If we try that on our `text`, the result is a list of strings that match and do not match the regular expression:

```
>>> re.split(r'(\W)', text)
['Father', ',', "'", ' ', 'father', ',', "'", ' ', 'where', ' ', 'are', ' ',
 'you', ' ', 'going', '?', '']
```

I'd like to group all the non-word characters together by adding + to the regex (see figure 15.11):

```
>>> re.split(r'(\W+)', text)
['Father', ' ', ' ', 'father', ' ', ' ', 'where', ' ', ' ', 'are', ' ', ' ', 'you', ' ', ' ', 'going',
 ' ', '?', ' ']
```

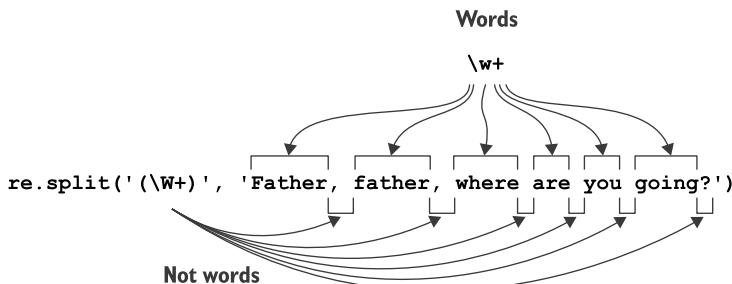


Figure 15.11 The `re.split()` function can use a captured regex to return both the parts that match the regex and those that do not.

That is so cool! Now we have a way to process each *actual* word and the bits in between them.

15.1.5 Writing the `fry()` function

Our next step is to write a function that will decide whether and how to modify *just one word*. That is, rather than thinking about how to handle all the text at once, we'll think about how to handle one word at a time. We can call this function `fry()`.

To help us think about how this function should work, let's start off by writing the `test_fry()` function and a stub for the actual `fry()` function that contains just the single command `pass`, which tells Python to do nothing. To get started on this, you can paste this into your program:

```
def fry(word):
    pass
```

pass is a way to do nothing. You might call it a “no-operation” or “NO-OP,” which kind of looks like “NOPE,” which is another way to remember that it does nothing. We’re just defining this `fry()` function as a placeholder so we can write the test.

```
def test_fry():
    assert fry('you') == "y'all"
    assert fry('You') == "Y'all"
    assert fry('fishing') == "fishin'"
    assert fry('Aching') == "Achin'"
    assert fry('swing') == "swing"
```

The word “you” should become “y’all.” Ensure the word’s capitalization is preserved.

This is a two-syllable “-ing” word that should be changed by dropping the final “g” for an apostrophe.

This is a one-syllable “-ing” word that should not be changed.

This is a two-syllable “-ing” word that starts with a vowel. It should likewise be changed.

The `test_fry()` function will pass in words we expect to be changed or not. We can’t check every word, so we’ll rely on spot-checking the major cases.

Now run `pytest friar.py` to see that, as expected, the test will fail:

```
===== FAILURES =====
test_fry
=====
def test_fry():
>     assert fry('you') == "y'all"
E     assert None == "y'all"
E     + where None = fry('you')
friar.py:47: AssertionError
===== 1 failed in 0.08 seconds =====
```

The first test is failing.
The result of `fry('you')` was `None`, which does not equal “`y'all`.”

Let’s change our `fry()` function to handle that string:

```
def fry(word):
    if word == 'you':
        return "y'all"
```

Now let’s run our tests again:

```
===== FAILURES =====
test_fry
=====
def test_fry():
>     assert fry('you') == "y'all"
E     assert fry('You') == "Y'all"
E     assert None == "Y'all"
E     + where None = fry('You')
friar.py:49: AssertionError
===== 1 failed in 0.16 seconds =====
```

Now the first test passes.
The second test fails because the “`You`” is capitalized.
The function returned `None` but should have returned “`Y'all`.”

Let’s handle those:

```
def fry(word):
    if word == 'you':
        return "y'all"
    elif word == 'You':
        return "Y'all"
```

If you run the tests now, you’ll see that the first two tests pass; however, I’m definitely not happy with that solution. There is already a good bit of duplicated code. Can we find a more elegant way to match both “`you`” and “`You`” and still return the correctly capitalized answer? Yes, we can!

```
def fry(word):
    if word.lower() == 'you':
        return word[0] + "'all"
```

Better still, we can write a regular expression! There is one difference between “you” and “You”—the “y” or “Y”—that we can represent using the character class ‘[yY]’ (see figure 15.12). This will match the lowercase version:

```
>>> re.match(' [yY] ou', 'you')
<re.Match object; span=(0, 3), match='you'>
```

Either
“y” or “Y”

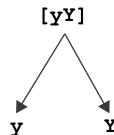


Figure 15.12 We can use a character class to match lower- and uppercase Y.

It will also match the capitalized version (see figure 15.13):

```
>>> re.match(' [yY] ou', 'You')
<re.Match object; span=(0, 3), match='You'>
```

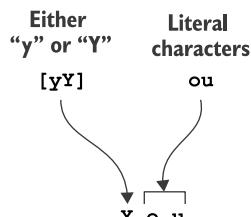


Figure 15.13 This regex will match
“you” and “You.”

Now we want to reuse the initial character (either “y” or “Y”) in the return value. We could *capture* it by placing it into parentheses. Try to rewrite your `fry()` function using this idea, and getting it to pass the first two tests again, before moving on:

```
>>> match = re.match('([yY])ou', 'You')
>>> match.group(1) + "all"
"Y'all"
```

The next step is to handle a word like “fishing”:

```
===== FAILURES =====
_____ test_fry _____
```

```

def test_fry():
    assert fry('you') == "y'all"
    assert fry('You') == "Y'all"
>    assert fry('fishing') == "fishin'"
```

The third test fails.

```
E    assert None == "fishin'"
E    +   where None = fry('fishing')
```

The return from fry('fishing')
was None, but the value
“fishin” was expected.

```
friar.py:52: AssertionError
=====
1 failed in 0.10 seconds =====
```

How can we identify a word that ends with “ing”? With the `str.endswith()` function:

```
>>> 'fishing'.endswith('ing')
True
```

A regular expression to find “ing” at the end of a string would use `$` (pronounced “dollar”) at the end of the expression to *anchor* the expression to the end of the string (see figure 15.14):

```
>>> re.search('ing$', 'fishing')
<re.Match object; span=(4, 7), match='ing'>
```

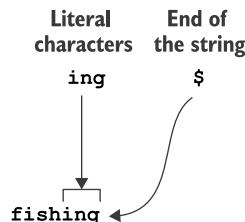


Figure 15.14 The dollar sign indicates the end of the word.

As shown in figure 15.15, we can use a string slice to get all the characters up to the last at index `-1` and then append an apostrophe.

```
f i s h i n g
————→ | -1
word[:-1] + "'" —————→ fishin'
```

Figure 15.15 Use a string slice to get all the letters up to the last one and add an apostrophe.

Add this to your `fry()` function and see how many tests you pass:

```
if word.endswith('ing'):
    return word[:-1] + "'"
```

Or you could use a group within the regex to capture the first part of the word (see figure 15.16):

```
>>> match = re.search('(.+)ing$', 'fishing')
>>> match.group(1) + "in"
"fishin'"
```

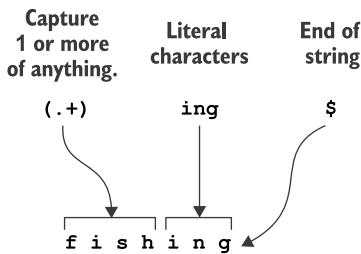


Figure 15.16 Using a capture group so we can access the matching string

You should be able to get results like this:

```
===== FAILURES =====
test_fry
```

```
def test_fry():
    assert fry('you') == "y'all"
    assert fry('You') == "Y'all"
    assert fry('fishing') == "fishin'"
    assert fry('Aching') == "Achin'"
>   assert fry('swing') == "swing"      <-- This test failed.
E   assert "swin'" == 'swing'        <-- The result of fry('swing')
E       - swin'                  <-- was "swin," but it should
E       ? ^                      <-- have been "swing."
E       + swing                 <-- Sometimes the test results will be able to highlight the
E       ? ^                      <-- exact point of failure. Here you are being shown that
                                there is an apostrophe ('') where there should be a "g."
friar.py:59: AssertionError
===== 1 failed in 0.10 seconds =====
```

We need a way to identify words that have two syllables. I mentioned before that we'll use a heuristic that looks for a vowel, '[aeiouy]', in the part of the word *before* the "ing" ending, as shown in figure 15.17. Another regex could do the trick:

Here we know there will be a match value, so we can use `match.group(1)` to get the first capture group, which will be anything immediately before "ing." In actual code, we should check that `match` is not `None` or we'd trigger an exception by trying to execute the `group` method on a `None`.

The `(.+)` will match and capture one or more of anything followed by the characters "ing." The return from `re.search()` will either be a `re.Match` object if the pattern was found or `None` to indicate it was not.

```
>>> match = re.search('(.+)ing$', 'fishing')      <-- We can use re.search()
>>> first = match.group(1)                         <-- on the first part of the
>>> re.search('[aeiouy]', first)                   <-- string to look for a
<<re.Match object; span=(1, 2), match='i'>          <-- vowel.

As the return from re.search() is a re.Match object, we
know there is a vowel in the first part, so the word
looks to have two syllables.
```

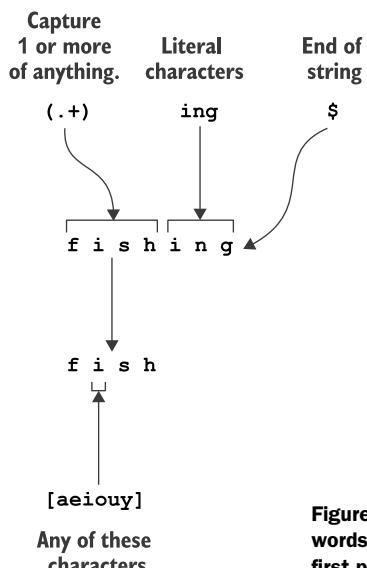


Figure 15.17 A possible way to find two-syllable words ending in "ing" is to look for a vowel in the first part of the word.

If the word matches this test, return the word with the final "g" replaced with an apostrophe; otherwise, return the word unchanged. I suggest you not proceed until you are passing all of `test_fry()`.

15.1.6 Using the `fry()` function

Now your program should be able to

- 1 Read input from the command line or a file
- 2 Read the input line by line
- 3 Split each line into words and non-words
- 4 `fry()` any individual word

The next step is to apply the `fry()` function to all the word-like units. I hope you can see a familiar pattern emerging—applying a function to all elements of a list! You can use a `for` loop:

```
Preserve the structure of the newlines in
args.text by using str.splitlines().
```

```
for line in args.text.splitlines():
    words = []
    for word in re.split(r'(\W+)', line.rstrip()):
        words.append(fry(word))
    print(''.join(words))
```

Create a `words` variable to hold the transformed words.

Split each line into words and non-words.

Add the fried word to the `words` list.

Print a new string of the joined words.

That (or something like it) should work well enough to pass the tests. Once you have a version that works, see if you can rewrite the `for` loop as a list comprehension and a `map()`.

Alrighty! Time to bear down and write this.

15.2 Solution

This reminds me of when Robin Hood's mate Friar Tuck was captured by the Sheriff of Nottingham. The Friar was sentenced to be boiled in oil, to which he replied "You can't boil me, I'm a friar!"

```
#!/usr/bin/env python3
"""Kentucky Friar"""

import argparse
import os
import re

# -----
def get_args():
    """Get command-line arguments"""
    parser = argparse.ArgumentParser(
        description='Southern fry text',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='text', help='Input text or file')

    args = parser.parse_args()

    if os.path.isfile(args.text):
        args.text = open(args.text).read()

    return args
```

If the argument is a file,
replace the text value with
the contents from the file.

Use the
`str.splitlines()`
method to
preserve the
line breaks in
the input text.

```
# -----
def main():
    """Make a jazz noise here"""
    args = get_args()           ← Get the command-line arguments.  
                            The text value will either be the  
                            command-line text or the  
                            contents of a file by this point.

    for line in args.text.splitlines():
        print('.join(map(fry, re.split(r'(\W+)', line.rstrip()))))') ←
```

Define a `fry()`
function that
will handle
one word.

```
# -----
def fry(word):
    """Drop the 'g' from '-ing' words, change 'you' to 'y'all'"""
```

Map the pieces of text split by the regular expression through the
`fry()` function, which will return the words modified as needed. Use
`str.join()` to turn that resulting list back into a string to print.

```

    ing_word = re.search('(.+)ing$', word)
    you = re.match('([Yy])ou$', word)

    if ing_word:
        prefix = ing_word.group(1)
        if re.search('[aeiouy]', prefix, re.IGNORECASE):
            return prefix + "in"
    elif you:
        return you.group(1) + "all"

    return word

```

Check if the search for “ing” returned a match.

Get the prefix (the bit before the “ing”), which is in group number 1.

Return the captured first character plus “all.”

Otherwise, return the word unaltered.

```

# -----
def test_fry():
    """Test fry"""

    assert fry('you') == "y'all"
    assert fry('You') == "Y'all"
    assert fry('fishing') == "fishin'"
    assert fry('Aching') == "Achin'"
    assert fry('swing') == "swing"

```

The tests for the `fry()` function

Check if the match for “you” succeeded.

Append “in” to the prefix and return it to the caller.

```

# -----
if __name__ == '__main__':
    main()

```

Search for “you” or “You” starting from the beginning of word. Capture the [yY] alternation in a group.

Search for “ing” anchored to the end of word. Use a capture group to remember the part of the string before the “ing.”

Perform a case-insensitive search for a vowel (plus “y”) in the prefix. If nothing is found, `None` will be returned, which evaluates to `False` in this Boolean context. If a match is returned, the `not None` value will evaluate to `True`.

15.3 Discussion

Again, there is nothing new in `get_args()`, so let’s just move to breaking the text into lines. In several previous exercises, I used a technique of reading an input file into the `args.text` value. If the input is coming from a file, there will be newlines separating each line of text. I suggested using a `for` loop to handle each line of input text returned by `str.splitlines()` to preserve the newlines in the output. I also suggested you start with a second `for` loop to handle each word-like unit returned by the `re.split()`:

```

for line in args.text.splitlines():
    words = []
    for word in re.split(r'(\W+)', line.rstrip()):
        words.append(fry(word))
    print(''.join(words))

```

That's five lines of code that could be written in two if we replace the second `for` with a list comprehension:

```
for line in args.text.splitlines():
    print(''.join([fry(w) for w in re.split(r'(\W+)', line.rstrip())]))
```

Or it could be slightly shorter using a `map()`:

```
for line in args.text.splitlines():
    print(''.join(map(fry, re.split(r'(\W+)', line.rstrip()))))
```

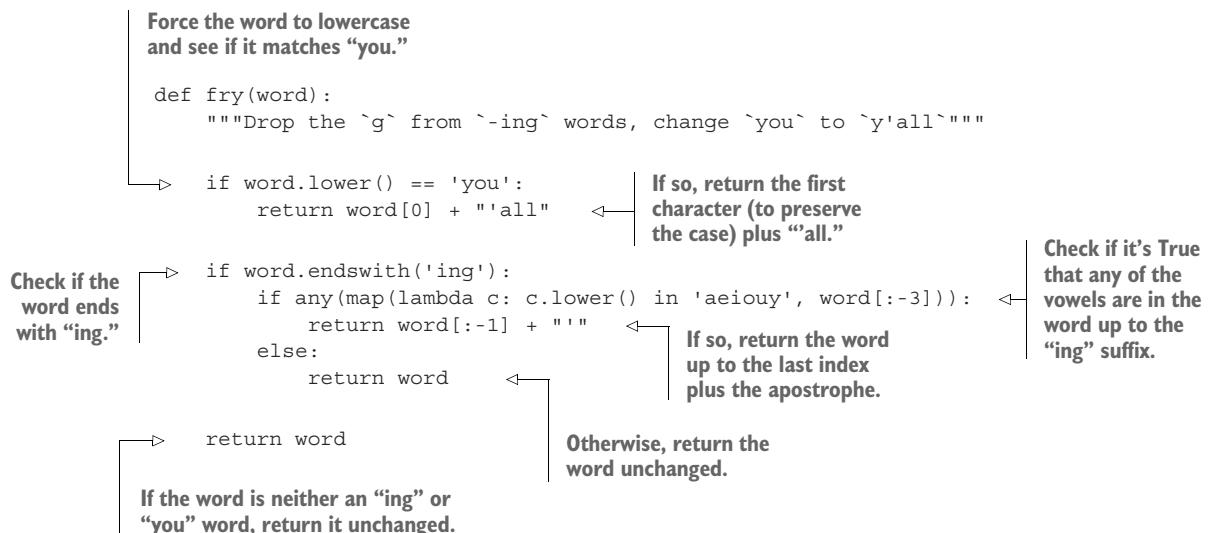
One other way to slightly improve readability is to use the `re.compile()` function to compile the regular expression. When you use the `re.split()` function inside the `for` loop, the regex must be compiled anew each iteration. By compiling the regex first, the compilation happens just once, so your code is (maybe just slightly) faster. More importantly, though, I think this is slightly easier to read, and the benefits are greater when the regex is more complicated:

```
splitter = re.compile(r'(\W+)')
for line in args.text.splitlines():
    print(''.join(map(fry, splitter.split(line.rstrip())))))
```

15.3.1 Writing the `fry()` function manually

You were not required, of course, to write a `fry()` function. However you wrote your solution, I hope you wrote tests for it!

The following version is fairly close to some of the suggestions I made earlier in the chapter. This version uses no regular expressions:



Let's take a moment to appreciate the `any()` function as it's one of my favorites. The preceding code uses a `map()` to check if each of the vowels exists in the portion of the word before the "ing" ending:

```
>>> word = "cooking"
>>> list(map(lambda c: (c, c.lower() in 'aeiouy'), word[:-3]))
[('c', False), ('o', True), ('o', True), ('k', False)]
```

The first character of "cooking" is "c," and it does not appear in the string of vowels. The next two characters ("o") do appear in the vowels, but "k" does not.

Let's reduce this to just the True/False values:

```
>>> list(map(lambda c: c.lower() in 'aeiouy', word[:-3]))
[False, True, True, False]
```

Now we can use `any` to tell us if *any* of the values are True:

```
>>> any([False, True, True, False])
True
```

It's the same as joining the values with `or`:

```
>>> False or True or True or False
True
```

The `all()` function returns True only if *all* the values are true:

```
>>> all([False, True, True, False])
False
```

That's the same as joining those values on `and`:

```
>>> False and True and True and False
False
```

If it's True that one of the vowels appears in the first part of the word, we have determined that this is (probably) a two-syllable word, and we can return the word with the final "g" replaced with an apostrophe. Otherwise, we return the unaltered word:

```
if any(map(lambda c: c.lower() in 'aeiouy', word[:-3])):
    return word[:-1] + "'"
else:
    return word
```

This approach works fine, but it's quite manual as we have to write quite a bit of code to find our patterns.

15.3.2 Writing the `fry()` function with regular expressions

Let's revisit the version of the `fry()` function that uses regular expressions:

The `re.match()` starts matching at the beginning of the given word, and it is looking for either an upper- or lowercase "y" followed by "ou" and then the end of the string (\$).

We use `re.search()` to look anywhere in the prefix for any of the vowels (plus "y") in a case-insensitive fashion. Remember that `re.match()` would start at the beginning of word, which is not what we want.

The pattern '(.+)ing\$' matches one or more of anything followed by "ing." The dollar sign anchors the pattern to the end of the string, so this is looking for a string that ends in "ing," but the string cannot just be "ing" as it has to have at least one of something before it. The parentheses capture the part before the "ing."

```
def fry(word):
    """Drop the 'g' from `-ing` words, change `you` to `y'all`"""

    ing_word = re.search('(.+)ing$', word)
    you = re.match('([Yy])ou$', word)
```

```
if ing_word:
    prefix = ing_word.group(1)
    if re.search('[aeiouy]', prefix, re.IGNORECASE):
        return prefix + "in"
elif you:
    return you.group(1) + "all"

return word
```

The prefix is the bit before the "ing" that we wrapped in parentheses. Because it is the first set of parentheses, we can fetch it with `ing_word.group(1)`.

If `ing_word` is `None`, that means it failed to match. If it is not `None` (so it is "truthy"), that means it is a `re.Match` object we can use.

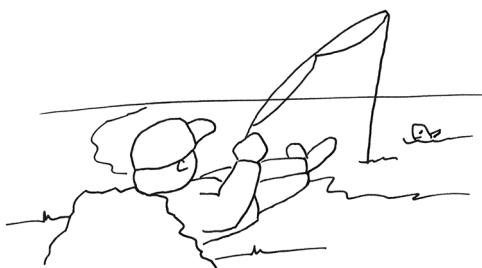
Return the prefix plus the string "in" so as to drop the final "g."

If `re.match()` for the "you" pattern fails, then "you" will be `None`. If it is not `None`, then it matched, and "you" is a `re.Match` object.

We used parentheses to capture the first character so as to maintain the case. That is, if the word was "You," we want to return "Y'all." Here we return that first group plus the string "all."

If the word matched neither a two-syllable "ing" pattern or the word "you," we return the word unchanged.

I've been using regexes for maybe 20 years, so this version seems much simpler to me than the manual version. You may feel differently. If you are completely new to regexes, trust me that they are so very worth the effort to learn. I absolutely would not be able to do much of my work without them.



15.4 Going further

- You could also replace "your" with "y'all's." For instance, "Where are your britches?" could become "Where are y'all's britches?"
- Change "getting ready" or "preparing" to "fixin'," as in "I'm getting ready to eat" to "I'm fixin' to eat." Also change the string "think" to "reckon," as in "I

think this is funny” to “I reckon this is funny.” You should also change “thinking” to “reckoning,” which then should become “reckonin’.” That means you either need to make two passes for the changes or find both “think” and “thinking” in the one pass.

- Make a version of the program for another regional dialect. I lived in Boston for a while and really enjoyed saying “wicked” all the time instead of “very,” as in “IT’S WICKED COLD OUT!”

Summary

- Regular expressions can be used to find patterns in text. The patterns can be quite complicated, like a grouping of non-word characters in between groupings of word characters.
- The `re` module has seriously handy functions like `re.match()` to find a pattern at the beginning of some text, `re.search()` to find a pattern anywhere inside some text, `re.split()` to break text on a pattern, and `re.compile()` to compile a regex so you can use it repeatedly.
- If you use capturing parentheses on the pattern for `re.split()`, the captured split pattern will be included in the returned values. This allows you to reconstruct the original string with the strings that are described by the pattern.

The scrambler: Randomly reordering the middles of words

Yuor brian is an azinamg cmiobiaontn of hdarware and sftraowe. You're rdineag tihgs right now eevn thgouh the wrdos are amses, but you're biran can mkae snese of it because the frsit and lsat ltrtees of each wrod hvae saeytd the smae. Yuor biran de'onst atlauclcy raed each lteetr of each wrod but rades wlohe wdors. The scrambeld wrdos difteienly solw you dwon, but y'roue not rlleay eevn tyinrg to ulsrmbance the lrttees, are you? It jsut hnaepps!



In this chapter, you will write a program called `scrambler.py` that will scramble each word of the text given as an argument. The scrambling should only work on words with four characters or more, and it should only scramble the letters in the middle of the word, leaving the first and last characters unchanged. The program should take an `-s` or `--seed` option (an `int` with default `None`) to pass to `random.seed()`.

It should handle text on the command line:

```
$ ./scrambler.py --seed 1 "foobar bazquux"  
faobor buuzaqx
```

Or text from a file:

```
$ cat ../inputs/spiders.txt  
Don't worry, spiders,  
I keep house  
casually.  
$ ./scrambler.py ../inputs/spiders.txt
```

```
D'not wrory, sdireps,
I keep hsueo
csalluay.
```

Figure 16.1 shows a string diagram to help you think about it.

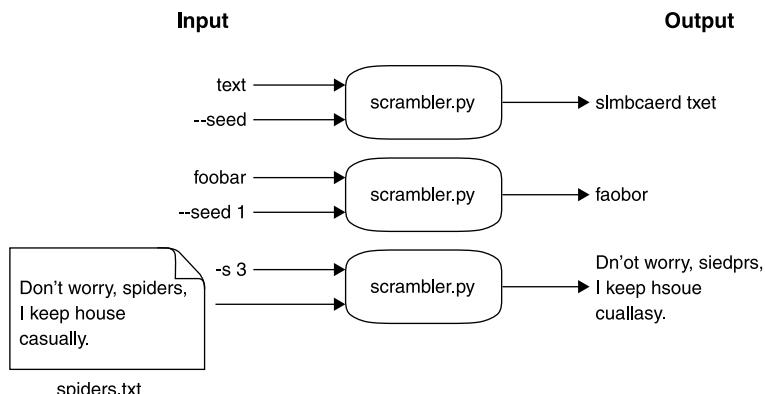


Figure 16.1 Our program will take input text from the command line or a file and will scramble the letters in words with four or more characters.

In this chapter you will

- Use a regular expression to split text into words
- Use the `random.shuffle()` function to shuffle a list
- Create scrambled versions of words by shuffling the middle letters while leaving the first and last letters unchanged

16.1 Writing *scrambler.py*

I recommend you start by using `new.py` `scrambler.py` to create the program in the `16_scrambler` directory. Alternatively, you can copy `template/template.py` to `16_scrambler/scrambler.py`. You can refer to previous exercises, like the one in chapter 5, to remember how to handle a positional argument that might be text or might be a text file to read.

When run with no arguments or the flags `-h` or `--help`, `scrambler.py` should present a usage statement:

```
$ ./scrambler.py -h
usage: scrambler.py [-h] [-s seed] text
```

Scramble the letters of words

```
positional arguments:
  text                  Input text or file
```

```
optional arguments:
  -h, --help            show this help message and exit
  -s seed, --seed seed  Random seed (default: None)
```

Once your program's usage statement matches this, change your `main()` definition as follows:

```
def main():
    args = get_args()
    print(args.text)
```

Then verify that your program can echo text from the command line:

```
$ ./scrambler.py hello
hello
```

Or from an input file:

```
$ ./scrambler.py ../inputs/spiders.txt
Don't worry, spiders,
I keep house
casually.
```

16.1.1 *Breaking the text into lines and words*

As in chapter 15, we want to preserve the line breaks of the input text by using `str.splitlines()`:

```
for line in args.text.splitlines():
    print(line)
```

If we are reading the `spiders.txt` haiku, this is the first line:

```
>>> line = "Don't worry, spiders,"
```

We need to break the line into words. In chapter 6 we used `str.split()`, but that approach leaves punctuation stuck to our words—both `worry` and `spiders` have commas:

```
>>> line.split()
['Don\'t', 'worry,', 'spiders,']
```

In chapter 15 we used the `re.split()` function with the regular expression `(\W+)` to split text on one or more non-word characters. Let's try that:

```
>>> re.split('(\W+)', line)
['Don', "'", 't', ' ', 'worry', ' ', ' ', 'spiders', ' ', '']
```

That won't work because it splits `Don't` into three parts: `Don`, `'`, and `t`.

Perhaps we could use `\b` to break on *word boundaries*. Note that we'd have to put an `r''` in front of the first quote, `r'\b'`, to denote that it is a “raw” string.

This still won't work because \b thinks the apostrophe is a word boundary and so splits the contracted word:

```
>>> re.split(r'\b', "Don't worry, spiders,")
['', 'Don', "", 't', ' ', 'worry', ' ', 'spiders', ',']
```

While searching the internet for a regex to split this text properly, I found the following pattern on a Java discussion board. It perfectly separates *words* from *non-words*:¹

```
>>> re.split("( [a-zA-Z] (?: [a-zA-Z]* [a-zA-Z] )? )", "Don't worry, spiders,")
['', "Don't", ' ', 'worry', ' ', 'spiders', ',']
```

The beautiful thing about regular expressions is that they are their own language—one that is used inside many other languages from Perl to Haskell. Let's dig into this pattern, shown in figure 16.2.

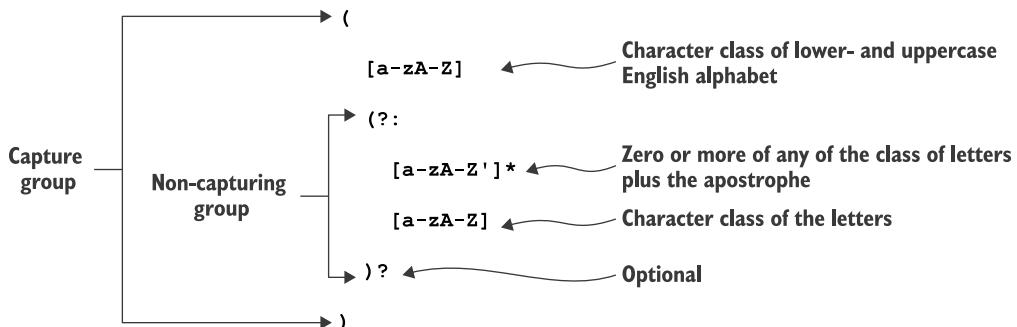


Figure 16.2 A regular expression that will find words that include an apostrophe

16.1.2 Capturing, non-capturing, and optional groups

In figure 16.2 you can see that groups can contain other groups. For instance, here is a regex that can capture the entire string “foobarbaz” as well as the substring “bar”:

```
>>> match = re.match('(foo(bar)baz)', 'foobarbaz')
```

Capture groups are numbered by the position of their left parenthesis. Since the first left parenthesis starts the capture starting at “f” and going to “z,” that is group 1:

```
>>> match.group(1)
'foobarbaz'
```

¹ I would like to stress that a significant part of my job is spent looking for answers both in the books I own but also on the internet!

The second left parenthesis starts just before the “b” and goes to the “r”:

```
>>> match.group(2)
'bar'
```

We can also make a group *non-capturing* by using the starting sequence `(?:`. If we use this sequence on the second group, we no longer capture the substring “bar”:

```
>>> match = re.match('(foo(?:bar)baz)', 'foobarbaz')
>>> match.groups()
('foobarbaz',)
```

Non-capturing groups are commonly used when you are grouping primarily for the purpose of making it optional by placing a ? after the closing parenthesis. For instance, we can make the “bar” optional and then match both “foobarbaz,”

```
>>> re.match('(foo(?:bar)?baz)', 'foobarbaz')
<re.Match object; span=(0, 9), match='foobarbaz'>
```

as well as “foobaz”:

```
>>> re.match('(foo(?:bar)?baz)', 'foobaz')
<re.Match object; span=(0, 6), match='foobaz'>
```

16.1.3 Compiling a regex

I mentioned the `re.compile()` function in chapter 15 as a way to incur the cost of compiling a regular expression just once. Whenever you use something like `re.search()` or `re.split()`, the regex engine must parse the `str` value you provide for the regex into something it understands and can use. This parsing step must happen *each time* you call the function. When you compile the regex and assign it to a variable, the parsing step is done before you call the function, which improves performance.

I especially like to use `re.compile()` to assign a regex to a meaningful variable name and/or reuse the regex in multiple places in my code. Because this regex is quite long and complicated, I think it makes the code more readable to assign it to a variable called `splitter`, which will help me to remember how it will be used:

```
>>> splitter = re.compile("([a-zA-Z](?:[a-zA-Z])*[a-zA-Z])?")
>>> splitter.split("Don't worry, spiders,")
['', "Don't", ' ', 'worry', ' ', 'spiders', ',']
```

16.1.4 Scrambling a word

Now that we have a way to process the *lines* and then *words* of the text, let's think about how we'll scramble the words by starting with just *one word*. You and I will need to use the same algorithm for scrambling the words in order to pass the tests, so here are the rules:

- If the word is three characters or shorter, return the word unchanged.
- Use a string slice to copy the characters, not including the first and last.
- Use the `random.shuffle()` method to mix up the letters in the middle.
- Return the new “word” by combining the first, middle, and last parts.



I recommend you create a function called `scramble()` that will do all this, and also create a test for it. Feel free to add this to your program:

```
def scramble(word) :           ← The pass is a no-op (no operation), so this function
    """Scramble a word"""\n    pass\n\n    ← Set random.seed() to a known value for testing.\n\n    def test_scramble():\n        """Test scramble"""\n        state = random.getstate()\n        random.seed(1)\n        assert scramble("a") == "a"\n        assert scramble("ab") == "ab"\n        assert scramble("abc") == "abc"\n        assert scramble("abcd") == "acbd"\n        assert scramble("abcde") == "acbde"\n        assert scramble("abcdef") == "aecbdf"\n        assert scramble("abcde'f") == "abcd'ef"\n        random.setstate(state)\n\n        ← Restore the state to the previous value.\n\n    ← The change we'll make by setting the random.seed() in the next line will be a global change. We'll want to restore the state after testing, so here we use random.getstate() to get the current state of the random module.\n\n    ← Words with three characters or fewer should be returned unchanged.\n\n    ← This word looks unchanged, but that's just because with the seed of 1 the shuffling didn't end up changing the middle characters.\n\n    ← Now it's more evident that the word is being scrambled.
```

Inside the `scramble()` function, we will have a word like “worry.” We can use list slices to extract part of a string. Since Python starts numbering at 0, we use 1 to indicate the *second* character:

```
>>> word = 'worry'\n>>> word[1]\n'o'
```

The last index of any string is -1:

```
>>> word[-1]
'Y'
```

To get a slice, we use the `list[start:stop]` syntax. Since the stop position is not included, we can get the middle like so:

```
>>> middle = word[1:-1]
>>> middle
'orr'
```

We can import `random` to get access to the `random.shuffle()` function. As with the `list.sort()` and `list.reverse()` methods, the argument will be shuffled *in place*, and the function will return `None`. That is, you might be tempted to write code like this:

```
>>> import random
>>> x = [1, 2, 3]
>>> shuffled = random.shuffle(x)
```

What is the value of `shuffled`? Is it something like `[3, 1, 2]`, or is it `None`?

```
>>> type(shuffled)
<class 'NoneType'>
```

The `shuffled` value now holds `None`, while the `x` list has been shuffled *in place* (see figure 16.3):

```
>>> x
[2, 3, 1]
```

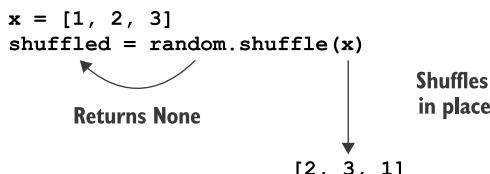


Figure 16.3 The return from `random.shuffle()` was `None`, so `shuffled` was assigned `None`.

If you've been following along, it turns out that we cannot shuffle the `middle` like this:

```
>>> random.shuffle(middle)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/Users/kyclark/anaconda3/lib/python3.7/random.py", line 278, in shuffle
    x[i], x[j] = x[j], x[i]
TypeError: 'str' object does not support item assignment
```

The `middle` variable is a `str`:

```
>>> type(middle)
<class 'str'>
```

The `random.shuffle()` function is trying to directly modify a `str` value in place, but `str` values in Python are *immutable*. One workaround is to make `middle` into a new list of the characters from `word`:

```
>>> middle = list(word[1:-1])
>>> middle
['o', 'r', 'r']
```

That is something we can shuffle:

```
>>> random.shuffle(middle)
>>> middle
['r', 'o', 'r']
```

Then it's a matter of creating a new string with the original first letter, the shuffled `middle`, and the last letter. I'll leave that for you to work out.

Use `pytest scrambler.py` to have Pytest execute the `test_scramble()` function to see if it works correctly. Run this command *after every change to your program*. Ensure that your program always compiles and runs properly. Only make one change at a time, and then save your program and run the tests.

16.1.5 Scrambling all the words

As in several previous exercises, we're now down to applying the `scramble()` function to all the words. Can you see a familiar pattern?

```
splitter = re.compile("([a-zA-Z](?:[a-zA-Z']*[a-zA-Z])?)")
for line in args.text.splitlines():
    for word in splitter.split(line):
        # what goes here?
```

We've talked about how to apply a function to each element in a sequence. You might try a `for` loop, a list comprehension, or maybe a `map()`. Think about how you can split the text into words, feed them to the `scramble()` function, and then join them back together to reconstruct the text.

Note that this approach will pass both the words and the non-words (the bits in between each word) to the `scramble()` function. You don't want to modify the non-words, so you'll need a way to check that the argument looks like a word. Maybe a regular expression?

That should be enough to go on. Write your solution and use the included tests to check your program.

16.2 Solution

To me, the program comes down to properly splitting the words and then figuring out the `scramble()` function. Then it's a matter of applying the function and reconstructing the text.

```
#!/usr/bin/env python3
"""Scramble the letters of words"""

import argparse
import os
import re
import random

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Scramble the letters of words',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='text', help='Input text or file') ←

    parser.add_argument('-s',
                        '--seed',
                        help='Random seed',
                        metavar='seed',
                        type=int,
                        default=None) ←
    The seed option is an int
    that defaults to None.

    args = parser.parse_args() ←
    Get the arguments so we
    can check the text value.

    if os.path.isfile(args.text): ←
        args.text = open(args.text).read().rstrip()

    return args ←
    Return the arguments
    to the caller.

# -----
def main():
    """Make a jazz noise here"""

    args = get_args() ←
    Get the
    command-line
    arguments.

    random.seed(args.seed) ←
    Use args.seed to set the
    random.seed() value. If args.seed
    is the default None, this is the
    same as not setting the seed.

    splitter = re.compile("(([a-zA-Z](?:[a-zA-Z']*[a-zA-Z])?)") ←
    Use str.splitlines() to
    preserve the line breaks
    in the input text.

    for line in args.text.splitlines():
        print(''.join(map(scramble, splitter.split(line)))) ←
    Use the splitter to break the line into a new
    list that map() will feed into the scramble()
    function. Join the resulting list on the empty
    string to create a new str to print.

    Save the
    compiled
    regex into
    a variable. →
```

```

Define a
function to
scramble()
a single
word.    # -----
        def scramble(word):
            """For words over 3 characters, shuffle the letters in the middle"""

            if len(word) > 3 and re.match(r'\w+', word):
                middle = list(word[1:-1])
                random.shuffle(middle)
                word = word[0] + ''.join(middle) + word[-1]

            return word

Shuffle the
middle
letters.      # -----  

# -----  

def test_scramble():
    """Test scramble"""

    random.seed(1)
    assert scramble("a") == "a"
    assert scramble("ab") == "ab"
    assert scramble("abc") == "abc"
    assert scramble("abcd") == "acbd"
    assert scramble("abcde") == "acbde"
    assert scramble("abcdef") == "aecbdf"
    assert scramble("abcde'f") == "abcd'ef"
    random.seed(None)

# -----  

if __name__ == '__main__':
    main()

```

Only scramble words with four or more characters if they contain word characters.

Copy the second through the second-to-last characters of the word into a new list called middle.

Set the word equal to the first character, plus the middle, plus the last character.

Return the word, which may have been altered if it met the criteria.

The test for the scramble() function

16.3 Discussion

There is nothing new in `get_args()`, so I trust you'll understand that code. Refer to chapter 5 if you want to revisit how to handle the `args.text` coming from the command line or from a file.

16.3.1 Processing the text

As mentioned earlier in the chapter, I often assign a *compiled* regex to a variable. Here I did it with the splitter:

```
splitter = re.compile("([a-zA-Z](?:[a-zA-Z']*[a-zA-Z])?)")
```

The other reason I like to use `re.compile()` is because I feel it can make my code more readable. Without it, I would have to write this:

```
for line in args.text.splitlines():
    print(''.join(map(scramble, re.split("([a-zA-Z](?:[a-zA-Z']*[a-zA-Z])?)", line))))
```

That ends up creating a line of code that is 86 characters wide, and the PEP 8 style guide (www.python.org/dev/peps/pep-0008/) recommends we “limit all lines to a maximum of 79 characters.” I find the following version much easier to read:

```
splitter = re.compile("([a-zA-Z](?:[a-zA-Z] *[a-zA-Z])?)")
for line in args.text.splitlines():
    print(''.join(map(scramble, splitter.split(line))))
```

You may still find that code somewhat confusing. Figure 16.4 shows the flow of the data:

- 1 First Python will split the string "Don't worry, spiders,".
- 2 The splitter creates a new list composed of words (that matched our regex) and non-words (the bits in between).
- 3 The map() function will apply the scramble() function to each element of the list.
- 4 The result of map() is a new list with the results of each application of the scramble() function.
- 5 The result of str.join() is a new string, which is the argument to print().

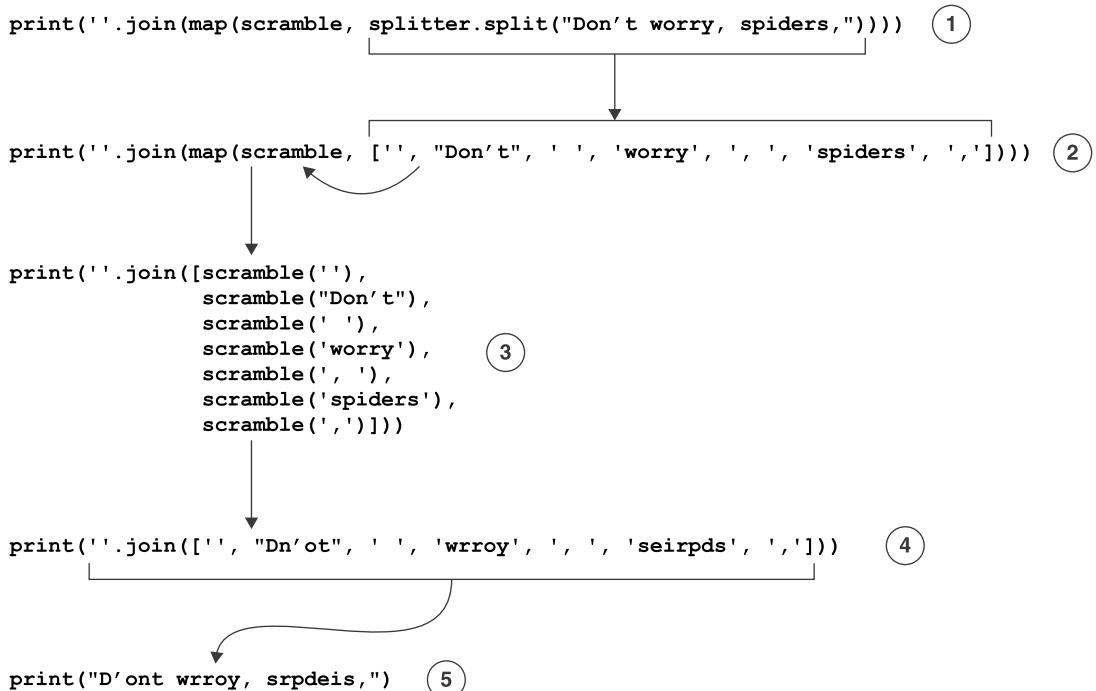


Figure 16.4 A visualization of how data moves through the map() function

A longer way to write this with a `for` loop might look like this:

```
Use str.splitlines() to preserve  
the original line breaks.  
  
for line in args.text.splitlines():  
    words = []  
    for word in splitter.split(line):  
        words.append(scramble(word))  
    print(''.join(words))  
  
Join the words on the empty string  
and pass the result to print().
```

Because the goal is to create a new list, this is better written as a list comprehension:

```
for line in args.text.splitlines():  
    words = [scramble(word) for word in splitter.split(line)]  
    print(''.join(words))
```

Or you could go in quite the opposite direction and replace all the `for` loops with `map()`:

```
print('\n'.join(  
    map(lambda line: ''.join(map(scramble, splitter.split(line))),  
        args.text.splitlines())))
```

This last solution reminds me of a programmer I used to work with who would jokingly say, “If it was hard to write, it should be hard to read!” It becomes somewhat clearer if you rearrange the code. Note that Pylint will complain about assigning a `lambda`, but I really don’t agree with that criticism:

```
scrambler = lambda line: ''.join(map(scramble, splitter.split(line)))  
print('\n'.join(map(scrambler, args.text.splitlines())))
```

Writing code that is correct, tested, and understandable is as much an art as it is a craft. Choose the version that you (and your teammates!) believe is the most readable.

16.3.2 Scrambling a word

Let’s take a closer look at my `scramble()` function. I wrote it in a way that would make it easy to incorporate into `map()`:

Check if the given word is one I ought to scramble. First, it must be longer than three characters. Second, it must contain one or more word characters because the function will be passed both “word” and “non-word” strings. If either check returns False, I will return the word unchanged. The `r'\w+'` is used to create a “raw” string. Note that the regex works fine with or without it being a raw string, but Pylint complains about an “invalid escape character” unless it is a raw string.

Copy the middle of the word to a new list called `middle`.

```
def scramble(word):  
    """For words over 3 characters, shuffle the letters in the middle"""  
    if len(word) > 3 and re.match(r'\w+', word):  
        middle = list(word[1:-1])
```

```

random.shuffle(middle)
word = word[0] + ''.join(middle) + word[-1]
return word

```

Return the word, which may or may not have been shuffled.

Reconstruct the word by joining together the first character, the shuffled middle, and the last character.

Shuffle the middle in place. Remember that this function returns None.

16.4 Going further

- Write a version of the program where the `scramble()` function sorts the middle letters into alphabetical order rather than shuffling them.
- Write a version that reverses each word rather than scrambles them.
- Write a program to *unscramble* the text. For this, you need to have a dictionary of English words, which I have provided as `inputs/words.txt.zip`. You will need to split the scrambled text into words and non-words, and then compare each “word” to the words in your dictionary. I recommend you start by comparing the words as anagrams (that is, they have the same composition/frequency of letters) and then using the first and last letters to positively identify the unscrambled word.

Summary

- The regex we used to split the text into words was quite complex, but it also gave us exactly what we needed. Writing the program without this piece would have been significantly more difficult. Regexes, while complex and deep, are wildly powerful black magic that can make your programs incredibly flexible and useful.
- The `random.shuffle()` function accepts a list, which is mutated in place.
- List comprehensions and `map()` can often lead to more compact code, but going too far can reduce readability. Choose wisely.



17

Mad Libs: *Using regular expressions*

When I was a wee lad, we used to play at Mad Libs for hours and hours. This was before computers, mind you, before televisions or radio or even paper! No, scratch that, we had paper. Anyway, point is we only had Mad Libs to play, and we loved it! And now you must play!

In this chapter, we'll write a program called `mad.py` that will read a file given as a positional argument and find all the placeholders in angle brackets, like `<verb>` or `<adjective>`. For each placeholder, we'll prompt the user for the part of speech being requested, like "Give me a verb" and "Give me an adjective." (Notice that you'll need to use the correct article, just as in chapter 2.) Each value from the user will then replace the placeholder in the text, so if the user says "drive" for the verb, then `<verb>` in the text will be replaced with `drive`. When all the placeholders have been replaced with inputs from the user, we'll print out the new text.

There is a `17_mad_libs/inputs` directory with some sample files you can use, but I also encourage you to create your own. For instance, here is a version of the "fox" text:

```
$ cd 17_mad_libs  
$ cat inputs/fox.txt  
The quick <adjective> <noun> jumps <preposition> the lazy <noun>.
```



When the program is run with this file as the input, it will ask for each of the place-holders and then print the silliness:

```
$ ./mad.py inputs/fox.txt
Give me an adjective: surly
Give me a noun: car
Give me a preposition: under
Give me a noun: bicycle
The quick surly car jumps under the lazy bicycle.
```

By default, this is an interactive program that will use the `input()` prompt to ask the user for their answers, but for testing purposes we will have an `-i` or `--inputs` option so the test suite can pass in all the answers and bypass the interactive `input()` calls:

```
$ ./mad.py inputs/fox.txt -i surly car under bicycle  
The quick surly car jumps under the lazy bicycle.
```

In this exercise, you will

- Learn to use `sys.exit()` to halt your program and indicate an error status
 - Learn about greedy matching with regular expressions
 - Use `re.findall()` to find all matches for a regex
 - Use `re.sub()` to replace found patterns with new text
 - Explore ways to write the solution without using regular expressions

17.1 Writing mad.py

To start off, create the program `mad.py` in the `17_mad_libs` directory using `new.py` or by copying `template/template.py` to `17_mad_libs/mad.py`. You would also do well to define the positional `file` argument as a readable text file using `type=argparse.FileType('rt')`. The `-i` or `--inputs` option should use `nargs='*' to define a list of zero or more str values.`

After this, your program should be able to produce a usage statement when given no arguments or the `-h` or `--help` flag:

If the given file argument does not exist, the program should error out:

```
$ ./mad.py blargh
usage: mad.py [-h] [-i [str [str ...]]] FILE
mad.py: error: argument FILE: can't open 'blargh': \
[Errno 2] No such file or directory: 'blargh'
```

If the text of the file contains no <> placeholders, the program should print a message and exit with an error value (something other than 0). Note that this error does not need to print a usage statement, so you don't have to use `parser.error()` as in previous exercises:

```
$ cat no_blanks.txt
This text has no placeholders.
$ ./mad.py no_blanks.txt
"no_blanks.txt" has no placeholders.
```

Figure 17.1 shows a string diagram to help you visualize the program.

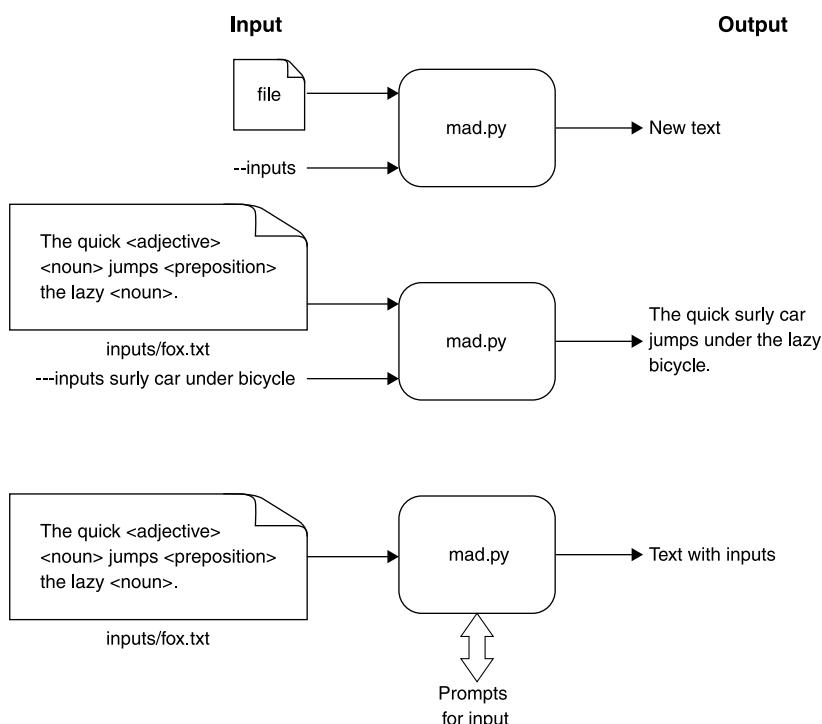


Figure 17.1 The Mad Libs program must have an input file. It may also have a list of strings for the substitutions or it will interactively ask the user for the values.

17.1.1 Using regular expressions to find the pointy bits

We've talked before about the possible dangers of reading an entire file into memory. Because we'll be parsing the text to find all the <...> bits in this program, we'll really need to read the whole file at once. We can do this by chaining the appropriate functions like so:

```
>>> text = open('inputs/fox.txt').read().rstrip()
>>> text
'The quick <adjective> <noun> jumps <preposition> the lazy <noun>.'
```

We're looking for patterns of text inside angle brackets, so let's use a regular expression. We can find a literal < character like so (see figure 17.2):

```
>>> import re
>>> re.search('<', text)
<re.Match object; span=(10, 11), match='<'>
```

The quick <adjective> <noun> jumps <preposition> the lazy <noun>.

Figure 17.2 Matching a literal less-than sign

Now let's find that bracket's mate. The . in a regular expression means "anything," and we can add a + after it to mean "one or more." I'll capture the match so it's easier to see:

```
>>> match = re.search('<.+>', text)
>>> match.group(1)
'<adjective> <noun> jumps <preposition> the lazy <noun>'
```

As shown in figure 17.3, that matched all the way to the end of the string instead of stopping at the first available >. It's common when you use * or + for zero, one, or more for the regex engine to be "greedy" on the *or more* part. The pattern matches beyond where we wanted, but it is technically matching exactly what we described. Remember that . means *anything*, and a right angle bracket (or greater-than sign) is

The quick <adjective> <noun> jumps <preposition> the lazy <noun>.

Figure 17.3 The plus sign to match one or more is a greedy match, matching as many characters as possible.

“anything.” It matches as many characters as possible until it finds the last right angle bracket to stop at, which is why this pattern is called “greedy.”

We can make the regex “non-greedy” by changing `+` to `+?` so that it matches the shortest possible string (see figure 17.4):

```
>>> re.search('<.+?>', text)
<re.Match object; span=(10, 21), match='<adjective>'>
```

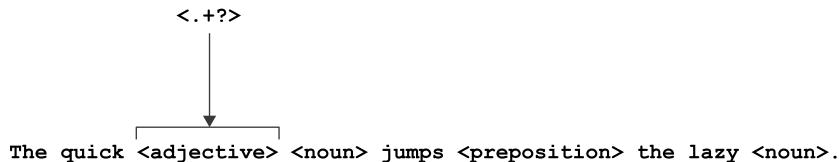


Figure 17.4 The question mark after the plus sign makes the regex stop at the shortest possible match.

Rather than using `.` for “anything,” it would be more accurate to say that we want to match one or more of anything *that is not either of the angle brackets*. The character class `[<>]` would match either bracket. We can negate (or complement) the class by putting a caret (`^`) as the first character, so we have `[^<>]` (see figure 17.5). That will match anything that is not a left or right angle bracket:

```
>>> re.search('<[^<>]+>', text)
<re.Match object; span=(10, 21), match='<adjective>'>
```

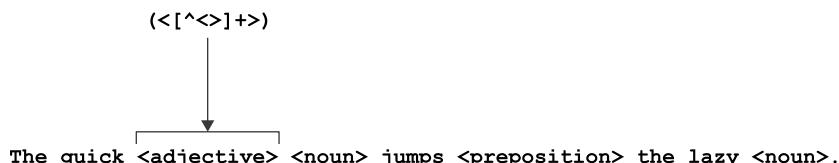


Figure 17.5 A negated character class to match anything other than the angle brackets

Why do we have both brackets inside the negated class? Wouldn’t the right bracket be enough? Well, I’m guarding against *unbalanced* brackets. With only the right bracket, it would match this text (see figure 17.6):

```
>>> re.search('<[^>]+>', 'foo <<bar> baz')
<re.Match object; span=(4, 10), match='<<bar>'>
```

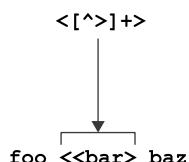


Figure 17.6 This regex leaves open the possibility of matching unbalanced brackets.

But with *both* brackets in the negated class, it finds the correct, balanced pair (see figure 17.7):

```
>>> re.search('<[^<>]+>', 'foo <<bar> baz')
<re.Match object; span=(5, 10), match='<bar>'>
```

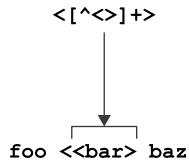


Figure 17.7 This regex finds the correctly balanced brackets and contained text.

We'll add two sets of parentheses `()`. The first will capture the *entire* placeholder pattern (see figure 17.8):

```
>>> match = re.search('(<([^\<>]+)>)', text)
>>> match.groups()
('<adjective>', 'adjective')
```

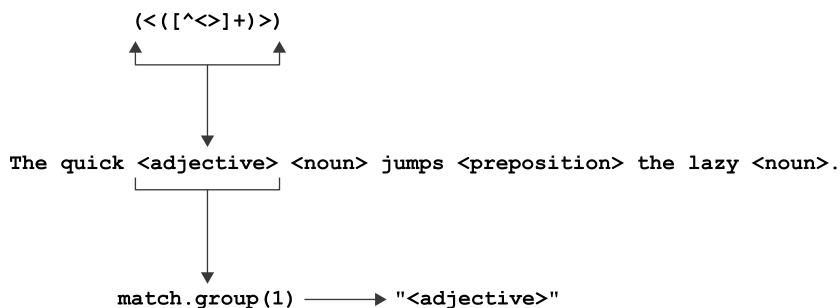


Figure 17.8 The outer parentheses capture the brackets and text.

The other is for the string *inside* the `<>` (see figure 17.9):

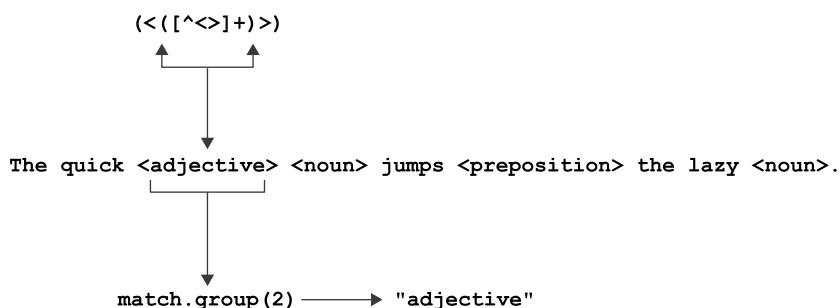


Figure 17.9 The inner parentheses capture just the text.

There is a very handy function called `re.findall()` that will return all matching text groups as a list of tuple values:

```
>>> from pprint import pprint
>>> matches = re.findall('<([<>]+)>', text)
>>> pprint(matches)
[('<adjective>', 'adjective'),
 ('<noun>', 'noun'),
 ('<preposition>', 'preposition'),
 ('<noun>', 'noun')]
```

Note that the capture groups are returned in the order of their opening parentheses, so the entire placeholder is the first member of each tuple, and the contained text is the second. We can iterate over this list, *unpacking* each tuple into variables (see figure 17.10):

```
>>> for placeholder, name in matches:
...     print(f'Give me {name}')
...
Give me adjective
Give me noun
Give me preposition
Give me noun
```

```
for placeholder, name in [('<adjective>', 'adjective')]:
    print(f'Give me {name}')
```

Figure 17.10 Since the list contains 2-tuples, we can unpack them into two variables in the `for` loop.

You should insert the correct article (“a” or “an,” as you did in chapter 2) to use as the prompt for `input()`.

17.1.2 Halting and printing errors

If we find there are no placeholders in the text, we need to print an error message. It’s common to print error messages to `STDERR` (standard error), and the `print()` function allows us to specify a `file` argument. We’ll use `sys.stderr`, just as we did in chapter 9. To do that, we need to import that module:

```
import sys
```

You may recall that `sys.stderr` is like an already open file handle, so there’s no need to `open()` it:

```
print('This is an error!', file=sys.stderr)
```

If there really are no placeholders, we should exit the program with an error value to indicate to the operating system that the program failed to run properly. The normal exit value for a program is 0, as in “zero errors,” so we need to exit with some `int` value that is *not* 0. I always use 1:

```
sys.exit(1)
```

One of the tests checks whether your program can detect missing placeholders and if your program exits correctly.

You can also call `sys.exit()` with a string value, in which case the string will be printed to `sys.stderr` and the program will exit with the value 1:

```
sys.exit('This will kill your program and print an error message!')
```

17.1.3 Getting the values

For each one of the parts of speech in the text, we need a value that will come either from the `--inputs` argument or directly from the user. If we have nothing for `--inputs`, we can use the `input()` function to get an answer from the user.

The `input()` function takes a `str` value to use as a prompt:

```
>>> value = input('Give me an adjective: ')
Give me an adjective: blue
```

And it returns a `str` value of whatever the user typed before pressing the Return key:

```
>>> value
'blue'
```

If, however, we have values for the inputs, we can use those and not bother with the `input()` function. I’m only making you handle the `--inputs` option for testing purposes. You can safely assume that you will always have the same number of inputs as you have placeholders (see figure 17.11).

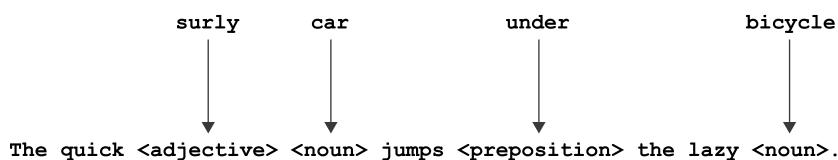


Figure 17.11 If given inputs from the command line, they will match up with the placeholders in the text.

For instance, you might have the following as the `--inputs` option to your program for the `fox.txt` example:

```
>>> inputs = ['surly', 'car', 'under', 'bicycle']
```

You need to remove and return the first string, “surly,” from inputs. The `list.pop()` method is what you need, but it wants to remove the *last* element by default:

```
>>> inputs.pop()
'bicycle'
```

The `list.pop()` method takes an optional argument to indicate the index of the element you want to remove. Can you figure out how to make that work? Be sure to read `help(list.pop)` if you’re stuck.

17.1.4 Substituting the text

When you have values for each of the placeholders, you will need to substitute them into the text. I suggest you look into the `re.sub()` (substitute) function, which will replace any text matching a given regular expression with some other value. I definitely recommend you read `help(re.sub)`:

```
sub(pattern, repl, string, count=0, flags=0)
    Return the string obtained by replacing the leftmost
    non-overlapping occurrences of the pattern in string by the
    replacement repl.
```

I don’t want to give away the ending, but you will need to use a pattern similar to the preceding to replace each `<placeholder>` with each value.

Note that it’s not a requirement that you use the `re.sub()` function to solve this. I challenge you, in fact, to try writing a solution that does not use the `re` module at all. Now go write the program, and use the tests to guide you!

17.2 Solution

Are you getting more comfortable with regular expressions? I know they are complicated, but really understanding them will help you more than you might expect.

```
#!/usr/bin/env python3
"""Mad Libs"""

import argparse
import re
import sys

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Mad Libs',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file',                         ← | The file argument should
                      metavar='FILE',                  | be a readable text file.
    )
    args = parser.parse_args()
    return args
```

```

        type=argparse.FileType('rt'),
        help='Input file')

parser.add_argument('-i',
                    '--inputs',
                    help='Inputs (for testing)',
                    metavar='input',
                    type=str,
                    nargs='*')

return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    inputs = args.inputs
    text = args.file.read().rstrip()           ←
    blanks = re.findall('<([^\>]+)>', text) ←

    if not blanks:
        sys.exit(f'{args.file.name} has no placeholders.')
    Check if there are no placeholders.

    tmp1 = 'Give me {} {}: '
    for placeholder, pos in blanks:
        article = 'an' if pos.lower()[0] in 'aeiou' else 'a'
        answer = inputs.pop(0) if inputs else input(tmp1.format(article, pos))
        text = re.sub(placeholder, answer, text, count=1)

        print(text)   ← Print the resulting text to STDOUT.

    # -----
    if __name__ == '__main__':
        main()

Replace the current placeholder text with the answer from the user. Use count=1 to ensure that only the first value is replaced. Overwrite the existing value of text so that all the placeholders will be replaced by the end of the loop.

If there are inputs, remove the first one for the answer; otherwise, use input() to prompt the user for a value.

Print a message to STDERR that the specified file contains no placeholders, and exit the program with a non-zero status to indicate an error to the operating system.

```

The --inputs option may have zero or more strings.

Open and read the input file, stripping off the trailing newline.

Use a regex to find all matches for a left angle bracket, followed by one or more of anything that is not a left or right angle bracket, followed by a right angle bracket. Use two capture groups to capture the entire expression and the text inside the brackets.

Choose the correct article based on the first letter of the name of the part of speech (pos): "an" for those starting with a vowel and "a" otherwise.

Iterate through the blanks, unpacking each tuple into variables.

Create a string template for the prompt to ask for input() from the user.

17.3 Discussion

We start off by defining our arguments well. The input file should be declared using `type=argparse.FileType('rt')` so that argparse will verify that the argument is a readable text file. The `--inputs` are optional, so we can use `nargs='*'` to indicate

zero or more strings. If no inputs are provided, the default value will be `None`, so be sure you don't assume it's a list and try doing list operations on a `None`.

17.3.1 Substituting with regular expressions

There is a subtle bug waiting for you in using `re.sub()`. Suppose we have replaced the first `<adjective>` with "blue" so that we have this:

```
>>> text = 'The quick blue <noun> jumps <preposition> the lazy <noun>.'
```

Now we want to replace `<noun>` with "dog," so we try this:

```
>>> text = re.sub('<noun>', 'dog', text)
```

Let's check on the value of `text` now:

```
>>> text
'The quick blue dog jumps <preposition> the lazy dog.'
```

Since there were two instances of the string `<noun>`, both got replaced with "dog," as shown in figure 17.12.

```
re.sub('<noun>', 'dog', 'The quick blue <noun> jumps <preposition> the lazy <noun>.')
```

Figure 17.12 The `re.sub()` function will replace all matches.

We must use `count=1` to ensure that only the first occurrence is changed (see figure 17.13):

```
>>> text = 'The quick blue <noun> jumps <preposition> the lazy <noun>.'
>>> text = re.sub('<noun>', 'dog', text, count=1)
>>> text
'The quick blue dog jumps <preposition> the lazy <noun>.'
```

```
re.sub('<noun>', 'dog', 'The quick blue <noun> jumps <preposition> the lazy <noun>.', count=1)
```

Figure 17.13 Use the `count` option to `re.sub()` to limit the number of replacements.

Now we can keep moving on to replace the other placeholders.

17.3.2 Finding the placeholders without regular expressions

I trust the explanation of the regex solution earlier in the chapter was sufficient. I find that solution fairly elegant, but it is certainly possible to solve this without using regexes. Here is how I might solve it manually.

First I need a way to search the text for <...>. I start off by writing a test that helps me imagine what I might give to my function and what I might expect in return for both good and bad values.

I decide to return `None` when the pattern is missing and to return a tuple of `(start, stop)` indices when the pattern is present:

```
def test_find_brackets():
    """Test for finding angle brackets"""
    assert find_brackets('') is None
    assert find_brackets('<>') is None
    assert find_brackets('<x>') == (0, 2)
    assert find_brackets('foo <bar> baz') == (4, 8)
```

Annotations for the test cases:

- `There is no text, so it should return None.`
- `There are angle brackets, but they lack any text inside, so this should return None.`
- `The pattern should be found at the beginning of a string.`
- `The pattern should be found further into the string.`

Now I need to write the code that will satisfy that test. Here is what I wrote:

```
def find_brackets(text):
    """Find angle brackets"""
    start = text.index('<') if '<' in text else -1
    stop = text.index('>') if start >= 0 and '>' in text[start + 2:] else -1
    return (start, stop) if start >= 0 and stop >= 0 else None
```

Annotations for the function logic:

- `Find the index of the left bracket if one is found in the text.`
- `Find the index of the right bracket if one is found starting two positions after the left.`
- `If both brackets were found, return a tuple of their start and stop positions; otherwise, return None.`

This function works well enough to pass the given tests, but it is not quite correct because it will return a region that contains unbalanced brackets:

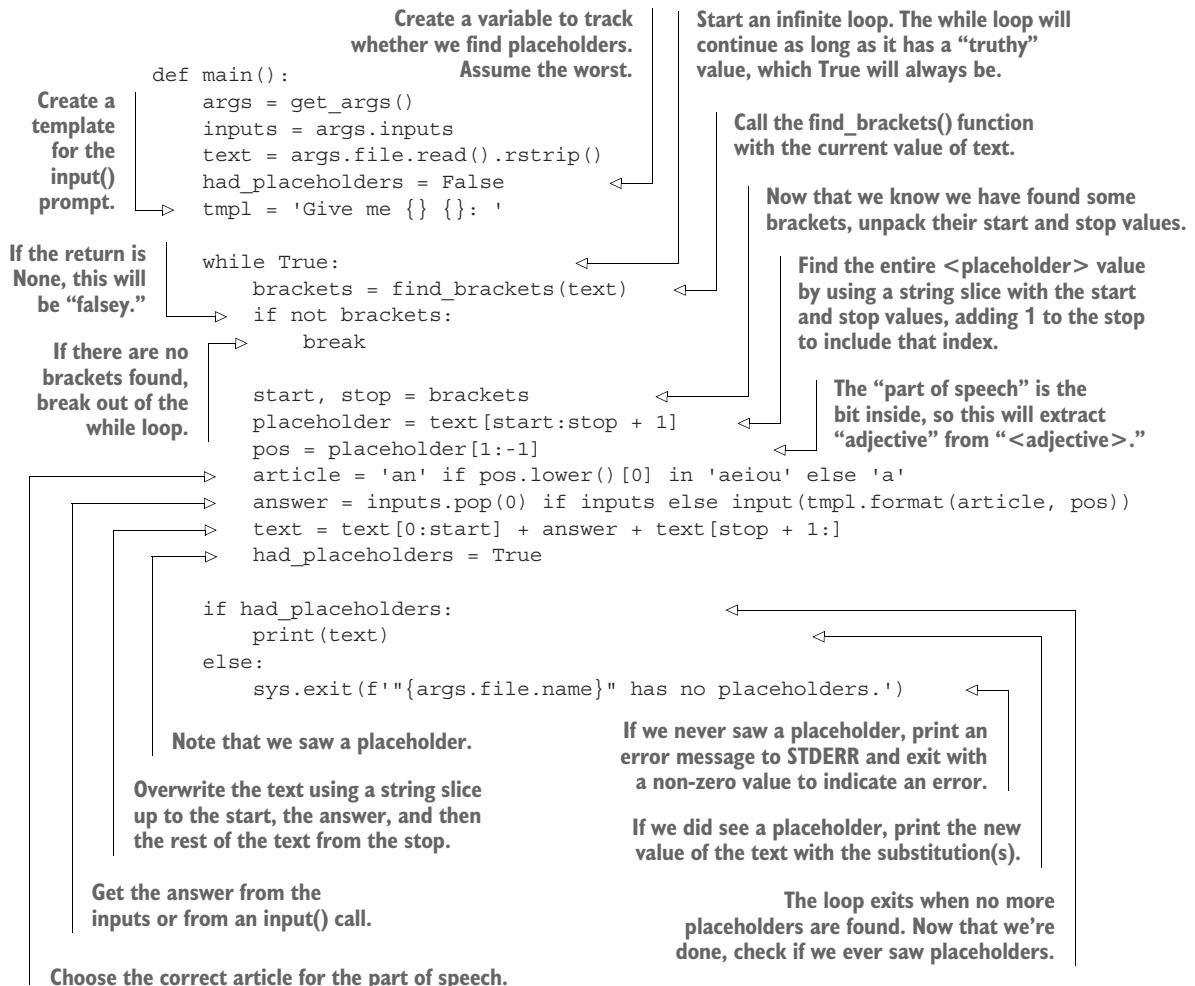
```
>>> text = 'foo <<bar> baz'
>>> find_brackets(text)
[4, 9]
>>> text[4:10]
'<<bar>'
```

That may seem unlikely, but I chose angle brackets to make you think of HTML tags like `<head>` and ``. HTML is notorious for being incorrect, maybe because it was hand generated by a human who messed up a tag or because some tool that generated the HTML had a bug. The point is that most web browsers have to be fairly relaxed in parsing HTML, and it would not be unexpected to see a malformed tag like `<<head>` instead of the correct `<head>`.

The regex version, on the other hand, specifically guards against matching unbalanced brackets by using the class `[^<>]` to define text that cannot contain any angle brackets. I could write a version of `find_brackets()` that finds only balanced brackets, but, honestly, it's just not worth it. This function points out that one of the strengths of the regex engine is that it can find a partial match (the first left bracket), see that it's unable to make a complete match, and start over (at the next left bracket). Writing this myself would be tedious and, frankly, not that interesting.

Still, this function works for all the given test inputs. Note that it only returns one set of brackets at a time. I will alter the text after I find each set of brackets, which will likely change the start and stop positions of any following brackets, so it's best to handle one set at a time.

Here is how I would incorporate it into the `main()` function:



17.4 Going further

- Extend your code to find all the HTML tags enclosed in `<...>` and `</...>` in a web page you download from the internet.
- Write a program that will look for unbalanced open/close pairs for parentheses `()`, square brackets `[]`, and curly brackets `{}`. Create input files that have balanced and unbalanced text, and write tests that verify your program identifies both.

Summary

- Regular expressions are almost like functions where we *describe* the patterns we want to find. The regex engine will do the work of trying to find the patterns, handling mismatches and starting over to find the pattern in the text.
- Regex patterns with * or + are “greedy” in that they match as many characters as possible. Adding a ? after them makes them “non-greedy” so that they match as *few* characters as possible.
- The `re.findall()` function will return a list of all the matching strings or capture groups for a given pattern.
- The `re.sub()` function will substitute a pattern in some text with new text.
- You can halt your program at any time using the `sys.exit()` function. If it’s given no arguments, the default exit value will be 0 to indicate no errors. If you wish to indicate there was an error, use any non-zero value such as 1. Or use a string value, which will be printed to `STDERR`, and a non-zero exit value will be used automatically.

Gematria: Numeric encoding of text using ASCII values

Gematria is a system for assigning a number to a word by summing the numeric values of each of the characters (<https://en.wikipedia.org/wiki/Gematria>). In the standard encoding (*Mispar hechrechi*), each character of the Hebrew alphabet is assigned a numeric value ranging from 1 to 400, but there are more than a dozen other methods for calculating the numeric value for the letters. To encode a word, these values are added together. Revelation 13:18 from the Christian Bible says, “Let the one who has insight calculate the number of the wild beast, for it is a man’s number, and its number is 666.” Some scholars believe that number is derived from the encoding of the characters representing Nero Caesar’s name and title and that it was used as a way of writing about the Roman emperor without naming him.

We will write a program called `gematria.py` that will numerically encode each word in a given text by similarly adding numeric values for the characters in each word. There are many ways we could assign these values. For instance, we could start by giving “a” the value 1, “b” the value 2, and so forth. Instead, we will use the ASCII table (<https://en.wikipedia.org/wiki/ASCII>) to derive a numeric value for English alphabet characters. For non-English characters, we could consider using a Unicode value, but this exercise will stick to ASCII letters.

The input text may be entered on the command line:

```
$ ./gematria.py 'foo bar baz'  
324 309 317
```



Or it could be in a file:

```
$ ./gematria.py ../inputs/fox.txt
289 541 552 333 559 444 321 448 314
```

Figure 18.1 shows a string diagram showing how the program should work.

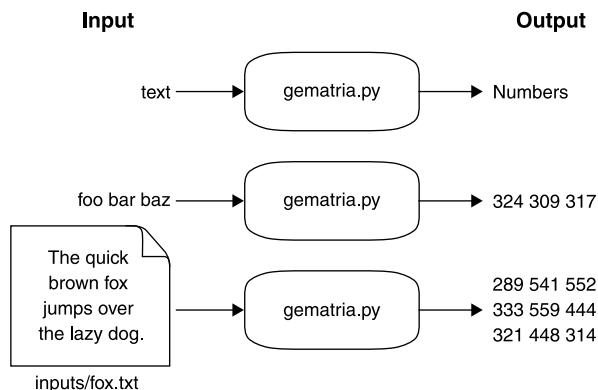


Figure 18.1 The gematria program will accept input text and will produce a numeric encoding for each word.

In this exercise, you will

- Learn about the `ord()` and `chr()` functions
- Explore how characters are organized in the ASCII table
- Understand character ranges used in regular expressions
- Use the `re.sub()` function
- Learn how `map()` can be written without `lambda`
- Use the `sum()` function and see how that relates to using `reduce()`
- Learn how to perform case-insensitive string sorting

18.1 Writing gematria.py

I will always recommend you start your programs in some way that avoids having to type all the boilerplate text. Either copy template/template.py to 18_gematria/gematria.py or use new.py gematria.py in the 18_gematria directory to create a starting point.

Modify the program until it prints the following usage statement if it's given no arguments or the `-h` or `--help` flag:

```
$ ./gematria.py -h
usage: gematria.py [-h] text

Gematria

positional arguments:
  text      Input text or file
```

```
optional arguments:
  -h, --help show this help message and exit
```

As in previous exercises, the input may come from the command line or from a file. I suggest you copy the code you used in chapter 5 to handle this, and then modify your `main()` function as follows:

```
def main():
    args = get_args()
    print(args.text)
```

Verify that your program will print text from the command line,

```
$ ./gematria.py 'Death smiles at us all, but all a man can do is smile back.'
Death smiles at us all, but all a man can do is smile back.
```

or from a file:

```
$ ./gematria.py ../inputs/spiders.txt
Don't worry, spiders,
I keep house
casually.
```

18.1.1 Cleaning a word

Let's discuss how a single word will be encoded, as it will affect how we will break the text in the next section. In order to be absolutely sure we are only dealing with ASCII values, let's remove anything that is not an upper- or lowercase English alphabet character or any of the Arabic numerals 0–9. We can define that class of characters using the regular expression `[A-Za-z0-9]`.

We can use the `re.findall()` function we used in chapter 17 to find all the characters in word that match this class. For instance, we should expect to find everything except the apostrophe in the word “Don’t” (see figure 18.2):

```
>>> re.findall('[A-Za-z0-9]', "Don't")
['D', 'o', 'n', 't']
```

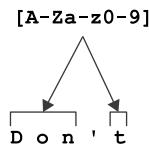


Figure 18.2 This character class only matches alphanumeric values.

If we put a caret (^) as the first character inside the class, like `[^A-Za-z0-9]`, we'll find anything that is *not* one of those characters. Now we would expect to match *only* the apostrophe (see figure 18.3):

```
>>> import re
>>> re.findall('[^A-Za-z0-9]', "Don't")
["'"]
```

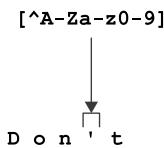


Figure 18.3 The caret will find the complement of the character class, so any non-alphanumeric character.

We can use the `re.sub()` function to replace any characters in that second class with the empty string. As you learned in chapter 17, this will replace *all* occurrences of the pattern unless we use the `count=n` option:

```
>>> word = re.sub('[^A-Za-z0-9]', '', "Don't")
>>> word
'Dont'
```

We will want to use this operation to clean each word that we'll encode, as shown in figure 18.4.

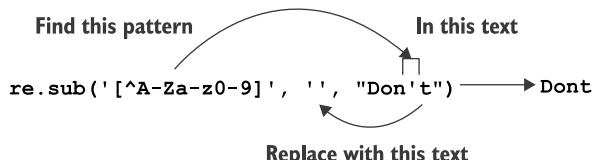


Figure 18.4 The `re.sub()` function will replace any text matching a pattern with another value.

18.1.2 Ordinal character values and ranges

We will encode a string like “Dont” by converting *each character* to a numeric value and then adding them together, so let’s first figure out how to encode a single character.

Python has a function called `ord()` that will convert a character to its “ordinal” value. For all alphanumeric values that we are using, this will be equal to the character’s position in the American Standard Code for Information Interchange (ASCII, pronounced like “as-kee”) table:

```
>>> ord('D')
68
>>> ord('o')
111
```

The `chr()` function works in reverse to convert a number to a character:

```
>>> chr(68)
'D'
>>> chr(111)
'o'
```

Following is the ASCII table. For simplicity's sake, I show "NA" ("not available") for the values up to index 31 as they are not printable.

\$./asciitbl.py							
0 NA	16 NA	32 SPACE	48 0	64 @	80 P	96 `	112 p
1 NA	17 NA	33 !	49 1	65 A	81 Q	97 a	113 q
2 NA	18 NA	34 "	50 2	66 B	82 R	98 b	114 r
3 NA	19 NA	35 #	51 3	67 C	83 S	99 c	115 s
4 NA	20 NA	36 \$	52 4	68 D	84 T	100 d	116 t
5 NA	21 NA	37 %	53 5	69 E	85 U	101 e	117 u
6 NA	22 NA	38 &	54 6	70 F	86 V	102 f	118 v
7 NA	23 NA	39 '	55 7	71 G	87 W	103 g	119 w
8 NA	24 NA	40 (56 8	72 H	88 X	104 h	120 x
9 NA	25 NA	41)	57 9	73 I	89 Y	105 i	121 y
10 NA	26 NA	42 *	58 :	74 J	90 Z	106 j	122 z
11 NA	27 NA	43 +	59 ;	75 K	91 [107 k	123 {
12 NA	28 NA	44 ,	60 <	76 L	92 \	108 l	124
13 NA	29 NA	45 -	61 =	77 M	93]	109 m	125 }
14 NA	30 NA	46 .	62 >	78 N	94 ^	110 n	126 ~
15 NA	31 NA	47 /	63 ?	79 O	95 _	111 o	127 DEL

NOTE I have included the `asciitbl.py` program in the `18_gematria` directory of the source code repository.

We can use a `for` loop to cycle through all the characters in a string:

```
>>> word = "Dont"
>>> for char in word:
...     print(char, ord(char))
...
D 68
o 111
n 110
t 116
```

Note that upper- and lowercase letters have different `ord()` values. This makes sense because they are two different letters:

```
>>> ord('D')
68
>>> ord('d')
100
```

We can iterate over the values from "a" to "z" by finding their `ord()` values:

```
>>> [chr(n) for n in range(ord('a'), ord('z') + 1)]
['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm',
'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z']
```

As you can see in the previous ASCII table, the letters "a" through "z" lie contiguously. The same is true for "A" to "Z" and "0" to "9," which is why we can use `[A-Za-z0-9]` as a regex.

Note that the uppercase letters have *lower* ordinal values than their lowercase versions, which is why you cannot use the range [a-Z]. Try this in the REPL and note the error you get:

```
>>> re.findall(' [a-Z] ', word)
```

If I execute the preceding function in the REPL, the last line of the error I see is this:

```
re.error: bad character range a-Z at position 1
```

You *can*, however, use the range [A-z]:

```
>>> re.findall(' [A-z] ', word)
['D', 'o', 'n', 't']
```

But note that “Z” and “a” are not contiguous:

```
>>> ord('Z'), ord('a')
(90, 97)
```

There are other characters in between them:

```
>>> [chr(n) for n in range(ord('Z') + 1, ord('a'))]
['[', '\\', ']', '^', '_', '`']
```

If we try to use that range on all the printable characters, you’ll see that it matches characters that are not letters:

```
>>> import string
>>> re.findall(' [A-z] ', string.printable)
['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm',
 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z',
 'A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M',
 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z',
 '[', '\\', ']', '^', '_', '`']
```

That is why it is safest to specify the characters we want as the three separate ranges, [A-Za-z0-9], which you may sometimes hear pronounced as “A to Z, a to z, zero to nine,” as it assumes you understand that there are two “a to z” ranges that are distinct according to their case.

18.1.3 Summing and reducing

Let’s keep reminding ourselves what the goal is here: convert all the characters in a word, and then sum those values. There is a handy Python function called `sum()` that will add a list of numbers:

```
>>> sum([1, 2, 3])
6
```



We can manually encode the string “Dont” by calling `ord()` on each letter and passing the results as a list to `sum()`:

```
>>> sum([ord('D'), ord('o'), ord('n'), ord('t'))])
405
```

The question is how to apply the function `ord()` to all the characters in a `str` and pass a list to `sum()`. You’ve seen this pattern many times now. What’s the first tool you’ll reach for? We can always start with our handy `for` loop:

```
>>> word = 'Dont'
>>> vals = []
>>> for char in word:
...     vals.append(ord(char))
...
>>> vals
[68, 111, 110, 116]
```

Can you see how to make that into a single line using a list comprehension?

```
>>> vals = [ord(char) for char in word]
>>> vals
[68, 111, 110, 116]
```

From there, we can move to a `map()`:

```
>>> vals = map(lambda char: ord(char), word)
>>> list(vals)
[68, 111, 110, 116]
```

Here I’d like to show that the `map()` version doesn’t need the `lambda` declaration because the `ord()` function expects a single value, which is exactly what it will get from `map()`. Here is a nicer way to write it:

```
>>> vals = map(ord, word)
>>> list(vals)
[68, 111, 110, 116]
```

To my eye, that is a really beautiful piece of code!

Now we can `sum()` that to get a final value for our word:

```
>>> sum(map(ord, word))
405
```

That is correct:

```
>>> sum([68, 111, 110, 116])
405
```

18.1.4 Using `functools.reduce`

If Python has a `sum()` function, you might suspect it also has a `product()` function to multiply a list of numbers together. Alas, this is not a built-in function, but it does represent a common idea of *reducing* a list of values into a single value.

The `reduce()` function from the `functools` module provides a generic way to reduce a list. Let's consult the documentation for how to use it:

```
>>> from functools import reduce
>>> help(reduce)
reduce(...)

    reduce(function, sequence[, initial]) -> value
```

Apply a function of two arguments cumulatively to the items of a sequence, from left to right, so as to reduce the sequence to a single value.
 For example, `reduce(lambda x, y: x+y, [1, 2, 3, 4, 5])` calculates $((((1+2)+3)+4)+5)$. If `initial` is present, it is placed before the items of the sequence in the calculation, and serves as a default when the sequence is empty.

This is another higher-order function that wants *another function* as the first argument, just like `map()` and `filter()`. The documentation shows us how to write our own `sum()` function:

```
>>> reduce(lambda x, y: x + y, [1, 2, 3, 4, 5])
15
```

If we change the `+` operator to `*`, we have a product:

```
>>> reduce(lambda x, y: x * y, [1, 2, 3, 4, 5])
120
```

Here is how you might write a function for this:

```
def product(vals):
    return reduce(lambda x, y: x * y, vals)
```

And now you can call it:

```
>>> product(range(1,6))
120
```

Instead of writing our own `lambda`, we can use any function that expects two arguments. The `operator.mul` function fits this bill:

```
>>> import operator
>>> help(operator.mul)
mul(a, b, /)
    Same as a * b.
```

So it would be easier to write this:

```
def product(vals):
    return reduce(operator.mul, vals)
```

Fortunately, the `math` module also contains a `prod()` function you can use:

```
>>> import math
>>> math.prod(range(1, 6))
120
```

If you think about it, the `str.join()` method also reduces a list of strings to a single `str` value. Here's how we can write our own:

```
def join(sep, vals):
    return reduce(lambda x, y: x + sep + y, vals)
```

I much prefer the syntax of calling this `join` over the `str.join()` function:

```
>>> join(', ', ['Hey', 'Nonny', 'Nonny'])
'Hey, Nonny, Nonny'
```

Whenever you have a list of values that you want to combine to produce a single value, consider using the `reduce()` function.

18.1.5 Encoding the words

That was a lot of work just to get to summing the ordinal values of the characters, but wasn't it fascinating to explore? Let's get back on track, though.

We can create a function to encapsulate the idea of converting a word into a numeric value derived from summing the ordinal values of the characters. I call mine `word2num()`, and here is my test:

```
def test_word2num():
    """Test word2num"""
    assert word2num("a") == "97"
    assert word2num("abc") == "294"
    assert word2num("ab'c") == "294"
    assert word2num("4a-b'c,") == "346"
```

Notice that my function returns a `str` value, not an `int`. This is because I want to use the result with the `str.join()` function that only accepts `str` values—so '`405`' instead of `405`:

```
>>> from gematria import word2num
>>> word2num("Don't")
'405'
```

To summarize, the `word2num()` function accepts a word, removes unwanted characters, converts the remaining characters to `ord()` values, and returns a `str` representation of the `sum()` of those values.

18.1.6 Breaking the text

The tests expect you to maintain the same line breaks as the original text, so I recommend you use `str.splitlines()` as in other exercises. In chapters 15 and 16, we used different regexes to split each line into “words,” a process sometimes called “tokenization” in programs that deal with natural language processing (NLP). If you write a `word2num()` function that passes the tests I’ve provided, then you can use `str.split()` to break a line on spaces because the function will ignore anything that is not a character or number. You are, of course, welcome to break the line into words using whatever means you like.

The following code will maintain the line breaks and reconstruct the text. Can you modify it to add the `word2num()` function so that it instead prints out encoded words as shown in figure 18.5?

```
def main():
    args = get_args()
    for line in args.text.splitlines():
        for word in line.split():
            # what goes here?
            print(' '.join(line.split()))
```

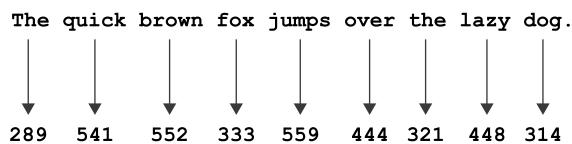
The quick brown fox jumps over the lazy dog.

 ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
 289 541 552 333 559 444 321 448 314

Figure 18.5 Each word of the text will be cleaned and encoded into a number.

The output will be one number for each word:

```
$ ./gematria.py ../inputs/fox.txt
289 541 552 333 559 444 321 448 314
```

Time to finish writing the solution. Be sure to use the tests! See you on the flip side.

18.2 Solution

I do enjoy the ideas of cryptography and encoding messages, and this program is (sort of) encrypting the input text, albeit in a way that cannot be reversed. Still, it’s fun to think of other ways you might process some text and transmogrify it to some other value.

```
#!/usr/bin/env python3
"""Gematria"""
```

```
import argparse
import os
import re
```

```

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Gematria',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('text', metavar='text', help='Input text or file')

    args = parser.parse_args()

    if os.path.isfile(args.text):
        args.text = open(args.text).read().rstrip()

    return args

```

Get the parsed command-line arguments.

The text argument is a string that might be a filename.

Check if the text argument is an existing file.

Overwrite the args.text with the contents of the file.

Return the arguments.


```

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    for line in args.text.splitlines():
        print(' '.join(map(word2num, line.split())))

```

Get the parsed arguments.

Split args.text on newlines to retain line breaks.

Split the line on spaces, map the result through word2num(), and then join that result on spaces.


```

# -----
def word2num(word):
    """Sum the ordinal values of all the characters"""

    return str(sum(map(ord, re.sub('[^A-Za-z0-9]', '', word))))

```

Define a function to convert a word to a number.

Use re.sub() to remove anything that's not an alphanumeric character. Map the resulting string through the ord() function, sum the ordinal values of the characters, and return a str representation of the sum.


```

# -----
if __name__ == '__main__':
    main()

```

Define a function to test the word2num() function.

18.3 Discussion

I trust you understand `get_args()`, as we've used this exact code several times now. Let's jump to the `word2num()` function.

18.3.1 Writing word2num()

I could have written the function like this:

```
def word2num(word):
    vals = []
    for char in re.sub('[^A-Za-z0-9]', '', word):
        vals.append(ord(char))
    return str(sum(vals))

    Initialize an empty list to
    hold the ordinal values.

    Iterate all the
    characters returned
    from re.sub().

    Convert the character to
    an ordinal value and
    append that to the
    values.

    Sum the values and return a
    string representation.
```

That's four lines of code instead of the one I wrote. I would at least rather use a list comprehension, which collapses three lines of code into one:

```
def word2num(word):
    vals = [ord(char) for char in re.sub('[^A-Za-z0-9]', '', word)]
    return str(sum(vals))
```

That could be written in one line, though it could be argued that readability suffers:

```
def word2num(word):
    return str(sum([ord(char) for char in re.sub('[^A-Za-z0-9]', '', word)]))
```

I still think the `map()` version is the most readable and concise:

```
def word2num(word):
    return str(sum(map(ord, re.sub('[^A-Za-z0-9]', '', word))))
```

Figure 18.6 shows how the three methods relate to each other.

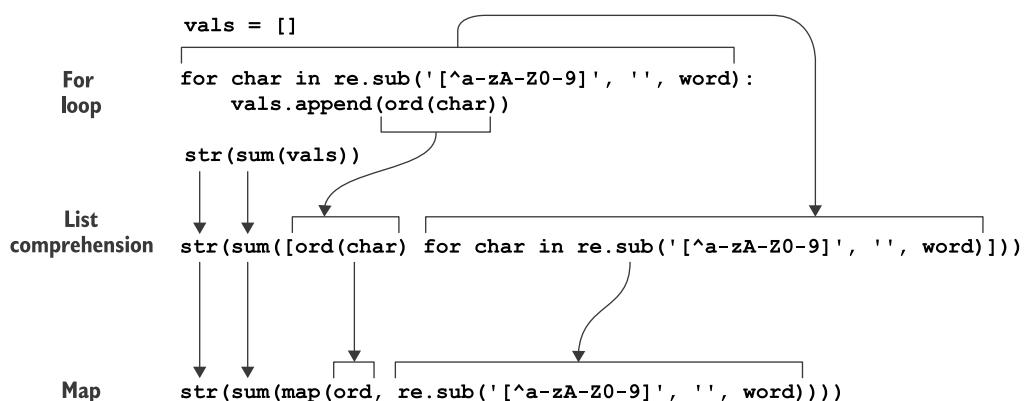


Figure 18.6 How the `for` loop, a list comprehension, and a `map()` relate to each other

Figure 18.7 will help you see how the data moves through the `map()` version with the string “Don’t.”

- 1 The `re.sub()` function will replace any character not in the character class with the empty string. This will turn a word like “Don’t” into “Dont” (without the apostrophe).
- 2 The `map()` will apply the given function `ord()` to each element of a sequence. Here that “sequence” is a `str`, so it will use each character of the word.
- 3 The result of `map()` is a new `list`, where each character from “Dont” is given to the `ord()` function.
- 4 The results of the calls to `ord()` will be a `list` of `int` values, one for each letter.
- 5 The `sum()` function will reduce a list of numbers to a single value by adding them together.
- 6 The final value from our function needs to be a `str`, so we use the `str()` function to turn the return from `sum()` into a string representation of the number.

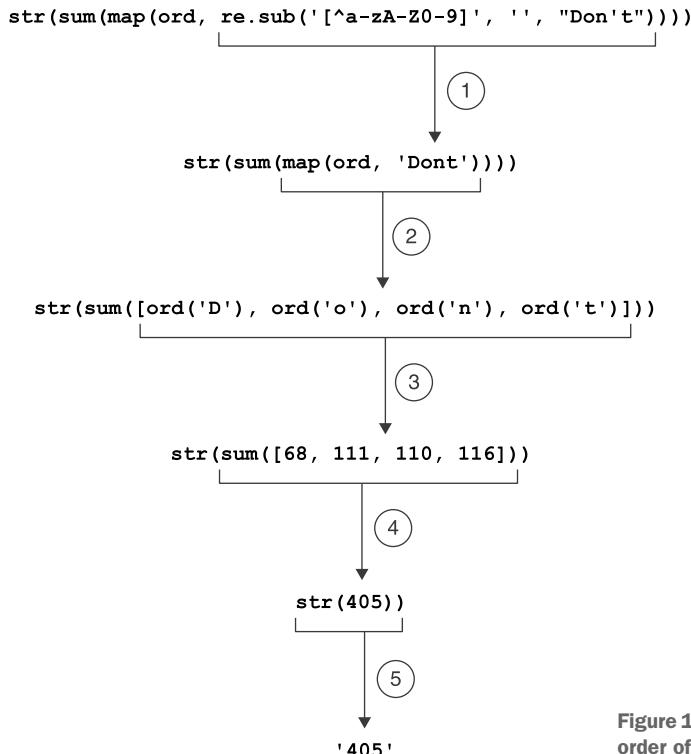


Figure 18.7 A representation of the order of operations for the functions

18.3.2 Sorting

The point of this exercise was less about the `ord()` and `chr()` functions and more about exploring regular expressions, function application, and how characters are represented inside programming languages like Python.

For instance, the sorting of strings is case sensitive because of the relative order of the `ord()` values of the characters (because the uppercase letters are defined earlier in the ASCII table than the lowercase values). Note that the words that begin with uppercase letters are sorted before those with lowercase letters:

```
>>> words = 'banana Apple Cherry anchovies cabbage Beets'
>>> sorted(words)
['Apple', 'Beets', 'Cherry', 'anchovies', 'banana', 'cabbage']
```

This is because all the uppercase ordinal values are lower than those of the lowercase letters. In order to perform a case sensitive sorting of strings, you can use `key=str.casefold`. The `str.casefold()` function will return “a version of the string suitable for caseless comparisons.” We are using the function’s name *without parentheses* here because we are passing *the function itself* as the argument for `key`:

```
>>> sorted(words, key=str.casefold)
['anchovies', 'Apple', 'banana', 'Beets', 'cabbage', 'Cherry']
```

If you add the parentheses, it will cause an exception. This is exactly the same way we pass functions as arguments to `map()` and `filter()`:

```
>>> sorted(words, key=str.casefold())
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: descriptor 'casefold' of 'str' object needs an argument
```

The option is the same with `list.sort()` if you prefer to sort the list in place:

```
>>> words.sort(key=str.casefold)
>>> words
['anchovies', 'Apple', 'banana', 'Beets', 'cabbage', 'Cherry']
```

Command-line tools like the `sort` program behave in the same way due to the same representation of characters. Given a file of these same words,

```
$ cat words.txt
banana
Apple
Cherry
anchovies
cabbage
Beets
```

the sort program on my Mac¹ will first sort the uppercase words and then the lowercase:

```
$ sort words
Apple
Beets
Cherry
anchovies
banana
cabbage
```

I have to read the sort manual page (via `man sort`) to find the `-f` flag to perform a case-insensitive sort:

```
$ sort -f words
anchovies
Apple
banana
Beets
cabbage
Cherry
```

18.3.3 Testing

I would like to take a moment to point out how often I use my own tests. Every time I write an alternative version of a function or program, I run my own tests to verify that I'm not accidentally showing you buggy code. Having a test suite gives me the freedom and confidence to extensively refactor my programs because I know I can check my work. If I ever find a bug in my code, I add a test to verify that the bug exists. Then I fix the bug and verify that it's handled. I know if I accidentally reintroduce that bug, my tests will catch it.

For the purposes of this book, I've tried to never write a program over 100 lines. It's common for programs to grow to thousands of lines of code spread over dozens of modules. I recommend you start writing and using tests, no matter how small you start. It's a good habit to establish early on, and it will only help you as you write longer code.

18.4 Going further

- Analyze text files to find other words that sum to the value 666. Are these particularly scary words?
- Given some text input, find the most frequently occurring value from `word2num()` and all the words that reduce to that value.
- Create a version using your own numeric values for each character. For instance, each letter could be encoded as its position in the alphabet so that “A”

¹ The GNU coreutils 8.30 version on one of my Linux machines will perform a case-insensitive sort by default. How does your `sort` work?

and “a” are 1, “B” and “b” are 2, and so on. Or you might decide to weigh each consonant as 1 and each vowel as –1. Create your own scheme, and write tests to ensure your program performs as you expect.

Summary

- The `ord()` function will return the Unicode code point of a character. For our alphanumeric values, the ordinal values correspond to their position in the ASCII table.
- The `chr()` function will return the character for a given ordinal value.
- You can use character ranges like `a-z` in regular expressions when ordinal values of the characters lie contiguously, such as in the ASCII table.
- The `re.sub()` function will replace matching patterns of text in a string with new values, such as replacing all non-characters with the empty string to remove punctuation and whitespace.
- A `map()` can be written using a function reference instead of a `lambda` if the function expects a single positional argument.
- The `sum()` function reduces a list of numbers using addition. You can manually write a version of this using the `functools.reduce()` function.
- To perform a case-insensitive sort of string values, use the `key=str.casefold` option with both the `sorted()` and `list.sort()` functions.

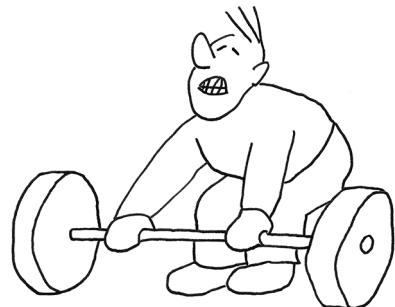
10

Workout of the Day: Parsing CSV files, creating text table output

Several years ago, I joined a workout group. We meet several times a week in our coach's unpaved driveway. We pick up and drop heavy things and run around trying to keep Death at bay for another day. I'm no paragon of strength and fitness, but it's been a nice way to exercise and visit with friends. One of my favorite parts of going is that our coach will write a "Workout of the Day" or "WOD" on the board. Whatever it says is what I do. It doesn't matter if I actually want to do 200 push-ups that day, I just get them done no matter how long it takes.¹

In that spirit, we'll write a program called `wod.py` to help us create a random daily workout that we have to do, no questions asked:

```
$ ./wod.py
Exercise      Reps
-----
Pushups       40
```



¹ See "More Isn't Always Better" by Barry Schwartz (<https://hbr.org/2006/06/more-isnt-always-better>). He notes that increasing the number of choices given to people actually creates more distress and feelings of dissatisfaction, whatever choice is made. Imagine an ice cream shop with three flavors: chocolate, vanilla, and strawberry. If you choose chocolate, you'll likely be happy with that choice. Now imagine that the shop has 60 flavors of ice cream, including 20 different fruit creams and sorbets and 12 different chocolate varieties from Rocky Road to Fudgetastic Caramel Tiramisu Ripple. Now when you choose a "chocolate" variety, you may leave with remorse about the 11 other kinds you could have chosen. Sometimes having no choice at all provides a sense of calm. Call it fatalism or whatnot.

Plank	38
Situps	99
Hand-stand pushups	5

NOTE Each time you run the program, you are required to perform all the exercises *immediately*. Heck, even just *reading* them means you have to do them. Like NOW. Sorry, I don't make the rules. Better get going on those sit-ups!

We'll choose from a list of exercises stored in a *delimited text file*. In this case, the “delimiter” is the comma, and it will separate each field value. Data files that use commas as delimiters are often described as *comma-separated values* or CSV files. Usually the first line of the file names the columns, and each subsequent line represents a row in the table:

```
$ head -3 inputs/exercises.csv
exercise,reps
Burpees,20-50
Situps,40-100
```

In this exercise, you will

- Parse delimited text files using the `csv` module
- Coerce text values to numbers
- Print tabular data using the `tabulate` module
- Handle missing and malformed data

This chapter and the next are meant to be a step up in how challenging they are. You will be applying many of the skills you've learned in previous chapters, so get ready!

19.1 Writing wod.py

You will be creating a program called `wod.py` in the `19_wod` directory. Let's start by taking a look at the usage that should print when it's run with `-h` or `--help`. Modify your program's parameters until it produces this:

```
$ ./wod.py -h
usage: wod.py [-h] [-f FILE] [-s seed] [-n exercises] [-e]

Create Workout Of (the) Day (WOD)

optional arguments:
  -h, --help            show this help message and exit
  -f FILE, --file FILE CSV input file of exercises (default:
                        inputs/exercises.csv)
  -s seed, --seed seed  Random seed (default: None)
  -n exercises, --num exercises
                        Number of exercises (default: 4)
  -e, --easy             Halve the reps (default: False)
```

Our program will read an input `-f` or `--file`, which should be a readable text file (default, `inputs/exercises.csv`). The output will be some `-n` or `--num` number of exercises

(default, 4). There might be an `-e` or `--easy` flag to indicate that the repetitions of each exercise should be cut in half. Since we'll be using the `random` module to choose the exercises, we'll need to accept an `-s` or `--seed` option (`int` with a default of `None`) to pass to `random.seed()` for testing purposes.

19.1.1 Reading delimited text files

We're going to use the `csv` module to parse the input file. This is a standard module that should already be installed on your system. You can verify that by opening a `python3` REPL and trying to import it. If this works, you're all set:

```
>>> import csv
```

We'll also look at two other modules that you probably will need to install:

- Tools from the `csvkit` module to look at the input file on the command line
- The `tabulate` module to format the output table

Run this command to install these modules:

```
$ python3 -m pip install csvkit tabulate
```

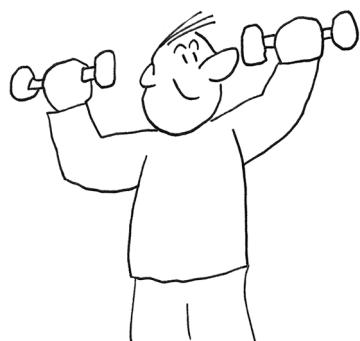
There is also a `requirements.txt` file, which is a common way to document the dependencies for a program. Instead of the previous command, you can install all the modules with this one:

```
$ python3 -m pip install -r requirements.txt
```

Despite having “`csv`” in the name, the `csvkit` module can handle just about any delimited text file. For instance, it's typical to use the tab (`\t`) character as a delimiter, too. The module includes many tools that you can read about in its documentation (<https://csvkit.readthedocs.io/en/1.0.3/>). I've included several delimited files in the `19_wod/inputs` directory that you can use to test your program.

After installing `csvkit`, you should be able to use `csvlook` to parse the `inputs/exercises.csv` file into a table structure showing the columns:

```
$ csvlook --max-rows 3 inputs/exercises.csv
| exercise | reps   |
| ----- | ----- |
| Burpees | 20-50 |
| Situps  | 40-100|
| Pushups | 25-75 |
| ...     | ...    |
```



The “`reps`” column of the input file will have two numbers separated by a dash, like `10-20` meaning “from 10 to 20 reps.” To select the final value for the `reps`, you will use

the `random.randint()` function to select an integer value between the low and high values. When run with a seed, your output should exactly match this:

```
$ ./wod.py --seed 1 --num 3
Exercise      Reps
-----
Pushups        32
Situps         71
Crunches       27
```

When run with the `--easy` flag, the reps should be halved:

```
$ ./wod.py --seed 1 --num 3 --easy
Exercise      Reps
-----
Pushups        16
Situps         35
Crunches       13
```

The `--file` option should default to the `inputs/exercises.csv` file, or we can indicate a different input file:

```
$ ./wod.py --file inputs/silly-exercises.csv
Exercise      Reps
-----
Hanging Chads  46
Squatting Chinups 46
Rock Squats    38
Red Barchettas 32
```

Figure 19.1 shows our trusty string diagram to help you think about it.

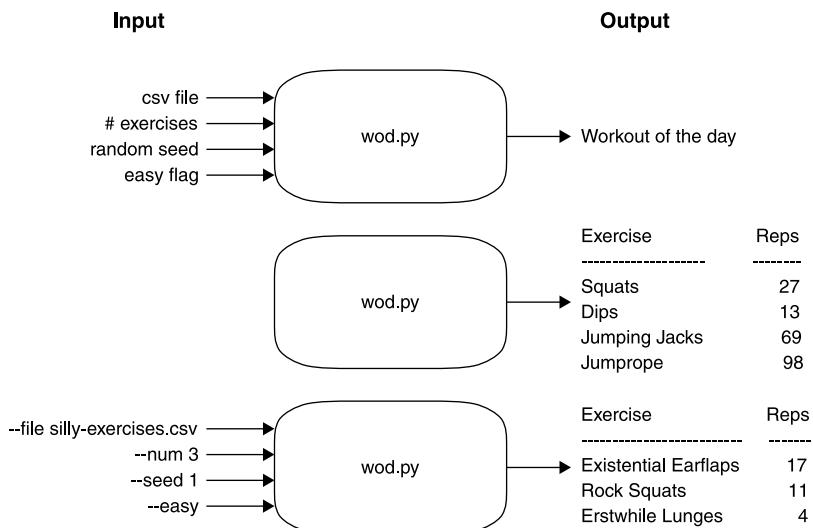


Figure 19.1 The WOD program will randomly select exercises and reps from a CSV file to create a table listing the workout of the day.

19.1.2 Manually reading a CSV file

First I'm going to show you how to manually parse each record from a CSV file into a list of dictionaries, and then I'll show you how to use the `csv` module to do this more quickly. The reason we want to make a dictionary from each record is so that we can get at the values for each exercise and the number of reps (repetitions, or how many times to repeat a given exercise). We're going to need to split the reps into low and high values so that we can get a range of numbers from which we'll randomly select the number of reps. Finally, we'll randomly select some exercises along with their reps to make a workout. Whew, just describing that was a workout!

Notice that reps is given as a range from a low number to a high number, separated by a dash:

```
$ head -3 inputs/exercises.csv
exercise,reps
Burpees,20-50
Situps,40-100
```

It would be convenient to read this as a list of dictionaries where the column names in the first line are combined with each line of data, like this:

```
$ ./manual1.py
[{'exercise': 'Burpees', 'reps': '20-50'},
 {'exercise': 'Situps', 'reps': '40-100'},
 {'exercise': 'Pushups', 'reps': '25-75'},
 {'exercise': 'Squats', 'reps': '20-50'},
 {'exercise': 'Pullups', 'reps': '10-30'},
 {'exercise': 'Hand-stand pushups', 'reps': '5-20'},
 {'exercise': 'Lunges', 'reps': '20-40'},
 {'exercise': 'Plank', 'reps': '30-60'},
 {'exercise': 'Crunches', 'reps': '20-30'}]
```

It may seem like overkill to use a dictionary for records that contain just two columns, but I regularly deal with records that contain dozens to *hundreds* of columns, and then field names are essential. A dictionary is really the only sane way to handle most delimited text files, so it's good to learn with a small example like this.

Let's look at the `manual1.py` code that will do this:

```
Use a for loop to read the
rest of the lines of fh.

Initialize records
as an empty list.

#!/usr/bin/env python3

from pprint import pprint

with open('inputs/exercises.csv') as fh:
    headers = fh.readline().rstrip().split(',')
    records = []
    for line in fh:
        record = {}
        for index, value in enumerate(line.split(',')):
            if index == 0:
                record['exercise'] = value
            else:
                record['reps'] = value
        records.append(record)

pprint(records)
```

Append the resulting dictionary to the records.

```

rec = dict(zip(headers, line.rstrip().split(',')))
records.append(rec)

pprint(records)
    ↪
Pretty-print the records.
    ↪

```

Strip and split the line of text into a list of field values. Use the `zip()` function to create a new list of tuples containing each of the headers paired with each of the values. Use the `dict()` function to turn this list of tuples into a dictionary.

Let's break this down a bit more. First we'll open() the file and read the first line:

```

>>> fh = open('exercises.csv')
>>> fh.readline()
'exercise,reps\n'

```

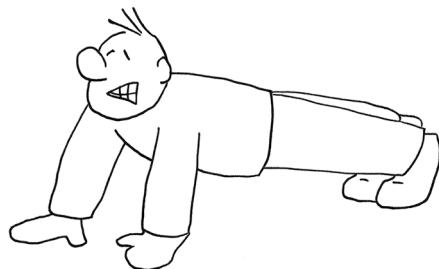
The line still has a newline stuck to it, so we can use the `str.rstrip()` function to remove that:

```

>>> fh = open('exercises.csv')
>>> fh.readline().rstrip()
'exercise,reps'

```

NOTE Note that I need to keep reopening this file for this demonstration, or each subsequent call to `fh.readline()` would read the next line of text.



Now let's use `str.split()` to split that line on the comma to get a list of strings:

```

>>> fh = open('exercises.csv')
>>> headers = fh.readline().rstrip().split(',')
>>> headers
['exercise', 'reps']

```

We can likewise read the next line of the file to get a list of the field values:

```

>>> line = fh.readline().rstrip().split(',')
>>> line
['Burpees', '20-50']

```

Next we use the `zip()` function to merge the two lists into one list where the elements of each list have been mated with their counterparts in the same positions. That might seem complicated, but think about the end of a wedding ceremony when the bride and groom turn around to face the assembled crowd. Usually they will hold hands and start walking down the aisle to leave the ceremony. Imagine three grooms-men ('G') and three bridesmaids ('B') left standing on their respective sides facing each other:

```

>>> groomsmen = 'G' * 3
>>> bridesmaids = 'B' * 3

```

If there are two lines each containing three people, then we end up with a single line containing three pairs:

```
>>> pairs = list(zip(groomsmen, bridesmaids))
>>> pairs
[('G', 'B'), ('G', 'B'), ('G', 'B')]
>>> len(pairs)
3
```

Or think of two lines of cars merging to exit a parking lot. It's customary for one car from one lane (say, "A") to merge into traffic, then a car from the other lane (say, "B"). The cars are combining like the teeth of a zipper, and the result is "A," "B," "A," "B," and so forth.

The `zip()` function will group the elements of the lists into tuples, grouping all the elements in the first position together, then the second position, and so on, as shown in figure 19.2. Note that this is another *lazy* function, so I will use `list` to coerce this in the REPL:

```
>>> list(zip('abc', '123'))
[('a', '1'), ('b', '2'), ('c', '3')]
```

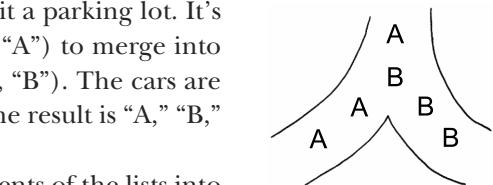
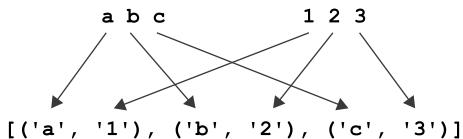


Figure 19.2 Zipping two lists creates a new list with pairs of elements.

The `zip()` function can handle more than two lists. Note that it will only create groupings for the shortest list. In the following example, the first two lists have four elements ("abcd" and "1234"), but the last has only three ("xyz"), so only three tuples are created:

```
>>> list(zip('abcd', '1234', 'xyz'))
[('a', '1', 'x'), ('b', '2', 'y'), ('c', '3', 'z')]
```

In our data, `zip()` will combine the header "exercise" with the value "Burpees" and then the header "reps" with the value "20–50" (see figure 19.3):

```
>>> list(zip(headers, line))
[('exercise', 'Burpees'), ('reps', '20-50')]
```

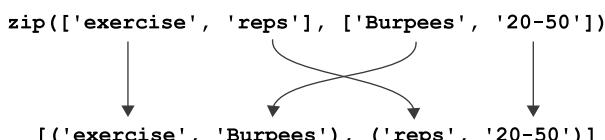


Figure 19.3 Zipping the headers and values together to create a list of tuples

That created a list of tuple values. Instead of `list()`, we can use `dict()` to create a dictionary:

```
>>> rec = dict(zip(headers, line))
>>> rec
{'exercise': 'Burpees', 'reps': '20-50'}
```

Recall that the `dict.items()` function will turn a dict into a list of tuple (key/value) pairs, so you can think of these two data structures as being fairly interchangeable:

```
>>> rec.items()
dict_items([('exercise', 'Burpees'), ('reps', '20-50')])
```

We can drastically shorten our code by replacing the `for` loop with a list comprehension:

```
with open('inputs/exercises.csv') as fh:
    headers = fh.readline().rstrip().split(',')
    records = [dict(zip(headers, line.rstrip().split(','))) for line in fh]
    pprint(records)
```

We still need to break out
the headers separately by
reading the first line.

This combines the three lines of the for loop
into a single list comprehension.

We can use `map()` to write equivalent code:

```
with open('inputs/exercises.csv') as fh:
    headers = fh.readline().rstrip().split(',')
    mk_rec = lambda line: dict(zip(headers, line.rstrip().split(',')))
    records = map(mk_rec, fh)
    pprint(list(records))
```

Flake8 will complain about assigning this lambda
expression. I generally write my code so as to produce no
warnings, but I do tend to disagree with this suggestion. I quite
like writing one-line functions using a lambda assignment.

In the next section, I'm going to show you how to use the `csv` module to handle much of this code, which may lead you to wonder why I bothered showing you how to handle this yourself. Unfortunately, I often have to handle data that is terribly formatted, such that the first line is not the header, or there are other rows of information between the header row and the actual data. When you've seen as many badly formatted Excel files as I have, you'll come to appreciate that you sometimes have no choice but to parse the file yourself.

19.1.3 **Parsing with the csv module**

Parsing delimited text files in this way is extremely common, and it would not make sense to write or copy this code every time you needed to parse a file. Luckily, the `csv` module is a standard module installed with Python, and it can handle all of this very gracefully.

Let's look at how our code can change if we use `csv.DictReader()` (see `using_csv1.py` in the repo):

```
#!/usr/bin/env python3
import csv
from pprint import pprint

with open('inputs/exercises.csv') as fh:
    reader = csv.DictReader(fh, delimiter=',')
    records = []
    for rec in reader:
        records.append(rec)

    pprint(records)
```

Import the csv module.

Create a `csv.DictReader()` that will create a dictionary for each record in the file. It zips the headers in the first line with the data values in the subsequent lines. It uses the delimiter to indicate the string value for splitting the columns of text.

Initialize an empty list to hold the records.

Use a for loop to iterate through each record returned by the reader.

The records will be a dictionary that is appended to the list of records.

The following code creates the same list of dict values as before, but with far less code. Note that each record is shown as an `OrderedDict`, which is a type of dictionary where the keys are maintained in their insertion order:

```
$ ./using_csv1.py
[OrderedDict([('exercise', 'Burpees'), ('reps', '20-50')]),
 OrderedDict([('exercise', 'Situps'), ('reps', '40-100')]),
 OrderedDict([('exercise', 'Pushups'), ('reps', '25-75')]),
 OrderedDict([('exercise', 'Squats'), ('reps', '20-50')]),
 OrderedDict([('exercise', 'Pullups'), ('reps', '10-30')]),
 OrderedDict([('exercise', 'Hand-stand pushups'), ('reps', '5-20')]),
 OrderedDict([('exercise', 'Lunges'), ('reps', '20-40')]),
 OrderedDict([('exercise', 'Plank'), ('reps', '30-60')]),
 OrderedDict([('exercise', 'Crunches'), ('reps', '20-30')])]
```

We can remove the entire for loop and use the `list()` function to coerce the reader to give us that same list. This code (in `using_csv2.py`) will print the same output:

```
with open('inputs/exercises.csv') as fh:
    reader = csv.DictReader(fh, delimiter=',')
    records = list(reader)
    pprint(records)
```

Open the file.

Create a `csv.DictReader()` to read `fh`, using the comma for the delimiter.

Use the `list()` function to coerce all the values from the reader.

Pretty-print the records.

19.1.4 Creating a function to read a CSV file

Let's try to imagine how we could write and test a function we might call `read_csv()` to read in our data. Let's start with a placeholder for our function and the `test_read_csv()` definition:

```
def read_csv(fh):
    """Read the CSV input"""
    pass

def test_read_csv():
    """Test read_csv"""
    text = io.StringIO('exercise,reps\nBurpees,20-50\nSitups,40-100') ←
    assert read_csv(text) == [('Burpees', 20, 50), ('Situps', 40, 100)] ←
```

Use `io.StringIO()` to create a mock file handle to wrap around a valid text that we might read from a file. The `\n` represents the newlines that break each line in the input data, and each line uses commas to separate the fields. We previously used `io.StringIO()` in the low-memory version of chapter 5's program.

Affirm that our imaginary `read_csv()` file would turn this text into a list of tuple values with the name of the exercise and the reps, which have been split into low and high values. Note that these values have been converted to integers.

Hey, we just did all that work to make a list of dict values, so why am I suggesting that we now create a list of tuple values? I'm looking ahead here to how we might use the `tabulate` module to print out the result, so just trust me here. This is a good way to go!

Let's go back to using `csv.DictReader()` to parse our file and think about how we can break the `reps` value into int values for the low and high:

```
reader = csv.DictReader(fh, delimiter=',')
exercises = []
for rec in reader:
    name, reps = rec['exercise'], rec['reps']
    low, high = 0, 0 # what goes here?
    exercises.append((name, low, high))
```

You have a couple of tools at your disposal. Imagine `reps` is this:

```
>>> reps = '20-50'
```

The `str.split()` function could break that into two strings, “20” and “50”:

```
>>> reps.split('-')
['20', '50']
```

How could you turn each of the `str` values into integers?

Another way you could go is to use a regular expression. Remember that `\d` will match a digit, so `\d+` will match one or more digits. (Refer back to chapter 15 to refresh your memory on `\d` as a shortcut to the character class of digits.) You can wrap that expression in parentheses to capture the “low” and “high” values:

```
>>> match = re.match('(\d+)-(\d+)', reps)
>>> match.groups()
('20', '50')
```

Can you write a `read_csv()` function that passes the previous `test_read_csv()`?

19.1.5 Selecting the exercises

By this point, I'm hoping you've got `get_args()` straight and your `read_csv()` passes the given test. Now we can start in `main()` with printing out the data structure:

```
Get the command-line arguments.
def main():
    args = get_args()    ←
    random.seed(args.seed)
    pprint(read_csv(args.file)) ←
Set the random.seed()
with the args.seed value. ←
Read the args.file (which will be an open file
handle) using the read_csv() function and print the
resulting data structure. Note that I've imported
the pprint() function for demonstration purposes.
```

If you run the preceding code, you should see this:

```
$ ./wod.py
[('Burpees', 20, 50),
 ('Situps', 40, 100),
 ('Pushups', 25, 75),
 ('Squats', 20, 50),
 ('Pullups', 10, 30),
 ('Hand-stand pushups', 5, 20),
 ('Lunges', 20, 40),
 ('Plank', 30, 60),
 ('Crunches', 20, 30)]
```

We will use the `random.sample()` function to select the `--num` of exercises indicated by the user. Add `import random` to your program and modify your `main` to match this:

```
def main():
    args = get_args()
    random.seed(args.seed) ←
    exercises = read_csv(args.file) ←
    pprint(random.sample(exercises, k=args.num)) ←
Always set your random
seed before calling
random functions. ←
Read the input file. ←
Randomly select the given
number of exercises.
```

Now instead of printing all the exercises, it should print a random sample of the correct number of exercises. In addition, your sampling should exactly match this output if you set the `random.seed()` value:

```
$ ./wod.py -s 1
[('Pushups', 25, 75),
 ('Situps', 40, 100),
 ('Crunches', 20, 30),
 ('Burpees', 20, 50)]
```

We need to iterate through the sample and select a single “reps” value using the `random.randint()` function. The first exercise is push-ups, and the range is between 25 and 75 reps:

```
>>> import random
>>> random.seed(1)
```

```
>>> random.randint(25, 75)
33
```

If `args.easy` is `True`, you will need to halve that value. Unfortunately, we cannot have a fraction of a rep:

```
>>> 33/2
16.5
```

You can use the `int()` function to truncate the number to the integer component:

```
>>> int(33/2)
16
```

19.1.6 Formatting the output

Modify your program until it can reproduce this output:

```
$ ./wod.py -s 1
[('Pushups', 56), ('Situps', 88), ('Crunches', 27), ('Burpees', 35)]
```

We will use the `tabulate()` function from the `tabulate` module to format this list of tuple values into a text table:

```
>>> from tabulate import tabulate
>>> wod = [('Pushups', 56), ('Situps', 88), ('Crunches', 27), ('Burpees', 35)]
>>> print(tabulate(wod))
-----
Pushups    56
Situps    88
Crunches   27
Burpees   35
-----
```

If you read `help(tabulate)`, you will see that there is a `headers` option where you can specify a list of strings to use for the headers:

```
>>> print(tabulate(wod, headers=['Exercise', 'Reps']))
Exercise      Reps
-----
Pushups        56
Situps        88
Crunches       27
Burpees       35
```

If you synthesize all these ideas, you should be able to pass the provided tests.

19.1.7 Handling bad data

None of the tests will give your program bad data, but I have provided several “bad” CSV files in the `19_wod/inputs` directory that you might be interested in figuring out how to handle:

- bad-headers-only.csv is well-formed but has no data. It only has headers.
- bad-empty.csv is empty. That is, it is a zero-length file that I created with `touch bad-empty.csv`, and it has no data at all.
- bad-headers.csv has headers that are capitalized, so “Exercise” instead of “exercise,” “Reps” instead of “reps.”
- bad-delimiter.tab uses the tab character (`\t`) instead of the comma (,) as the field delimiter.
- bad-reps.csv contains reps that are not in the format x-y or which are not numeric or integer values.

Once your program passes the given tests, trying running it on the “bad” files to see how your program breaks. What should your program do when there is no usable data? Should your program print error messages when it encounters bad or missing values, or should it quietly ignore errors and only print the usable data? These are all real-world concerns that you will encounter, and it’s up to you to decide what your program will do. After the solution, I will show you ways I might deal with these files.

19.1.8 Time to write

OK, enough lollygagging. Time to write this program. You must do 10 push-ups every time you find a bug!

Here are a few hints:

- Use `csv.DictReader()` to parse the input CSV files.
- Break the `reps` field on the - character, coerce the low/high values to `int` values, and then use `random.randint()` to choose a random integer in that range.
- Use `random.sample()` to select the correct number of exercises.
- Use the `tabulate` module to format the output into a text table.

19.2 Solution

How did that go for you? Did you manage to modify your program to gracefully handle all the bad input files?

```
#!/usr/bin/env python3
"""Create Workout Of (the) Day (WOD)"""

import argparse
import csv
import io
import random
from tabulate import tabulate

# -----
def get_args():
    """Get command-line arguments"""
    # Import the tabulate
    # function we will use to
    # format the output table.
    # -----
```

```

parser = argparse.ArgumentParser(
    description='Create Workout Of (the) Day (WOD)',
    formatter_class=argparse.ArgumentDefaultsHelpFormatter)

parser.add_argument('-f',
                    '--file',
                    help='CSV input file of exercises',
                    metavar='FILE',
                    type=argparse.FileType('rt'),      ←
                    default='exercises.csv')

parser.add_argument('-s',
                    '--seed',
                    help='Random seed',
                    metavar='seed',
                    type=int,
                    default=None)

parser.add_argument('-n',
                    '--num',
                    help='Number of exercises',
                    metavar='exercises',
                    type=int,
                    default=4)

parser.add_argument('--easy',
                    '--easy',
                    help='Halve the reps',
                    action='store_true')

args = parser.parse_args()

if args.num < 1:
    parser.error(f"--num '{args.num}' must be greater than 0") ←

return args

```

Randomly select a value for reps that is in the provided range.

If args.easy is "truthy," cut the reps in half.

Ensure that args.num is a positive value.

Randomly sample the given number of exercises. The result will be a list of tuples that each contain three values, which can be unpacked directly into the variables name and low and high values.

Initialize wod as an empty list.

Read the input file into a list of exercises.

Append a tuple containing the name of the exercise and the reps to the wod.

```

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    random.seed(args.seed)
    wod = []
    exercises = read_csv(args.file) ←

    for name, low, high in random.sample(exercises, k=args.num):
        reps = random.randint(low, high)
        if args.easy:
            reps = int(reps / 2)
        wod.append((name, reps)) ←

```

```

print(tabulate(wod, headers=('Exercise', 'Reps')))

# -----
def read_csv(fh):
    """Read the CSV input"""

    exercises = []
    for row in csv.DictReader(fh, delimiter=','):
        low, high = map(int, row['reps'].split('-'))
        exercises.append((row['exercise'], low, high))

    return exercises

# -----
def test_read_csv():
    """Test read_csv"""

    text = io.StringIO('exercise,reps\nBurpees,20-50\nSitups,40-100')
    assert read_csv(text) == [('Burpees', 20, 50), ('Situps', 40, 100)]

# -----
if __name__ == '__main__':
    main()

```

Define a function to read an open CSV file handle.

Use the tabulate() function to format the wod into a text table using the appropriate headers.

Initialize exercises to an empty list.

Append a tuple containing the name of the exercise with the low and high values.

Return the list of exercises to the caller.

Verify that read_csv() can handle valid input data.

Create a mock file handle containing valid sample data.

Split the “reps” column on the dash, turn those values into integers, and assign to low and high variables.

Iterate through the file handle using the csv.DictReader() to create a dictionary combining the column names from the first row with the field values from the rest of the file. Use the comma as the field delimiter.

19.3 Discussion

Almost half the lines of the program are found within the `get_args()` function! Even though there's nothing new to discuss, I really want to point out how much work is being done to validate the inputs, provide defaults, create the usage statement, and so forth. Let's dig into the program, starting with the `read_csv()` function.

19.3.1 Reading a CSV file

Earlier in the chapter, I left you with one line where you needed to split the `reps` column and convert the values to integers. Here is one way:

```

def read_csv(fh):
    exercises = []
    for row in csv.DictReader(fh, delimiter=','):
        low, high = map(int, row['reps'].split('-'))
        exercises.append((row['exercise'], low, high))

    return exercises

```

Split the reps field on the dash, map the values through the int() function, and assign to low and high.

The annotated line works as follows. Assume a reps value like so:

```
>>> '20-50'.split('-')
['20', '50']
```

We need to turn each of those into an int value, which is what the int() function will do. We could use a list comprehension:

```
>>> [int(x) for x in '20-50'.split('-')]
[20, 50]
```

But the map() is much shorter and easier to read, in my opinion:

```
>>> list(map(int, '20-50'.split('-')))
[20, 50]
```

Since that produces exactly two values, we can assign them to two variables:

```
>>> low, high = map(int, '20-50'.split('-'))
>>> low, high
(20, 50)
```

19.3.2 Potential runtime errors

This code makes many, many assumptions that will cause it to fail miserably when the data doesn't match the expectations. For instance, what happens if the reps field contains no dash? It will produce one value:

```
>>> list(map(int, '20'.split('-')))
[20]
```

That will cause a *runtime* exception when we try to assign one value to two variables:

```
>>> low, high = map(int, '20'.split('-'))
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: not enough values to unpack (expected 2, got 1)
```

What if one or more of the values cannot be coerced to an int? It will cause an exception, and, again, you won't discover this until you run the program with bad data:

```
>>> list(map(int, 'twenty-thirty'.split('-')))
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: invalid literal for int() with base 10: 'twenty'
```

What happens if there is no reps field in the record, as is the case when the field names are capitalized?

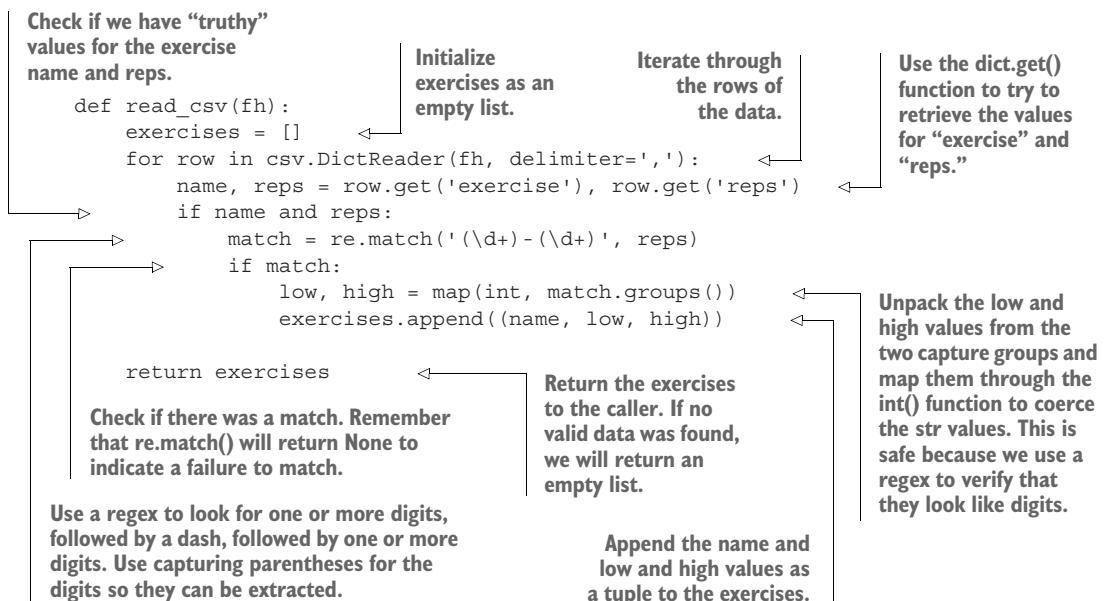
```
>>> rec = {'Exercise': 'Pushups', 'Reps': '20-50'}
```

Then the dictionary access `rec['reps']` will cause an exception:

```
>>> list(map(int, rec['reps'].split('-')))
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'reps'
```

The `read_csv()` function seems to work just fine as long as we pass it well-formed data, but the real world does not always give us clean datasets. An unfortunately large part of my job, in fact, is finding and correcting errors like this.

Earlier in the chapter, I suggested you might use a regular expression to extract the low and high values from the `reps` field. A regex has the advantage of inspecting the entire field, ensuring that it looks correct. Here is a more robust way to implement `read_csv()`:



19.3.3 Using `pandas.read_csv()` to parse the file

Many people familiar with statistics and data science will likely know the Python module called `pandas`, which mimics many ideas from the R programming language. I specifically chose the function name `read_csv()` because this is similar to a built-in function in R called `read.csv`, which was in turn used as the model for the `pandas.read_csv()` function. Both R and `pandas` tend to think of the data in delimited/CSV files in terms of a “data frame”—a two-dimensional object that allows you to deal with columns and rows of data.

To run the `using_pandas.py` version, you’ll need to install `pandas` like so:

```
$ python3 -m pip install pandas
```

Now you can try running this program:

```
import pandas as pd

df = pd.read_csv('inputs/exercises.csv')
print(df)
```

You'll see this output:

```
$ ./using_pandas.py
      exercise      reps
0        Burpees    20-50
1        Situps     40-100
2       Pushups     25-75
3       Squats      20-50
4      Pullups     10-30
5 Hand-stand pushups    5-20
6        Lunges     20-40
7        Plank      30-60
8      Crunches    20-30
```

Learning how to use pandas is far beyond the scope of this book. Mostly I just want you to be aware that this is a very popular way to parse delimited text files, especially if you intend to run statistical analyses over various columns of the data.

19.3.4 Formatting the table

Let's look at the `main()` function I included in the solution. You may notice a runtime exception waiting to happen:

```
def main():
    args = get_args()
    random.seed(args.seed)
    wod = []
    exercises = read_csv(args.file)

    for name, low, high in random.sample(exercises, k=args.num):
        reps = random.randint(low, high)
        if args.easy:
            reps = int(reps / 2)
        wod.append((name, reps))

    print(tabulate(wod, headers=['Exercise', 'Reps']))
```

This line will fail if args.num is greater than the number of elements in exercises, such as if read_csv() returns None or an empty list.

If you test the given solution with the `bad-headers-only.csv` file, you will see this error:

```
$ ./wod.py -f inputs/bad-headers-only.csv
Traceback (most recent call last):
  File "./wod.py", line 93, in <module>
    main()
  File "./wod.py", line 62, in main
    for name, low, high in random.sample(exercises, k=args.num):
  File "/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/random.py", line 363, in sample
```

```

raise ValueError("Sample larger than population or is negative")
ValueError: Sample larger than population or is negative

```

A safer way to handle this is to check that `read_csv()` returns enough data to pass to `random.sample()`. We have a couple of possible errors:

- No usable data was found in the input file.
- We are trying to sample too many records from the file.

Here is a possible way to handle these problems. Remember that calling `sys.exit()` with a string value will cause the program to print the message to `sys.stderr` and exit with a value of 1 (which is an error value):

```

def main():
    """Make a jazz noise here"""

    args = get_args()
    random.seed(args.seed)
    exercises = read_csv(args.file) ←
        | Read the input file into exercises.
        | The function should only return a
        | list, possibly empty.

    if not exercises: ←
        | Check if exercises is "falsey," such as an empty list.
        | sys.exit(f'No usable data in --file "{args.file.name}"')

    num_exercises = len(exercises)
    if args.num > num_exercises: ←
        | Check if we are trying to sample too many records.
        | sys.exit(f"--num '{args.num}' > exercises '{num_exercises}'")

    wod = []
    for name, low, high in random.sample(exercises, k=args.num):
        reps = random.randint(low, high)
        if args.easy:
            reps = int(reps / 2)
        wod.append((name, reps)) ←
            | Continue after we verify that we have enough valid data.

    print(tabulate(wod, headers=('Exercise', 'Reps')))


```

The version in `solution2.py` has these updated functions and gracefully handles all the bad input files. Note that I moved the `test_read_csv()` function to the `unit.py` file because it became much longer as I tested with various bad inputs.

You can run `pytest -xv unit.py` to run the unit tests. Let's inspect `unit.py` to see a more rigorous testing scheme:

```

import io
from wod import read_csv ←
    | Remember that you can import your own functions from
    | your own modules into other programs. Here we are
    | bringing in our read_csv() function. If we had instead used
    | import wod, we could call wod.read_csv().

The original, valid input
    | def test_read_csv():
    |     """Test read_csv"""
    |     good = io.StringIO('exercise,reps\nBurpees,20-50\nnSitups,40-100')
    |     assert read_csv(good) == [('Burpees', 20, 50), ('Situps', 40, 100)]

```

```

Testing with no data at all → no_data = io.StringIO('')
assert read_csv(no_data) == []

headers_only = io.StringIO('exercise,reps\n')
assert read_csv(headers_only) == []

→ bad_headers = io.StringIO('Exercise,Reps\nBurpees,20-50\nSitups,40-100')
assert read_csv(bad_headers) == []

→ bad_numbers = io.StringIO('exercise,reps\nBurpees,20-50\nSitups,forty-100')
assert read_csv(bad_numbers) == [('Burpees', 20, 50)]

no_dash = io.StringIO('exercise,reps\nBurpees,20\nSitups,40-100')
assert read_csv(no_dash) == [('Situps', 40, 100)] ←

tabs = io.StringIO('exercise\treps\nBurpees\t20-40\nSitups\t40-100') ←
assert read_csv(tabs) == []

A string ("forty") that cannot be coerced by int() to a numeric value

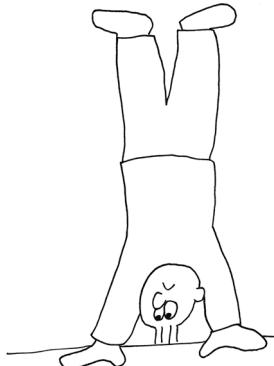
Well-formed file (correct headers and delimiter), but no data
Well-formed data with correct headers, but using a tab for the delimiter
A "reps" value ("20") missing a dash

```

The headers are capitalized, but only lowercase headers are expected.

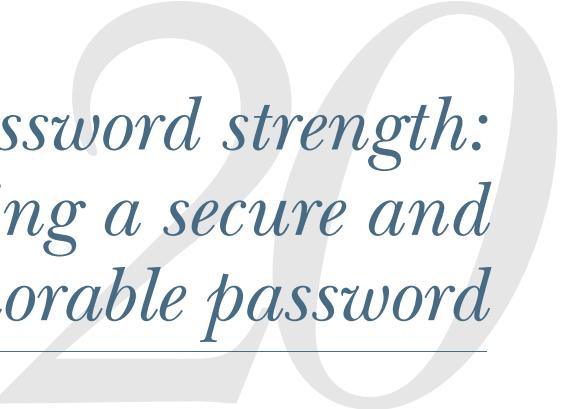
19.4 Going further

- Add an option to use a different delimiter, or guess that the delimiter is a tab if the input file extension is “.tab” as in the bad-delimiter.tab file.
- The `tabulate` module supports many table formats, including plain, simple, grid, pipe, orgtbl, rst, mediawiki, latex, latex_raw, and latex_booktabs. Add an option to choose a different `tabulate` format using these as the valid choices. Choose a reasonable default value.



Summary

- The `csv` module is useful for parsing delimited text data such as CSV and tab-delimited files.
- Text values representing numbers must be coerced to numeric values using `int()` or `float()` in order to be used as numbers inside your program.
- The `tabulate` module can be used to create text tables to format tabular output.
- Great care must be taken to anticipate and handle bad and missing data values. Tests can help you imagine all the ways in which your code might fail.



Password strength: Generating a secure and memorable password

It's not easy to create passwords that are both difficult to guess and easy to remember. An XKCD comic describes an algorithm that provides both security and recall by suggesting that a password be composed of "four random common words" (<https://xkcd.com/936/>). For instance, the comic suggests that the password composed of the words "correct," "horse," "battery," and "staple" would provide "~44 bits of entropy" which would require around 550 years for a computer to guess, given 1,000 guesses per second.

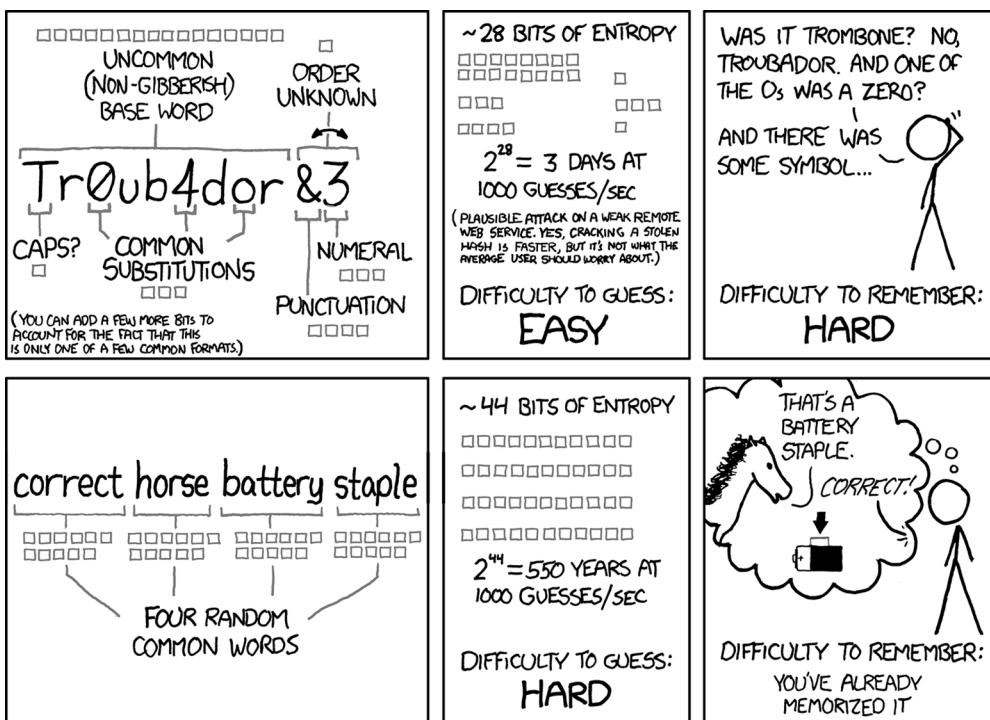
We're going to write a program called password.py that will create passwords by randomly combining words from some input files. Many computers have a file that lists thousands of English words, each on a separate line. On most of my systems, I can find this at /usr/share/dict/words, and it contains over 235,000 words! As the file can vary by system, I've added a version to the repo so that we can use the same file. This file is a little large, so I've compressed to inputs/words.txt.zip. You should unzip it before using it:

```
$ unzip inputs/words.txt.zip
```

Now we should both have the same inputs/words.txt file so that this is reproducible for you:

```
$ ./password.py ../inputs/words.txt --seed 14
CrotalLeavesMeeredLogy
NatalBurrelTizzyOddman
UnbornSignerShodDehort
```

Hmm, maybe those aren't going to be the easiest to remember! Perhaps instead we should be a bit more judicious about the source of our words? We're drawing from



THROUGH 20 YEARS OF EFFORT, WE'VE SUCCESSFULLY TRAINED EVERYONE TO USE PASSWORDS THAT ARE HARD FOR HUMANS TO REMEMBER, BUT EASY FOR COMPUTERS TO GUESS.

(Image used with permission from xkcd.com.)

a pool of over 200,000 words, but the average speaker tends to use somewhere between 20,000 and 40,000 words.

We can generate more memorable passwords by drawing from an actual piece of English text, such as the US Constitution. Note that to use a piece of input text in this way, we will need to remove any punctuation, as we have done in previous exercises:

```
$ ./password.py --seed 8 ../inputs/const.txt
DulyHasHeadsCases
DebtSevenAnswerBest
ChosenEmitTitleMost
```

Another strategy for generating memorable words could be to limit the pool of words to the more interesting parts of speech, like nouns, verbs, and adjectives taken from texts like novels or poetry. I've included a program I wrote called `harvest.py` that uses a natural language processing library in Python called spaCy (<https://spacy.io>) that will extract those parts of speech into files that we can use as input to our program. If

you want to use this program on your own input files, you'll need to be sure you first install the module:

```
$ python3 -m pip install spacy
```

I ran the `harvest.py` program on some texts and placed the outputs into directories in the `20_password` directory of the source repo. For instance, here is the output drawing from nouns found in the US Constitution:

```
$ ./password.py --seed 5 const/nouns.txt
TaxFourthYearList
TrialYearThingPerson
AidOrdainFifthThing
```

And here we have passwords generated using only verbs found in *The Scarlet Letter* by Nathaniel Hawthorne:

```
$ ./password.py --seed 1 scarlet/verbs.txt
CrySpeakBringHold
CouldSeeReplyRun
WearMeanGazeCast
```

And here are some generated from adjectives extracted from William Shakespeare's sonnets:

```
$ ./password.py --seed 2 sonnets/adjs.txt
BoldCostlyColdPale
FineMaskedKeenGreen
BarrenWiltFemaleSeldom
```

Just in case that does not result in a strong enough password, we will also provide a `--133t` flag to further obfuscate the text by

- 1** Passing the generated password through the `ransom.py` algorithm from chapter 12
- 2** Substituting various characters with a given table, as we did in `jump_the_five.py` from chapter 4
- 3** Adding a randomly selected punctuation character to the end

Here is what the Shakespearean passwords look like with this encoding:

```
$ ./password.py --seed 2 sonnets/adjs.txt --133t
B0LDco5TLYColdp@l3,
f1n3M45K3dK3eNGR33N[
B4rReNW1LTFeM4l3seldoM/
```

In this exercise, you will

- Take a list of one or more input files as positional arguments
- Use a regular expression to remove non-word characters
- Filter words by some minimum length requirement
- Use sets to create unique lists

- Generate a given number of passwords by combining some given number of randomly selected words
- Optionally encode text using a combination of algorithms we've previously written

20.1 Writing `password.py`

Our program should be written in the 20_password directory and will be called `password.py`. It will create some `--num` number of passwords (default, 3) each by randomly choosing some `--num_words` number of words (default, 4) from a unique set of words from one or more input files. As it will use the `random` module, the program will also accept a random `--seed` argument, which should be an integer value with a default of `None`. The words from the input files will need to be a `--min_word_len` minimum length (default, 3) up to a `--max_word_len` maximum length (default, 6) after removing any non-characters.

As always, our first priority is to sort out the inputs to the program. Do not move ahead until your program can produce this usage with the `-h` or `--help` flags and can pass the first eight tests:

```
$ ./password.py -h
usage: password.py [-h] [-n num_passwords] [-w num_words] [-m minimum]
                   [-x maximum] [-s seed] [-l]
                   FILE [FILE ...]

Password maker

positional arguments:
  FILE                  Input file(s)

optional arguments:
  -h, --help            show this help message and exit
  -n num_passwords, --num num_passwords
                        Number of passwords to generate (default: 3)
  -w num_words, --num_words num_words
                        Number of words to use for password (default: 4)
  -m minimum, --min_word_len minimum
                        Minimum word length (default: 3)
  -x maximum, --max_word_len maximum
                        Maximum word length (default: 6)
  -s seed, --seed seed  Random seed (default: None)
  -l, --l33t            Obfuscate letters (default: False)
```

The words from the input files will be title cased (first letter uppercase, the rest lowercase), which we can achieve using the `str.title()` method. This makes it easier to see and remember the individual words in the output. Note that we can vary the number of words included in each password as well as the number of passwords generated:

```
$ ./password.py --num 2 --num_words 3 --seed 9 sonnets/*
QueenThenceMasked
GullDeemdEven
```

The `--min_word_len` argument helps to filter out shorter, less interesting words like “a,” “I,” “an,” “of,” and so on, while the `--max_word_len` argument prevents the passwords from becoming unbearably long. If you increase these values, the passwords change quite drastically:

```
$ ./password.py -n 2 -w 3 -s 9 -m 10 -x 20 sonnets/*
PerspectiveSuccessionIntelligence
DistillationConscienceCountenance
```

The `--133t` flag is a nod to “leet”-speak, where `31337 H4X0R` means “ELITE HACKER”.¹ When this flag is present, we’ll encode each of the passwords in two ways. First, we’ll pass the word through the `ransom()` algorithm we wrote in chapter 12:

```
$ ./ransom.py MessengerRevolutionImportune
MESSENG3RReV0LUTI0N IMPORTUNE
```

Then, we’ll use the following substitution table to substitute characters in the same way we did in chapter 4:

```
a => @
A => 4
O => 0
t => +
E => 3
I => 1
S => 5
```

To cap it off, we’ll use `random.choice()` to select one character from `string.punctuation` to add to the end:

```
$ ./password.py --num 2 --num_words 3 --seed 9 --min_word_len 10 --max_word_len
20 sonnets/* --133t
p3RsPeC+1Vesucces5i0niN+3lL1Genc3$D1s+iLl@+ioNconsc1eNc3coun+eN@Nce^
```

Figure 20.1 shows a string diagram that summarizes the inputs.

20.1.1 Creating a unique list of words

Let’s start off by making our program print the name of each input file:

```
def main():
    args = get_args()
    random.seed(args.seed)           ← Always set random.seed() right
                                    away as it will globally affect all
                                    actions by the random module.

    for fh in args.files:          ← Iterate through the file arguments.
        print(fh.name)             ← Print the name of the file.
```

¹ See the “Leet” Wikipedia page (<https://en.wikipedia.org/wiki/Leet>) or the Cryptii translator <https://cryptii.com/>.

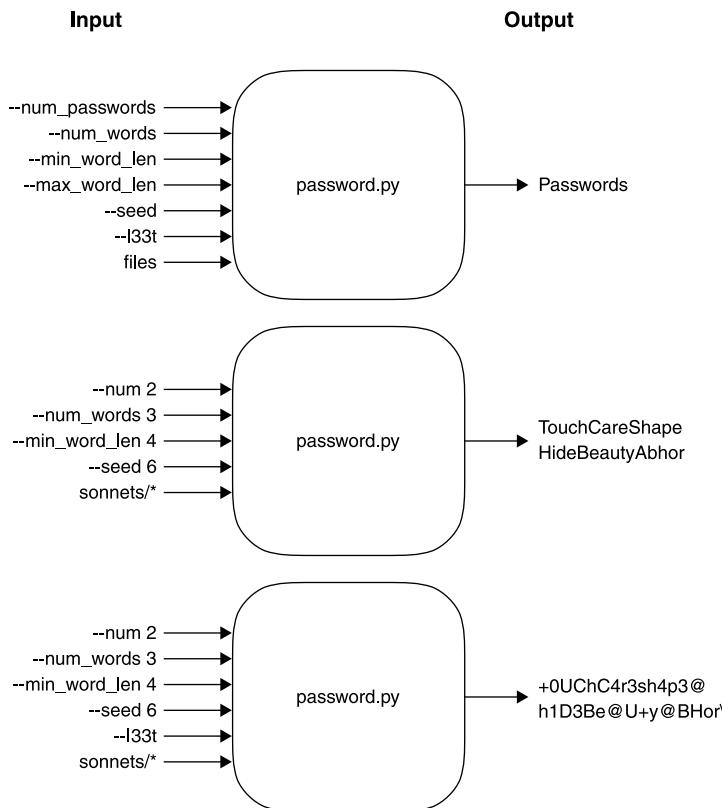


Figure 20.1 Our program has many possible options but requires only one or more input files. The output will be unbreakable passwords.

Let's test it with the `words.txt` file:

```
$ ./password.py ../inputs/words.txt
../inputs/words.txt
```

Now let's try it with some of the other inputs:

```
$ ./password.py scarlet/*
scarlet/adjs.txt
scarlet/nouns.txt
scarlet/verbs.txt
```

Our first goal is to create a unique list of words we can use for sampling. So far we've used lists to keep ordered collections of things like strings and numbers. The elements in a list do not have to be *unique*, though. We've also used dictionaries to create key/value pairs, and the keys of a dictionary *are* unique. Since we don't care about the values, we could set each key of a dictionary equal to some arbitrary value, like 1:

```

def main():
    args = get_args()
    random.seed(args.seed)
    words = {}

    for fh in args.file:
        for line in fh:
            for word in line.lower().split():
                words[word] = 1

    print(words)

```

Create an empty dict to hold the unique words.

Iterate through the files.

Iterate through the lines of the file.

Lowercase the line and split it on spaces into words.

Set the key words[word] equal to 1 to indicate we saw it. We're only using a dict to get the unique keys. We don't care about the values, so you could use whatever value you like.

If you run this on the US Constitution, you should see a fairly large list of words (some output elided here):

```
$ ./password.py ../inputs/const.txt
{'we': 1, 'the': 1, 'people': 1, 'of': 1, 'united': 1, 'states,'': 1, ...}
```

I can spot one problem, in that the word 'states,' has a comma attached to it. If we try in the REPL with the first bit of text from the Constitution, we can see the problem:

```
>>> 'We the People of the United States,'.lower().split()
['we', 'the', 'people', 'of', 'the', 'united', 'states,']
```

How can we get rid of the punctuation?

20.1.2 Cleaning the text

We've seen several times that splitting on spaces leaves punctuation, but splitting on non-word characters can break contracted words like "Don't" in two. We'd like a function that will `clean()` a word.

First let's imagine the test for it. Note that in this exercise, I'll put all my unit tests into a file called `unit.py`, which I can run with `pytest -xv unit.py`.

Here is the test for our `clean()` function:

```

def test_clean():
    assert clean('') == ''           ← It's always good to test your functions on nothing,
    assert clean("states,") == 'states'   ← just to make sure it does something sane.
    assert clean("Don't") == 'Dont'      ← The function should
                                         remove punctuation at
                                         the end of a string.

                                         The function should not split a
                                         contracted word in two.

```

I would like to apply this to all the elements returned by splitting each line into words, and `map()` is a fine way to do that. We often use a `lambda` when writing `map()`, as in figure 20.2.

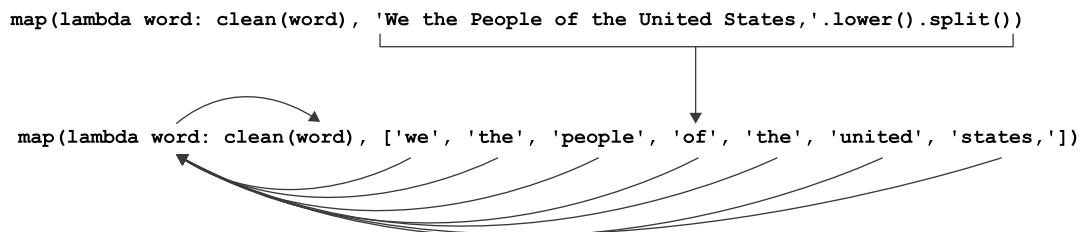


Figure 20.2 Writing `map()` using a `lambda` to accept each word from splitting a string

We don't actually need to write a `lambda` for `map()` here because the `clean()` function expects a single argument, as shown in figure 20.3.

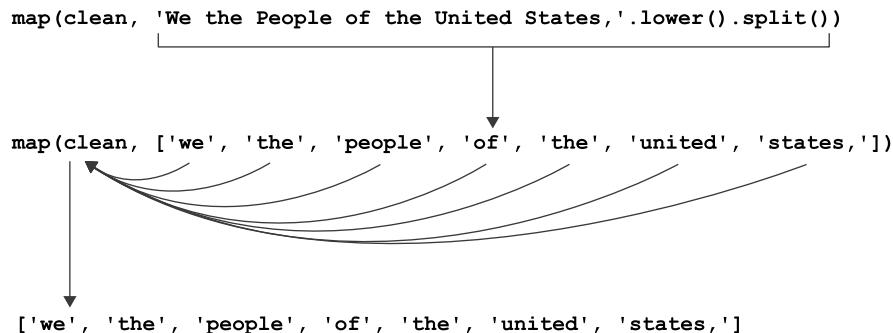


Figure 20.3 Writing the `map()` without the `lambda` because the function expects a single value

See how it integrates with the code:

```

def main():
    args = get_args()
    random.seed(args.seed)
    words = {}
    for fh in args.file:
        for line in fh:
            for word in map(clean, line.lower().split()):
                words[word] = 1
    print(words)
  
```

Use `map()` to apply the `clean()` function to the results of splitting the line on spaces.
No `lambda` is required because `clean()` expects a single argument.

If we run that on the US Constitution again, we can see that 'states' has been fixed:

```
$ ./password.py ../inputs/const.txt
{'we': 1, 'the': 1, 'people': 1, 'of': 1, 'united': 1, 'states': 1, ...}
```

I'll leave it to you to write a `clean()` function that will satisfy that test. You might use a list comprehension, a `filter()`, or maybe a regular expression. The choice is yours, so long as it passes the test.

20.1.3 Using a set

There is a better data structure than a `dict` to use for our purposes here. It's called a `set`, and you can think of it as being like a unique list or just the keys of a `dict`. Here is how we could change our code to use a `set` to keep track of *unique* words:

```
def main():
    args = get_args()
    random.seed(args.seed)
    words = set()           ← Use the set() function to
                           create an empty set.

    for fh in args.file:
        for line in fh:
            for word in map(clean, line.lower().split()):
                words.add(word)   ← Use set.add() to
                           add a value to a set.

    print(words)
```

If you run this code now, you will see slightly different output, where Python shows you a data structure in curly brackets (`{}`) that will make you think of a `dict`, but you'll notice that the contents look more like a list (as pointed out in figure 20.4):

```
$ ./password.py ../inputs/const.txt
{'', 'impartial', 'imposed', 'jared', 'levying', ...}
```

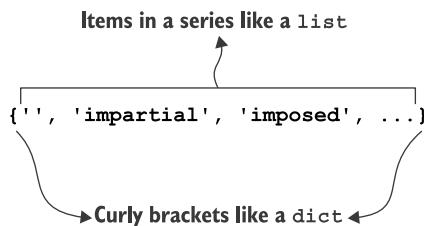


Figure 20.4 A set looks like a cross between a dictionary and a list.

We're using sets here because they so easily allow us to keep a unique list of words, but sets are much more powerful than this. For instance, you can find the shared values between two lists by using `set.intersection()`:

```
>>> nums1 = set(range(1, 10))
>>> nums2 = set(range(5, 15))
>>> nums1.intersection(nums2)
{5, 6, 7, 8, 9}
```

You can read `help(set)` in the REPL or in the documentation online to learn about all the amazing things you can do with sets.

20.1.4 Filtering the words

If we look again at the output we have, we'll see that the empty string is the first element:

```
$ ./password.py ../inputs/const.txt
{'', 'impartial', 'imposed', 'jared', 'levying', ...}
```

We need a way to filter out unwanted values like strings that are too short. In chapter 14 we looked at the `filter()` function, which is a higher-order function that takes two arguments:

- A function that accepts one element and returns `True` if the element should be kept or `False` if the element should be excluded
- Some “iterable” (like a list or `map()`) that produces a sequence of elements to be filtered

In our case, we want to accept only words that have a length greater than or equal to the `--min_word_len` argument, and less than or equal to `--max_word_len`. In the REPL, we can use a `lambda` to create an anonymous function that accepts a word and makes these comparisons. The result of that comparison is either `True` or `False`. Only words with a length from 3 to 6 are allowed, so this has the effect of removing short, uninteresting words. Remember that `filter()` is lazy, so I have to coerce it using the `list` function in the REPL to see the output:

```
>>> shorter = ['', 'a', 'an', 'the', 'this']
>>> min_word_len = 3
>>> max_word_len = 6
>>> list(filter(lambda word: min_word_len <= len(word) <= max_word_len, shorter))
['the', 'this']
```

This `filter()` will also remove longer words that would make our passwords cumbersome:

```
>>> longer = ['that', 'other', 'egalitarian', 'disequilibrium']
>>> list(filter(lambda word: min_word_len <= len(word) <= max_word_len, longer))
['that', 'other']
```

One way we could incorporate the `filter()` is to create a `word_len()` function that encapsulates the preceding `lambda`. Note that I defined it inside `main()` in order to create a *closure*, because I want to reference the values of `args.min_word_len` and `args.max_word_len`:

```
def main():
    args = get_args()
    random.seed(args.seed)
    words = set()

    def word_len(word):
        <-- This function will return True
        if the length of the given word
        is in the allowed range.
        return args.min_word_len <= len(word) <= args.max_word_len
```

```

for fh in args.file:
    for line in fh:
        for word in filter(word_len, map(clean, line.lower().split())):
            words.add(word)
print(words)

```

We can use `word_len` (without the parentheses!) as the function argument to `filter()`.

We can again try our program to see what it produces:

```
$ ./password.py ../inputs/const.txt
{'measures', 'richard', 'deprived', 'equal', ...}
```

Try it on multiple inputs, such as all the nouns, adjectives, and verbs from *The Scarlet Letter*:

```
$ ./password.py scarlet/*
{'walk', 'lose', 'could', 'law', ...}
```



20.1.5 Titlecasing the words

We used the `line.lower()` function to lowercase all the input, but the passwords we generate will need each word to be in “Title Case,” where the first letter is uppercase and the rest of the word is lowercase. Can you figure out how to change the program to produce this output?

```
$ ./password.py scarlet/*
{'Dark', 'Sinful', 'Life', 'Native', ...}
```

Now we have a way to process any number of files to produce a unique list of title-cased words that have non-word characters removed and have been filtered to remove the ones that are too short or long. That’s quite a lot of power packed into a few lines of code!

20.1.6 Sampling and making a password

We’re going to use the `random.sample()` function to randomly choose `--num` number of words from our set to create an unbreakable, yet memorable, password. We’ve talked before about the importance of using a random seed to test that our “random” selections are reproducible. It’s also quite important that the items from which we sample always be ordered in the same way so that the same selections are made. If we use the `sorted()` function on a set, we get back a sorted list, which is perfect for using with `random.sample()`.

We can add this line to the code from before:

```
words = sorted(words)
print(random.sample(words, args.num_words))
```

Now when I run the program with *The Scarlet Letter* input, I will get a list of words that might make an interesting password:

```
$ ./password.py scarlet/*
['Lose', 'Figure', 'Heart', 'Bad']
```

The result of `random.sample()` is a list that you can join on the empty string in order to make a new password:

```
>>> ''.join(random.sample(words, num_words))
'TokenBeholdMarketBegin'
```

You will need to create the number of passwords indicated by the user, similar to how we created some number of insults in chapter 9. How will you do that?

20.1.7 l33t-ify

The last piece of our program involves creating an `l33t()` function that will obfuscate the password. The first step is to convert the password with the same algorithm we wrote for `ransom.py`. I'm going to create a `ransom()` function for this, and here is the test that is in `unit.py`:

```
def test_ransom():
    state = random.getstate()           ← Save the current
    random.seed(1)                     ← global state.
    assert ransom('Money') == 'moNeY'   ← Set random.seed()
    assert ransom('Dollars') == 'DOLLaRs'← to a known value
    random.setstate(state)             ← for the test.
                                         ← Restore the state.
```

I'll leave it to you to create the function that satisfies this test.

NOTE You can run `pytest -xv unit.py` to run the unit tests. The program will import the various functions from your `password.py` file to test. Open `unit.py` and inspect it to understand how this happens.

Next I will replace some of the characters according to the following table. I recommend you revisit chapter 4 to see how you did that:

a =>	@
A =>	4
O =>	0
t =>	+
E =>	3
I =>	1
S =>	5

I wrote an `l33t()` function that combines `random()` with the preceding substitution and then adds a punctuation character by appending `random.choice(string.punctuation)`.

Here is the `test_l33t()` function you can use to write your function. It works almost identically to the previous test, so I shall eschew commentary:

```
def test_l33t():
    state = random.getstate()
    random.seed(1)
    assert l33t('Money') == 'moNeY{'
    assert l33t('Dollars') == 'DOLL4r5`'
    random.setstate(state)
```

20.1.8 Putting it all together

Without giving away the ending, I'd like to say that you need to be *really careful* about the order of operations that include the random module. My first implementation would print different passwords given the same seed when I used the `--l33t` flag. Here was the output for plain passwords:

```
$ ./password.py -s 1 -w 2 sonnets/*
EagerCarcanet
LilyDial
WantTempest
```

I would have expected the *exact same passwords*, only encoded. Here is what my program produced instead:

```
$ ./password.py -s 1 -w 2 sonnets/* --l33t
3@G3RC@rC@N3+{
m4dness5iNcoN5+4n+|
MouTh45s15T4nCe^
```

The first password looks OK, but what are those other two? I modified my code to print both the original password and the l33ted one:

```
$ ./password.py -s 1 -w 2 sonnets/* --l33t
3@G3RC@rC@N3+{ (EagerCarcanet)
m4dness5iNcoN5+4n+| (MadnessInconstant)
MouTh45s15T4nCe^ (MouthAssistance)
```

The random module uses a global state to make each of its “random” choices. In my first implementation, I was modifying this state after choosing the first password by immediately modifying the new password with the `l33t()` function. Because the `l33t()` function also uses random functions, the state was altered for the next password. My solution was to first generate *all* the passwords and then alter them using the `l33t()` function, if necessary.

Those are all the pieces you should need to write your program. You have the unit tests to help you verify the functions, and you have the integration tests to ensure your program works as a whole.

20.2 Solution

I hope you will use your program to generate your passwords. Be sure to share them with your author, especially the ones to your bank account and favorite shopping sites!

```
#!/usr/bin/env python3
"""Password maker, https://xkcd.com/936/"""

import argparse
import random
import re
import string
```

```
# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Password maker',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file',
                        metavar='FILE',
                        type=argparse.FileType('rt'),
                        nargs='+',
                        help='Input file(s)')

    parser.add_argument('-n',
                        '--num',
                        metavar='num_passwords',
                        type=int,
                        default=3,
                        help='Number of passwords to generate')

    parser.add_argument('-w',
                        '--num_words',
                        metavar='num_words',
                        type=int,
                        default=4,
                        help='Number of words to use for password')

    parser.add_argument('-m',
                        '--min_word_len',
                        metavar='minimum',
                        type=int,
                        default=3,
                        help='Minimum word length')

    parser.add_argument('-x',
                        '--max_word_len',
                        metavar='maximum',
                        type=int,
                        default=6,
                        help='Maximum word length')

    parser.add_argument('-s',
                        '--seed',
                        metavar='seed',
                        type=int,
                        help='Random seed')

    parser.add_argument('-l',
                        '--l33t',
                        action='store_true',
                        help='Obfuscate letters')

    return parser.parse_args()
```

```

Set the random.seed() to the given
value or the default None, which is
the same as not setting the seed.

# -----
def main():
    args = get_args()
    random.seed(args.seed)
    words = set()

Iterate
through
each open
file handle.

Title-case the word
before adding it
to the set.

Create an empty set
to hold all the unique
words we'll extract
from the texts.

Iterate through each word generated by splitting
the lowercased line on spaces, removing non-
word characters with the clean() function, and
filtering for words of an acceptable length.

Create a word_len() function
for filter() that returns True
if the word's length is in the
allowed range and False
otherwise.

for fh in args.file:
    for line in fh:
        for word in filter(word_len, map(clean, line.lower().split())):
            words.add(word.title())

Iterate through each line
of text in the file handle.

words = sorted(words)
passwords = [
    ''.join(random.sample(words, args.num_words)) for _ in range(args.num)

] See if the args.l33t
flag is True.

if args.l33t:
    passwords = map(l33t, passwords)

print('\n'.join(passwords))

Use a list comprehension with a range to create the
correct number of passwords. Since I don't need the
actual value from range, I can use _ to ignore the value.

Use the sorted() function to order
words into a new list.

Define a
function
to l33t()
a word.

Define a function
to clean() a word.

Use the regular expression to
substitute the empty string for
anything that is not an English
alphabet character.

Use the ransom()
function to randomly
capitalize letters.

Make a translation
table/dict for character
substitutions.

Use the str.translate() function to perform the substitutions
and append a random piece of punctuation.

# -----
def clean(word):
    """Remove non-word characters from word"""

    return re.sub('[^a-zA-Z]', '', word)

# -----
def l33t(text):
    """l33t"""

    text = ransom(text)
    xform = str.maketrans({
        'a': '@', 'A': '4', 'O': '0', 'T': '+', 'E': '3', 'I': '1', 'S': '5'
    })
    return text.translate(xform) + random.choice(string.punctuation)

Use the str.translate() function to perform the substitutions
and append a random piece of punctuation.

```

```

# -----
def ransom(text):
    """Randomly choose an upper or lowercase letter to return"""
    return ''.join(
        map(lambda c: c.upper() if random.choice([0, 1]) else c.lower(), text))

# -----
if __name__ == '__main__':
    main()

Return a new string created by randomly
upper- or lowercasing each letter in a word.

```

Define a function for the
ransom() algorithm
from chapter 12.

20.3 Discussion

I hope you found this program challenging and interesting. There wasn't anything new in `get_args()`, but, again, about half the lines of code are found just in this function. I feel this is indicative of just how important it is to correctly define and validate the inputs to a program!

Now, let's get on with talking about the auxiliary functions.

20.3.1 Cleaning the text

I chose to use a regular expression to remove any characters that are outside the set of lower- and uppercase English characters:

```

def clean(word):
    """Remove non-word characters from word"""
    return re.sub('[^a-zA-Z]', '', word)

```

The `re.sub()` function will substitute any text matching the pattern (the first argument) found in the given text (the third argument) with the value given by the second argument.

Recall from chapter 18 that we can write the character class `[a-zA-Z]` to define the characters in the ASCII table bounded by those two ranges. We can then *negate* or complement that class by placing a caret (^) as the *first character* inside that class, so `[^a-zA-Z]` can be read as “any character not matching a to z or A to Z.”

It's perhaps easier to see it in action in the REPL. In the following example, only the letters “AbCd” will be left from the text “Alb*C!d4”:

```

>>> import re
>>> re.sub('[^a-zA-Z]', '', 'Alb*C!d4')
'AbCd'

```

If the only goal were to match ASCII letters, it would be possible to solve it by looking for membership in `string.ascii_letters`:

```

>>> import string
>>> text = 'Alb*C!d4'
>>> [c for c in text if c in string.ascii_letters]
['A', 'b', 'C', 'd']

```

A list comprehension with a guard can also be written using `filter()`:

```
>>> list(filter(lambda c: c in string.ascii_letters, text))
['A', 'b', 'C', 'd']
```

Both of the non-regex versions seem like more effort to me. Additionally, if the function ever needed to be changed to allow, say, numbers and a few specific pieces of punctuation, the regular expression version becomes significantly easier to write and maintain.

20.3.2 A King's ransom

The `ransom()` function was taken straight from the `ransom.py` program in chapter 12, so there isn't too much to say about it except, hey, look how far we've come! What was the idea for an entire chapter is now a single line in a much longer and more complicated program:

```
def ransom(text):
    """Randomly choose an upper or lowercase letter to return"""
    return ''.join(
        map(lambda c: c.upper() if random.choice([0, 1]) else c.lower(), text))
```

Join the resulting list from the `map()` on the empty string to create a new string.

Use `map()` to iterate through each character in the text and select either the upper- or lowercase version of the character based on a “coin toss,” using `random.choice()` to select between a “truthy” value (1) or a “falsey” value (0).

20.3.3 How to l33t()

The `l33t()` function builds on `ransom()` and then adds a text substitution that is straight out of chapter 4. I like the `str.translate()` version of that program, so I used it again here:

```
def l33t(text):
    """l33t"""
    text = ransom(text)
    xform = str.maketrans({
        'a': '@', 'A': '4', 'O': '0', 't': '+', 'E': '3', 'I': '1', 'S': '5'
    })
    return text.translate(xform) + random.choice(string.punctuation)
```

Randomly capitalize the given text.

Make a translation table from the given dict that describes how to modify one character to another. Any characters not listed in the keys of this dict will be ignored.

Use the `str.translate()` method to make all the character substitutions. Use `random.choice()` to select one additional character from `string.punctuation` to append to the end.

20.3.4 Processing the files

To use these functions, we need to create a unique set of all the words in our input files. I wrote this bit of code with an eye both on performance and on style:

```
words = set()
for fh in args.file:
    for line in fh:
```

Iterate through each open file handle.

Read the file handle line by line with a for loop, not with a method like `fh.read()`, which will read the entire contents of the file at once.

```

for word in filter(word_len, map(clean, line.lower().split())):
    words.add(word.title())

```

Title-case the word before adding it to the set.

Reading this code requires starting at the end where I split line.lower() on spaces. Each word from str.split() goes into clean(), which then must pass through the filter() function.

Figure 20.5 shows a diagram of that for line.

- 1 line.lower() will return a lowercase version of line.
- 2 The str.split() method will break the text on whitespace to return words.
- 3 Each word is fed into the clean() function to remove any character that is not in the English alphabet.
- 4 The cleaned words are filtered by the word_len() function.
- 5 The resulting word has been transformed, cleaned, and filtered.

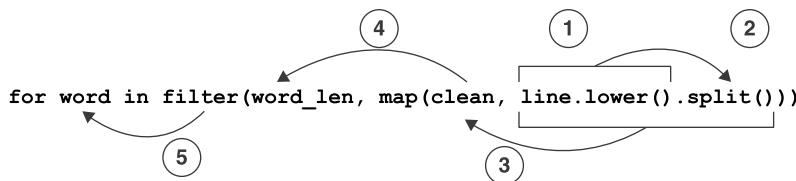


Figure 20.5 A visualization of the order of operations for the various functions

If you don't like the `map()` and `filter()` functions, you might rewrite the code like so:

```

words = set()
for fh in args.files:
    for line in fh:
        for word in line.lower().split():
            word = map(clean)
            if args.min_word_len <= len(word) <= args.max_word_len:
                words.add(word.title())

```

Iterate through each open file handle.

Iterate through each line of the file handle.

Iterate through each “word” from splitting the lowercased line on spaces.

Remove unwanted characters.

Add the title-cased word to the set.

Check if the word is an acceptable length.

However you choose to process the files, at this point you should have a complete set of all the unique, title-cased words from the input files.

20.3.5 Sampling and creating the passwords

As noted earlier, it's vital to sort the words for our tests so that we can verify that we are making consistent choices. If you only wanted random choices and didn't care about testing, you would not need to worry about sorting—but then you'd also be a morally

deficient person for not testing, so perish the thought! I chose to use the `sorted()` function, as there is no other way to sort a set:

```
words = sorted(words)
```

There is no `set.sort()` function. Sets are ordered internally by Python. Calling `sorted()` on a set will create a new, sorted list.

We need to create a given number of passwords, and I thought it might be easiest to use a `for` loop with a `range()`. In my code, I used `for _ in range(...)` just as in chapter 9 because I don't need to know the value each time through the loop. The underscore (`_`) is a way to indicate that you are ignoring the value. It's fine to say `for i in range(...)` if you want, but some linters might complain if they see that your code declares the variable `i` but never uses it. That could legitimately be a bug, so it's best to use the `_` to show that you mean to ignore this value.

Here is the first way I wrote the code that led to the bug I mentioned earlier, where different passwords would be chosen even when I used the same random seed. Can you spot the bug?

Iterate through the `args.num` of passwords to create.

Each password will be based on a random sampling from words, and I will choose the value given in `args.num_words`. The `random.sample()` function returns a list of words that I `str.join()` on the empty string to create a new string.

```
for _ in range(args.num):
    password = ''.join(random.sample(words, args.num_words))
    print(l33t(password) if args.l33t else password)
```

If the `args.l33t` flag is True, we'll print the l33t version of the password; otherwise, I'll print the password as is. This is the bug! Calling `l33t()` here modifies the global state used by the `random` module, so the next time I call `random.sample()`, I get a different sample.

The solution is to separate the concerns of generating the passwords and possibly modifying them:

```
passwords = [
    ''.join(random.sample(words, args.num_words)) for _ in range(args.num)
]

if args.l33t:
    passwords = map(l33t, passwords)

print('\n'.join(passwords))
```

Use a list comprehension to iterate through `range(args.num)` to generate the correct number of passwords.

If the `args.leet` flag is True, use the `l33t()` function to modify the passwords.

Print the passwords joined on newlines.

20.4 Going further

- The substitution part of the `l33t()` function changes every available character, which perhaps makes the password too difficult to remember. It would be better to modify only maybe 10% of the password, much like how we changed the input strings in chapter 10's Telephone exercise.

- Create programs that combine other skills you've learned. Like maybe a lyrics generator that randomly selects lines from files of songs by your favorite bands, then encodes the text as in chapter 15, then changes all the vowels to one vowel as in chapter 8, and then SHOUTS IT OUT as in chapter 5?

Summary

- A set is a unique collection of values. Sets can interact with other sets to create differences, intersections, unions, and more.
- Changing the order of operations using the random module can change the output of a program because the global state of the random module may be affected.
- Short, tested functions can be composed to create more complicated, tested programs. Here we combined many ideas from previous exercises in concise, powerful expressions.



21

Tic-Tac-Toe: Exploring state

One of my favorite movies is the 1983 release *War Games* starring Matthew Broderick, whose character, David, plays a young hacker who enjoys cracking into computer systems ranging from his school's grade book to a Pentagon server that has the potential to launch intercontinental ballistic missiles. Central to the plot is the game of Tic-Tac-Toe, a game so simple that it usually ends in a draw between the two players.

In the movie, David engages Joshua, an artificial intelligence (AI) agent, who is capable of playing lots of nice games like chess. David would rather play the game Global Thermonuclear War with Joshua. Eventually David realizes that Joshua is using the simulation of a war game to trick the US military into initiating a nuclear first strike against the Soviet Union. Understanding the mutually assured destruction (MAD) doctrine, David asks Joshua to play himself at Tic-Tac-Toe so that he can explore the futility of games that can never result in victory. After hundreds or thousands of rounds all ending in draws, Joshua concludes that “the only winning move is not to play,” at which point Joshua stops trying to destroy the Earth and suggests instead that they could play “a nice game of chess.”



I assume you already know the game of Tic-Tac-Toe, but we'll review briefly in case your childhood missed countless games of this with your friends. The game starts out with a 3-by-3 square grid. There are two players who take turns marking first X and then O in the cells. A player wins by placing their mark in any three squares in a straight line, horizontally, vertically, or diagonally. This is usually impossible, as each player will generally use their moves to block a potential win by their opponent.

We will spend the last two chapters writing Tic-Tac-Toe. We will explore ideas for representing and tracking program *state*, which is a way of thinking about how the pieces of a program change over time. For instance, we'll start off with a blank board, and the first player to go is X. Play alternates between the X and O, and after each round two cells on the board will have been taken by the two players. We'll need to keep track of these moves and more, so that, at any moment, we always know the state of the game.

If you recall, the hidden state of the `random` module proved to be a problem in chapter 20, where an early solution we explored produced inconsistent results depending on the order of the operations that used the module. In this exercise, we're going to think about ways to make the state of our game, and any changes to it, explicit.

In this chapter, we'll write a program that plays just one turn of the game; then in the next chapter we'll expand the program to handle a full game. This version of the program will be given a string that represents the state of the playing board at any time during a game. The default is the empty board at the beginning of the game, before either player has made a move. The program may also be given one move to add to that board. It will print a picture of the board and report if there is a winner after making the move.

For this program, we need to track at least two ideas in our state:

- The board, identifying which player has marked which squares of the grid
- The winner, if there is one

For the next version, we'll write an interactive version of the game where we will need to track and update several more items in the state through a complete game of Tic-Tac-Toe.

In this exercise, you will

- Consider how to use elements like strings and lists to represent aspects of a program's state
- Enforce the rules of a game in code, such as preventing a player from playing in a cell that has already been taken
- Use a regular expression to validate the initial board
- Use `and` and `or` to reduce combinations of Boolean values to a single value
- Use lists of lists to find a winning board
- Use the `enumerate()` function to iterate a list with the index and value

21.1 Writing tictactoe.py

You will create a program called tictactoe.py in the 21_tictactoe directory. As usual, I would recommend you start the program using new.py or template.py. Let's discuss the parameters for the program.

The initial state of the board will come from a `-b` or `--board` option that describes which cells are occupied by which players. Since there are nine cells, we'll use a string that is nine characters long, composed only of the characters X and O, or the period (.) to indicate that the cell is open. The default board will be a string of nine dots. When you display the board, you will either display the player's mark in a cell or the cell's number, from one to nine. In the next version of the game, this number will be used by the player to identify a cell for their move. As there is no winner for the default board, the program should print "No winner":

```
$ ./tictactoe.py
-----
| 1 | 2 | 3 |
-----
| 4 | 5 | 6 |
-----
| 7 | 8 | 9 |
-----
No winner.
```

The `--board` option will describe which cells should be marked for which player, where the positions in the string describe the different cells, ascending from 1 to 9. In the string X.O..O..X, the positions 1 and 9 are occupied by "X" and positions 3 and 6 by "O" (see figure 21.1).

Here is how that grid would be rendered by the program:

```
$ ./tictactoe.py -b X.O..O..X
-----
| x | 2 | o |
-----
| 4 | 5 | o |
-----
| 7 | 8 | x |
-----
No winner.
```

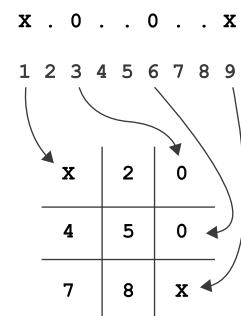


Figure 21.1 The board is nine characters describing the nine cells of the board.

We can additionally modify the given `--board` by passing a `-c` or `--cell` option of 1–9 and a `-p` or `--player` option of "X" or "O." For instance, we can mark the first cell as "X" like so:

```
$ ./tictactoe.py --cell 1 --player X
-----
| x | 2 | 3 |
-----
```

```
| 4 | 5 | 6 |
-----
| 7 | 8 | 9 |
-----
No winner.
```

The winner, if any, should be declared with gusto:

```
$ ./tictactoe.py -b XXXOO.....
-----
| X | X | X |
-----
| O | O | 6 |
-----
| 7 | 8 | 9 |
-----
X has won!
```

As usual, we'll use a test suite to ensure that our program works properly. Figure 21.2 shows the string diagram.

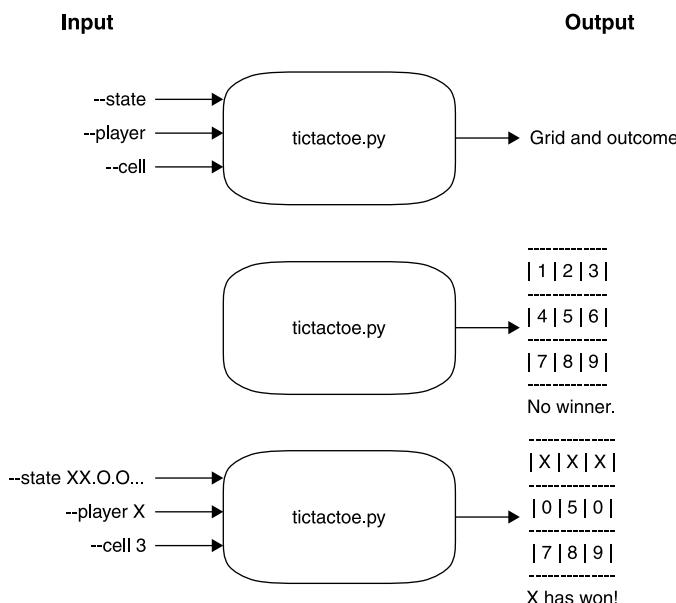


Figure 21.2 Our Tic-Tac-Toe program will play one turn of the game using a board, player, and cell. It should print the board and winner.

21.1.1 Validating user input

There's a fair bit of input validation that needs to happen. The --board needs to ensure that any argument is exactly 9 characters and is composed only of X, O, and .:

```
$ ./tictactoe.py --board XXXOOO...
usage: tictactoe.py [-h] [-b board] [-p player] [-c cell]
tictactoe.py: error: --board "XXXOOO..." must be 9 characters of ., X, O
```

Likewise, the --player can only be X or O:

```
$ ./tictactoe.py --player A --cell 1
usage: tictactoe.py [-h] [-b board] [-p player] [-c cell]
tictactoe.py: error: argument -p/--player: \
invalid choice: 'A' (choose from 'X', 'O')
```

And the --cell can only be an integer value from 1 to 9:

```
$ ./tictactoe.py --player X --cell 10
usage: tictactoe.py [-h] [-b board] [-p player] [-c cell]
tictactoe.py: error: argument -c/--cell: \
invalid choice: 10 (choose from 1, 2, 3, 4, 5, 6, 7, 8, 9)
```

Both --player and --cell must be present together, or neither can be present:

```
$ ./tictactoe.py --player X
usage: tictactoe.py [-h] [-b board] [-p player] [-c cell]
tictactoe.py: error: Must provide both --player and --cell
```

Lastly, if the --cell specified is already occupied by an X or an O, the program should error out:

```
$ ./tictactoe.py --player X --cell 1 --board X..O.....
usage: tictactoe.py [-h] [-b board] [-p player] [-c cell]
tictactoe.py: error: --cell "1" already taken
```

I would recommend you put all this error checking into `get_args()` so that you can use `parser.error()` to throw the errors and halt the program.

21.1.2 Altering the board

The initial board, once validated, describes which cells are occupied by which player. This board can be altered by adding the --player and --cell arguments. It may seem silly to not just pass in the already altered --board, but this is necessary practice for writing the interactive version.

If you represent board as a `str` value, like 'XX.O.O..X', and you need to change cell 3 to an X, for instance, how will you do that? For one thing, cell 3 is not found at *index* 3 in the given board—the index is *one less* than the cell number. The other issue is that a `str` is immutable. Just as in chapter 10's Telephone program, you'll need to figure out a way to modify one character in the board value.

21.1.3 Printing the board

Once you have a board, you'll need to format it with ASCII characters to create a grid. I recommend you make a function called `format_board()` that takes the board string as an argument and returns a `str` that uses dashes (-) and vertical pipes (|) to create a table. I have provided a `unit.py` file that contains the following test for the default, unoccupied grid:

```
def test_board_no_board():
    """makes default board"""

    board = """
-----  

| 1 | 2 | 3 |
-----  

| 4 | 5 | 6 |
-----  

| 7 | 8 | 9 |
-----  

""".strip()

    assert format_board('.' * 9) == board
```

Use triple quotes because the string has embedded newlines. The final `str.strip()` call will remove the trailing newline used to format the code.

If you multiply a string by an integer value, Python will repeat the given string that number of times. Here we create a string of nine dots as the input to `format_board()`. We expect the return should be an empty board as formatted here.

Now try formatting a board with some other combination. Here's another test I wrote that you might like to use, but feel free to write your own:

```
def test_board_with_board():
    """makes board"""

    board = """
-----  

| 1 | 2 | 3 |
-----  

| O | X | X |
-----  

| 7 | 8 | 9 |
-----  

""".strip()

    assert format_board('...OXX...') == board
```

The given board should have the first and third rows open and the second row with "OXX."

It would be impractical to test every possible combination for the board. When you're writing tests, you'll often have to rely on spot-checking your code. Here I am checking the empty board and a non-empty board. Presumably if the function can handle these two arguments, it can handle any others.

21.1.4 Determining a winner

Once you have validated the input and printed the board, your last task is to declare a winner if there is one. I chose to write a function called `find_winner()` that returns either `X` or `O` if one of those is the winner, or returns `None` if there is no winner. To test

this, I wrote out every possible winning board, to test my function with values for both players. You are welcome to use this test:

```
def test_winning():
    """test winning boards"""

    wins = [('PPP.....'), ('...PPP...'), ('.....PPP'), ('P..P..P..'),
            ('.P..P..P.'), ('..P..P..P'), ('P...P...P'), ('..P..P..P.')]
    Check for both players, X and O.

    for player in 'XO':
        other_player = 'O' if player == 'X' else 'X'
        Iterate through each of the winning combinations.

        for board in wins:
            board = board.replace('P', player)
            dots = [i for i in range(len(board)) if board[i] == '.']
            mut = random.sample(dots, k=2)
            test_board = ''.join([
                other_player if i in mut else board[i]
                for i in range(len(board))
            ])
            Change all the P (for "player") values in the given board to the player that we're checking.

            assert find_winner(test_board) == player
            Find the indexes of the open cells (indicated by a dot).
            Assert that find_winner() will determine that this board wins for the given player.

This is a list of the board indexes that, if occupied by the same player, would win.

Determine which is the opposite player from X or O.

Randomly sample two open cells. We will mutate these, so I call them mut.

Alter the board to change the two selected mut cells to other_player.
```

I also wanted to be sure I would not falsely claim that a losing board is winning, so I also wrote the following test to ensure that `None` is returned when there is no winner:

```
def test_losing():
    """test losing boards"""

    losing_board = list('XXOO.....')
    Run 10 tests.

    for _ in range(10):
        random.shuffle(losing_board)
        assert find_winner(''.join(losing_board)) is None
No matter how this board is arranged, it cannot win, as there are only two marks for each player.

Shuffle the losing board into another configuration.

Assert that, no matter how the board is arranged, we will still find no winner.
```

If you choose the same function names as I did, you can run `pytest -xv unit.py` to run the unit tests I wrote. If you wish to write different functions, you can create your own unit tests either inside your `tictactoe.py` file or in another unit file.

After printing the board, be sure to print “[Winner] has won!” or “No winner” depending on the outcome. All righty, you have your orders, so get marching!

21.2 Solution

We’re taking baby steps towards the full, interactive game in the next chapter. Right now we need to cement some basics on how just one turn will be played. It’s good to make iterations of difficult programs, where you start as simply as possible and slowly add features to build a more complex idea.

```

#!/usr/bin/env python3
"""Tic-Tac-Toe"""

import argparse
import re

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Tic-Tac-Toe',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('-b',
                        '--board',
                        help='The state of the board',
                        metavar='board',
                        type=str,
                        default='.' * 9)

    parser.add_argument('-p',
                        '--player',
                        help='Player',
                        choices='XO',
                        metavar='player',
                        type=str,
                        default=None)

    parser.add_argument('-c',
                        '--cell',
                        help='Cell 1-9',
                        metavar='cell',
                        type=int,
                        choices=range(1, 10),
                        default=None)

    args = parser.parse_args()

    if any([args.player, args.cell]) and not all([args.player, args.cell]):
        parser.error('Must provide both --player and --cell')

    if not re.search('^[.XO]{9}$', args.board):
        parser.error(f'--board "{args.board}" must be 9 characters of ., X, O')

    if args.player and args.cell and args.board[args.cell - 1] in 'XO':
        parser.error(f"--cell '{args.cell}' already taken")

    return args

```

Use a regular expression to check that --board is comprised of exactly nine valid characters.

The --board will default to nine dots. If you use the multiplication operator (*) with a string value and an integer (in any order), the result is the string value repeated that many times. So “! * 9” will produce ‘.....’!

The --player must be either X or O, which can be validated using choices.

The --cell must be an integer from 1 to 9, which can be validated with type=int and choices=range(1, 10), remembering that the upper bound (10) is not included.

The combination of any() and all() is a way to test that both arguments are present or neither is.

If both --player and --cell are present and valid, verify that the cell in the board is not currently occupied.

Modify the board if both cell and player are “truthy.” Since the arguments are validated in `get_args()`, it’s safe to use them here. That is, I won’t accidentally assign an index value that is out of range because I have taken the time to check that the cell value is acceptable.

```
# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    board = list(args.board)

    if args.player and args.cell:
        board[args.cell - 1] = args.player

    print(format_board(board))
    winner = find_winner(board)
    print(f'{winner} has won!' if winner else 'No winner.')
```

Look for a winner in the board.

Since we may need to alter the board, it’s easiest to convert it to a list.

Since the cells start numbering at 1, subtract 1 from the cell to change the correct index in board.

Define a function to format the board. The function does not print() the board because that would make it hard to test. The function returns a new string value that can be printed or tested.

Print the outcome of the game. The `find_winner()` function returns either X or O if one of the players has won, or None to no indicate no winner.

```
# -----
def format_board(board):
    """Format the board"""

    cells = [str(i) if c == '.' else c for i, c in enumerate(board, 1)]
    bar = '-----'
    cells_tmpl = '| {} | {} | {} | '
    return '\n'.join([
        cells_tmpl.format(*cells[:3]), bar,
        cells_tmpl.format(*cells[3:6]), bar,
        cells_tmpl.format(*cells[6:]), bar
    ])
```

Iterate through the cells in the board and decide whether to print the player, if the cell is occupied, or the cell number, if it is not.

Define a function that returns a winner or the value None if there is no winner. Again, the function does not print() the winner but only returns an answer that can be printed or tested.

The return from the function is a new string created by joining all the lines of the grid on newlines.

```
# -----
def find_winner(board):
    """Return the winner"""

    winning = [[0, 1, 2], [3, 4, 5], [6, 7, 8], [0, 3, 6], [1, 4, 7],
               [2, 5, 8], [0, 4, 8], [2, 4, 6]]
```

Iterate through both players, X and O.

Iterate through each winning combination of cells, unpacking them into the variables i, j, and k.

There are eight winning boards, which are defined as eight lists of the cells that need to be occupied by the same player. Note that I chose here to represent the actual zero-offset index values and not the 1-based values I expect from the user.

```

combo = [board[i], board[j], board[k]]
if combo == [player, player, player]:
    return player
else:
    return None

```

If that is True, return the player. If this is never True for any of the combinations, we exit the function without returning a value, and so `None` is returned by default.

21.2.1 Validating the arguments and mutating the board

Most of the validation can be handled by using argparse effectively. Both the `--player` and `--cell` options can be handled by the `choices` option. It's worth taking time to appreciate the use of `any()` and `all()` in this code:

```

if any([args.player, args.cell]) and not all([args.player, args.cell]):
    parser.error('Must provide both --player and --cell')

```

We can play with these functions in the REPL. The `any()` function is the same as using `or` in between Boolean values:

```

>>> True or False or True
True

```

If `any` of the items in a given list is “truthy,” the whole expression will evaluate to `True`:

```

>>> any([True, False, True])
True

```

If `cell` is a non-zero value, and `player` is not the empty string, they are both “truthy”:

```

>>> cell = 1
>>> player = 'X'
>>> any([cell, player])
True

```

The `all()` function is the same as using `and` in between all the elements in a list, so *all* of the elements need to be “truthy” in order for the whole expression to be `True`:

```

>>> cell and player
'X'

```

Why does that return `X`? It returns the last “truthy” value, which is the `player` value, so if we reverse the arguments, we’ll get the `cell` value:

```

>>> player and cell
1

```

If we use `all()`, it evaluates the truthiness of anding the values, which will be `True`:

```
>>> all([cell, player])
True
```

We are trying to figure out if the user has provided only *one* of the arguments for `--player` and `--cell`, because we need both or we want neither. So we pretend `cell` is `None` (the default) but `player` is `X`. It's true that `any()` of those values is "truthy":

```
>>> cell = None
>>> player = 'X'
>>> any([cell, player])
True
```

But it's not true that they *both* are:

```
>>> all([cell, player])
False
```

So when we and those two expressions, they return `False`,

```
>>> any([cell, player]) and all([cell, player])
False
```

because that is the same as saying this:

```
>>> True and False
False
```

The default for `--board` is provided as nine dots, and we can use a regular expression to verify that it's correct:

```
>>> board = '.' * 9
>>> import re
>>> re.search('^[.XO]{9}$', board)
<re.Match object; span=(0, 9), match='.....'>
```

Our regular expression creates a character class composed of the dot `(.)`, "X," and "O" by using `[.XO]`. The `{9}` indicates that there must be exactly 9 characters, and the `^` and `$` characters anchor the expression to the beginning and end of the string, respectively (see figure 21.3).

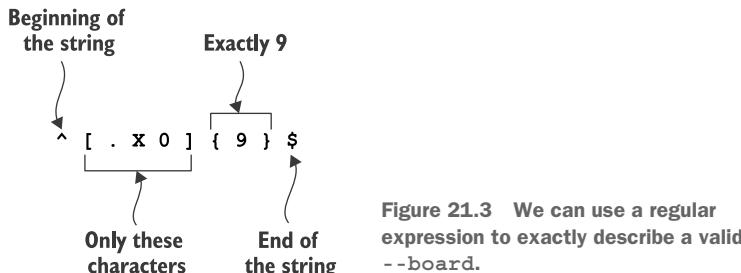


Figure 21.3 We can use a regular expression to exactly describe a valid `--board`.

You could manually validate this using the magic of `all()` again:

- Is the length of board exactly 9 characters?
- Is it true that each of the characters is one of those allowed?

Here is one way to write it:

```
>>> board = '...XXXOOO'
>>> len(board) == 9 and all([c in '.XO' for c in board])
True
```

The `all()` part is checking this:

```
>>> [c in '.XO' for c in board]
[True, True, True, True, True, True, True, True]
```

Since each character `c` (“cell”) in `board` is in the allowed set of characters, all the comparisons are `True`. If we change one of the characters, a `False` will show up:

```
>>> board = '...XXXOOA'
>>> [c in '.XO' for c in board]
[True, True, True, True, True, True, True, True, False]
```

Any `False` value in an `all()` expression will return `False`:

```
>>> all([c in '.XO' for c in board])
False
```

The last piece of validation checks if the --cell being set to --player is already occupied:

```
if args.player and args.cell and args.board[args.cell - 1] in 'XO':
    parser.error(f"--cell \"{args.cell}\" already taken")
```

Because --cell starts counting from 1 instead of 0, we must subtract 1 when we use it as an index into the --board argument. Given the following inputs, the first cell has been set to X, and now O wants the same cell:

```
>>> board = 'X.....'
>>> cell = 1
>>> player = 'O'
```

We can ask if the value in `board` at `cell - 1` has already been set:

```
>>> board[cell - 1] in 'XO'
True
```

Or you could instead check if that position is *not* a dot:

```
>>> boards[cell - 1] != '.'
True
```

It's rather exhausting to validate all the inputs, but this is the only way to ensure that the game is played properly.

In the `main()` function, we might need to alter the board of the game if there are arguments for both cell and player. I decided to make `board` into a `list` precisely because I might need to alter it in this way:

```
if player and cell:
    board[cell - 1] = player
```

21.2.2 Formatting the board

Now it's time to create the grid. I chose to create a function that returns a string value that I could test rather than directly printing the grid. Here is my version:

```
def format_board(board):
    """Format the board"""

    cells = [str(i) if c == '.' else c for i, c in enumerate(board, start=1)] ←
    bar = '-----'
    cells_tmpl = '| {} | {} | {} | '
    return '\n'.join([
        bar,
        cells_tmpl.format(*cells[:3]), bar, ←
        cells_tmpl.format(*cells[3:6]), bar,
        cells_tmpl.format(*cells[6:]), bar
    ])
    The asterisk, or “splat” (*), is shorthand to  
expand the list returned by the list slice operation  
into values that the str.format() function can use.
```

I used a list comprehension to iterate through each position and character of board using the `enumerate()` function. Because I would rather start counting from index position 1 than 0, I used the `start=1` option. If the character is a dot, I want to print the position as the cell number; otherwise, I print the character, which will be X or O.

The “splat” syntax of `*cell[:3]` is a shorter way of writing the code, like so:

```
return '\n'.join([
    bar,
    cells_tmpl.format(cells[0], cells[1], cells[2]), bar,
    cells_tmpl.format(cells[3], cells[4], cells[5]), bar,
    cells_tmpl.format(cells[6], cells[7], cells[8]), bar
])
```

The `enumerate()` function returns a list of tuples that include the index and value of each element in a list (see figure 21.4). Since it's a lazy function, I must use the `list()` function in the REPL to view the values:

```
>>> board = 'XX.O.O...'
>>> list(enumerate(board))
[(0, 'X'), (1, 'X'), (2, '.'), (3, 'O'), (4, '.'), (5, 'O'), (6, '.'), (7, '.'), (8, '.')]
```

In this instance, I would rather start counting at 1, so I can use the `start=1` option:

```
>>> list(enumerate(board, start=1))
[(1, 'X'), (2, 'X'), (3, '.'), (4, 'O'), (5, '.'), (6, 'O'), (7, '.'), (8, '.'), (9, '.')]
```

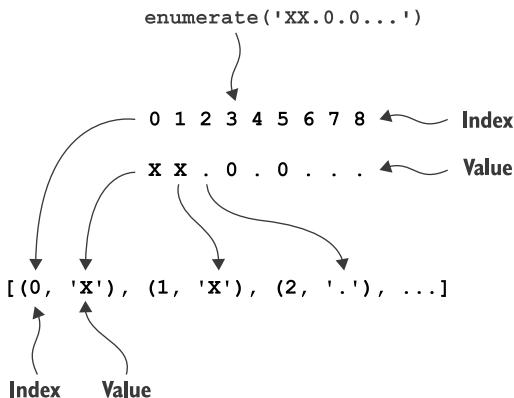


Figure 21.4 The `enumerate()` function will return the index and value of items in a series. By default, the initial index is 0.

This list comprehension could alternatively be written as a `for` loop:

```

cells = [] ← Initialize an empty list to hold the cells.
for i, char in enumerate(board, start=1): ← Unpack each tuple of the index (starting at 1) and value of each character in board into the variables i (for "integer") and char.
    cells.append(str(i) if char == '.' else char) ← If the char is a dot, use the string version of the i value; otherwise, use the char value.
  
```

Figure 21.5 illustrates how `enumerate()` is unpacked into `i` and `char`.

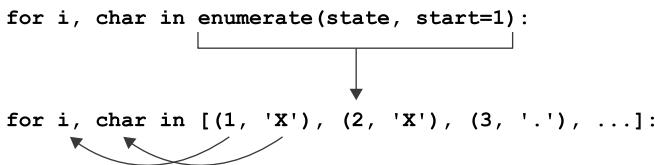


Figure 21.5 The tuples containing the indexes and values returned by `enumerate()` can be assigned to two variables in the `for` loop.

This version of `format_board()` passes all the tests found in `unit.py`.

21.2.3 Finding the winner

The last major piece to this program is determining if either player has won by placing three of their marks in a row horizontally, vertically, or diagonally.

```

def find_winner(board):
    """Return the winner"""
    winning = [[0, 1, 2], [3, 4, 5], [6, 7, 8], [0, 3, 6], [1, 4, 7],
               [2, 5, 8], [0, 4, 8], [2, 4, 6]] ← There are eight winning positions—the three horizontal rows, the three vertical columns, and the two diagonals—so I decided to create a list where each element is also a list that contains the three cells in a winning configuration.
  
```

```

for player in ['X', 'O']:
    for i, j, k in winning:
        combo = [board[i], board[j], board[k]]
        if combo == [player, player, player]:
            return player

```

It's typical to use `i` as a variable name for "integer" values, especially when their life is rather brief, as here. When more similar names are needed in the same scope, it's also common to use `j`, `k`, `l`, etc. You may prefer to use names like `cell1`, `cell2`, and `cell3`, which are more descriptive but also longer to type. The unpacking of the cell values is exactly the same as the unpacking of the tuples in the previous `enumerate()` code (see figure 21.6).

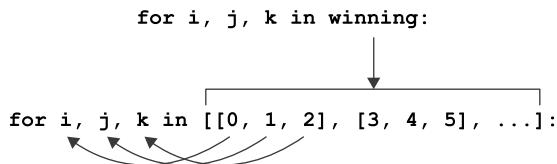


Figure 21.6 As with the unpacking of the `enumerate()` tuples, each list of three elements can be unpacked into three variables in the `for` loop.

The rest of the code checks if either `X` or `O` is the only character at each of the three positions. I worked out half a dozen ways to write this, but I'll just share this one alternate version that uses two of my favorite functions, `all()` and `map()`:

```

Iterate through each combination
of cells in winning.
for combo in winning:
    group = list(map(lambda i: board[i], combo))
    for player in ['X', 'O']:
        if all(x == player for x in group):
            return player

```

Use map() to get the value
of board at each position in
the combination.

Check for
each player,
X and O.

See if all the values in the
group are equal to the
given player.

If so, return
that player.

If a function has no explicit `return` or never executes a `return`, as would be the case here when there is no winner, Python will use the `None` value as the default `return`. We'll interpret `None` to mean there is no winner when we print the outcome of the game:

```

winner = find_winner(board)
print(f'{winner} has won!' if winner else 'No winner.')

```

That covers this version of the game that plays just one turn of Tic-Tac-Toe. In the next chapter, we'll expand these ideas into an interactive version that starts with a blank board and dynamically requests user input to play the game.

21.3 Going further

- Write a game that will play one hand of a card game like Blackjack (Twenty-one) or War.

Summary

- This program uses a `str` value to represent the Tic-Tac-Toe board with nine characters representing X, O, or . to indicate a taken or empty cell. We sometimes convert that to a `list` to make it easier to modify.
- A regular expression is a handy way to validate the initial board. We can declaratively describe that it should be a string exactly nine characters long composed only of the characters ., X, and O.
- The `any()` function is like chaining `or` between multiple Boolean values. It will return `True` if *any* of the values is “truthy.”
- The `all()` function is like using `and` between multiple Boolean values. It will return `True` only if every one of the values is “truthy.”
- The `enumerate()` function will return the list index and value for each element in an iterable like a `list`.

Tic-Tac-Toe redux: An interactive version with type hints

In this last exercise, we’re going to revisit the Tic-Tac-Toe game from the previous chapter. That version played one turn of the game by accepting an initial `--board` and then modifying it if there were also valid options for `--player` and `--cell`. It printed the one board and the winner, if any. We’re going to extend those ideas into a version that will always start from an empty board and will play as many turns as needed to complete a game, ending with a winner or a draw.

This program will be different from all the other programs in this book because it will accept no command-line arguments. The game will always start with a blank “board” and with the X player going first. It will use the `input()` function to interactively ask each player, X and then O, for a move. Any invalid move, such as choosing an occupied or non-existing cell, will be rejected. At the end of each turn, the game will decide to stop if it determines there is a win or a draw.

In this chapter you will

- Use and break out of an infinite loop
- Add type hints to your code
- Explore tuples, named tuples, and typed dictionaries
- Use `mypy` to analyze code for errors, especially misuse of types



22.1 Writing `itictactoe.py`

This is the one program where I won't provide an integration test. The program doesn't take any arguments, and I can't easily write tests that will interact dynamically with the program. This also makes it difficult to show a string diagram, because the output of the program will be different depending on the moves you make. Still, figure 22.1 is an approximation of how you could think of the program starting with no inputs and then looping until some outcome is determined, or the player quits.

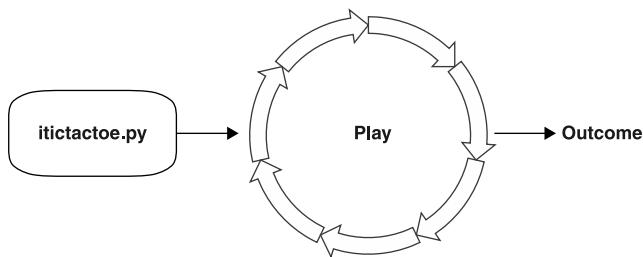


Figure 22.1 This version of Tic-Tac-Toe accepts no arguments and will play in an infinite loop until some conclusion like a win, draw, or forfeit.

I encourage you to start off by running the `solution1.py` program to play a few rounds of the game. The first thing you may notice is that the program clears the screen of any text and shows you an empty board, along with a prompt for the X player's move. I'll type 1 and press Enter:

```

-----
| 1 | 2 | 3 |
-----
| 4 | 5 | 6 |
-----
| 7 | 8 | 9 |
-----
Player X, what is your move? [q to quit]: 1

```

Then you will see that cell 1 is now occupied by X, and the player has switched to O:

```

-----
| X | 2 | 3 |
-----
| 4 | 5 | 6 |
-----
| 7 | 8 | 9 |
-----
Player O, what is your move? [q to quit]:

```

If I choose 1 again, I am told that cell is already taken:

```

-----
| X | 2 | 3 |
-----
| 4 | 5 | 6 |

```

```
-----  
| 7 | 8 | 9 |  
-----  
Cell "1" already taken  
Player O, what is your move? [q to quit]:
```

Note that the player is still O because the previous move was invalid. The same happens if I put in some value that cannot be converted to an integer:

```
-----  
| x | 2 | 3 |  
-----  
| 4 | 5 | 6 |  
-----  
| 7 | 8 | 9 |  
-----  
Invalid cell "biscuit", please use 1-9  
Player O, what is your move? [q to quit]:
```

Or if I enter an integer that is out of range:

```
-----  
| x | 2 | 3 |  
-----  
| 4 | 5 | 6 |  
-----  
| 7 | 8 | 9 |  
-----  
Invalid cell "10", please use 1-9  
Player O, what is your move? [q to quit]:
```

You should be able to reuse many of the ideas from chapter 21's version of the game to validate the user input.

If I play the game to a conclusion where one player gets three in a row, it prints the winning board and proclaims the victor:

```
-----  
| x | o | 3 |  
-----  
| 4 | x | 6 |  
-----  
| 7 | o | x |  
-----  
X has won!
```

22.1.1 Tuple talk

In this version, we'll write an interactive game that always starts with an empty grid and plays as many rounds as necessary to reach a conclusion with a win or a draw. The idea of "state" in the last game was limited to the board—which players were in which cells. This version requires us to track quite a few more variables in our game state:

- The cells of the board, like `..XO..X.O`
- The current player, either `X` or `O`
- Any error, such as the player entering a cell that is occupied or that does not exist or a value that cannot be converted to a number
- Whether the user wishes to quit the game early
- Whether the game is a draw, which happens when all the cells of the grid are occupied but there is no winner
- The winner, if any, so we know when the game is over

You don't need to write your program exactly the way I wrote mine, but you still may find yourself needing to keep track of many items. A `dict` is a natural data structure for that, but I'd like to introduce a new data structure called a "named tuple," as it plays nicely with Python's type hints, which will figure prominently in my solution.

We've encountered tuples throughout the exercises. They've been returned by something like `match.groups()` when a regular expression contains capturing parentheses, like in chapters 14 and 17; when using `zip` to combine two lists, like in chapter 19; or when using `enumerate()` to get a list of index values and elements from a `list`. A `tuple` is an immutable `list`, and we'll explore how that immutability can prevent us from introducing subtle bugs into our programs.

You create a tuple whenever you put commas between values:

```
>>> cell, player
(1, 'X')
```

It's most common to put parentheses around them to make it more explicit:

```
>>> (cell, player)
(1, 'X')
```

We could assign this to a variable called `state`:

```
>>> state = (cell, player)
>>> type(state)
<class 'tuple'>
```

We index into a tuple using list index values:

```
>>> state[0]
1
>>> state[1]
'X'
```

Unlike with a list, we cannot change any of the values inside the tuple:

```
>>> state[1] = 'O'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'tuple' object does not support item assignment
```

It's going to be inconvenient remembering that the first position is the `cell` and the second position is the `player`, and it will get much worse when we add all the other fields. We could switch to using a dict so that we can use strings to access the values of state, but dictionaries are mutable, and it's also easy to misspell a key name.

22.1.2 Named tuples

It would be nice to combine the safety of an immutable tuple with named fields, which is exactly what we get with the `namedtuple()` function. First, you must import it from the `collections` module:

```
>>> from collections import namedtuple
```

The `namedtuple()` function allows us to describe a new class for values. Let's say we want to create a class that describes the idea of `State`. A class is a group of variables, data, and functions that together can be used to represent some idea. The Python language itself, for example, has the `str` class, which represents the idea of a sequence of characters that can be contained in a variable that has some `len` (length), and which can be converted to uppercase with `str.upper()`, can be iterated with a `for` loop, and so forth. All these ideas are grouped into the `str` class, and we've used `help(str)` to read the documentation for that class inside the REPL.

The class name is the first argument we pass to `namedtuple()`, and the second argument is a list of the field names in the class. It's common practice to capitalize class names:

```
>>> State = namedtuple('State', ['cell', 'player'])
```

We've just created a new type called `State`!

```
>>> type(State)
<class 'type'>
```

Just as there is a function called `list()` to create a list type, we can now use the `State()` function to create a named tuple of the type `State` that has two named fields, `cell` and `player`:

```
>>> state = State(1, 'X')
>>> type(state)
<class '__main__.State'>
```

We can still access the fields with index values, like any list or tuple:

```
>>> state[0]
1
>>> state[1]
'X'
```

But we can also use their names, which is much nicer. Notice that there are no parentheses at the end, as we are accessing a field, not calling a method:

```
>>> state.cell
1
>>> state.player
'X'
```

Because `state` is a tuple, we cannot mutate the value once it has been created:

```
>>> state.cell = 1
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: can't set attribute
```

This is actually *good* in many instances. It's often quite dangerous to change your data values once your program has started. You should use tuples or named tuples whenever you want a list- or dictionary-like structure that cannot be accidentally modified.

There is a problem, however, in that there's nothing to prevent us from instantiating a `state` with the fields out of order *and of the wrong types*—`cell` should be an `int`, and `player` should be a `str`!

```
>>> state2 = State('O', 2)
>>> state2
State(cell='O', player=2)
```

In order to avoid that, you can use the field names, so that their order no longer matters:

```
>>> state2 = State(player='O', cell=2)
>>> state2
State(cell=2, player='O')
```

Now you have a data structure that looks like a dict but has the immutability of a tuple!

22.1.3 Adding type hints

We still have a big problem in that there's nothing preventing us from assigning a `str` to the `cell`, which ought to be an `int`, and vice versa for `int` and `player`:

```
>>> state3 = State(player=3, cell='X')
>>> state3
State(cell='X', player=3)
```

Starting in Python 3.6, the `typing` module allows you to add *type hints* to describe the data types for variables. You should read PEP 484 (www.python.org/dev/peps/pep-0484/) for more information, but the basic idea is that we can use this module to describe the appropriate types for variables and type signatures for functions.

I'm going to improve our `State` class by using the `NamedTuple` class from the `typing` module as the base class. First we need to import from the `typing` module the classes we'll need, such as `NamedTuple`, `List`, and `Optional`, the last of which describes a type that could be `None` or some other class like a `str`:

```
from typing import List, NamedTuple, Optional
```

Now we can specify a `State` class with named fields, types, and even default values to represent the initial state of the game where the board is empty (all dots) and player X goes first. Note that I decided to store the board as a list of characters rather than a `str`:

```
class State(NamedTuple):
    board: List[str] = list('.' * 9)
    player: str = 'X'
    quit: bool = False
    draw: bool = False
    error: Optional[str] = None
    winner: Optional[str] = None
```

We can use the `State()` function to create a new value that's set to the initial state:

```
>>> state = State()
>>> state.board
['.', '.', '.', '.', '.', '.', '.', '.', '.']
>>> state.player
'X'
```

You can override any default value by providing the field name and a value. For instance, we could start the game off with player O by specifying `player='O'`. Any field we don't specify will use the default:

```
>>> state = State(player='O')
>>> state.board
['.', '.', '.', '.', '.', '.', '.', '.', '.']
>>> state.player
'O'
```

We get an exception if we misspell a field name, like `playre` instead of `player`:

```
>>> state = State(playre='O')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: __new__() got an unexpected keyword argument 'playre'
```

22.1.4 Type verification with Mypy

As nice as all the above is, *Python will not generate a runtime error if we assign an incorrect type*. For instance, I can assign `quit` a `str` value of `'True'` instead of the `bool` value `True`, and nothing at all happens:

```
>>> state = State(quit='True')
>>> state.quit
'True'
```

The benefit of type hints comes from using a program like Mypy to check our code. Let's place all this code into a small program called `typehints.py` in the repo:

```
#!/usr/bin/env python3
""" Demonstrating type hints """

from typing import List, NamedTuple, Optional

class State(NamedTuple):
    board: List[str] = list('. ' * 9)
    player: str = 'X'
    quit: bool = False
    draw: bool = False
    error: Optional[str] = None
    winner: Optional[str] = None

state = State(quit='False') ←
print(state) ←
```

Annotations for the State class:

- A callout points to the `quit` field with the text: "quit is defined as a bool, which means it should only allow values of True and False."
- A callout points to the assignment of `'False'` to `quit` with the text: "We are assigning the str value 'True' instead of the bool value True, which might be an easy mistake to make, especially in a very large program. We'd like to know this type of error will be caught!"

The program will execute *with no errors*:

```
$ ./typehints.py
State(board=['.', '.', '.', '.', '.', '.', '.', '.', '.'], player='X', \
quit='False', draw=False, error=None, winner=None)
```

But the Mypy program will report the error of our ways:

```
$ mypy typehints.py
typehints.py:16: error: Argument "quit" to "State" has incompatible type
      "str"; expected "bool"
Found 1 error in 1 file (checked 1 source file)
```

If I correct the program like so,

```
#!/usr/bin/env python3
""" Demonstrating type hints """

from typing import List, NamedTuple, Optional

class State(NamedTuple):
    board: List[str] = list('. ' * 9)
    player: str = 'X'
    quit: bool = False
    draw: bool = False
    error: Optional[str] = None
    winner: Optional[str] = None
```

Annotation for the corrected quit field:

- A callout points to the `quit` field with the text: "Again, quit is a bool value."

```
state = State(quit=True)
print(state)
```

We have to assign an actual
bool value in order to pass
muster with Mypy.

now Mypy will be satisfied:

```
$ mypy typehints2.py
Success: no issues found in 1 source file
```

22.1.5 Updating immutable structures

If one of the advantages of using NamedTuples is their *immutability*, how will we keep track of changes to our program? Consider our initial state of an empty grid with the player X going first:

```
>>> state = State()
```

Imagine X takes cell 1, so we need to change board to X..... and the player to O. We can't directly modify state:

```
>>> state.board=list('X.....')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: can't set attribute
```

We could use the `State()` function to create a new value to overwrite the existing state. That is, since we can't change anything *inside* the state variable, we could instead point state to an entirely new value. We did this in the second solution in chapter 8, where we needed to change a `str` value, because they are also immutable in Python.

To do this, we can copy all the current values that haven't changed and combine them with the changed values:

```
>>> state = State(board=list('X.....'), player='O', quit=state.quit, \
    draw=state.draw, error=state.error, winner=state.winner)
```

The `namedtuple._replace()` method, however, provides a much simpler way to do this. Only the values we provide are changed, and the result is a new State:

```
>>> state = state._replace(board=list('X.....'), player='O')
```

We overwrite our state variable with the return from `state._replace()`, just as we have repeatedly overwritten string variables with new values:

```
>>> state
State(board=['X', '.', '.', '.', '.', '.', '.', '.', '.', '.'], player='O', \
    quit=False, draw=False, error=None, winner=None)
```

This is much more convenient than having to list all the fields—we only need to specify the fields that have changed. We are also prevented from accidentally modifying

any of the other fields, and we are likewise prevented from forgetting or misspelling any fields or setting them to the wrong types.

22.1.6 Adding type hints to function definitions

Now let's look at how we can add type hints to our function definitions. For an example, we can modify our `format_board()` function to indicate that it takes a parameter called `board`, which is a list of string values, by adding `board: List[str]`. Additionally, the function returns a `str` value, so we can add `-> str` after the colon on the `def` to indicate this, as in figure 22.2.

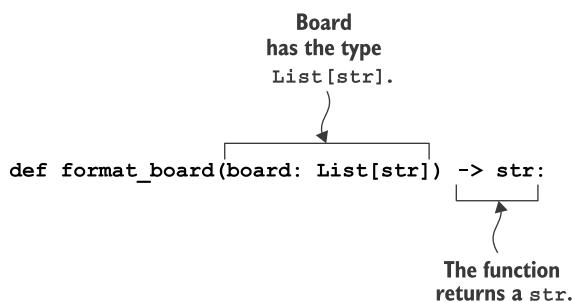


Figure 22.2 Adding type hints to describe the type of the parameter and the return value

The annotation for `main()` indicates that the `None` value is returned, as shown in figure 22.3.

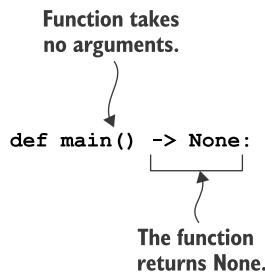


Figure 22.3 The `main()` function accepts no parameters and returns `None`.

What's really terrific is that we can define a function that takes a value of the type `State`, and Mypy will check that this kind of value is actually being passed (see figure 22.4).

Try playing my version of the game and then writing your own that behaves similarly. Then take a look at how I wrote an interactive solution that incorporates these ideas of data immutability and type safety.

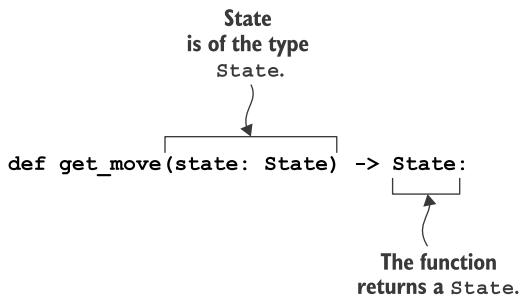


Figure 22.4 We can use custom types in type hints. This function takes and returns a value of the type `State`.

22.2 Solution

This is the last program! I hope that writing the simpler version in the previous chapter gave you ideas for making this work. Did the type hints and unit tests also help?

```
#!/usr/bin/env python3
""" Interactive Tic-Tac-Toe using NamedTuple """
from typing import List, NamedTuple, Optional
```

Import the classes we'll need from the `typing` module.

```
class State(NamedTuple):
    board: List[str] = list('.' * 9)
    player: str = 'X'
    quit: bool = False
    draw: bool = False
    error: Optional[str] = None
    winner: Optional[str] = None
```

Declare a class that is based on the `NamedTuple` class. Define field names, types, and defaults for the values this class can hold.

Start an infinite loop. When we have a reason to stop, we can break out of the loop.

Print a special sequence that most terminals will interpret as a command to clear the screen.

```
# -----
def main() -> None:
    """Make a jazz noise here"""

    state = State()
```

Instantiate the initial state as an empty grid and the first player as X.

Print the current state of the board.

```
while True:
    print("\u033[H\u033[J")
    print(format_board(state.board))
```

Print any errors, such as the user not choosing a valid cell.

```
    if state.error:
        print(state.error)
```

If there is a winner, proclaim the victor and break out of the loop.

```
    elif state.winner:
        print(f'{state.winner} has won!')
        break
```

Get the next move from the player. The `get_move()` function accepts a `State` type and returns one too. We overwrite the existing `state` variable each time through the loop.

```
    state = get_move(state)
```

```

if state.quit:
    print('You lose, loser!')
    break
elif state.draw:
    print("All right, we'll call it a draw.")
    break

```

If the user has decided to withdraw from the game prematurely, insult them, and break from the loop.

If so, replace the quit value of the state with True and return with the new state. Note that no other values in the state are modified.

Copy the player from the state, since we'll refer to it several times in the function body.

```

# -----
def get_move(state: State) -> State:
    """Get the player's move"""

    player = state.player
    cell = input(f'Player {player}, what is your move? [q to quit]: ')

```

Define a `get_move()` function that takes and returns a `State` type.

Check if the user entered a value that can be converted to a digit using `str.isdigit()` and if the integer version of the value is in the valid range.

Use the `input()` function to ask the player for their next move. Tell them how to quit the game early so they don't have to use `Ctrl-C` to interrupt the program.

```

if cell == 'q':
    return state._replace(quit=True)

```

First check if the user wants to quit.

```

if not (cell.isdigit() and int(cell) in range(1, 10)):
    return state._replace(error=f'Invalid cell "{cell}", please use 1-9')

```

See if the board is open at the indicated cell.

```

cell_num = int(cell)
if state.board[cell_num - 1] in 'XO':
    return state._replace(error=f'Cell "{cell}" already taken')

```

After we have verified that cell is a valid integer value, convert it to an integer.

If not, return an updated state that has an error. Note that the current state and player remain unchanged so that the same player has a retry with the same board until they provide valid input.

board = state.board
board[cell_num - 1] = player
return state._replace(board=board,
 player='O' if player == 'X' else 'X',
 winner=find_winner(board),
 draw='.' not in board,
 error=None)

Return a new state value with the new board value, the current player switched to the other player, and if there is a winner or a draw.

Use the cell value to update the board with the current player.

Copy the current board because we need to modify it and `state.board` is immutable.

If not, return an updated state with an error. Again, nothing else about the state is changed, so we retry the round with the same player and state.

```
# -----
def format_board(board: List[str]) -> str:
    """Format the board"""

    cells = [str(i) if c == '.' else c for i, c in enumerate(board, 1)]
    bar = '-----'
    cells_tmpl = '| {} | {} | {} |'
    return '\n'.join([
        bar,
        cells_tmpl.format(*cells[:3]), bar,
        cells_tmpl.format(*cells[3:6]), bar,
        cells_tmpl.format(*cells[6:]), bar
    ])
```

The only change from the previous version of this function is the addition of type hints. The function accepts a list of string values (the current board) and returns a formatted grid of the board state.

```
# -----
def find_winner(board: List[str]) -> Optional[str]:
    """Return the winner"""

    winning = [[0, 1, 2], [3, 4, 5], [6, 7, 8], [0, 3, 6], [1, 4, 7],
               [2, 5, 8], [0, 4, 8], [2, 4, 6]]

    for player in ['X', 'O']:
        for i, j, k in winning:
            combo = [board[i], board[j], board[k]]
            if combo == [player, player, player]:
                return player

    return None
```

This is also the same function as before, but with type hints. The function accepts the board as a list of strings and returns an optional string value, which means it could also return None.

```
# -----
if __name__ == '__main__':
    main()
```

22.2.1 A version using TypedDict

New to Python 3.8 is the `TypedDict` class, which looks very similar to a `NamedTuple`. Let's look at how using this as the base class changes parts of our program. One crucial difference is that you cannot (yet) set default values for the fields:

```
#!/usr/bin/env python3
""" Interactive Tic-Tac-Toe using TypedDict """

from typing import List, Optional, TypedDict
```

Import `TypedDict` instead of `NamedTuple`.

```
class State(TypedDict):
    board: str
    player: str
    quit: bool
    draw: bool
    error: Optional[str]
    winner: Optional[str]
```

Base State on a `TypedDict`.

We have to set our initial values when we instantiate a new state:

```
def main() -> None:
    """Make a jazz noise here"""

    state = State(board='.' * 9,
                  player='X',
                  quit=False,
                  draw=False,
                  error=None,
                  winner=None)
```

Syntactically, I prefer using `state.board` with the named tuple rather than the dictionary access of `state['board']`:

```
while True:
    print("\033[H\033[J")
    print(format_board(state['board']))

    if state['error']:
        print(state['error'])
    elif state['winner']:
        print(f"{state['winner']} has won!")
        break

    state = get_move(state)

    if state['quit']:
        print('You lose, loser!')
        break
    elif state['draw']:
        print('No winner.')
        break
```

Beyond the convenience of accessing the fields, I prefer the read-only nature of the `NamedTuple` to the mutable `TypedDict`. Note how in the `get_move()` function, we can change the state:

```
def get_move(state: State) -> State:
    """Get the player's move"""

    player = state['player']
    cell = input(f'Player {player}, what is your move? [q to quit]: ')

    if cell == 'q':
        state['quit'] = True
        return state
    else:
        state['error'] = f'Invalid cell "{cell}", please use 1-9'
        return state

    if state['board'][cell_num - 1] in 'XO':
        state['error'] = f'Cell {cell} is already occupied'
```

Here we are directly modifying the `TypedDict`, whereas the `NamedTuple` version used `state._replace()` to return an entirely new state value.

Another place where the state is directly modifiable. You may prefer this approach.

```

state['error'] = f'Cell "{cell}" already taken'
return state

board = list(state['board'])
board[cell_num - 1] = player

return State(
    board=''.join(board),
    player='O' if player == 'X' else 'X',
    winner=find_winner(board),
    draw='.' not in board,
    error=None,
    quit=False,
)

```

In my opinion, a NamedTuple has nicer syntax, default values, and immutability over the TypedDict version, so I prefer it. Regardless of which you choose, the greater lesson I hope to impart is that we should try to be explicit about the “state” of the program and when and how it changes.

22.2.2 Thinking about state

The idea of program state is that a program can remember changes to variables over time. In the previous chapter, our program accepted a given --board and possible values for --cell and --player that might alter the board. Then the game printed a representation of the board. In this chapter’s interactive version, the board always begins as an empty grid and changes with each turn, which we modeled as an infinite loop.

It is common in programs like this to see programmers use *global variables* that are declared at the top of the program outside of any function definitions so that they are *globally* visible throughout the program. While common, it’s not considered a best practice, and I would discourage you from ever using globals unless you can see no other way. I would suggest, instead, that you stick to using small functions that accept all the values required and return a single type of value. I would also suggest you use data structures like typed, named tuples to represent program state, and that you guard the changes to state very carefully.



22.3 Going further

- Incorporate spicier insults. Maybe bring in the Shakespearean generator?
- Write a version that allows the user to start a new game without quitting and restarting the program.
- Write other games like Hangman.

Summary

- Type hints allow you to annotate variables as well as function parameters and return values with the types of the values.
- Python itself will ignore type hints at runtime, but Mypy can use type hints to find errors in your code before you ever run it.
- A `NamedTuple` behaves a bit like a dictionary and a bit like an object but retains the immutability of tuples.
- Both `NamedTuple` and `TypedDict` allow you to create a novel type with defined fields and types that you can use as type hints to your own functions.
- Our program used a `NamedTuple` to create a complex data structure to represent the state of our program. The state included many variables, such as the current board, the current player, any errors, the winner, and so on, each of which was described using type hints.
- While it is difficult to write integration tests for an interactive program, we can still break a program into small functions (such as `format_board()` or `get_winner()`) for which we write and run unit tests.

Epilogue

Well, that's the whole book. We came a long way, from writing the crow's nest program in chapter 2 to chapter 22's interactive Tic-Tac-Toe game, incorporating a custom class based on named tuples and using type hints. I hope you can see now how much you can do with Python's strings, lists, tuples, dictionaries, sets, and functions. I especially hope I've convinced you that, above all, you should always write programs that are

- *Flexible*, by taking command-line arguments
- *Documented*, by using something like `argparse` to parse your arguments and produce usage statements
- *Tested*, by writing both *unit* tests for your functions and *integration* tests for your program as a whole

The people using your programs will really appreciate knowing how to use your program and how to make it behave differently. They'll also appreciate that you took the time to verify that your program is correct. Let's be honest, though. The person most likely to be using and modifying your programs will be you, several months from now. I've heard it said that "documentation is a love letter to your future self." All this work you put into making your programs good will be very appreciated by you when you come back to your code.

Now that you've worked through all the exercises and seen how to use the tests I've written, I challenge you to go back to the beginning and read the `test.py` programs. If you intend to adopt test-driven development, you may find that you can steal many ideas and techniques from those programs.

Further, each chapter included suggestions for how to extend the ideas and exercises presented. Go back and think about how you can use ideas you learned later in the book to improve or extend earlier programs. Here are some ideas:

- Chapter 2 (The crow’s nest)—Add an option to randomly select a greeting other than “Hello” from a list like “Hello,” “Hola,” “Salut,” and “Ciao.”
- Chapter 3 (Going on a picnic)—Allow the program to take one or more options and incorporate those into the output with the correct articles for each item joined on the Oxford comma.
- Chapter 7 (Gashlycrumb)—Download *The Devil’s Dictionary* by Ambrose Bierce from Project Gutenberg. Write a program that will look up a word’s definition if it appears in the text.
- Chapter 16 (The scrambler)—Use the scrambled text as the basis for encrypting messages. Force the scrambled words to uppercase, remove all the punctuation and spaces, and then format the text into “words” of five characters followed by a space, with no more than five per line. Pad the end so that the text completely fills the last line. Can you make sense of the output?
- new.py—I first wrote a program to create a new program when I was a green-horn Perl hacker. My new-pl program would add a random quote from the poetry of William Blake (yes, really—I also went through phases with the Brontës and Dickinson). Alter your version of new.py to add a random quote or joke or to customize it in some way for your programs.

I hope you’ve had as much fun writing the programs as I’ve had creating and teaching them. I want you to feel you now have dozens of programs and tests with ideas and functions you can steal to create even more programs.

All the best to you in your coding adventures!



appendix

Using argparse

Often, getting the right data into your program is a real chore. The `argparse` module makes it much easier to validate arguments from users and to generate useful error messages when they provide bad input. It's like your program's "bouncer," only allowing the right kinds of values into the program. Defining the arguments properly with `argparse` is the crucial first step to making the programs in this book work.

For instance, chapter 1 discusses a very flexible program that can extend warm salutations to an optionally named entity, such as the "World" or "Universe":

```
$ ./hello.py  
Hello, World!  
$ ./hello.py --name Universe  
Hello, Universe!
```

When the program runs with no input values,
it will use "World" for the entity to greet.

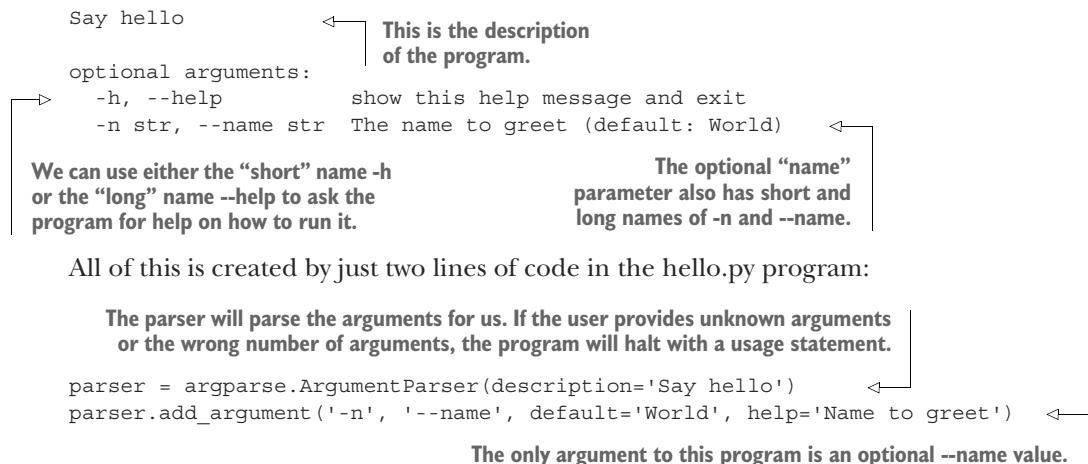
The program can take an optional
--name value to override the default.

The program will respond to the `-h` and `--help` flags with helpful documentation:

```
$ ./hello.py -h  
usage: hello.py [-h] [-n str]
```

The argument to the program is `-h`,
which is the "short" flag to ask for help.

This line shows a summary of all the options the
program accepts. The square brackets [] around
the arguments show that they are optional.



NOTE You do not need to define the -h or --help flags. Those are generated automatically by argparse. In fact, you should never try to use those for other values because they are almost universal options that most users will expect.

The argparse module helps us define a parser for the arguments and generates help messages, saving us loads of time and making our programs look professional. Every program in this book is tested on different inputs, so you'll really understand how to use this module by the end. I recommend you look over the argparse documentation (<https://docs.python.org/3/library/argparse.html>).

Now let's dig further into what this module can do for us. In this appendix, you will

- Learn how to use argparse to handle positional parameters, options, and flags
- Set default values for options
- Use type to force the user to provide values like numbers or files
- Use choices to restrict the values for an option

A.1 Types of arguments

Command-line arguments can be classified as follows:

- *Positional arguments*—The order and number of the arguments is what determines their meaning. Some programs might expect, for instance, a filename as the first argument and an output directory as the second. Positional arguments are generally required (not optional) arguments. Making them optional is difficult—how would you write a program that accepts two or three arguments where the second and third ones are independent and optional? In the first version of hello.py in chapter 1, the name to greet was provided as a positional argument.
- *Named options*—Most command-line programs define a *short* name like -n (one dash and a single character) and a *long* name like --name (two dashes and a word) followed by some value, like the name in the hello.py program. Named

options allow arguments to be provided in any order—their *position* is not relevant. This makes them the right choice when the user is not required to provide them (they are *options*, after all). It's good to provide reasonable default values for options. When we changed the required positional name argument of hello.py to the optional --name argument, we used “World” for the default so that the program could run with no input from the user. Note that some other languages, like Java, might define long names with a single dash, like -jar.

- **Flags**—A Boolean value like “yes”/“no” or True/False is indicated by something that starts off looking like a named option, but there is no value after the name; for example, the -d or --debug flag to turn on debugging. Typically the presence of the flag indicates a True value for the argument, and its absence would mean False, so --debug turns *on* debugging, whereas its absence means it is off.

A.2 Using a template to start a program

It's not easy to remember all the syntax for defining parameters using argparse, so I've created a way for you to write new programs from a template that includes this plus some other structure that will make your programs easier to read and run.

One way to start a new program is to use the new.py program. From the top level of the repository, you can execute this command:

```
$ bin/new.py foo.py
```

Alternatively, you could copy the template:

```
$ cp template/template.py foo.py
```

The resulting program will be identical no matter how you create it, and it will have examples of how to declare each of the argument types outlined in the previous section. Additionally, you can use argparse to validate the input, such as making sure that one argument is a number while another argument is a file.

Let's look at the help generated by our new program:

<p>Optional arguments can be left out, so you should provide reasonable default values for them.</p> <p>This is the description of the entire program.</p> <pre>\$./foo.py -h usage: foo.py [-h] [-a str] [-i int] [-f FILE] [-o] str</pre> <p>→ Rock the Casbah</p> <p>positional arguments:</p> <p>str</p> <p>optional arguments:</p> <p>-h, --help</p>	<p>The -h and --help arguments are always present when you use argparse; you do not need to define them.</p> <p>Every program should respond to -h and --help with a help message.</p> <p>This program defines one positional parameter, but you could have many more. You'll see how to define those shortly.</p> <p>A positional argument</p>	<p>This a brief summary of the options that are described in greater detail below.</p> <p>show this help message and exit</p>
---	--	---

```

→ -a str, --arg str      A named string argument (default: )
→ -i int, --int int     A named integer argument (default: 0)
-f FILE, --file FILE   A readable file (default: None)
-o, --on                A boolean flag (default: False)

```

The `-i` or `--int` option must be an integer value. If the user provides “one” or “4.2,” these will be rejected.

The `-a` or `--arg` option accepts some text, which is often called a “string.”

The `-o` or `--on` is a flag. Notice how the `-f` FILE description specifies that a “FILE” value should follow the `-f`, but for this flag no value follows the option. The flag is either present or absent, and so it’s either True or False, respectively.

The `-f` or `--file` option must be a valid, readable file.

A.3 Using argparse

The code to generate the preceding usage is found in a function, called `get_args()`, that looks like the following:

```

def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Rock the Casbah',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('positional',
                        metavar='str',
                        help='A positional argument')

    parser.add_argument('-a',
                        '--arg',
                        help='A named string argument',
                        metavar='str',
                        type=str,
                        default='')

    parser.add_argument('-i',
                        '--int',
                        help='A named integer argument',
                        metavar='int',
                        type=int,
                        default=0)

    parser.add_argument('-f',
                        '--file',
                        help='A readable file',
                        metavar='FILE',
                        type=argparse.FileType('r'),
                        default=None)

    parser.add_argument('-o',
                        '--on',
                        help='A boolean flag',
                        action='store_true')

    return parser.parse_args()

```

You are welcome to put this code wherever you like, but defining and validating the arguments can sometimes get rather long. I like to separate this code out into a function I call `get_args()`, and I always define this function first in my program. That way I can see it immediately when I'm reading the source code.

The `get_args()` function is defined like this:

```
def get_args():    ←
    """Get command-line arguments"""\n\n← The def keyword defines a new function, and the arguments to the
← function are listed in the parentheses. Even though the get_args()
← function takes no arguments, the parentheses are still required.
```

The triple-quoted line after the function `def` is the “docstring,” which serves as a bit of documentation for the function. Docstrings are not required, but they are good style, and Pylint will complain if you leave them out.

A.3.1 Creating the parser

The following snippet creates a parser that will deal with the arguments from the command line. To “parse” here means to derive some meaning from the order and syntax of the bits of text provided as arguments:

```
Call the argparse.ArgumentParser()
function to create a new parser.\n\nparser = argparse.ArgumentParser(←
    description='Argparse Python script', ←
    formatter_class=argparse.ArgumentDefaultsHelpFormatter) ←\n\n← A short summary of your
← program's purpose.\n\n← The formatter_class argument tells argparse
← to show the default values in usage.
```

You should read the documentation for `argparse` to see all the other options you can use to define a parser or the parameters. In the REPL, you can start with `help(argparse)`, or you could look up the docs on the internet at <https://docs.python.org/3/library/argparse.html>.

A.3.2 Creating a positional parameter

The following line will create a new *positional* parameter:

```
The lack of leading dashes makes this a
positional parameter, not the name "positional." ←\n\nparser.add_argument('positional', ←
    metavar='str', ←
    help='A positional argument') ←\n\n← Provide a hint to the user about
← the data type. By default, all
← arguments are strings.\n\n← A brief description of the parameter for the usage
```

Remember that the parameter is not positional because the *name* is “positional.” That’s just there to remind you that it *is* a positional parameter. `argparse` interprets the string ‘`positional`’ as a positional parameter *because the name does not start with any dashes*.

A.3.3 Creating an optional string parameter

The following line creates an *optional* parameter with a short name of `-a` and a long name of `--arg`. It will be a `str` with a default value of '' (the empty string).

```
parser.add_argument('-a',           ← The short name
                   '--arg',        ← The long name
                   help='A named string argument',   ← Brief description
                                           metavar='str',          for the usage
                                           type=str,             ← The actual Python data type
                                           default='',           (note the lack of quotes
                                                               around str)
```

Type hint for usage The short name
 The default value The long name
 The actual Python data type
 (note the lack of quotes
 around str)

NOTE You can leave out either the short or long name in your own programs, but it's good form to provide both. Most of the tests in this book will test your programs using both short and long option names.

If you wanted to make this a required, named parameter, you would remove the `default` and add `required=True`.

A.3.4 Creating an optional numeric parameter

The following line creates an option called `-i` or `--int` that accepts an `int` (integer) with a default value of 0. If the user provides anything that cannot be interpreted as an integer, the `argparse` module will stop processing the arguments and will print an error message and a short usage statement.

```
parser.add_argument('-i',           ← The short name
                   '--int',        ← The long name
                   help='A named integer argument',   ← A brief description for
                                           metavar='int',          the usage statement
                                           type=int,             ← A Python data type that the string must be converted
                                           default=0)           to. You can also use float for a floating point value
                                                               (a number with a fractional component like 3.14).
```

A type hint for the usage statement The short name
 The default value The long name
 A Python data type that the string must be converted to. You can also use float for a floating point value
 (a number with a fractional component like 3.14).

One of the big reasons to define numeric arguments in this way is that `argparse` will convert the input to the correct type. All values coming from the command are strings, and it's the job of the program to convert each value to an actual numeric value. If you tell `argparse` that the option should be `type=int`, it will have already been converted to an actual `int` value when you ask the parser for the value.

If the value provided by the user cannot be converted to an `int`, the value will be rejected. Note that you can also use `type=float` to accept and convert the input to a floating-point value. That saves you a lot of time and effort.

A.3.5 Creating an optional file parameter

The following line creates an option called `-f` or `--file` that will only accept a valid, readable file. This argument alone is worth the price of admission, as it will save you oodles of time validating the input from your user. Note that pretty much every exercise

that has a file as input will have tests that pass *invalid* file arguments to ensure that your program rejects them.

```

parser.add_argument('-f',      # The short name
                   '--file',   # The long name
                   help='A readable file',    # A brief usage statement
                   metavar='FILE',
                   type=argparse.FileType('r'),
                   default=None)           # Says that the argument must name a readable (r) file
  
```

The diagram shows the `parser.add_argument` call with several annotations:

- A type suggestion**: Points to the `'FILE'` in `metavar='FILE'`.
- The default value**: Points to `default=None`.
- The short name**: Points to `'-f'`.
- The long name**: Points to `--file`.
- A brief usage statement**: Points to `help='A readable file'`.
- Says that the argument must name a readable (r) file**: Points to `type=argparse.FileType('r')`.

The person running the program is responsible for providing the location of the file. For instance, if you created the `foo.py` program in the top level of the repository, there will be a `README.md` file there. We could use that as the input to our program, and it would be accepted as a valid argument:

```

$ ./foo.py -f README.md foo
str_arg = ""
int_arg = "0"
file_arg = "README.md"
flag_arg = "False"
positional = "foo"
  
```

If we provide a bogus `--file` argument, like “blaragh,” we will get an error message:

```

$ ./foo.py -f blaragh foo
usage: foo.py [-h] [-a str] [-i int] [-f FILE] [-o] str
foo.py: error: argument -f/--file: can't open 'blaragh': \
[Errno 2] No such file or directory: 'blaragh'
  
```

A.3.6 Creating a flag option

The flag option is slightly different in that it does not take a value like a string or integer. Flags are either present or not, and they *usually* indicate that some idea is True or False.

You’ve already seen the `-h` and `--help` flags. They are not followed by any values. They either are present, in which case the program should print a “usage” statement, or they are absent, in which case the program should not. For all the exercises in this book, I use flags to indicate a `True` value when they are present and `False` otherwise, which we can represent using `action='store_true'`.

For instance, `new.py` shows an example of this kind of a flag called `-o` or `--on`:

```

parser.add_argument('-o',      # Short name
                   '--on',    # Long name
                   help='A boolean flag',    # Brief usage statement
                   action='store_true')
  
```

The diagram shows the `parser.add_argument` call with several annotations:

- Short name**: Points to `'-o'`.
- Long name**: Points to `--on`.
- Brief usage statement**: Points to `help='A boolean flag'`.
- What to do when this flag is present. When it is present, we use the value True for on. The default value will be False when the flag is not present.**: Points to `action='store_true'`.

What to do when this flag is present. When it is present, we use the value True for on. The default value will be False when the flag is not present.

It's not always the case that a "flag" like this should be interpreted as `True` when present. You could instead use `action='store_false'`, in which case `on` would be `False` when the flag is present, and the default value would be `True`. You could also store one or more constant values when the flag is present.

Read the `argparse` documentation for the various ways you can define this parameter. For the purposes of this book, we will only use a flag to turn "on" some behavior.

A.3.7 *Returning from get_args*

The final statement in `get_args()` is `return`, which returns the result of having the `parser` object parse the arguments. That is, the code that calls `get_args()` will receive the result of this expression:

```
return parser.parse_args()
```

This expression could fail because `argparse` finds that the user provided invalid arguments, such as a string value when it expected a `float` or perhaps a misspelled file-name. If the parsing succeeds, we will be able to access all the values the user provided from inside our program.

Additionally, the values of the arguments will be of the *types* that we indicated. That is, if we indicated that the `--int` argument should be an `int`, then when we ask for `args.int`, it will already be an `int`. If we define a file argument, we'll get an *open file handle*. That may not seem impressive now, but it's really enormously helpful.

If you refer to the `foo.py` program we generated, you'll see that the `main()` function calls `get_args()`, so the `return` from `get_args()` goes back to `main()`. From there, we can access all the values we just defined using the names of the positional parameters or the long names of the optional parameters:

```
def main():
    args = get_args()
    str_arg = args.arg
    int_arg = args.int
    file_arg = args.file
    flag_arg = args.on
    pos_arg = args.positional
```

A.4 *Examples using argparse*

Many of the program tests in this book can be satisfied by learning how to use `argparse` effectively to validate the arguments to your programs. I think of the command line as the boundary of your program, and you need to be judicious about what you let into your program. You should always expect and defend against every argument being wrong.¹ Our `hello.py` program in chapter 1 is an example of a single, positional argument and then a single, optional argument. Let's look at some more examples of how you can use `argparse`.

¹ I always think of the kid who will type "fart" for every input.

A.4.1 A single positional argument

This is the first version of chapter 1’s hello.py program, which requires a single argument specifying the name to greet:

```
#!/usr/bin/env python3
"""A single positional argument"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='A single positional argument',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('name', metavar='name', help='The name to greet') ←

    return parser.parse_args()

# -----
def main():
    """Make a jazz noise here"""

    args = get_args()
    print('Hello, ' + args.name + '!')

# -----
if __name__ == '__main__':
    main()
```

The name parameter does not start with dashes, so this is a **positional** parameter. The metavar will show up in the help to let the user know what this argument is supposed to be.

Whatever is provided as the first positional argument to the program will be available in the args.name slot.

This program will not print the “Hello” line if it’s not provided exactly one argument. If given nothing, it will print a brief usage statement about the proper way to invoke the program:

```
$ ./one_arg.py
usage: one_arg.py [-h] name
one_arg.py: error: the following arguments are required: name
```

If we provide more than one argument, it complains again. Here “Emily” and “Bronte” are two arguments because spaces separate arguments on the command line. The program complains about getting a second argument that has not been defined:

```
$ ./one_arg.py Emily Bronte
usage: one_arg.py [-h] name
one_arg.py: error: unrecognized arguments: Bronte
```

Only when we give the program exactly one argument will it run:

```
$ ./one_arg.py "Emily Bronte"
Hello, Emily Bronte!
```

While it may seem like overkill to use argparse for such a simple program, it shows that argparse can do quite a bit of error checking and validation of arguments for us.

A.4.2 Two different positional arguments

Imagine you want two *different* positional arguments, like the *color* and *size* of an item to order. The color should be a str, and the size should be an int value. When you define them positionally, the order in which you declare them is the order in which the user must supply the arguments. Here we define *color* first, and then *size*:

```
#!/usr/bin/env python3
"""Two positional arguments"""

import argparse

# -----
def get_args():
    """Get args"""

    parser = argparse.ArgumentParser(
        description='Two positional arguments',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('color',
                        metavar='color',
                        type=str,
                        help='The color of the garment')

    parser.add_argument('size',
                        metavar='size',
                        type=int,
                        help='The size of the garment')

    return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    print('color =', args.color)
    print('size =', args.size)

if __name__ == '__main__':
    main()
```

This will be the first of the positional arguments because it is defined first. Notice that metavar has been set to 'color' instead of 'str' as it's more descriptive of the kind of string we expect—one that describes the “color” of the garment.

This will be the second of the positional arguments. Here metavar='size', which could be a number like 4 or a string like 'small', so it's still ambiguous.

The “color” argument is accessed via the name of the color parameter.

The “size” argument is accessed via the name of the size parameter.

Again, the user must provide exactly two positional arguments. Entering no arguments triggers a short usage statement:

```
$ ./two_args.py
usage: two_args.py [-h] color size
two_args.py: error: the following arguments are required: color, size
```

Just entering one argument won't cut it either. We are told that "size" is missing:

```
$ ./two_args.py blue
usage: two_args.py [-h] color size
two_args.py: error: the following arguments are required: size
```

If we give it two strings, like "blue" for the color and "small" for the size value will be rejected because it needs to be an integer value:

```
$ ./two_args.py blue small
usage: two_args.py [-h] color size
two_args.py: error: argument size: invalid int value: 'small'
```

If we give it two arguments, the second of which can be interpreted as an int, all is well:

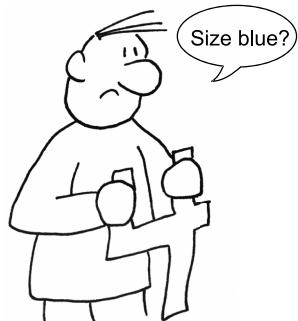
```
$ ./two_args.py blue 4
color = blue
size = 4
```

Remember that *all* the arguments coming from the command line are strings. The command line doesn't require quotes around blue or the 4 to make them strings the way that Python does. On the command line, everything is a string, and all arguments are passed to Python as strings.



When we tell argparse that the second argument needs to be an int, argparse will attempt to convert the string '4' to the integer 4. If you provide 4.1, that will be rejected too:

```
$ ./two_args.py blue 4.1
usage: two_args.py [-h] str int
two_args.py: error: argument int: invalid int value: '4.1'
```



Positional arguments require the user to remember the correct order of the arguments. If we mistakenly switch around str and int arguments, argparse will detect invalid values:

```
$ ./two_args.py 4 blue
usage: two_args.py [-h] COLOR SIZE
two_args.py: error: argument SIZE: invalid int
value: 'blue'
```

Imagine, however, a case of two strings or two numbers that represent two *different* values, like a car’s make and model or a person’s height and weight. How could you detect that the arguments are reversed?

Generally speaking, I only ever create programs that take exactly one positional argument or one or more of *the same thing*, like a list of files to process.

A.4.3 **Restricting values using the choices option**

In our previous example, there was nothing stopping the user from providing *two integer values*:

```
$ ./two_args.py 1 2
color = 1
size = 2
```

The 1 is a string. It may look like a number to you, but it is actually the *character '1'*. That is a valid string value, so our program accepts it.

Our program would also accept a “size” of -4, which clearly is not a valid size:

```
$ ./two_args.py blue -4
color = blue
size = -4
```

How can we ensure that the user provides both a valid color and size? Let’s say we only offer shirts in primary colors. We can pass in a list of valid values using the *choices* option.

In the following example, we restrict the color to “red,” “yellow,” or “blue.” Additionally, we can use `range(1, 11)` to generate a list of numbers from 1 to 10 (11 isn’t included!) as the valid sizes for our shirts:

```
#!/usr/bin/env python3
"""Choices"""

import argparse

# -----
def get_args():
    """get args"""

    parser = argparse.ArgumentParser(
        description='Choices',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('color',
                        metavar='str',
                        help='Color',
                        choices=['red', 'yellow', 'blue']) ←

    parser.add_argument('size',
                        metavar='size',
```

The choices option takes a list of values. argparse stops the program if the user fails to supply one of these.

```

        type=int,
        choices=range(1, 11),
        help='The size of the garment') ←
    return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    print('color =', args.color)
    print('size =', args.size) ←

# -----
if __name__ == '__main__':
    main()

```

The user must choose from the numbers 1–10 or argparse will stop with an error.

If our program makes it to this point, we know that args.color will definitely be one of those values and that args.size is an integer value in the range of 1–10. The program will never get to this point unless both arguments are valid.

Any value not present in the list will be rejected, and the user will be shown the valid choices. Again, no value is rejected:

```

$ ./choices.py
usage: choices.py [-h] color size
choices.py: error: the following arguments are required: color, size

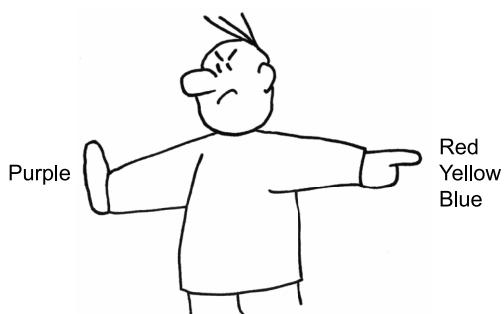
```

If we provide “purple,” it will be rejected because it is not in the choices we defined. The error message that argparse produces tells the user the problem (“invalid choice”) and even lists the acceptable colors:

```

$ ./choices.py purple 1
usage: choices.py [-h] color size
choices.py: error: argument color: \
invalid choice: 'purple' (choose from 'red', 'yellow', 'blue')

```



Likewise with a negative size argument:

```

$ ./choices.py red -1
usage: choices.py [-h] color size

```

```
choices.py: error: argument size: \
invalid choice: -1 (choose from 1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

Only when both arguments are valid may we continue:

```
$ ./choices.py red 4
color = red
size = 4
```

That's really quite a bit of error checking and feedback that you never have to write. The best code is code you don't write!

A.4.4 Two of the same positional arguments

If we were writing a program that adds two numbers, we could define them as two positional arguments, like `number1` and `number2`. But since they are the same kinds of arguments (two numbers that we will add), it might make more sense to use the `nargs` option to tell `argparse` that you want exactly two of a thing:

```
#!/usr/bin/env python3
"""nargs=2"""

import argparse

# -----
def get_args():
    """get args"""

    parser = argparse.ArgumentParser(
        description='nargs=2',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('numbers',
                        metavar='int',
                        nargs=2,
                        type=int,
                        help='Numbers')
    return parser.parse_args()

# -----
def main():
    """main"""

    args = get_args()
    n1, n2 = args.numbers
    print(f'{n1} + {n2} = {n1 + n2}')  

# -----
if __name__ == '__main__':
    main()
```

The `nargs=2` will require exactly two values.

Each value must be parsable as an integer value, or the program will error out.

Since we defined that there are exactly two values for `numbers`, we can copy them into two variables.

Because these are actual int values, the result of + will be numeric addition and not string concatenation.

The help indicates we want two numbers:

```
$ ./nargs2.py
usage: nargs2.py [-h] int int
nargs2.py: error: the following arguments are required: int
```

When we provide two good integer values, we get their sum:

```
$ ./nargs2.py 3 5
3 + 5 = 8
```

Notice that argparse converts the n1 and n2 values to actual integer values. If you change the type=int to type=str, you'll see that the program will print 35 instead of 8 because the + operator in Python both adds numbers and concatenates strings!

```
>>> 3 + 5
8
>>> '3' + '5'
'35'
```



A.4.5 One or more of the same positional arguments

You could expand your two-number adding program into one that sums as many numbers as you provide. When you want *one or more* of some argument, you can use nargs='+' :

```
#!/usr/bin/env python3
"""nargs=+"""

import argparse

# -----
def get_args():
    """get args"""

    parser = argparse.ArgumentParser(
        description='nargs=+',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('numbers',
                        metavar='int',
                        nargs='+',
                        type=int,
                        help='Numbers')

    return parser.parse_args()

# -----
def main():
    """main"""


The + will make nargs accept
one or more values.

The int means that all the
values must be integer
values.
```

```
# -----
def main():
    """main"""
```

```

args = get_args()
numbers = args.numbers
print('{} = {}'.format(' + '.join(map(str, numbers)), sum(numbers)))
# -----
if __name__ == '__main__':
    main()

```

numbers will be a list with at least one element.

Don't worry if you don't understand this line. You will by the end of the book.

Note that this will mean `args.numbers` is always a list. Even if the user provides just one argument, `args.numbers` will be a list containing that one value:

```

$ ./nargs+.py 5
5 = 5
$ ./nargs+.py 1 2 3 4
1 + 2 + 3 + 4 = 10

```

You can also use `nargs='*'` to indicate *zero or more* of an argument, and `nargs='?'` means *zero or one* of the argument.

A.4.6 File arguments

So far you've seen how you can specify that an argument should be of a type like `str` (which is the default), `int`, or `float`. There are also many exercises that require a file as input, and for that you can use the type of `argparse.FileType('r')` to indicate that the argument must be a *file* that is *readable* (the '`r`' part).

If, additionally, you want to require that the file be *text* (as opposed to a *binary* file), you would add a '`t`'. These options will make more sense after you've read chapter 5.

Here is an implementation in Python of the command `cat -n`, where `cat` will *concatenate* a readable text file, and the `-n` says to *number* the lines of output:

```

#!/usr/bin/env python3
"""Python version of `cat -n`"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Python version of `cat -n`',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file',
                        metavar='FILE',
                        type=argparse.FileType('rt'), ←
                        help='Input file')

    return parser.parse_args()


```

The argument will be rejected if it does not name a valid, readable text file.

```
# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    for i, line in enumerate(args.file, start=1):
        print(f'{i:6} {line}', end='')

# -----
```

The value of args.file is an open file handle that we can directly read. Again, don't worry if you don't understand this code. We'll talk all about file handles in the chapters.

```
# -----
if __name__ == '__main__':
    main()
```

When we define an argument as type=int, we get back an actual int value. Here, we define the file argument as a FileType, so we receive an *open file handle*. If we had defined the file argument as a string, we would have to manually check if it were a file and then use open() to get a file handle:

```
#!/usr/bin/env python3
"""Python version of `cat -n`, manually checking file argument"""

import argparse
import os
```

```
# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Python version of `cat -n`',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('file', metavar='str', type=str, help='Input file')
    args = parser.parse_args()

    if not os.path.isfile(args.file):
        parser.error(f'{args.file} is not a file')

    args.file = open(args.file)
    return args
```

Intercept the arguments.

Check if the file argument is not a file.

Replace the file with an open file handle.

Print an error message and exit the program with a non-zero value.

```
# -----
def main():
    """Make a jazz noise here"""

    args = get_args()

    for i, line in enumerate(args.file, start=1):
        print(f'{i:6} {line}', end='')
```

```
# -----
if __name__ == '__main__':
    main()
```

With the `FileType` definition, you don't have to write any of this code.

You can also use `argparse.FileType('w')` to indicate that you want the name of a file that can be opened for *writing* (the '`w`'). You can pass additional arguments specifying how to open the file, like the encoding. See the documentation for more information.

A.4.7 Manually checking arguments

It's also possible to manually validate arguments before we `return` from `get_args()`. For instance, we can define that `--int` should be an `int`, but how can we require that it must be between 1 and 10?

One fairly simple way to do this is to manually check the value. If there is a problem, you can use the `parser.error()` function to halt execution of the program, print an error message along with the short usage statement, and then exit with an error value:

```
#!/usr/bin/env python3
"""Manually check an argument"""

import argparse

# -----
def get_args():
    """Get command-line arguments"""

    parser = argparse.ArgumentParser(
        description='Manually check an argument',
        formatter_class=argparse.ArgumentDefaultsHelpFormatter)

    parser.add_argument('-v',
                        '--val',
                        help='Integer value between 1 and 10',
                        metavar='int',
                        type=int,
                        default=5)

    args = parser.parse_args()
    if not 1 <= args.val <= 10:
        parser.error(f"--val '{args.val}' must be between 1 and 10")

    return args
```

Parse the arguments.

Check if the args.int value is not between 1 and 10.

If we get here, everything was OK, and the program will continue as normal.

Call `parser.error()` with an error message. The error message and the brief usage statement will be shown to the user, and the program will immediately exit with a non-zero value to indicate an error.

```

args = get_args()
print(f'val = "{args.val}"')

# -----
if __name__ == '__main__':
    main()

```

If we provide a good `--val`, all is well:

```
$ ./manual.py -v 7
val = "7"
```

If we run this program with a value like 20, we get an error message:

```
$ ./manual.py -v 20
usage: manual.py [-h] [-v int]
manual.py: error: --val "20" must be between 1 and 10
```

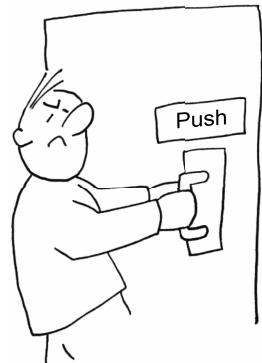
It's not possible to tell here, but the `parser.error()` also caused the program to exit with a non-zero status. In the command-line world, an exit status of 0 indicates “zero errors,” so anything not 0 is considered an error. You may not realize yet just how wonderful that is, but trust me. It is.

A.4.8 Automatic help

When you define a program’s parameters using `argparse`, the `-h` and `--help` flags will be reserved for generating help documentation. You do not need to add these, nor are you allowed to use these flags for other purposes.

I think of this documentation as being like a door into your program. Doors are how we get into buildings and cars and such. Have you ever come across a door that you can’t figure out how to open? Or one that requires a “PUSH” sign when clearly the handle is designed to “pull”? The book *The Design of Everyday Things* by Don Norman (Basic Books, 2013) uses the term *affordances* to describe the interfaces that objects present to us that do or do not inherently describe how we should use them.

The usage statement of your program is like the handle of the door. It should let users know exactly how to use it. When I encounter a program I’ve never used, I either run it with no arguments or with `-h` or `--help`. I *expect* to see some sort of usage statement. The only alternative would be to open the source code itself and study how to make the program run and how I can alter it, and this is a truly unacceptable way to write and distribute software!



When you start creating a new program with `new.py foo.py`, this is the help that will be generated:

```
$ ./foo.py -h
usage: foo.py [-h] [-a str] [-i int] [-f FILE] [-o] str

Rock the Casbah

positional arguments:
  str                  A positional argument

optional arguments:
  -h, --help            show this help message and exit
  -a str, --arg str     A named string argument (default: )
  -i int, --int int     A named integer argument (default: 0)
  -f FILE, --file FILE  A readable file (default: None)
  -o, --on              A boolean flag (default: False)
```

Without writing a single line of code, you have

- An executable Python program
- A variety of command-line arguments
- A standard and useful help message

This is the “handle” to your program, and you don’t have to write a single line of code to get it!

Summary

- Positional parameters typically are required parameters. If you have two or more positional parameters representing different ideas, it would be better to make them named options.
- Optional parameters can be named, like `--file fox.txt` where `fox.txt` is the value for the `--file` option. It is recommended that you always define a default value for options.
- `argparse` can enforce many argument types, including numbers like `int` and `float`, or even files.
- Flags like `--help` do not have an associated value. They are (usually) considered `True` if present and `False` if not.
- The `-h` and `--help` flags are reserved for use by `argparse`. If you use `argparse`, your program will automatically respond to these flags with a usage statement.

index

Symbols

\$ character 250, 361
: (colon) 113
" (single quotes) 42, 157
"" (double quotes) 42, 157
[] (square brackets) 43, 59, 78,
80, 89, 125, 293, 385
[<>] character class 285
{ } (curly brackets) 48, 52, 78–79,
113, 124–125, 293, 339
* (asterisk) character 108, 363
* (multiplication operator) 358
. (period) character 235, 284,
353, 361
^ (caret) character 254, 285,
297, 346
- (hyphen character) 40
+ operator 176
+= sign 88–89, 203
= (equal sign) 41, 46, 57
>>> prompt 41, 87
| (pipe operator) 108, 356

A

abuse.py (Dial-a-Curse
program) 150–164
constructing insults
162–163
defining adjectives and
nouns 155–156
defining arguments 159
exit values 160–161
final program 157–159
formatting output 156–157

importing and seeding ran-
dom module 154–155
parser.error() function 160
random.seed() function 161
range() function 162
STDERR (standard
error) 160–161
taking random samples and
choices 156
throwaway variables 162
validating arguments 153–154
writing program 151–157
add1() function 145, 189–190
affordances 403
AI (artificial intelligence) 351
algorithm design 207–224
counting 209–210
creating ordinal value 211–212
final program 216–217
generating verses 221–222
making verses 213–215,
218–221
printing verses 215, 222
verse() function 215
writing program 208–215
all() function 358, 360, 365
alpha characters 169
anonymous function 144, 241
any() function 358, 360
append() function 61
apples.py (Apples and Bananas
program) 128–149
altering strings 130–132
str.replace() method 131
str.translate() method
131–132

defining parameters 134–135
final program 133–134
refactoring with tests 149
replacing vowels 135–149
iterating through every
character 135–136
list comprehensions
140–142
list comprehensions with
functions 142–144
map() function 144–147
map() function with named
functions 147
regular expressions 148–149
str.replace() method 136
str.translate() method
137–139
argparse module 385–404
creating flag option 391–392
creating optional file
parameter 390–391
creating optional numeric
parameter 390
creating optional string
parameter 390
creating parser 389
creating positional
parameter 389
returning from get_args()
function 392
types of arguments 386–387
use cases 392–404
automatic help 403–404
file arguments 400–402
manually checking
arguments 402–403

argparse module (*continued*)
 one or more of the same positional arguments 399–400
 restricting values using choices option 396–398
 single positional argument 393–394
 two different positional arguments 394–396
 two of the same positional arguments 398–399
 using template.py to start programs 387–388
 argparse.ArgumentParser()
 function 389
 arguments
abuse.py 153–154, 159
crowsnest.py 39–40, 50–51
 file 400–402
gashlycrumb.py 123–124
hello.py 24–25
howler.py 102
 manually checking 402–403
picnic.py 73
 positional
 one or more of the same 399–400
 single 393–394
 two different 394–396
 two of the same 398–399
tictactoe.py 360–363
 types of 386–387
wc.py 115
 artificial intelligence (AI) 351
 ASCII values 295–310
 breaking text 304
 cleaning words 297–298
 encoding words 303
 final program 304–305
 functools.reduce()
 function 302–303
 ordinal character values and ranges 298–300
 sorting 308–309
 summing and reducing 300–301
 testing 309
 word2num() function 306–307
 writing program 296–304
 assert statements 60, 83
 asterisk (*) character 108, 363

B

binary decisions 191, 198
 Black tool 28–29
 Boolean values 202
bottles.py (Bottles of Beer program) 178–194
 counting down 180–181, 189
 final program 187–188
 iterating through verses 191–193
 other solutions 194
 test-driven development 189–190
verse() function 190–191
 using 186–187
 writing 181–182
 writing tests for 182–186
 writing program 179–187
 breaking text 270–271, 304

C

capture groups 232–236
 caret (^) character 254, 285, 297, 346
 chmod (change mode) command 20
 choices option 132, 134, 360, 386, 396–398
choose() function 197–199, 202, 204
chr() function 296, 298
 class, defined 42
clean() function 337–338, 348
 cleaning text 297–298, 337–339, 346–347
 colon (:) 113
 comma-separated values files.
See CSV (comma-separated values) files
 comment lines 16
 adding shebang line 18–19
 regular expressions 242–243
 conditional branching 47–48, 70–71
 consonants pattern 233
copy (cp) command 33
 counting down 180–181, 189
cp (copy) command 33
crowsnest.py (Crow’s Nest program) 35–54
 classifying first character of words 51–52
 concatenating strings 41–42
 conditional branching 47–48
 defining arguments 39–40, 50–51
 final program 49–50
 getting individual characters of strings 43–44
 main() function 51
 printing results 52
 REPL (Read-Evaluate-Print-Loop) 44
 starting with new.py 37–38
 string comparisons 45–47
 string formatting 48
 string methods 44–45
 test-driven development 52–53
 variable types 42
 writing and testing little by little 38–39
 writing program 49
crowsnest.py file 38
 CSV (comma-separated values) files 311–330
 final program 323–325
 formatting output 322
 formatting table 328–330
 handling bad data 322–323
 parsing
 with csv module 318–319
 with pandas.read_csv()
 function 327–328
 potential runtime errors 326–327
 reading 313–314, 325–326
 creating function to read 320–321
 manually 315–318
 selecting exercises 321–322
 writing program 312–323
 csv module 312, 315, 318–319
csv.DictReader() function 319–320, 323, 325
 csvkit module 313
 curly brackets ({}) 48, 52, 78–79, 113, 124–125, 293, 339

D

declarative programming 149
def (defines) 28, 181
 delimited text file 312
 deterministic approach 169
Dial-a-Curse program. *See abuse.py (Dial-a-Curse program)*

dict() function 78–79, 124
 dict.get() method 80, 84, 87, 126, 327
 dictionaries 77–82
 accessing values 80
 creating 78–79
 defining parameters 85
 dictionary
 comprehensions 125
 looking up items in 124–126
 methods 81–82
 processing items in series 86–90
 turning for loop into list comprehension 89–90
 using for loop to build new list 89
 using for loop to build new string 88
 using for loop to print() each character 86–88
 using str.translate() function 90
 using for encoding 85
 dict.items() method 81, 318
 dict.keys() method 81
 dict.pop() function 82
 dict.values() method 81
 domain-specific language (DSL) 148
 double quotes ("") 42, 157
 double-under methods 60
 DSL (domain-specific language) 148
 dynamically typed language 48

E

enumerate() function 352, 363–364
 env command 18–19
 equal sign (=) 46, 57
 error checking 287–288
 hello.py 28–29
 executability, making programs executable 20
 exit values 160–161
 extend() function 62

F

fh (file handle) 94–95
 fh.read() method 94–95, 103, 111

fh.readline() function 315–316
 fh.seek(0) method 95
 fh.write() method 223
 file arguments 400–402
 file handle (fh) 94–95
 file parameter (argparse) 390–391
 files
 choosing output file handle 104
 printing output 104
 reading 93–97
 reading input from file or command line 103–104
 writing 97–99
 FileType 401–402
 filter() function 45, 147, 227, 240–241, 339–340, 348
 find_brackets() function 292–293
 find_winner() function 356–357, 359
 flag option 31, 391–392
 flags 387
 Flake8 program 28
 float type 82, 167
 floating-point values 237
 foo() function 143
 foo.py program 391
 for loops
 building new lists with 89
 building new strings with 88
 printing each character with 86–88
 reading files using 115–117
 turning into list
 comprehensions 89–90
 format_board() function 356
 formatting
 CSV output 322
 game board 363–364
 list items 73–74
 output 156–157
 regular expressions 242–243
 strings 48, 112–114
 text tables 328–330
 friar.py (Kentucky Friar program) 248–267
 final program 262–263
 fry() function
 using 261–262
 writing manually 264–265
 writing with regular expressions 256–261, 266
 negated shorthand classes 254–255
 re.split() with captured regex 255–256
 shorthand classes 252
 splitting text 251
 writing program 250–262
 fry() function
 using 261–262
 writing manually 264–265
 writing with regular expressions 256–261, 266
 f-strings 40, 49, 52, 113, 159
 functions
 adding type hints to definition 376
 using 186–187
 writing 181–182
 writing tests for 182–186
 functools module 302
 functools.reduce()
 function 302–303

G

gashlycrumb.py (Gashlycrumb Timies program) 118–127
 dictionary comprehensions 125
 dictionary lookups 126
 final program 122–123
 handling arguments 123–124
 reading input file 124–125
 writing program 119–122
 gematria.py (Gematria program) 295–310
 breaking text 304
 cleaning words 297–298
 encoding words 303
 final program 304–305
 functools.reduce()
 function 302–303
 ordinal character values and ranges 298–300
 sorting 308–309
 summing and reducing 300–301
 testing 309
 word2num() function 306–307
 writing program 296–304

`get_args()` function 39, 49–50, 57, 71, 84–85, 100, 102, 104, 119, 121, 123, 134–135, 154, 157, 167–168, 179, 189, 200, 209, 263, 277, 305, 321, 325, 346, 355, 359, 388–389, 402
adding 27–29
defining arguments with 50–51
returning from 392
`getoutput()` function 40
global changes 198
guard statement 240

H

hashable values 161
`hello.py` (Hello, World! program) 15–16
adding `get_args()` function 27–29
adding help messages 22–24
adding `main()` function 26–27
adding parameters 22–24
comment lines 16
error checking 28–29
making argument optional 24–25
making program executable 20
\$PATH variable 20–22
shebang line 18–19
style checking 28–29
testing 17–18, 26, 29–30
help messages
adding 22–24
higher-order function. *See HOF (higher-order function)*
HOF (higher-order function) 147, 240, 302
\$HOME variable 18, 21
`howler.py` (Howler program) 92
defining arguments 102
files
choosing output file handle 104
printing output 104
reading 93–97
reading input from file or command line 103–104
writing 97–99
final program 101–102
low-memory version 104–106
writing program 99–101
hyphen character (-) 40

I

if expressions 47, 52, 142, 202
if/elif/else statements, conditional branching with 70–71
if/else statements 47, 70
immutable strings 130, 165
immutable values 275
imperative methods 225
import statements 154
import sys command 101
in-del (insertion-deletions) 177
index method 64
indexing 63
indexing lists 63–64
input handles 94
input() function 127, 288
input-output (io) module 105
int type 82, 167
int values 153, 166
int() function 322, 326–327
integer values 153, 365
integration tests 186, 383
io (input-output) module 105
io.StringIO() function 105–106, 320
IPython 41, 59, 78
iterator object 68
`ittictactoe.py` (Interactive Tic-Tac-Toe program) 367–382
final program 377–381
state 381
tuples 369–372
type hints 372–373
adding to function definitions 376
type verification with Mypy 373–375
updating immutable structures 375–376
version using TypedDict 379–381
writing program 368–376

J

joining
lists 70
strings 41
`jump.py` (Jump the Five program) 76–91
dictionaries 77–82
accessing values 80
creating 78–79

defining parameters 85
methods 81–82
processing items in series 86–90
using for encoding 85
final program 84–85
writing program 82–84
Jupyter Notebook 41, 59, 78

K

Kentucky Friar program. *See friar.py (Kentucky Friar program)*

L

l33t() function 342–343, 345, 347, 349
lambda keyword 144, 241
lazy functions 140, 180, 317
len() function 45, 60, 63, 72, 81, 112, 116
lines of code (LOC) 189
line.split() method 114, 116
linters 28
list comprehensions
replacing vowels 140–142
turning for loops into 89–90
with functions 142–144
list context 136
list variable 55
list() function 59, 68, 140, 171, 180, 220, 245, 319, 363
list.append() function 60, 140, 142, 202–203, 220
list.extend() method 62, 217, 220
list.index() method 64
list.insert() method 62
list.pop() method 65, 289
list.remove() method 66
list.reverse() method 67, 69, 214, 274
lists 59–70
adding many elements to 61–63
adding one element to 60–61
building with for loops 89
finding elements in 64–65
formatting items 73–74
indexing 63–64
iterating 111
joining 70
mutability of 69–70

lists (*continued*)

- mutating 176–177
- printing items 74
- removing elements from 65–67
- reversing 67–69
- slicing 64
- sorting items 67–69, 73
- `list.sort()` method 67–69, 274, 308
- LOC (lines of code) 189

M

- MAD (mutually assured destruction) doctrine 351
- `mad.py` (Mad Libs program) 281–294
- final program 289–290
 - finding placeholders 284–287
 - finding placeholders without regular expressions 291–293
 - getting values 288–289
 - halting and printing errors 287–288
 - limiting replacements 291
 - substituting text 289
 - writing program 282–289
- main namespace 50
- `main()` function 26–27, 32, 49–51, 72, 85, 104, 120, 136, 143, 158, 160, 168, 181–182, 189, 244, 270, 293, 297, 328
- `make directory (mkdir)` command 21
- `make` program 29
- `make test` command 29, 36, 49, 56, 119, 186
- `Makefile` 29
- `map()` function 45, 144, 146, 179, 186–188, 192, 196, 199, 222, 262, 265, 276, 296, 301, 306, 318, 337, 340, 347, 365
- `ransom.py` 204
- replacing vowels 144–147
- with named functions 147
- `MapReduce` 205
- `match.groups()` function 234, 255
- `Mispar hechrechi` 295
- `mkdir` (make directory) command 21

- multiplication operator (*) 358
- mutable elements 221
- mutating
- game board 360–363
 - lists 69–70, 176–177
 - strings 165–177
 - calculating number of mutations 168–169
 - final program 173–174
 - mutation space 169
 - selecting characters to mutate 169–172
 - writing program 167–173
- `text` 197–198
- updating immutable structures 375–376
- `-mutations` option 166–167
- `-mutations` parameter 174–175
- mutually assured destruction (MAD) doctrine 351
- `Mypy` 373–375
-
- N**
- `-name` option 24–25, 31
- `name` parameter 23
- `-name` value 385
- named options 386
- `nargs` option 71, 398
- natural language processing (NLP) 304
- negative index numbers 43
- `new_char()` function 142–144, 147
- newline character 98–99
- `new.py` file 30, 38, 56
- `new.py`, starting new program with 30–33, 37–38
- NLP (natural language processing) 304
- non-capturing groups 272
- non-deterministic selection 170
- non-greedy regex 285
- no-op (no operation) 273
- numeric encoding of text 295–310
- breaking text 304
- cleaning words 297–298
- encoding words 303
- final program 304–305
- `functools.reduce` 302–303
- ordinal character values and ranges 298–300

- sorting 308–309
- summing and reducing 300–301
- testing 309
- `word2num()` function 306–307
- writing program 296–304
- numeric parameter (`argparse`) 390

O

-
- open file handles 110, 115, 392
- `open()` method 94, 104–105, 135, 215, 401
- operating system (`os`) module 93
- `operator.mul` function 302
- optional arguments 23, 25
- optional parameters 25
- `ord()` function 138, 212, 296, 298–299, 301, 307
- `OrderedDict` 319
- ordinal value 138
- `os` (operating system) module 93
- `os.path.basename()` method 94
- `os.path.dirname()` method 93
- `os.path.isfile()` function 98, 103, 110, 168
- output handles 94
- overwritten files 98

P

-
- `pandas.read_csv()`
- function 327–328
- parallel operations 205
- parameters
- `apples.py` 134–135
 - `argparse`
 - file 390–391
 - numeric 390
 - positional 389
 - string 390
 - `jump.py` 85
 - `parser.add_argument()` function 50
- `parser.error()` function 151, 154, 157–158, 160, 167, 179, 187, 189, 209, 216, 283, 355, 402–403
- pass statement 182

password.py (Secure Password Generator program) 331–350
 cleaning text 337–339, 346–347
 creating unique list of words 335–337
 filtering words 340–341
 final program 343–346
`l33t()` function 342, 347
 processing files 347–348
`ransom()` function 347
 sampling and creating passwords 341–342, 348–349
 title-casing words 341
 using sets 339
 writing program 334–343
`$PATH` variable 20–22
 PCRE (Perl-Compatible Regular Expressions) 252
 period (.) character 232, 235, 284, 353, 361
 Perl-Compatible Regular Expressions (PCRE) 252
 picnic.py (Picnic List program) 55–75
 defining arguments 73
 final program 71–72
 lists 59–70
 adding many elements to 61–63
 adding one element to 60–61
 conditional branching with `if/elif/else` statements 70–71
 finding elements in 64–65
 formatting items 73–74
 indexing 63–64
 joining 70
 mutability of 69–70
 printing items 74
 removing elements from 65–67
 slicing 64
 sorting items 67–69, 73
 starting new program 56–58
 writing program 58–59
 pip module 28
 pipe operator (`|`) 108, 356
 placeholders
 finding with regular expressions 284–287

finding without regular expressions 291–293
 positional arguments 23, 25, 32, 39, 50, 386
 one or more of the same 399–400
 single 393–394
 two different 394–396
 two of the same 398–399
 positional parameters 25, 389
`pprint` module 124
`pprint()` function 124–125
`pprint.pprint()` function 138
`print()` function 42, 59, 74, 84, 86, 88, 98, 116, 121, 124, 140, 157, 159, 210, 222, 278, 287
`printf()` function 113
 printing
 characters with for loops 86–88
 errors 287–288
 game board 356
 list items 74
 output 52, 104, 215, 222
 REPL (Read-Evaluate-Print-Loop) 44
`product()` function 302
 pseudo-random events 150
 PyCharm 15, 29
`python3` 18–19, 243

R

random events
 capitalizing text 195
 comparing methods 204–205
 creating new strings 198–199
 final program 199–200
 flipping coin 198
 iterating through elements in sequence 200–202
`list.append()` function 202–203
`map()` function 204
 mutating text 197–198
 using list comprehensions 203
 using strings instead of lists 203
 writing function to choose letter 202
 writing program 197–199

mutating strings 165–177
 calculating number of mutations 168–169
 final program 173–174
 mutation space 169
 selecting characters to mutate 169–172
 using lists instead of strings 176–177
 writing program 167–173
 reordering middles of words 268–280
 breaking text into lines and words 270–271
 capturing, non-capturing, and optional groups 271–272
 compiling regexes 272
 final program 276–277
 processing text 277–279
 scrambling all words 275
 scrambling words 273–275, 279–280
 writing program 269–275
 word generation 150–164
 constructing insults 162–163
 defining adjectives and nouns 155–156
 defining arguments 159
 exit values 160–161
 final program 157–159
 formatting output 156–157
 importing and seeding random module 154–155
`parser.error()` function 160
`random.seed()` function 161
`range()` function 162
`STDERR` (standard error) 160–161
 taking random samples and choices 156
 throwaway variables 162
 validating arguments 153–154
 writing program 151–157
 random module 161, 166, 313, 343, 352
 random seeds 198
`random.choice()` function 153, 156–157, 163, 170, 174–175, 335
`random.choices()` function 156
`random.getstate()` function 273
`random.randint()` function 314, 321, 323

`random.random()` function 170
`random.sample()` function 156–157, 159, 163, 171, 174, 176, 321, 323, 329, 341–342
`random.seed()` function 154, 156–158, 161, 168, 175, 197, 200, 268, 273, 276, 313, 321, 335, 342, 345
`random.shuffle()` function 269, 273–275
`range()` function 140, 146, 153, 156–157, 159, 162, 180, 186–187, 193, 208–209, 211, 349
`ransom()` function 335, 342, 345–347
`ransom.py` (Ransom Note program) 195
 comparing methods 204–205
 creating new strings 198–199
 final program 199–200
 flipping coin 198
 iterating through elements in sequence 200–202
`list.append()` function 202–203
`map()` function 204
 mutating text 197–198
 using list
 comprehensions 203
 using strings instead of lists 203
 writing function to choose letter 202
 writing program 197–199
 Read the Fine Manual (RTFM) 44
`read()` method 94, 96, 135
`read_csv()` function 320–321, 325, 327, 329
`re.compile()` function 264, 272
`reduce()` function 296, 302–303
 refactoring 149
`re.findall()` function 282, 287, 297
 regular expressions 148–149
 captured 255–256
 compiling 272
`friar.py` 248–267
 final program 262–263
 `fry()` function 256–262, 264–266
 negated shorthand
 classes 254–255
`re.split()` with captured regex 255–256

shorthand classes 252
 splitting text 251
 writing program 250–262
`mad.py` 281–294
 final program 289–290
 finding placeholders 284–287
 finding placeholders without regular expressions 291–293
 getting values 288–289
 halting and printing errors 287–288
 substituting text 289
 substituting with regular expressions 291
 writing program 282–289
`rhymer.py` 225–247
 commenting 242–243
 creating output 238
 creating rhyming strings 244–245
 final program 238–239
 formatting 242–243
 stemmer() function 228–229, 240–241, 243, 245
 truthiness 236–237
 using 229–232
 using capture groups 232–236
 writing `rhymer.py` 227–238
`re.Match` object 260
`re.match()` function 230, 232, 234, 239, 250–251, 266, 272, 327
`re.Match.groups()` method 232
`Repl.it` 20, 29
`requirements.txt` file 313
`re.search()` function 230, 232, 250, 252, 260, 266, 272
`re.split()` function 250, 255–256, 263–264, 270
`re.sub()` function 148, 282, 291, 296, 298, 305, 307, 346
 return statement 182
`reversed()` function 68, 180, 214, 220
 reversing lists 67–69
`rhymer.py` (Rhyme program) 225–247
 commenting 242–243
 creating output 238
 creating rhyming strings 244–245
 final program 238–239

formatting 242–243
`stemmer()` function 240–241
 `rhymer.py` 228–229
 using outside `rhymer.py` 243
 writing without regular expressions 245

truthiness 236–237
 using capture groups 232–236
 using regular expressions 229–232
 writing program 227–238
`round()` function 168
 round-tripping 91
`rstrip()` method 135
 RTFM (Read the Fine Manual) 44

S

`scramble()` function 273, 275, 277–278
`scrambler.py` (Scrambler program) 268–280
 breaking text into lines and words 270–271
 capturing, non-capturing, and optional groups 271–272
 compiling regexes 272
 final program 276–277
 processing text 277–279
 scrambling all words 275
 scrambling words 273–275, 279–280
 writing program 269–275
 secure password
 generation 331–350
 cleaning text 337–339, 346–347
 creating unique list of words 335–337
 filtering words 340–341
 final program 343–346
`l33t()` function 342, 347
 processing files 347–348
`ransom()` function 347
 sampling and creating passwords 341–342, 348–349
 title-casing words 341
 using sets 339
 writing program 334–343
 Secure Password Generator program. *See* `password.py` (Secure Password Generator program)

sep argument 86
 set() function 339
 set.add() function 339
 set.intersection() function 339
 sets 339
 set.sort() function 349
 shebang (#!) line 18–19
 shorthand classes 252, 254–255
 shuffled value 274
 single quotes ("") 42, 157
 slice notation 43
 slices 173
 slicing lists 64
 SNP (single nucleotide polymorphisms) 177
 SNV (single nucleotide variations) 177
 sorted() function 68–69, 174, 245, 345, 349
 sorting
 list items 67–69, 73
 strings 308–309
 spaCy library 332
 square brackets ([]]) 43, 59, 78, 80, 89, 125, 293, 385
 state 351–366, 381
 altering board 355
 determining winner 356–357, 364–365
 final program 357–365
 mutating board 360–363
 printing board 356
 validating arguments 360–363
 validating user input 355
 writing *tictactoe.py* 353–357
 static typing 48
 STDERR (standard error) 160–161, 287, 290, 293
 STDIN (standard in) 110
 STDOUT (standard out) 97, 101, 108, 110, 130, 157, 160, 208, 246, 290
 choosing output file handle 104
 writing files 97–99
 stemmer() function 228–229, 236, 238, 244
 rhymers.py 228–229, 240–241
 using outside rhymers.py 243
 writing without regular expressions 245
 step value 181
 stop value 43, 180, 292

str class 42, 44–45, 60, 63, 85, 90, 96, 134, 172, 183, 197, 303, 320, 394
 str.casefold() function 308
 str.endswith() function 259
 str.format() method 48, 52, 72, 112–113, 157, 363
 string module 169
 string parameter (argparse) 390
 string slices 173
 string.ascii_letters 169
 string.ascii_lowercase 241
 string.punctuation 169, 335
 strings
 altering 130–132
 str.replace() method 131
 str.translate() method 131–132
 building with for loops 88
 comparisons of 45–47
 concatenating 41–42
 creating new 198–199
 formatting 48, 112–114
 getting individual characters of strings 43–44
 methods for 44–45
 mutating 172–173, 175–176
 using instead of lists 203
 str.isupper() function 45
 str.join() method 70, 89, 136, 147, 157, 163, 192, 210, 215, 222, 230, 278, 303, 349
 str.lower() method 49
 str.maketrans() function 132, 137
 str.replace() method 84, 90–91, 131, 171, 175
 altering strings 131
 replacing vowels 136
 str.rstrip() method 97, 103, 124, 133, 315–316
 str.split() method 112, 155, 157, 162, 251, 270, 316, 348
 str.splitlines() function 261, 263, 270, 276
 str.splitlines() method 251, 262
 str.strip() function 356
 str.title() method 334
 str.translate() function 90–91, 139, 345, 347
 str.translate() method 90
 altering strings 131–132
 replacing vowels 137–139
 str.upper() function 44, 49, 97
 sum() function 296, 300, 302, 307
 sys.exit() function 160, 282, 288, 329
 sys.stderr 287
 sys.stdout file 101–102, 208

T

tab character 323
 tabulate module 312–313, 320, 323, 330
 tabulate() function 322, 325
 TDD (test-driven development) 53, 182
telephone.py (Telephone program) 165–177
 calculating number of mutations 168–169
 final program 173–174
 mutating strings 172–173, 175–176
 mutation space 169
 selecting characters to mutate 169–172
 non-deterministic selection 170
 randomly sampling characters 170–172
 using lists instead of strings 176–177
 writing program 167–173
template.py, starting new programs with 387–388
 test-driven development 52–53, 189–190
 testing
 apples.py 149
 gematria.py 309
 hello.py 17–18, 26, 29–30
 writing and testing little by little 38–39
 writing tests for functions 182–186
test.py 17–18, 30
 TextIOWrapper class 96
 Tic-Tac-Toe program. *See* *tictactoe.py* (Interactive Tic-Tac-Toe program)
tictactoe.py (Tic-Tac-Toe program) 351–366
 altering board 355
 determining winner 356–357, 364–365
 final program 357–365

`tictactoe.py` (Tic-Tac-Toe program) (*continued*)
 formatting board 363–364
 mutating board 360–363
 printing board 356
 validating arguments 360–363
 validating user input 355
 writing program 353–357

TIMTOWTDI (There Is More Than One Way To Do It) 135

title-casing 341

truthiness 236–237

tuples 369–372

`twelve_days.py` (Twelve Days of Christmas program) 207–224
 counting 209–210
 creating ordinal value 211–212
 final program 216–217
 generating verses 221–222
 making verses 213–215, 218–221
 printing verses 215, 222
`verse()` function 215
 writing program 208–215

two equal signs (==) 46

type hints 372–373
 adding to function definitions 376
 type verification with `Mypy` 373–375

type() function 42, 45, 60, 78, 80, 97

TypedDict 379–381

U

unbalanced brackets 285
 unit tests 183, 186, 383
`unit.py` file 356
 usage statement 23
 user input validation 355

V

variables
`$PATH` variable 20–22
 throwaway 162
 types of 42
 verbose output 17
`verse()` function 187–189, 193, 210–211, 217, 221
`twelve_days.py` 215
 using 186–187
 writing 181–182
 writing tests for 182–186
`VS Code` 15, 29

W

`wc.py` (Word Count program) 107–117
 defining arguments 115
 defining file inputs 110
 final program 114–115
 formatting results 112–114
 iterating lists 111
 reading files using for loops 115–117
 what's being counted 111–112
 writing program 109–114
 which command 19–20
 wildcard 108
`wod.py` (Workout of the Day program) 311–330
 final program 323–325
 formatting text table output 322, 328–330
 handling bad data 322–323
 parsing delimited text files with `csv` module 318–319
 with `pandas.read_csv()` function 327–328
potential runtime errors 326–327
 reading delimited text files 313–314, 325–326
 creating function to read 320–321
 manually 315–318

selecting exercises 321–322
 writing program 312–323
`wod.read_csv()` function 329
Word Count program. *See wc.py* (Word Count program)
word variable 41–44, 48, 51, 116
`word2num()` function 303–307, 309
`write()` method 96, 98
 writing programs 15–34, 353–357
`abuse.py` 151–157
`bottles.py` 179–187
`crowsnest.py` 49
`friar.py` 250–262
`gashlycrumb.py` 119–122
`gematria.py` 296–304
`hello.py` 15–30
`howler.py` 99–101
`itictactoe.py` 368–376
`jump.py` 82–84
`mad.py` 282–289
`password.py` 334–343
`picnic.py` 58–59, 71
`ransom.py` 197–199
`rhymer.py` 227–238
`scrambler.py` 269–275
 starting new program with `new.py` 30–33
 starting new program with `template.py` 33
`telephone.py` 167–173
`twelve_days.py` 208–215
`wc.py` 109–114
`wod.py` 312–323
 writing and testing little by little 38–39
`wt` (writing text) 104

Y

`YAPF` 28–29

Z

zero-offset indexing 63
`zip()` function 316–317

Tiny Python Projects

Ken Youens-Clark

Who says learning to program has to be boring? The 21 activities in this book teach Python fundamentals through puzzles and games. Not only will you be entertained with every exercise, but you'll learn about text manipulation, basic algorithms, and lists and dictionaries as you go. It's the ideal way for any Python newbie to gain confidence and experience.

The projects are tiny, but the rewards are big: each chapter in **Tiny Python Projects** challenges you with a new Python program, including a password creator, a word rhymer, and a Shakespearean insult generator. As you complete these entertaining exercises, you'll graduate from a Python beginner to a confident programmer—and you'll have a good time doing it!

What's Inside

- Write command-line Python programs
- Manipulate Python data structures
- Use and control randomness
- Write and run tests for programs and functions
- Download testing suites for each project

For readers with beginner programming skills.

Ken Youens-Clark is a Senior Scientific Programmer at the University of Arizona. He has an MS in Biosystems Engineering and has been programming for over 20 years.

To download their free eBook in PDF, ePUB, and Kindle formats,
owners of this book should visit
www.manning.com/books/tiny-python-projects



Free eBook

See first page

“*Tiny Python Projects* is a gentle, amusing introduction to Python that will firm up several key concepts while occasionally making you snicker.”

—Amanda Debler
Schaeffler Technologies

“Knowledge meets humor meets succinctness. A gem.”

—Mafinar Khan, theScore

“Learning based on doing small projects is effective, and that’s why this book is perfect.”

—Marcin Sek, e-Xim IT

“Excellent pick for those who want to improve coding skills with Python.”

—José Apablaza, Steadfast

ISBN: 978-1-61729-751-9



MANNING

\$39.99 / Can \$52.99 [INCLUDING eBook]